

'Types Transformation in Spark'

```
-----
1. Map           --> Realtime
2. FlatMap       --> Realtime
3. Filter        --> Realtime
4. join          --> Realtime
5. groupByKey    --> Realtime
6. reduceByKey   --> Realtime
7. aggregateByKey
8. mapPartition
9. mapPartitionWithIndex
10. coalesce      --> Realtime
11. repartition  --> Realtime
12. cogroup
13. union
14. union all
15. distinct
16. sortBy
17. intersect
18. cartesian
```

Key-value

```
1. aggregateByKey
2. reduceByKey
3. groupByKey
4. sortByKey
5. join
6. cogroup
```

-----28 th May -----

```
val edata = sc.textFile("file:///home/cloudera/emp.txt")
val ddata = sc.textFile("file:///home/cloudera/dept.txt")

val edata_pair = edata.map{ x =>
val w = x.split(",")
val eno = w(0).toInt
val ename = w(1)
val sal = w(2).toInt
val gendar = w(3)
val dno = w(4).toInt
(dno,sal)
}

val ddata_pair = ddata.map { x=>
val w = x.split(",")
val dno = w(0).toInt
val dname = w(1)
val dloc = w(2)
(dno,dloc)
}

val edata_pair_join_ddata_pair = edata_pair.join(ddata_pair)

val result = edata_pair_join_ddata_pair.map{ x =>
val dno = x._1
val sal = x._2._1.toInt()
val loc = x._2._2
val avgsal = sal.sum/sal.size
(dno,loc,avgsal)
}
```

```
val result = edata_pair_join_ddata_pair.map { x =>
val dno = x._1
val sal = x._2._1
val loc = x._2._2
(dno,sal,loc)
}
```

----DISTINCT TRANSFORMATION

```
scala> val data1= sc.parallelize(1 to 10)
data1: org.apache.spark.rdd.RDD[Int] = ParallelCollectionRDD[19] at parallelize at <console>:27

scala>
data2: org.apache.spark.rdd.RDD[Int] = ParallelCollectionRDD[20] at parallelize at <console>:27

scala> val result = data1.union(data2)
result: org.apache.spark.rdd.RDD[Int] = UnionRDD[21] at union at <console>:31

scala> result.collect()
res11: Array[Int] = Array(1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17, 18, 19, 20)
```

'* In Spark Union transformation combine 2 data sets and allow duplicate values*'

```
scala> val result = data1.union(data2).distinct()
result: org.apache.spark.rdd.RDD[Int] = MapPartitionsRDD[25] at distinct at <console>:31

scala> result.collect()
res12: Array[Int] = Array(4, 16, 8, 12, 20, 13, 1, 17, 9, 5, 14, 6, 18, 10, 2, 19, 15, 11, 3, 7)
```

---Its Wide transformation

```
res15: String =
(4) MapPartitionsRDD[33] at distinct at <console>:31 []
| ShuffledRDD[32] at distinct at <console>:31 []
+- (4) MapPartitionsRDD[31] at distinct at <console>:31 []
| UnionRDD[30] at union at <console>:31 []
| ParallelCollectionRDD[19] at parallelize at <console>:27 []
| ParallelCollectionRDD[20] at parallelize at <console>:27 []

scala> val resultsort = result.sortBy( x => x,false)
resultsort: org.apache.spark.rdd.RDD[Int] = MapPartitionsRDD[38] at sortBy at <console>:33

scala> resultsort.collect()
res16: Array[Int] = Array(20, 19, 18, 17, 16, 15, 14, 13, 12, 11, 10, 9, 8, 7, 6, 5, 4, 3, 2, 1)

scala> val resultsort = result.sortBy( x => x,true)
resultsort: org.apache.spark.rdd.RDD[Int] = MapPartitionsRDD[43] at sortBy at <console>:33

scala> resultsort.collect()
res17: Array[Int] = Array(1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17, 18, 19, 20)
```

'Intersect '

```
scala> val data1= sc.parallelize(1 to 10)
data1: org.apache.spark.rdd.RDD[Int] = ParallelCollectionRDD[44] at parallelize at <console>:27

scala> val data2 = sc.parallelize (5 to 20)
data2: org.apache.spark.rdd.RDD[Int] = ParallelCollectionRDD[45] at parallelize at <console>:27

scala>
```

```
scala> val result = data1.intersection(data2)
result: org.apache.spark.rdd.RDD[Int] = MapPartitionsRDD[51] at intersection at <console>:31
```

```
scala> result.collect().foreach(println)
6
8
10
7
9
5
```

```
scala>
```

```
'CROSS PRODUCT '
```

```
scala> val x = sc.parallelize(List(1,2,3,4,5),2)
x: org.apache.spark.rdd.RDD[Int] = ParallelCollectionRDD[60] at parallelize at <console>:27
```

```
scala> val y = sc.parallelize(List('a','b','c','d','e','f'),2)
y: org.apache.spark.rdd.RDD[Char] = ParallelCollectionRDD[61] at parallelize at <console>:27
```

```
scala> x.cartesian(y).collect().foreach(println)
```

```
(1,a)
(1,b)
(1,c)
(2,a)
(2,b)
(2,c)
(1,d)
(1,e)
(1,f)
(2,d)
(2,e)
(2,f)
(3,a)
(3,b)
(3,c)
(4,a)
(4,b)
(4,c)
(5,a)
(5,b)
(5,c)
(3,d)
(3,e)
(3,f)
(4,d)
(4,e)
(4,f)
(5,d)
(5,e)
(5,f)
```

```
scala>
```

```
'Cogroup ' -- If Join and group operates on Same column then use Cogroup
```

```
select dno,avg(sal)
from emp e join dept d
on e.deptno = d.deptno
group by deptno;
```

Data Should be in Key-value pairs

```
val a = sc.parallelize(List((1,'a'),(2,'a'),(3,'a'),(4,'a'),(2,'c'),(3,'c'),(3,'a')),2)
```

```
val b = sc.parallelize(List((1, 'b'), (2, 'b'), (3, 'b'), (4, 'b')), 2)
```

```
'Actions'
```

```
-----
```

1. **collect**
2. **saveAsSequenceFile** -- Recently introduced
3. **saveAsObjectFile**
4. **count**

```
'count'
scala> val x = List(1,2,3,4,5)
x: List[Int] = List(1, 2, 3, 4, 5)

scala> val r = sc.parallelize(x);
r: org.apache.spark.rdd.RDD[Int] = ParallelCollectionRDD[75] at parallelize at <console>:29
```

```
scala> r.count
res23: Long = 5
```

```
scala> r.first
res24: Int = 1
```

```
scala> r.take(2);
res28: Array[Int] = Array(1, 2)
```

```
'Difference between reduce and reduceByKey'
```

```
reduceByKey is a transformation
recude is an action
```

If you **are using** reduceByKey your **data** should be **in Key-value**
 If you **are using** reduce action **key,value is not** mandatory

Reduce :

```
select sum(sal),avg(sal),count(sal) from emp ;
```

```
val rrd1 = sc.parallelize(List(10,20,30,40,50,60,70,80,90),2)
```

```
rrd1.sum // All Partitions data will be collected into local, sum executed at local no parallel  

procession
```

```
rrd1.reduce(_+_ ) // It executed at cluster all separatlely for each partition
```

```
-----29 May
```

Spark **SQL** API Enables spark core functionality **as** well **as** sequel api functionality.

Rules **for** Spark **SQL** API .

1. Spark **SQL** Api Requireds **Sql** context **object** - SQLContext

2. **Data** should be **in** structured **format**
3. Proper **Schema** for your **data**

SQLContext + Structured **Data** + Proper **Schema** = **Data** Frame

Data Frame **is** a **SQL Table**

'How to create a Data frame'

```
select dno,loc,avg(sal),max(sal),min(sal) from emp e join dept d
where e.dno=d.dno
group by dno,dloc
```

```
import org.apache.spark.sql.SQLContext
```

```
val sqlContext = new SQLContext(sc)
```

```
val emp = sc.textFile("file:///home/cloudera/emp.txt")
val dept = sc.textFile("file:///home/cloudera/dept.txt")
```

```
case class Employee (eno:Int,ename:String,sal:Int,gendar:String,dno:Int)
case class Department (dno:Int,dname:String,dloc:String)
```

```
val edata = emp.map{ x =>
val w = x.split(",")
val eno = w(0).toInt
val ename=w(1)
val sal = w(2).toInt
val gendar = w(3)
val dno =w(4).toInt
Employee (eno,ename,sal,gendar,dno)
}
```

```
val ddata = dept.map{ x =>
val w = x.split(",")
val dno = w(0).toInt
val dname=w(1)
val dloc = w(2)
Department (dno,dname,dloc)
}
```

```
--//converting RDD to DataFrame
```

```
import sqlContext.implicits._
```

```
val edf = edata.toDF
val ddf = ddata.toDF
```

```
edf.show()
ddf.show()
```

```
edf.registerTempTable("empview")
ddf.registerTempTable("deptview")
```

```
val eresult= sqlContext.sql("select d.dno,d.dloc,avg(sal) as AVG_SAL ,max(sal) MAX_SAL
,min(sal) MIN_SAL,count(*) COUNT_SAL from empview e join deptview d on e.dno=d.dno group by
d.dno,d.dloc")
```

```
val eresult= sqlContext.sql("select d.dno,d.dloc from empview e join deptview d on e.dno=d.dno")
```

In Spark 1.x --> SQL Context allows only DRL (Data Retrieval only select statement)
 In Spark 2.x --> we can write DDL Statements , DML Statements

Spark 1.x (1.6.0/1.6.2/1.6.3)

Spark2.x (2.1.0 / 2.3x)

1. Spark core Context (RDD)
 2. SQLContext (SQL Api, supports DRL)
 context , SSC

1. In Spark 2.x only 1 context Spark Session
 Spark Session -- Spark Context,SQL

3. Hive Context (Spark + Hive)
 (Integration,
 Support DDL,DML,DRL

)
 4. Spark StreamingContext
 (
 SSC , Supports DStream
)

'Spark + Hive Integration using Hive Context '

1. Copy

-----31st May -----

'How to enable Hive context in Spark 2.x' --***** V Imp

Method :
 enablehivesupport

Warehouse Directory :

cp /usr/lib/hive/conf/hive-site.xml /usr/lib/spark/

```
val spark= SparkSession
    .builder()
    .appName("Spark Hive Example")
    .config("spark.sql.warehouse.dir","warehouseLocation")
    .enablehivesupport
    .getOrCreate()
```

'Why DataFrames are very Powerful'

DataFrame = RDD + Catalyst Optimizer + DAG + in Memory

DataSet = RDDs + Catalyst Optimizer + CPU Caches

CBO = Cost Based Optimizer

Catalyst Optimizer internally uses CBO

Catalyst Optimizer -- Read more about it

Catalog represents a serializer .

serialization **and** De-serialization -- How to read and write a data.

'Different serialization available in Spark '

Spark supports Different types **of** serialization

1. **Java** serialization --default
2. Kyro serialization -- Advance to Java
3. avro serialization -- Read and write data of avro types
4. **Sequence** serialization -- Read and write data of sequence serialization
5. Parquet serialization --Read and write parquet serialization

How **to** enable

```
val conf = new SparkConf()
    .set("spark-serializer","org.apache.spark.serializer.KryoSerializer")
```

It re-commendable **to use** kryo serialization **for** performance

'#Fetching the Data from RDS using Spark DataFrames'

RDS

cloud

```
mysql -----*****-----mysql
                        |
                        |
                        |
                        Spark
```

mysql --> mysql via AWS Virtual private cloud

Spark cannot **read** the **data from** cloud because **of 2** reasons

1. Firewalls
2. DNS

Spark **read data from local** mysql

Services --> Database --> RDS (Relational Data Services)

--> Launch DB Instance

--> MySQL

* **Check** the **check** box --> only enable options eligible for free usage type

-> **Next**

--> Check If Free Usage

Settings

1. DB Instance Identifier -- Unique Name for the database (mysqldb)
2. username **and** password

Configure Advanced settings

1. **Default** VPC --> Keep it default
2. **Public**
 - yes --> Access outside of the cloud
 - no --> Access outside of the cloud
3. Database **option**

Database Name : testdb

Port **No** : 3304

Launch

Install Client Machines

1. **SQL** WorkBench

Download generic **package for all** systems

copy **End** Point information

Enabling -- In bound and outbound connection

--> Click Security Group
 --> Inbound
 --> Add Rule

Type : MySQL/Aurora
 Source : Anywhere

Save

--> Outbound type

Edit --> Add Rule --> MySQL/Aurora
 Desitination --> Anywhere

Step **no** : 3 **Open SQL** WorkBench

Give DB Name

MySQL JDBC Jar Download -- Platform Independent

Manage Driver --> add Driver

Copy URL **and replace with** AWS information

'1-June'

Reading the **data from** RDS **using** Spark **SQL** Engine

IntelliJ

File Menu

Project structured

Modules

Dependencies

Right Most corner

Jars

mysql

'4-June'

'How to Handle csv file format, JSON file format and xml file formats using spark sql'

Spark packages (**2.x**)

csv
 xml
 json
 html
 rcv
 etc


```
import spark.implicits._ --> To convert RDD To DataFrame
```

```
1.x
```

```
//registerTempTable --> Available upto that session
```

```
2.x
```

```
createOrReplaceTempView == registerTempTable
```

```
createOrReplaceGlobalTempView --> Available even after closing the spark sql
```

```
'Handling csv file '
```

```
import org.apache.spark.sql.Session
```

```
object SparkCsv {
```

```
  def main(args: Array[String]): Unit = {
```

```
    val spark = SparkSession.builder.master("local[*]").appName("SparkCsv").getOrCreate()
    //local represents local machine & star represents no of resources i.e utilising all resources
    //getOrCreate => creating a application or using already existed application
    //appName => creating an application with name "csvExample_1"
```

```
    //creating 2 context sparkContext and sqlContext
```

```
    val sc = spark.sparkContext
```

```
    val sqlContext = spark.sqlContext
```

```
    import spark.implicits._
```

```
    //Creating DataFrame using SQLContext Reading data)
```

```
    val df = sqlContext.read.format("com.databricks.spark.csv").option("header", "true").option(
      "inferSchema", "true").load("C:\\Users\\NaraVish\\Desktop\\#Personal\\#Interview
      Documents\\filformatsinspark\\Police_Department_Incidents.csv")
```

```
    //option("header", "true") => Use first line of all files as Header
```

```
    //option("inferSchema", "true") => Automatically infer data types
```

```
    df.show()
```

```
    //df.select("Category").distinct.collect().foreach(println)
```

```
    df.createOrReplaceTempView("sfpd") //Creating Temp table
```

```
    // sqlContext.sql("select Category from sfpd").collect().foreach(println)
```

```
    //top 10 results
```

```
    //sqlContext.sql("SELECT Resolution , count(Resolution) as rescount FROM sfpd group by
    Resolution order by rescount desc limit 10").collect().foreach(println)
```

```
    val t = sqlContext.sql("select Category, count(Category) as catcount from sfpd group by
    Category order by catcount desc limit 10")
```

```
    t.show()
```

```
    t.map(t=> "column 0: " + t(0)).collect().foreach(println)
```

```
    spark.stop()
```

```
  }
```

```
}
```

```
'Handling of Json files using spark SQL'
```

```
Json is inbuilt feature of spark
```

```
sqlContext.read.json("Json_path")
```

```
import org.apache.spark.sql.Session
```

```
object JsonExample {
```

```
def main(args: Array[String]): Unit = {

    val spark = SparkSession.builder.master("local[*]").appName("JsonExample").getOrCreate()
    val sc = spark.sparkContext
    val sqlContext = spark.sqlContext

    //Converting RDD to Data Frame
    import spark.implicits._
    import spark.sql
    val df = sqlContext.read.json("C:\\Users\\NaraVish\\Desktop\\#Personal\\#Interview
Documents\\filformatsinspark\\world_bank.json")
    df.printSchema() //printing schema
    df.createOrReplaceTempView("jsondata_one")
    df.show()

    //val result = sqlContext.sql("select url,totalamt,abc.* from jsondata_one " + "lateral
view explode(theme_namecode) as abc")
    val result = sqlContext.sql("select _id from jsondata_one")
    result.show(10)

    //    println(result)

    //result.write.format("com.databricks.spark.csv").option("header","true").save(
"C:\\Users\\sonirai\\Desktop\\Hadoop GV\\Spark\\SparkSQL\\datasets\\jsontocsv")
    spark.stop()
}
}
```

'How to handle XML Files in Spark SQL '

```
import org.apache.spark.sql.SparkSession
object XMLExample {
```

```
    def main(args: Array[String]): Unit = {

        val spark = SparkSession.builder.master("local[*]").appName("xmlExample_1").getOrCreate()
        val sc = spark.sparkContext
        val sqlContext = spark.sqlContext
        val df = sqlContext.read.format("com.databricks.spark.xml").option("rootTag","books").option
("rowTag","book").load("C:\\Users\\NaraVish\\Desktop\\#Personal\\#Interview
Documents\\filformatsinspark\\books.xml")
        df.show()
        //    df.printSchema()
        //df.createOrReplaceTempView("jsondata_one")
        //sqlContext.sql("select url,totalamt,abc.code,abc.name from jsondata_one " +
        //"lateral view explode(theme_namecode) as abc").show(5)
        //    result.write.format("com.databricks.spark.csv").option("header","true").save(
        "C:\\Kalyan\\POC\\Spark\\spark_datasets\\json_data\\jsontocsv")
        spark.stop()
    }
}
```

Major Sources of Data Frame :

1. RDS
2. CSV
3. JSON
4. XML

Spark Core **is** faster **than** Spark **SQL** Because Spark **Sql** another layer **on** Spark Core

DATA frame = **Schema** + Structured + RDD + Cost Based Optimzer

Data Set = **Schema** + Structured + RDD + Cost Based Optimzer + CPU Caching

Spark **SQL** writes the **data to** Driver Node.


Two issues **with** above

1. Memory Issues
2. Storage issue **in** Driver Node

Data Frame **is** subset **of Data Set**.

CPU Cache :

NVDR



Data Set Using one more serializer encoder (along **with Java or** Kyro)
It **is** 10 times faster **than** kyro serializer

4040 :

Data Frame DSL : **Domain Specific Language**

Read About Encoder serializer

RDDs :

1. Functional Programming
2. **Type** Safe

DataFrames

1. Relational
2. Catalyst Query Optimization
3. Tungsten direct/packed RAM
4. JIT Code Generation
5. Sorting/shuffling **without** De-serialization

CPU Cache **is** also called Tungsten **in-memory in** Spark

Data Frame Operations

1. First

It **returns** the **first row**
df.**first()** --> Action in RDD internally

2. Take --> It randomly collects the data

It you want **to** display **first** n **number of rows use** take
df.**take(8)**

3. Head()

Returns the first n rows in the form of rows and array Format.

```
df.head(5)
```

4. collect

Returns an array that contains all of rows in this DataFrame

Its similar to head(all), if the data is too large, Its not re-commendable

```
df.collect()
```

Take randomly takes the record with ordering , head first sorts and show the record so degrade the performance

5. count

Returns the number of rows in the DataFrame

```
df.count()
```

6. show

Displays the table on scree

```
df.show(5)
```

7. printSchema

'DSL' --> Domain Specific Language

Is the Native Language of Data Frame. Its is faster.

Its only avaiable for Retrieval purpose

```
import org.apache.spark.sql.SQLContext

val sqlContext = new SQLContext(sc)

val emp = sc.textFile("file:///home/cloudera/emp.txt")
```

```
val dept = sc.textFile("file:///home/cloudera/dept.txt")
```

```
case class Employee (eno:Int,ename:String,sal:Int,gendar:String,dno:Int)
case class Department (dno:Int,dname:String,dloc:String)
```

```
val edata = emp.map{ x =>
val w = x.split(",")
val eno = w(0).toInt
val ename=w(1)
val sal = w(2).toInt
val gendar = w(3)
val dno =w(4).toInt
Employee (eno,ename,sal,gendar,dno)
}
```

```
val ddata = dept.map{ x =>
val w = x.split(",")
val dno = w(0).toInt
val dname=w(1)
val dloc = w(2)
Department (dno,dname,dloc)
}
```

```
--//converting RDD to DataFrame
```

```
import sqlContext.implicits._
```

```
val edf = edata.toDF
val ddf = ddata.toDF
```

```
ddf.show()
edf.collect()
edf.take(10)
edf.count()
edf.first()
```

```
---DSL
```

```
edf.select($"eno",$"ename",$"sal",$"gendar",$"dno").show()
edf.select($max("sal").show()
```

```
*** msck repair table
```

```
val dataset = Seq(1,2,3).toDS()
dataset.show()
```

```
case class Person(name:String , age :Int)
```

```
val personDS = Seq (Person("Max",33),Person("Adam",32),Person("Muller",62)).toDS()
personDS.show()
```

```
---
val rdd = sc.parallelize(Seq((1,"Spark"),(2,"databricks")))
val integerDS= rdd.toDS()
integerDS.show()
```

```
--- converting DF to Data Set
```

```
case class Company (name : String,foundingYear :Int,numEmployees :Int)
```

```
val inputSeq = Seq(Company("ABC",1998,310),Company("XYZ",1998,310),Company("NOP",1998,310))
val df = sc.parallelize(inputSeq).toDF()
```

```
val companyDS = df.as[Company]
```

```
companyDS.show()
```

```
--- 4 steps in 1.x
```

```
import org.apache.spark.sparkContext
import org.apache.spark.SQLContext
```

```
val sc = new sparkContext
```

```
val sqlContext= new SQLContext(sc)
```

```
--**Download this file winutils.exe**
```

```
C : \
create Directory winutils
    create bin Directory bin
        copy winutils.exe
```

Spark Important

1. What **is** Spark
2. Why Spark **is** faster **than** Error
3. **Difference between** Hadoop **and** Spark
4. What **is** RDD
5. What **is** lenience
6. How Fault tolerance works **in** Spark
7. RDD - RESILIENT DISTRIBUTED DATATASET
8. What **is** RDD **and** what **is not** RDD
9. RDD Property
 - a. List **of** Partititions
 - b. Compute Functions
 - c. **If** you **call function in** RDD it will perform **for all** elements **not to specific**
 - d. List **of** Dependencies
 - e. Main **function of** RDD **is** keep track **of operation not** the **data**.
 - f. logical plan **!=** Actual plan
10. Dependencies
 - a. Narrow Dependencies
 - b. Wide Dependencies
11. Intermediate **data is** cached but **not** the RDDs itself

Google --> Apache kafka

kafka.apache.org

Download

0.10.2.1

scala 2.11 --> kafka_2.11.0.10.2.1.tgz --> Download this

**** msck repair

Apache Kafka

```
Web Services 2000-03
JMS (Java Messaging Services) 2005-07
    -- only 1 queue but not distributed
TIBCO / Web Methods / Web Speher / Active MQ / Rabbit MQ -- >2008 (Web brokers)
    --n queues but only 1 queue is distributed
Apache Kafka
    -- Uses a queue as well as broker which uses n queues and all are distributed
```

Apache Kafka **is** Independent System , its **not** dependent **on** Hadoop.
Its runs **without** Hadoop. But we can integrate Kafka **to** Hadoop **or** Kafka **to** spark.

Apache Flume Architecture :

1. Flume **is** part **of** Hadoop
2. Apache flume **is** a **real time log** processing (**Data** Ingestion) technique , it **contains** Components **like**
 - a. Agents
 - b. source
 - c. sink
 - d. channel
 - e. Event
 - f. Interceptor
 - g. channel selector

```
Source      --> Aggent --> Events E1,E2 -- Sink
                -----channel-----
```

** Learn Flume Architecture **in** Detail

Drawback : **if** we **are** selecting multiple Sourcess **and** multiple sinks **then** flume agent will fail
There **is no** backup **for** Source,sink **and** channel.

If we doing more sensitive **data then** flume **is not** usefull.

1.7 Flume has secondary channel

Kafka **is less** security.

Other Systems :

1. Event Hub (Kafka + Security) -- Microsoft
2. AWS - kinesis (Kafka + Security)

Kafka Architecture :

1. Producer --no limits for sources
2. consumer --no limits for targets
3. Kafka **Cluster**
 - a. Kafka Brokers
 - i. Topic (logical, phycally it called partitions)

Each Topic **contains** N **number of** Partitions
Kafka maintains **order using index** called offset.
Kafka Replication

Producers

1. IOT
2. WebApp
3. NoSQL DBs
4. RDBMS
5. Sensor **Data** (Telemetric System)

Producers

1. IOT
2. WebApp
3. NoSQL DBs
4. Spark
5. Hadoop
6. Filesystem
7. cloud
8. RDBMS

'Communication between Producers to Kafka broker'

By Defaultly communication **between** producer **to** Kafka broker **is** synchronous communication

sync communication : **Any** system that sends acknowledge **to** source **is** sync system

Its **not** re-commendable **as** it increases throughput

'Asynchrononus communication'

There **is no** communication **between** producer **and** kafka broker.

Mostly it will be Asynchrononus .

Realtime Asynchrononus communication **is** used .

Disadvantage **with** Kafka :

Very **Less** Security
dependent **on** ZOOKEEPER

In Latest version **of** Kafka **1.x** there **is no** Dependencies **on** zookeer
(bootstap servers)

kafka_2.11-1.1.0.tgz

```
[cloudera@quickstart ~]$ cd kafka_2.11-1.1.0
```

```
[cloudera@quickstart kafka_2.11-1.1.0]$ ls
```

```
bin  config  libs  LICENSE  NOTICE  site-docs
```

```
[cloudera@quickstart kafka_2.11-1.1.0]$ ls -l
```

```
total 52
```

```
drwxr-xr-x 3 cloudera cloudera 4096 Mar 23 15:54 bin
```

```
drwxr-xr-x 2 cloudera cloudera 4096 Mar 23 15:54 config
```

```
drwxr-xr-x 2 cloudera cloudera 4096 Jun 11 20:26 libs
```

```
-rw-r--r-- 1 cloudera cloudera 28824 Mar 23 15:51 LICENSE
```

```
-rw-r--r-- 1 cloudera cloudera 336 Mar 23 15:51 NOTICE
```

```
drwxr-xr-x 2 cloudera cloudera 4096 Mar 23 15:54 site-docs
```

```
[cloudera@quickstart kafka_2.11-1.1.0]$
```

In Kafka **Bin** Directory **contains** list **of all** kafka services

Producer **start &** stop

borker **start and** stop

consumer **start and** stop etc

By Defaultly **all** the services available **in bin** directory takes **default** properties

In config directory **contains default** properties **like** server.properties , zookeeper.properties, producer.properties & consumer.properties

-----14th June-----

server properties --> Retension policy (Default 24*7 = 168 Hours)

Topic **is** replicated **and** it **is** taken care **of cluster**

IN ZOOKEEPER.PROPERITES **file contains 2** imp properties

1. **Data** Directory (dataDir)
Stores Zooker **log** information ,
Default path of log directory **is** /tmp/zookeeper
2. clientPort
The **Default** port **no of** zookeeper **is** 2181

dataDir=/tmp/zookeeper
the port **at** which the clients will **connect**
clientPort=**2181**

QuorumPeerMain --> This means your zookeeper Running (Internally using 2181)

How **to** kill the service := kill -9 Process_id

1. Change port **no in** zookeeper properties
2. root **user**
start zookeeper server
kafka/bin
./zookeeper-server-start.sh ../config/zookeeper.properties
3.
server.properties **file contains 4** major **parameter**
 1. broker.id
 2. broker ip address (listener)
 3. broker **log** Directory
 4. zookeeper ip address **and** port **no**

9092 is default port **no of** a kafka broker

4. **Start** zookeeper service

./kafka-server-start.sh ../config/server.properties &

5. jps

kafka services will be started

- 6.

```
[cloudera@quickstart bin]$ cd
[cloudera@quickstart ~]$ cd kafka_2.11-1.1.0
[cloudera@quickstart kafka_2.11-1.1.0]$ cd bin/
[cloudera@quickstart bin]$
```

```
./kafka-topics.sh --create --zookeeper localhost:2182 --partitions 2 --replication-factor 1
--topic test_20180613
```

```
./kafka-topics.sh --list --zookeeper localhost:2182
```

```
[cloudera@quickstart bin]$ pwd
/usr/lib/zookeeper/bin
```

```
[cloudera@quickstart bin]$ ./zkCli.sh -server localhost:2182
Connecting to localhost:2182
2018-06-12 20:18:53,474 [myid:] - INFO [main:Environment@100] - Client environment:zookeeper.
version=3.4.5-cdh5.12.0--1, built on 06/29/2017 11:30 GMT
2018-06-12 20:18:53,479 [myid:] - INFO [main:Environment@100] - Client environment:host.name=
quickstart.cloudera
2018-06-12 20:18:53,479 [myid:] - INFO [main:Environment@100] - Client environment:java.version
=1.7.0_67
2018-06-12 20:18:53,484 [myid:] - INFO [main:Environment@100] - Client environment:java.vendor=
Oracle Corporation
2018-06-12 20:18:53,484 [myid:] - INFO [main:Environment@100] - Client environment:java.home=/
usr/java/jdk1.7.0_67-cloudera/jre
2018-06-12 20:18:53,484 [myid:] - INFO [main:Environment@100] - Client environment:java.class.
path=/usr/lib/zookeeper/bin/./build/classes:/usr/lib/zookeeper/bin/./build/lib
/*.*jar:/usr/lib/zookeeper/bin/./lib/slf4j-log4j12.jar:/usr/lib/zookeeper/bin/./lib/slf4j-log4j1
2-1.7.5.jar:/usr/lib/zookeeper/bin/./lib/slf4j-api-1.7.5.jar:/usr/lib/zookeeper/bin/./lib/netty
-3.10.5.Final.jar:/usr/lib/zookeeper/bin/./lib/log4j-1.2.16.jar:/usr/lib/zookeeper/bin/./lib/jl
ine-2.11.jar:/usr/lib/zookeeper/bin/./zookeeper-3.4.5-cdh5.12.0.jar:/usr/lib/zookeeper/bin/./sr
c/java/lib/*.*jar:/usr/lib/zookeeper/bin/./conf:
2018-06-12 20:18:53,484 [myid:] - INFO [main:Environment@100] - Client
environment:java.library.path=/usr/java/packages/lib/amd64:/usr/lib64:/lib64:/lib:/usr/lib
2018-06-12 20:18:53,484 [myid:] - INFO [main:Environment@100] - Client
environment:java.io.tmpdir=/tmp
2018-06-12 20:18:53,484 [myid:] - INFO [main:Environment@100] - Client
environment:java.compiler=<NA>
2018-06-12 20:18:53,484 [myid:] - INFO [main:Environment@100] - Client environment:os.name=Linux
2018-06-12 20:18:53,484 [myid:] - INFO [main:Environment@100] - Client environment:os.arch=amd64
2018-06-12 20:18:53,486 [myid:] - INFO [main:Environment@100] - Client
environment:os.version=2.6.32-573.el6.x86_64
2018-06-12 20:18:53,486 [myid:] - INFO [main:Environment@100] - Client
environment:user.name=cloudera
2018-06-12 20:18:53,486 [myid:] - INFO [main:Environment@100] - Client
environment:user.home=/home/cloudera
2018-06-12 20:18:53,486 [myid:] - INFO [main:Environment@100] - Client
environment:user.dir=/usr/lib/zookeeper/bin
2018-06-12 20:18:53,488 [myid:] - INFO [main:ZooKeeper@438] - Initiating client connection,
connectString=localhost:2182 sessionTimeout=30000
watcher=org.apache.zookeeper.ZooKeeperMain$MyWatcher@5cd4927f
Welcome to ZooKeeper!
```

```
2018-06-12 20:18:53,639 [myid:] - INFO
[main-SendThread(localhost:2182):ClientCnxn$SendThread@975] - Opening socket connection to
server localhost/127.0.0.1:2182. Will not attempt to authenticate using SASL (unknown error)
JLine support is enabled
2018-06-12 20:18:53,679 [myid:] - INFO
[main-SendThread(localhost:2182):ClientCnxn$SendThread@852] - Socket connection established,
initiating session, client: /127.0.0.1:33792, server: localhost/127.0.0.1:2182
2018-06-12 20:18:53,706 [myid:] - INFO
[main-SendThread(localhost:2182):ClientCnxn$SendThread@1235] - Session establishment complete
on server localhost/127.0.0.1:2182, sessionId = 0x163f6fde80d0004, negotiated timeout = 30000
```

WATCHER::

```
WatchedEvent state:SyncConnected type:None path:null
[zookeeper: localhost:2182(CONNECTED) 0]
```

*/

```
[zk: localhost:2182(CONNECTED) 2] ls brokers/
Command failed: java.lang.IllegalArgumentException: Path must start with / character
[zk: localhost:2182(CONNECTED) 3] ls /brokers
[seqid, topics, ids]
[zk: localhost:2182(CONNECTED) 4] ls /brokers/ids
[0]
[zk: localhost:2182(CONNECTED) 5] ls /brokers/topics
[test_20180613]
[zk: localhost:2182(CONNECTED) 6] ls /brokers/seqid
[]
```

--Start Producer as normal user

There are two types of producers in kafka

1. console producer (Its default producer)
2. custom producer (End User create custom producer by using producer api option)

'Step 5'

```
./kafka-console-producer.sh --broker-list localhost:9092 --topic test_20180613
```

'Step 6'

Two Types of consumer

1. console consumer (Its Default consumer)
2. custom consumer (End user create custom consumer by using api consumer api)

```
./kafka-console-consumer.sh --bootstrap-server localhost:9092 --topic test_20180613
--from-beginning
```

-----14th June-----

```
./zookeeper-server-start.sh ../config/zookeeper.properties &
```

```
cp server.properties server_2.properties
```

```
350 gedit server.properties
351 cp server.properties server_1.properties
352 gedit server.properties
353 gedit server_1.properties
354 gedit server_2.properties
355 mkdir /tmp/kb0
356 mkdir /tmp/kb1
357 jps
358 exit
359 ls
360 ./zookeeper-server-start.sh ../config/zookeeper.properties &
361 jps
```

```
./kafka-server-start.sh ../config/server_1.properties &
368 ./kafka-server-start.sh ../config/server_2.properties &
369 jps
370 ./kafka-server-start.sh ../config/server_1.properties &
371 ./kafka-server-start.sh ../config/server_2.properties &
```

```
[zk: localhost:2182(CONNECTED) 0] ls /brokers
[seqid, topics, ids]
[zk: localhost:2182(CONNECTED) 1] ls /brokers/ids
[2, 1, 0]
```

```
[zk: localhost:2182(CONNECTED) 2]
```

'Command to Identify who is the leader and who is slave'

```
/kafka-topics.sh --desc --zookeeper localhost:2182 --topic test_20180614
```

```
0 --> Means Leader
```

```
1 --> Follower 1
```

```
2 --> Follower 2
```

Isr means **in**-sync-replica

Partitions wise Leader **is** selected **not** topic wise

```
-----
```

Custom Producer

MySQL DB --> Custom Producer --> Kafka Broker --> Console
(producer api)

```
import java.util.Properties;
import java.sql.*;
import kafka.javaapi.producer.Producer; // send method
import kafka.producer.KeyedMessage; // Serializer
import kafka.producer.ProducerConfig; // configuration where is my broker list

public class JdbcProducer{
public static void main(String[] args) throws ClassNotFoundException,SQLException{
    Properties props = new Properties();
    props.put("zk.connect","localhost:2182");
    props.put("serializer.class","kafka.serializer.StringEncoder");
    props.put("metadata.broker.list","localhost:9092");
    ProducerConfig config = new ProducerConfig(props);
    Producer producer = new Producer(config);
    try{
        class.forName("com.mysql.jdbc.Driver")
        Connection con = DriverManager.getConnection(
            "jdbc:mysql://localhost:3306/test","root","root");

        // test is the DB name
        Statement stmt = con.createStatement();
        ResultSet rs = stmt.executeQuery("select * from emp");
        while(rs.next())
            producer.send(new KeyedMessage("test",rs.getString(1)+" "+rs.getString(
                2)));
        con.close()
    }catch (Exception e){
        System.out.println(e)
    }
}
```

Key Serializer

Value Serializer

we need **to** copy mysql jar **to** kafka/lib

'How to start kafka '

```
./kafka-server-start.sh ../config/server.properties &
```

```
./kafka-topics.sh --create --zookeeper localhost:2182 --partitions 2 --replication-factor 1
--topic testodbc
```

```
./kafka-topics.sh --list --zookeeper localhost:2182
```

```
./kafka-console-consumer.sh --bootstrap-server localhost:9092 --topic testodbc --from-beginning
```

'Custom Consumer'

Major Consumer

```
Kafka + Spark
Kafka + Storm
Kafka + Filesystem
```

'Integration of Kafka with Spark'

Kafka	Spark
(Near Real Time)	(Batch + Inmemory)

```
Spark Streaming
(Micro Batch)
```

Spark Streaming internally **using** a concept **of** Micro batch processing

Batch --> Static and Historical Data , not transactional Data (Days - Months - Years)

Micro Batch --> Not a static data , -- Almost Transactional Data

```
-- Seconds
-- Minutes
-- Rarely Hours
```

Spark Streaming internally **using** a concept **of** D-Streams

D-Streams --Discretized Data Streams

D-Streams internally working **on** RDDs **only**.

Spark Streaming always expects 2 threads

1. Spark Streaming Context **Object** (ssc)
2. Spark Context

P -- > Kafka --> Spark Streaming Context --> D Streams --> Spark Context

D-Streams works Inmemory **and** reside **in** Worker Node

```
spark-shell --master local[2]
```

```
import org.apache.spark._
import org.apache.spark.streaming._
import org.apache.spark.streaming.StreamingContext._
```

```
val ssc = new StreamingContext(sc,Seconds(10))
```

```
val lines = ssc.socketTextStream("localhost",9999)
val words = lines.flatMap(x => x.split(" "))
val pairs = words.map(x => (x,1))
val res = pairs.reduceByKey(_+_ )
res.print()
ssc.start
```

```
nc -lk 9999
```

```
nc --> net cat Server
-lk --> localhost
9999 --> port no
```

'Integration of Kafka to Spark '

```

                                KF
Console                        --> (topic : Input) --> SSC          -->      SC
Producer                      |
                                |
--Console Consumer <-- Output  <-- (topic :output )             <--      SC Custom Producer
```

1. Start the ZooKeeper

```
kafka/bin
./zookeeper-server-start.sh ../config/zookeeper.properties
```

2. Start Kafka Broker

```
kafka/bin
./zookeeper-server-start.sh ../config/server.properties
```

3. Create topics

```
./kafka-topics.sh --create -zookeeper localhost:2182 --partitions 1 -replication-factor 1
topic-input
./kafka-topics.sh --create -zookeeper localhost:2182 --partitions 1 -replication-factor 1
topic-output
```

4. Start console producer

```
kafka/bin
./kafka-console-producer.sh --broker-list localhost:9092 --topic input
```

'* Kafka Utils establish between kafka broker to Spark Streaming context*'

5.

'*RDDs --> Partitions --> Records*'

```
spark-shell --master local[2]
```

```
import org.apache.spark._
import org.apache.spark.streaming._
import org.apache.spark.streaming.StreamingContext._
```

```
val ssc = new StreamingContext(sc,Seconds(10))
val lines = ssc.socketTextStream("localhost",9999)
val words = lines.flatMap(x => x.split(" "))
val pairs = words.map(x => (x,1))
val res = pairs.reduceByKey(_+_ )
res.print()
ssc.start
```

```
nc -lk 9999
```

```
-----
kafka - Spark Streaming Example
```

- 1) please **start** zookeeper
`./bin/zookeeper-server-start.sh config/zookeeper.properties &`
- 2) please **start** kafka-Broker
`./bin/kafka-server-start.sh config/server.properties &`
- 3) spark-shell --master local[2]
- 4) **create** two topics mytopic,results
`./bin/kafka-topics.sh --create --zookeeper localhost:2182 --partitions 1 --replication-factor 1 --topic mytopic`

`./bin/kafka-topics.sh --create --zookeeper localhost:2182 --partitions 1 --replication-factor 1 --topic results`

`./bin/kafka-topics.sh --list --zookeeper localhost:2182`
- 5) please **start** console producer **with** mytopic
`./bin/kafka-console-producer.sh --broker-list localhost:9092 --topic mytopic`
- 6) **create** another **to write** the spark streaming **output data**(topic name **result**)
- 7) please **start** console consumer **to read** the **data from result** topic.
`./bin/kafka-console-consumer.sh --bootstrap-server localhost:9092 --topic results --from-beginning`

```
spark-shell --master local[2]
8) import org.apache.spark.streaming.StreamingContext
import org.apache.spark.streaming.Seconds
val ssc = new StreamingContext(sc, Seconds(5))
import org.apache.spark.streaming.kafka.KafkaUtils
val kafkaStream = KafkaUtils.createStream(ssc,"localhost:2182",
"spark-streaming-consumer-group",Map("mytopic" ->5))
val lines = kafkaStream.map(x => x._2.toUpperCase)
val words = lines.flatMap(x => x.split(" "))
val pairs = words.map(x => (x,1))
val res = pairs.reduceByKey(_+_ )
import org.apache.kafka.clients.producer.ProducerConfig
import java.util.HashMap
import org.apache.kafka.clients.producer.KafkaProducer
import org.apache.kafka.clients.producer.ProducerRecord
res.foreachRDD(rdd =>
  rdd.foreachPartition(partition =>
    partition.foreach{
      case (w:String,cnt:Int)=>{
        val x = w+"\t"+cnt
        val props = new HashMap[String,Object]()
        props.put(ProducerConfig.BOOTSTRAP_SERVERS_CONFIG,"localhost:9092")
        props.put(ProducerConfig.VALUE_SERIALIZER_CLASS_CONFIG,
"org.apache.kafka.common.serialization.StringSerializer")
        props.put(ProducerConfig.KEY_SERIALIZER_CLASS_CONFIG,
"org.apache.kafka.common.serialization.StringSerializer")

import org.apache.kafka.clients.producer.KafkaProducer
import org.apache.kafka.clients.producer.ProducerRecord
val producer = new KafkaProducer[String,String](props)
val message = new ProducerRecord[String,String]("results",null,x)
producer.send(message)

      }
    }
  )
)

8)please start streaming context
ssc.start
```

Accumulator --21st June
 Broadcast **Variable** --21st Jun
 Spark Memory Management -- 22rd June

Sqoop (Mon **and** Tue) --25th,26th June
 Flume (Wed **and** Thurs) --27th,28th June
 Pig (Friday , Mond, Tue) --29th,2nd and 3rd July
 MR (Wed **to** Friday) -- 4th,5th,6th July
 Oozie (Mon **&** Tue) --9th July , 10th July
 Hbase (Wed **&** Thurs) --12th ,12th July
 MongoDB (Frid **and** Mon) -- 13th,16th July

Hive SerDe **&** UDF (Tues) --17th July
 GIT **and** Bitbucket (Wed) --18th July
 Challenging Prod (Thursday **&** Friday) --19th & 20th July

'Shard variables in Spark'

In Spark two types **of** shared variables

1. Broadcast variables
2. Accumulators

1. Broadcast Variables

Broadcast variables **are** immutable

Broadcast variables **are** readonly(immutable)

Broadcast variables **are** fault tolerant.

Broadcast variables should fit **in** Mememory

Broadcast variables **are** distributed **to** the **cluster**

Broadcast variable **are not** re-recommendable production (Same concept available **in** MR, We **call** it has distributed cache (**Map Join**))

How **to create** Broadcast **variable**

If you want **to create** Broadcast **variable** we require **2** methods

1. **sc.Broadcast**
2. **convert static table to a HashMap table using** a method called **collectAsMap**

sc.Broadcast method **to initialize** the Broadcast **variable in** Executor **in-memory**

collectAsMap Method converts a scala **object to or** RDD **to** a HashMap **Table**

```
val hoods = Seq((1,"Mission"),(2,"Soma"))
val checkins = Seq((234,1),(567,2))
val hoodsRDD=sc.parallelize(hoods)
val checkRDD=sc.parallelize(checkins)
val broadcastedHoods=sc.broadcast(hoodsRDD.collectAsMap)

val checkinsWithHoods = checkRDD.mapPartitions({
row =>
row.map(x => (x._1,x._2,broadcastedHoods.value.getOrElse(x._2,-1))) // -1 means descending order
},preservesPartitioning=true)// to make it as narrow transformation

checkinsWithHoods.take(5)
```

// Accumulator

Accumulator **are** broadcast **variable**. It will **collect** performance counter **of each and every** worker

node basic statistics of executor performance.
and Submitted back to driver node.

Major Accumulator will be taken care by administrator.

It will take performance counters (map Reduce : Counting with Counters= Accumators in Spark)
performance counters means stats about the executor

Accumulator collects performance counter for each executor to Driver node

'How to create an Accumulator'

In Spark Core Context a method called Accumulator and it expected Initial value of counter.

```
scala> val accum= sc.Accumulator[Int] = 0
<console>:1: error: ';' expected but '=' found.
      val accum= sc.Accumulator[Int] = 0
                                   ^
```

```
scala> val accum= sc.accumulator(0)
accum: org.apache.spark.Accumulator[Int] = 0
```

```
scala> sc.parallelize(Array(1,2,3,4)).foreach( x => accum +=x)
```

```
scala> accum.value
res3: Int = 10
```

'Memory Management is Spark'

In Spark Memory is splitted into 3 core Components

1. Reserved Memory
2. User Memory
3. Spark Memory
 - Further Divided into 2 splits
 - a. Executor Memory
 - b. Storage Memory

1. Reserved Memory

To Start a Spark Application we required a certain amount of Memory and that memory is called as reserved memory
and default size of reserved Mememory is 300 MB

2. User Memory

The Developer or End User start creating RDDs , Partitions , Transformation , Actions internally using a memory block called User Memory.

This Mememory belong to Driver Node.

25% of Whole Memory (After subtraction of Reserved Memory 300)

3. Spark Memory -- Worker Node

Remaining 75% Allocated to Spark

```
Spark 1.x :      Executor = Storage = 50% of Spark Memory
Spark 2.x :      ?
```

Further Divided into 2 splits
a. Executor Memory -- only to execute the job
b. Storage Memory --Store the intermediate data

Checkpoint Service ?

```
RDD.cache()
RDD.persist()
```

Both Represents Storage Memory

Cache	Persist
Default it is stored in Storage Memory -- RAM	1. MEMORY_ONLY 2. DISK_ONLY (Worker node Disk)
MEMORY_ONLY configurable	3. MEMORY_AND_DISK (70% in Memeory and 30% in Disk) --this is 4. MEMORY_AND_DISK_SER 5. MEMORY_ONLY_SER

spark.**exit()** --> All Data will be

RDD.cache() --> It can store in any one of the worker node

cache Never **use in Real time**

2 more options in latest **release**

```
MEMORY_AND_DISK_2
MEMORY_ONLY_2
```

```
uncache()
unpersist()
```

Drawback of Spark 1.x is storage 50% is a bottleneck

In Spark

-----END OF SPARK-----

Scala Programs

```
scala_practice (1).txt
```

Details

Activity

```
scala_practice (1).txt
```

Sharing Info

k

d

J

R

+8

General Info

Type

Text

Size

12 KB (12,342 bytes)

Storage used

0 bytesOwned by someone else

Location

Scala

Owner

kumar K

Modified
 24 Mar 2018 by Soni Rai
 Created
 24 Mar 2018
 Description
 Add a description
 Download permissions
 Viewers can download

Scala Practice Exercises

=====

```
Ex:1) val l = List(10,20,30,40,50,60,70,80,90,100)
      //find elements greater than 40
      val result = l.filter(x => x>40)
```

```
val l = 1 to 10 toList
val x = 1 to 10 toArray
val y = 1 to 100 by 2
```

```
Ex:2) val names = List("abc","def","ghi")
      //convert all elements to uppercase
      val result = names.map(x => x.toUpperCase)
      val result = names.map(_.toUpperCase)
      names.
```

```
Ex:3) val x = List(10,20,30,40,50)
      //sum of all elements in the list using for loop
      var tot = 0
      for(i<-x)
        tot+=i
      tot
```

```
Ex:4) val lst = List(1,2,3,4,5,6,7,8,9,10)
      //4th element to remaining elements
      val r1 = lst.slice(3,lst.size)
      //1st element to 5th element
      val r2 = lst.slice(0,5)
      //3rd element to 8th element
      val r3 = lst.slice(2,8)
```

```
Ex:5) val x = List(10,20,30)
      5::x
      x::35 //not possible but its possible in ListBuffers

      val x = 1 to 5 toList
      val y = 6 to 10 toList
      x::y
      x::y Or x++y
```

```
Ex:6) val x = List(10,20,30,40,50,60)
      val y = List(70,80,90)
      x++y
      val z = y++x
```

```
Ex:7) val x = List(10,20,30,40)
      x++List(50)
```

```
Ex:8) val t=("Ravi",35,"Mtech","Hyd")
      val name= t._1
      val age = t._2
      val qual = t._3
      val loc = t._4
```

```
Ex:9) var x= Map("x" -> "abc", "y"->"def")
      x("x")
```

```
x+=(z->"ghi")
x
```

```
Ex:10) // Transformations Map, FlatMap and Filter
val x = List(10,20,30,40,50)
val y = x.map(x => x+100)
val z = y.filter(x => x>120)
(or) val r = x.map(x => x+100).filter(v => v>120)
```

```
Ex:11) val name = List("rAvI","rani","VaNi","VeNu")
val result = name.map{ x =>
  val w = x.trim()
  val fc = w.substring(0,1).toUpperCase
  val rc = w.substring(1).toLowerCase
  fc+rc
}
```

```
Ex:12) val sal = List(10000,20000,30000,40000,50000)
// net salary tax 10% and hra 20%
val nets = sal.map{x =>
  val tax = x*10/100
  val hra = x*20/100
  val net = x+hra-tax
  net
}
```

```
Ex:13) val name = List("rAvI","rani","VaNi","VeNu")
val result = name.map{ x =>
  val w = x.trim()
  val fc = w.slice(0,1).toUpperCase
  val rc = w.slice(1,w.size).toLowerCase
  fc+rc
}
```

```
Ex:14) // Diff between map and flatmap
val l = List(List(1,2,3),List(3,4),List(1,3,5,6),List(1,2,3))
// using map transformation
val r = l.map(x => x.sum)
val res = r.sum
// using flatMap transformation
val res = l.flatMap(x => x).sum
```

```
Ex:15) val lines = List("I love hadoop","I Love Spark","I love Spark and Hadoop","Spark is Great")
// ((I,1), (love,1), (hadoop,1), (I,1), (Love,1), (Spark,1), (I,1), (love,1), (Spark,1), (
and,1), (Hadoop,1), (Spark,1), (is,1), (Great,1))
val rmap = lines.map(x => x.split(" "))
val rflatMap = lines.flatMap(x => x.split(" ")).map(x => (x,1))
```

```
Ex:16) val recs = List("101,Anil,100000,m,11","102,amala,50000,f,12","103,giri,60000,m,11",
"104,girija,90000,f,13","105,Mani,10000,m,12")
// select sex,sum(sal) from emp group by sex;
recs.foreach(println)
val arr = recs.map(x => x.split(","))
val pair = arr.map(x => (x(3),x(2).toInt))
```

```
Ex:17) val data = List("100,200,300","800,200,300,400,500","10000,30000","900,1000,5000")
```

```
Ex:18) val l = List(List(1,2,3),List(3,4),List(1,2,3,4),List(1,2,3,4,5))
val result = l.filter(x => x.size>2)
```

```
Ex:19) val recs = List("101,Anil,100000,m,11","102,amala,50000,f,12","103,giri,60000,m,11",
"104,girija,90000,f,13","105,Mani,10000,m,12")
// select * from emp where gendar = "m";
val m = recs.filter(x => x.contains("m"))

val males = recs.filter{ x =>
```

```

    val gendar = x.split(",").toLowerCase
    gendar == "m"
  }

```

```
males.foreach(println)
```

```
val recs = List("101,Anil,100000,m,11","102,amala,50000,f,12","103,giri,60000,m,11",
"104,girija,90000,f,13","105,Mani,10000,m,12")
```

```
val result = recs.map{x=>
val w = x.split(",")
val dno = w(4).toInt
val ename = w(1)
val sal = w(2).toInt
val gendar = w(3)
(dno,(ename,sal,gendar))
}
```

Ex:20) //Conditional Transformations

```
val lst = List(10000,30000,90000,20000,60000,80000)
lst.sum
lst.size
val avg = lst.sum/lst.size
val result = lst.map(x => if(x>=avg) "Above AVG Sal" else "Below AVG Sal")
```

```
val recs = List("101,Anil,100000,m,11","102,amala,50000,f,12","103,giri,60000,m,11",
"104,girija,90000,f,13","105,Mani,10000,m,12")
```

```
val result = recs.map{x=>
val w = x.split(",")
val sal = w(2).toInt
val ssum = sal.sum
val scnt = sal.size
val savg = ssum/scnt
val r = w.map(x => if(x>=savg) "Above AVG Sal" else "Below AVG Sal")
r
}
```

Ex:21) val a =100
val b =250
val c =150

```
if(a>b) a else if(b>c)b else if(a>c)a else c
```

Ex:22) val dno = List(11,12,13,11,11,11,12,13,12,12,13,14,15,11)
val dname = dno.map{x =>
if(x==11) "Marketing" else if(x==12) "HR" else if(x==13) "Finance" else "Others"
}

Ex:23) val gendar = "m"
val result = gendar match{
case "m" => "Male"
case "f" => "Female"
case other => "Unknown"
}

Ex:24) val gendar ="f"
val result = (gendar == "m") match{
case true => "Male"
case false => "Female"
}

Ex:25) val sal = List(10000,30000,90000,20000,60000,80000)
val r =sal.map{ x =>
var grade = " "
if (x>=70000) grade = "A" else if(x>=50000) grade ="B" else if(x>=30000) grade ="C" else
grade ="D"
grade

```
}

```

```
Ex:26) val sal = List(10000,30000,90000,20000,60000,80000)
      val res = sal.map{x =>
        var grade = " "
        (x>=70000) match{
        case true => grade = "A"
        case false => (x>=50000) match{
        case true => grade = "B"
        case false => (x>=30000) match{
        case true => grade = "C"
        case false => grade = "D"
        }
        }
        }
      }
      grade
    }
```

```
Ex:27) //using match and if combination
      val sal = List(10000,30000,90000,20000,60000,80000)
      val res = sal.map{x =>
        var grade = " "
        (x>=70000) match {
        case true => grade = "A"
        case false => if(x>=50000) grade="B" else if(x>=30000) grade="C" else grade = "D"
        }
      }
      grade
    }
```

```
Ex:28) val records = List("101,Anil,40000,m,11","102,aMaLA,80000,F,12","103,ManI,10000,m,13",
"104,GIRI,45000,m,14","105,SuReSH,60000,f,12","106,SIRI,90000,M,15")
      //name --->first char is upper and remaining are lower.
      //sal ---> calculate net salary(net =sal+hra-tax)hra=20% and tax=10%
      //sal --->sal grades like A,B,C,D
      //gendar ---> m/M and f/F convert as Male and Female
      //dno ---> 11-Marketing,12-HR,13-Finance,others

      val results = records.map{x =>
        val w = x.split(",")
        val id = w(0)
        val name = w(1).trim()
        var sal = w(2).toInt
        var gendar = w(3)
        val dno = w(4).toInt
        val fc = name.slice(0,1).toUpperCase
        val rc = name.slice(1,name.size).toLowerCase
        val newname = fc+rc
        gendar = if(gendar.toUpperCase=="M") "Male" else "Female"
        val tax = sal*0.1
        val hra = sal*0.2
        val net = sal+hra-tax
        var grade = " "
        if(net>=70000) grade="A" else if(net>=50000) grade = "B" else if(net>=30000) grade = "C"
        else "D"
        grade
        val dname = dno match{
        case 11 => "Marketing"
        case 12 => "HR"
        case 13 => "Finance"
        case other => "Others"
        }
        val newList = List(id,newname,sal.toString,hra.toString,tax.toString,net.toString,grade,
        gendar,dno.toString,dname)
        newList.mkString(",")
      }
    }
```

```
Ex:29) val emp = Array("101,aaaa,30000,m,11","102,bbbb,50000,f,12","103,hhhh,60000,m,11",
"104,qqqq,80000,f,11")
//select gendar,sum(sal) from emp group by gendar;
val pair1 = emp.map{x =>
val w = x.split(",")
val gendar = w(3)
val sal = w(2).toInt()
(gendar,sal)
}
```

```
Ex:30) val emp = Array("101,aaaa,30000,m,11","102,bbbb,50000,f,12","103,hhhh,60000,m,11",
"104,qqqq,80000,f,11")
//select dno,sum(sal) from emp group by gendar;
val pair1 = emp.map{x =>
val w = x.split(",")
val dno = w(4)
val sal = w(2).toInt()
(dno,sal)
}
```

```
Ex:31) val emp = Array("101,aaaa,30000,m,11","102,bbbb,50000,f,12","103,hhhh,60000,m,11",
"104,qqqq,80000,f,11")
//select dno,gend,sum(sal) from emp group by gendar;
val pair1 = emp.map{x =>
val w = x.split(",")
val dno = w(4)
val gend = w(3)
val sal = w(2).toInt()
val mykey = (dno,gend)
(mykey,sal)
}
```

Ex:32) //Making the records into structures

i) Tuple

ii) case class

```
emp = Array("101,aaaa,30000,m,11","102,bbbb,50000,f,12","103,hhhh,60000,m,11",
"104,qqqq,80000,f,11")
emp.foreach(println)
```

```
val recs = emp.map{ x =>
val w = x.split(",")
val id = w(0).toInt
val name = w(1)
val sal = w(2).toInt
val gend = w(3)
val dno = w(4).toInt
(id,name,sal,gend,dno)
}
```

```
recs.foreach(println)
```

```
val pair4 = recs.map(x => (x._4,x._3.toInt))
```

```
//select sum(sal) from recs;
```

```
val tsum = recs.map(x => x._3).sum
```

```
Ex:33) val textdata = " I Love Spark "
val word = textdata.trim()
val s = word.split(" ")
val w = s.filter(x => x!="")
val result = w.mkString(" ")
```

```
Ex:34) val post = List("I Love Spark ", "you Love hadoop", "hadoop and spark are great
big data systems")
val result = post.map{x =>
```

```

    val w = x.trim().split(" ")
    val words = w.filter(x => x!="")
    words.mkString(" ")
  }
  result.foreach(println)

```

```

Ex:35) //case class
case class Samp(a:Int,b:Int,c:Int)
val s1 = Samp(10,20,30)
val s2 = Samp(1,2,3)
val s3 = Samp(100,200,300)
val s = List(s1,s2,s3)

```

```

Ex:36) //case class
case class Emp(id:Int,name:String,sal:Int,gendar:String,dname:String)
val e = emp.map{x =>
  val w = x.split(",")
  val id = w(0).toInt
  val name = w(1)
  val sal = w(2).toInt
  val gendar = w(3)
  val dno = w(4).toInt
  val dname = dno match{
    case 11 => "Marketing"
    case 12 => "HR"
    case 13 => "Finance"
    case other => "Others"
  }
  val result = Emp(id,name,sal,gendar,dname)
  result
}
e.foreach(println)
val pair5 = e.map(x => (x.dname,x.sal))

```

```

Ex:37) //Functions
def f:String = "Hello World"
f

```

```

Ex:38) def f ={
  val x = "hello"
  val y = x.toUpperCase
  y
}

```

```

Ex:39) def fx(a:Int):Int = a+100
fx(10)

```

```

Ex:40) case class Empl(id:Int,name:String,sal:Int,sex:String,dno:Int,
dname:String,grade:String)

```

```

def FirstUpper(x:String):String={
  val w = x.trim()
  val fc = w.slice(0,1).toUpperCase
  val rc = w.slice(1,w.size).toLowerCase
  val name = fc+rc
  name
}

def gend(x:String):String={
  if(x.toUpperCase == "M") "Male" else "Female"
}

def grade(x:Int):String={
  if(x>=70000) "A" else if(x>=50000) "B" else if(x>=30000) "C" else "D"
}

```



```

def dept(x:Int):String={
  val dname = x match{
    case 11 => "Marketing"
    case 12 => "HR"
    case 13 => "Finance"
    case other => "Others"
  }
  dname
}

def toEmp(line:String):Empl={
  val w = line.split(",")
  val id = w(0).toInt
  val name = FirstUpper(w(1))
  val sal = w(2).toInt
  val gendar = gend(w(3))
  val dno = w(4).toInt
  val dname = dept(dno)
  val grad = grade(sal)
  val e = Empl(id,name,sal,gendar,dno,dname,grad)
  e
}

toEmp("201,ANiL,80000,m,11")

val emps = emp.map(x =>toEmp(x))
emps.foreach(println)

```

```

Ex:41) def isMale(x:String):Boolean={
  x.toUpperCase == "M"
}
val lst = List("m","M","f","F","M","f","M")
val males = lst.filter(x => isMale(x))
val females = lst.filter(x => !isMale(x))

val m = emp.filter(x =>isMale(x.split(",")(3)))
val f = emp.filter(x =>!isMale(x.split(",")(3)))

val m = emps.filter(x => isMale(x.gendar.slice(0,1)))
val f = emps.filter(x => !isMale(x.gendar.slice(0,1)))

```

```

Ex:42) def power(x:Int,n:Int):Int={
  if(n>=1) x*power(x,n-1)
  else 1
}

```

```

Ex:43) def fact(x:Int):Int={
  if(x>1) x*fact(x-1)
  else 1
}

```

--- Spark Core

```

Ex:1) val r1 = List((11,10000),(11,20000),(12,30000),(12,40000),(13,50000))
val r2 = List((11,"hyd"),(12,"bang"),(13,"hyd"))
val rdd1 = sc.parallelize(r1)
val rdd2 = sc.parallelize(r2)
val j = rdd1.join(rdd2)
j.collect.foreach(println)
val citysalpair = j.map{x =>
  val city = x._2._2
  val sal = x._2._1
  (city,sal)
}

```

```
citysalpair.collect.foreach(println)
val result = citysalpair.reduceByKey(_+_ )
result.collect.foreach(println)
```

```
Ex:2) val e = List((11,30000,10000),(11,40000,20000),(12,50000,30000),
(13,60000,20000),(12,80000,30000))
val ee = sc.parallelize(e)
ee.collect.foreach(println)
rdd2.collect.foreach(println)
val j2 = ee.join(rdd2)//error because both structures in key and value.
val e3 = ee.map{x =>
    val dno = x._1
    val sal = x._2
    val bonus = x._3
    (dno,(sal,bonus))
}
val j3 = e3.join(rdd2)
j3.collect.foreach(println)
val pair = j3.map{x =>
    val sal= x._2._1._1
    val bonus = x._2._1._2
    val tot = sal+bouus
    val city = x._2._2
    (city,tot)
}
pair.collect.foreach(println)
val result2 = pair.reduceByKey(_+_ )
result2.collect.foreach(println)
```

```
Ex:3) emp file
cat > emp
101,aaaa,70000,m,12
102,bbbb,90000,f,12
103,cccc,10000,m,11
104,dddd,40000,m,12
105,eeee,70000,f,13
106,ffff,80000,f,13
107,gggg,90000,m,14
108,hhhh,10000,f,14
109,iiii,30000,m,11
110,jjjj,60000,f,14
111,kkkk,90000,m,15
112,llll,10000,m,15
dept file
cat > dept
11,marketing,hyd
12,hr,del
13,finance,hyd
14,admin,del
15,accounts,hyd
//move the files from lfs to hdfs.
hadoop fs -mkdir /sparkcore
hadoop fs -put emp /sparkcore
hadoop fs -put dept /sparkcore
hadoop fs -ls /saprkcore
val emp=sc.textFile("/user/cloudera/sparkcore/emp")
val dept = sc.textFile("/user/cloudera/sparkcore/dept")
emp.collect.foreach(println)
dept.collect.foreach(println)
val e = emp.map{ x =>
    val w = x.split(",")
    val dno = w(4).toInt
    val id = w(0)
    val name = w(1)
    val sal = w(2).toInt
    val sex = w(3)
```

```

    val info = id+", "+name+", "+sal+", "+sex
    (dno,info)
    // (dno, (id,name,sal,gendar))
  }

```

```

e.collect.foreach(println)
val d = dept.map{ x =>
  val w = x.split(",")
  val dno = w(0).toInt
  val dname = w(1)
  val city = w(2)
  val info = dname+", "+city
  (dno,info)
}
d.collect.foreach(println)
val ed = e.join(d)
ed.collect.foreach(println)
val ed2 = ed.map{x =>
  val einfo = x._2._1
  val dinfo = x._2._2
  val info = einfo+", "+dinfo
  info
}
ed2.collect.foreach(println)
ed2.saveAsTextFile("/user/cloudera/sparkcore/res1")

```

```

Ex:4) emp
dept
val ednosal = emp.map{ x =>
  val w = x.split(",")
  val dno = w(4)
  val sal = w(2).toInt
  (dno,sal)
}
ednosal.collect.foreach(println)
val dnocity = dept.map{ x =>
  val w = x.split(",")
  val dno = w(0)
  val city = w(2)
  (dno,city)
}
dnocity.collect.foreach(println)

val edj = ednosal.join(dnocity)
edj.collect.foreach(println)

val result = edj.map{ x =>
  val city = x._2._2
  val sal = x._2._1
  (city,sal)
}
val res = result.reduceByKey(_+_ )
res.collect.foreach(println)

//how to get no of partitions?
r1.partitions.size
Ex:5) cat > comment
I love spark

```

```

I love hadoop
I love hadoop and spark
Spark is super speed
  hadoop fs -put comment /sparkcore
  val data = sc.textFile("/user/cloudera/sparkcore/comment")
  data.collect.foreach(println)
  val word = data.flatMap(" ")
  val pair = word.map(x => (x,1))
  val result = pair.reduceByKey(_+_ )
  result.collect.foreach(println)

```

```

Ex:6) cat > comment2
spark      Spark      spark
hadoop     spark Hadoop  HADOOP
HADOOP     Hadoop

```

```

def rmSpace(x:String):String={
  "  I   love      SpARk   "

  val line =x.trim()
  val w = line.split(" ")
  val words = w.filter(x => x!="").map(x => x.toLowerCase)
  words.mkString(" ")
}

val lines = sc.textFile("/user/cloudera/sparkcore/comment2")
val data = lines.map(x => rmSpace(x))
val arr = data.flatMap(x => x.split(" "))
val pair = arr.map(x => (x,1))
val result = pair.reduceByKey(_+_ )
result.collect.foreach(println)

```

```

Ex:7) val data = sc.textFile("/user/cloudera/sparkcore/emp")
      data.collect
      //select dno,sex,sum(sal) from emp
      //group by dno,sex;
      val result = data.map{ x =>
        val w = x.split(",")
        val dno =w(4)
        val sex = w(3)
        val sal = w(2).toInt
        val mkey=(dno,sex)
        (mkey,sal)
      }
      result.collect.foreach(println)
      val fina_res = result.reduceByKey(_+_ )
      final_res.collect.foreach(println)
      //no.of unique blocks(for files) = no.of partitions(for ram)

```

```

Ex:8)
val data1= sc.textFile("/user/cloudera/sparkcore/emp",2)
//data1.partitions.size
val r1 = data1.map(x => x.split(","))
val r2 = r1.map(x => (x(3), x(2).toInt))
//r2.persist
val res1 = r2.reduceByKey(_+_ )
val res2 = r2.reduceByKey(Math.max(_,_))
val res3 = r2.reduceByKey(Math.min(_,_))
res1.saveAsTextFile("/user/cloudera/sparkcore/res1")
res2.saveAsTextFile("/user/cloudera/sparkcore/res2")
res3.saveAsTextFile("/user/cloudera/sparkcore/res3")
val tres1 = res1.map(x => x._1+"\t"+x._2)
tres1.collect
tres1.saveAsTextFile("/user/cloudera/sparkcore/tres1")

```

```

Ex:9) val data = sc.textFile("/user/cloudera/sparkcore/emp")
      val arr = data.map(x => x.split(","))
      val pair1 = arr.map(x => (x(4),x(2).toInt))

```

```

pair1.collect
val grp = pair1.groupByKey()
grp.collect
val res = grp.map{ x =>
  val dno = x._1
  val cb = x._2
  val tot = cb.sum
  val cnt = cb.size
  val avg = tot/cnt
  val max = cb.max
  val min = cb.min
  val r = dno+", "+tot+", "+cnt+", "+avg+", "+max+", "+min
  r
}
res.collect.foreach(println)

```

```

val pairs = sc.parallelize(Array(("a", 3), ("a", 1), ("b", 7), ("b", 2), ("a", 5)))
//0 is initial value, _+ inside partition, + between partitions
val resaggregateByKey = pairs.aggregateByKey(0)(_+_ ,*_ )
resaggregateByKey.collect().foreach(println)

```

```

val a = sc.parallelize(List((1, 'a'), (2, 'a'), (3, 'a'), (4, 'a'), (1, 'c'), (2, 'c'), (3, 'c'), (4, 'c')), 2)
val b = sc.parallelize(List((1, 'b'), (2, 'b'), (3, 'b'), (4, 'b')), 2)
a.cogroup(b).sortByKey(false).collect().foreach(println)

```

```

Ex:10) data.collect
arr.collect
arr.persist
val pair2 = arr.map(x => ((x(4), x(3)), x(2).toInt))
pair2.collect
val grp2 = pair2.groupByKey()
grp2.collect.foreach(println)
val result2 = grp2.map{x =>
  val k = x._1
  val dno = k._1
  val sex = k._2

```

```

    val cb = x._2
    val tot = cb.sum
    val cnt = cb.size
    val avg = tot/cnt
    val max = cb.max
    val min = cb.min
    (dno,sex,tot,cnt,avg,max,min)
  }
result2.collect.foreach(println)

```

Ex:11) `select sum(sal),avg(sal),count(*),max(sal),min(sal) from emp;`

```

data.collect
arr.collect
val sal = arr.map(x => x(2).toInt)
sal.collect

val tot = sal.sum
val cnt = sal.count
val avg:Int = (tot/cnt).toInt
val max = sal.max
val min = sal.min

```

```

or
val tot = sal.reduce(_+_ )
val cnt = sal.count
val avg = tot/cnt
val max = sal.reduce(Math.max(_,_))
val min = sal.reduce(Math.min(_,_))

```

Ex:12) `val rrd1 = sc.parallelize(List(10,20,30,40,50,60,70,80,90),2)`

partition1 ---> `List(10,20,30,40,50)`

partition2 ---> `List(60,70,80,90)`

`rrd1.sum` // all partitions data will be collected into local,
sum executed at local (no parallel processing)

`rrd1.reduce(_+_)`

the above operation executed at cluster level separately for
each partition

partition1 result =150

partition2 result =300

these independent results will be collected into any one slave of spark cluster.

`List(150,300)=450`

Ex:13) `//select sum(sal) from emp where sex='m'`

```

def isMale(x:String)={
  val w = x.split(",")
  val sex = w(3).toLowerCase
  sex == "m"
}
val males = data.filter(x => isMale(x))
val sals = males.map(x => x.split(",")(2).toInt)
sals.reduce(_+_ )
//select sum(sal) from emp where sex='f'
//select max(sal) from emp where sex='m'
//select min(sal) from emp where sex='m'

```

14) `scala> val parallel = sc.parallelize(1 to 9, 3)`

`parallel.mapPartitionsWithIndex((index: Int, it: Iterator[Int]) =>`

```
it.toList.map(x => index + ", "+x).iterator).collect
```

--- Spark Transformation

```
*****Spark-Core*****
*****Sprark-Template*****
import org.apache.spark.sql.SparkSession
object ${NAME} {
def main(args: Array[String]) {
  //val spark = SparkSession.builder.master("local[*]").appName("${NAME}").config(
  "spark.sql.warehouse.dir", "/home/hadoop/work/warehouse").enableHiveSupport().getOrCreate()
  val spark = SparkSession.builder.master("local[*]").appName("${NAME}").getOrCreate()
  val sc = spark.sparkContext
  val sqlContext = spark.sqlContext
  import spark.implicits._
  import spark.sql
spark.stop()
}
}
*****
```

1) aggregateByKeyTransformation

```
-----
package SparkCore

/**
 * Created by Kalyan on 5/26/2017.
 */

import org.apache.spark.sql.SparkSession

object aggregateByKeyTransformation {
def main(args: Array[String]) {
  //val spark = SparkSession.builder.master("local[*]").appName("aggregateByKeyTransformation"
  ).config("spark.sql.warehouse.dir", "/home/hadoop/work/warehouse").enableHiveSupport().
  getOrCreate()
  val spark = SparkSession.builder.master("local[*]").appName("aggregateByKeyTransformation").
  getOrCreate()
  val sc = spark.sparkContext
  val sqlContext = spark.sqlContext
  import spark.implicits._
  import spark.sql
  //Example for reduceByKey
  // val pairs = sc.parallelize(Array(("a", 3), ("a", 1), ("b", 7), ("a", 5)))
  // val resReduceByKey = pairs.reduceByKey(_+_ )
  // resReduceByKey.collect().foreach(println)
  //Example-1 for aggregateByKey
  val pairs = sc.parallelize(Array(("a", 3), ("a", 1), ("b", 7), ("b", 2), ("a", 5)))
  //0 is initial value, _+ inside partition, _+ between partitions
  val resaggregateByKey = pairs.aggregateByKey(0)(_+_ ,_*_)
  resaggregateByKey.collect().foreach(println)
  //Example-2 for aggregateByKey
  // import scala.collection.mutable.HashSet
  // //the initial value is a void Set. Adding an element to a set is the first
  // //+_ Join two sets is the ++
  // val sets = pairs.aggregateByKey(new HashSet[Int])(_+_ , _++)
  spark.stop()
}
}
```

2) cartesianTransformation

```
-----
package SparkCore
```

```
/**
```

```

* Created by Kalyan on 5/26/2017.
*/
import org.apache.spark.sql.Session

object cartesianTransformation {
  def main(args: Array[String]) {
    //val spark = SparkSession.builder.master("local[*]").appName("cartesianTransformation").
    config("spark.sql.warehouse.dir", "/home/hadoop/work/warehouse").enableHiveSupport().
    getOrCreate()
    val spark = SparkSession.builder.master("local[*]").appName("cartesianTransformation").
    getOrCreate()
    val sc = spark.sparkContext
    val sqlContext = spark.sqlContext
    import spark.implicits._
    import spark.sql
    val x = sc.parallelize(List(1,2,3,4,5),2)
    val y = sc.parallelize(List('a','b','c','d','e','f'), 2)
    x.cartesian(y).collect().foreach(println)
    spark.stop()
  }
}

```

3)coalesceTransformation

```
package SparkCore
```

```

/**
 * Created by Kalyan on 5/26/2017.
 */
import org.apache.spark.sql.Session

object coalesceTransformation {
  def main(args: Array[String]) {
    //val spark = SparkSession.builder.master("local[*]").appName("coalesceTransformation").
    config("spark.sql.warehouse.dir", "/home/hadoop/work/warehouse").enableHiveSupport().
    getOrCreate()
    val spark = SparkSession.builder.master("local[*]").appName("coalesceTransformation").
    getOrCreate()
    val sc = spark.sparkContext
    val sqlContext = spark.sqlContext
    import spark.implicits._
    import spark.sql
    val x = sc.parallelize(Array(1,2,3,4,5),3)
    val y = x.coalesce(2,false)
    println(y.getNumPartitions)
    spark.stop()
  }
}

```

4)cogroupTransformation

```
package SparkCore
```

```

/**
 * Created by Kalyan on 5/26/2017.
 */
import org.apache.spark.sql.Session

object cogroupTransformation {
  def main(args: Array[String]) {
    //val spark = SparkSession.builder.master("local[*]").appName("cogroupTransformation").
    config("spark.sql.warehouse.dir", "/home/hadoop/work/warehouse").enableHiveSupport().
    getOrCreate()
    val spark = SparkSession.builder.master("local[*]").appName("cogroupTransformation").
    getOrCreate()
    val sc = spark.sparkContext

```



```

    val sqlContext = spark.sqlContext
    import spark.implicits._
    import spark.sql
    val a = sc.parallelize(List((1,'a'),(2,'a'),(3,'a'),(4,'a'),(1,'c'),(2,'c'),(3,'c'),(4,'c')
    )), 2)
    val b = sc.parallelize(List((1,'b'),(2,'b'),(3,'b'),(4,'b')), 2)
    a.cogroup(b).sortByKey(false).collect().foreach(println)
    spark.stop()
  }
}

```

5)distinctTransformation

```
package SparkCore
```

```
/**
```

```
 * Created by Kalyan on 5/26/2017.
```

```
*/
```

```
import org.apache.spark.sql.SparkSession
```

```
object distinctTransformation {
```

```
  def main(args: Array[String]) {
```

```
    //val spark = SparkSession.builder.master("local[*]").appName("distinctTransformation").
```

```
    config("spark.sql.warehouse.dir", "/home/hadoop/work/warehouse").enableHiveSupport().
```

```
    getOrCreate()
```

```
    val spark = SparkSession.builder.master("local[*]").appName("distinctTransformation").
```

```
    getOrCreate()
```

```
    val sc = spark.sparkContext
```

```
    val sqlContext = spark.sqlContext
```

```
    import spark.implicits._
```

```
    import spark.sql
```

```
    //example-1 with files
```

```
    System.setProperty("hadoop.home.dir", "C:\\winutils");
```

```
    val data = sc.textFile("D:\\Sprak-Training\\spark_datasets\\duplicate_data.csv")
```

```
    val result = data.distinct()
```

```
    result.collect().foreach(println)
```

```
    //example-2 using parallelize
```

```
//    val data1 = sc.parallelize(1 to 10)
```

```
//    val data2 = sc.parallelize(5 to 20)
```

```
//    val result1 = data1.union(data2).distinct().sortBy(x =>x,false)
```

```
//    result1.collect().foreach(println)
```

```
    spark.stop()
```

```
  }
```

```
}
```

6)filterTransformation

```
package SparkCore
```

```
/**
```

```
 * Created by Kalyan on 5/26/2017.
```

```
*/
```

```
import org.apache.spark.sql.SparkSession
```

```
object filterTransformation {
```

```
  def main(args: Array[String]) {
```

```
    //val spark = SparkSession.builder.master("local[*]").appName("filterTransformation").config
```

```
    ("spark.sql.warehouse.dir", "/home/hadoop/work/warehouse").enableHiveSupport().getOrCreate()
```

```
    val spark = SparkSession.builder.master("local[*]").appName("filterTransformation").
```

```
    getOrCreate()
```

```
    val sc = spark.sparkContext
```

```
    val sqlContext = spark.sqlContext
```

```
    import spark.implicits._
```

```
    import spark.sql
```

```
    System.setProperty("hadoop.home.dir", "C:\\winutils");
```

```
    val dataset = sc.textFile("D:\\Sprak-Training\\spark_datasets\\sample_data.csv")
```

```

    val result = dataset.map(x=>(x.split(", ")))
    val filtereddata = result.map(x=>((x(2).toInt))).filter(x=>x<25)
    filtereddata.collect().foreach(println)
    spark.stop()
  }
}

```

7) flatMapTransformation

```
package SparkCore
```

```

/**
 * Created by Kalyan on 5/26/2017.
 */
import org.apache.spark.sql.SparkSession

object flatMapTransformation {
  def main(args: Array[String]) {
    //val spark = SparkSession.builder.master("local[*]").appName("MapTransformation").config(
    "spark.sql.warehouse.dir", "/home/hadoop/work/warehouse").enableHiveSupport().getOrCreate()
    val spark = SparkSession.builder.master("local[*]").appName("MapTransformation").getOrCreate()
    val sc = spark.sparkContext
    val sqlContext = spark.sqlContext
    import spark.implicits._
    import spark.sql
    System.setProperty("hadoop.home.dir", "C:\\winutils");
    val dataRDD = sc.textFile("D:\\Sprak-Training\\spark_datasets\\favourite_animals.csv")
    val res = dataRDD.flatMap(x => x.split(","))
    res.collect().foreach(println)
    spark.stop()
  }
}

```

8) groupByKeyTransformation

```
package SparkCore
```

```

/**
 * Created by Kalyan on 5/26/2017.
 */
import org.apache.spark.sql.SparkSession

object groupByKeyTransformation {
  def main(args: Array[String]) {
    //val spark = SparkSession.builder.master("local[*]").appName("groupByKeyTransformation").
    config("spark.sql.warehouse.dir", "/home/hadoop/work/warehouse").enableHiveSupport().
    getOrCreate()
    val spark = SparkSession.builder.master("local[*]").appName("groupByKeyTransformation").
    getOrCreate()
    val sc = spark.sparkContext
    val sqlContext = spark.sqlContext
    import spark.implicits._
    import spark.sql
    System.setProperty("hadoop.home.dir", "C:\\winutils");
    val data = sc.textFile("D:\\Sprak-Training\\spark_datasets\\sample_baby_names.csv")
    val header = data.first()
    val rows = data.filter(row => row != header).map(x=>x.split(","))
    val namesToCountries = rows.map(x=>(x(2), x(1)))
    namesToCountries.groupByKey().collect().foreach(println)
    spark.stop()
  }
}

```

9) intersectionTransformation

```
package SparkCore
```

```
/**
 * Created by Kalyan on 5/26/2017.
 */
import org.apache.spark.sql.SparkSession

object intersectionTransformation {
  def main(args: Array[String]) {
    //val spark = SparkSession.builder.master("local[*]").appName("intersectionTransformation").
    config("spark.sql.warehouse.dir", "/home/hadoop/work/warehouse").enableHiveSupport().
    getOrCreate()
    val spark = SparkSession.builder.master("local[*]").appName("intersectionTransformation").
    getOrCreate()
    val sc = spark.sparkContext
    val sqlContext = spark.sqlContext
    import spark.implicits._
    import spark.sql
    System.setProperty("hadoop.home.dir", "C:\\winutils");
    val data1 = sc.parallelize(1 to 10)
    val data2 = sc.parallelize(5 to 20)
    val result = data1.intersection(data2)
    result.collect().foreach(println)
    spark.stop()
  }
}
```

```
10)joinTransformation
```

```
-----
package SparkCore
```

```
/**
 * Created by Kalyan on 5/26/2017.
 */
import org.apache.spark.sql.SparkSession

object joinTransformation {
  def main(args: Array[String]) {
    //val spark = SparkSession.builder.master("local[*]").appName("joinTransformation").config(
    "spark.sql.warehouse.dir", "/home/hadoop/work/warehouse").enableHiveSupport().getOrCreate()
    val spark = SparkSession.builder.master("local[*]").appName("joinTransformation").
    getOrCreate()
    val sc = spark.sparkContext
    val sqlContext = spark.sqlContext
    import spark.implicits._
    import spark.sql
    val x = sc.parallelize(Array(("a",1), ("b",2)))
    val y = sc.parallelize(Array(("a",3), ("a",4), ("b",5)))
    //Inner Join
    val innerjoin = x.join(y)
    println("Inner Join :" + innerjoin.collect().mkString(", "))
    //Left Outer Join
    val leftjoin = x.leftOuterJoin(y)
    println("Left Outer Join :" + leftjoin.collect().mkString(", "))
    //Right Outer Join
    val rightjoin = x.rightOuterJoin(y)
    println("Right Outer Join :" + rightjoin.collect().mkString(", "))
    //Full Outer Join
    val fulljoin = x.rightOuterJoin(y)
    println("Full Outer Join :" + fulljoin.collect().mkString(", "))
    spark.stop()
  }
}
```

```
11)mapPartitionsTransformation
```

```
-----
```

```
package SparkCore
```

```
/**
 * Created by Kalyan on 5/26/2017.
 */
import org.apache.spark.sql.SparkSession

object mapPartitionsTransformation {
  def main(args: Array[String]) {
    //val spark = SparkSession.builder.master("local[*]").appName("mapPartitionsTransformation")
    //    .config("spark.sql.warehouse.dir", "/home/hadoop/work/warehouse").enableHiveSupport().
    //    getOrCreate()
    val spark = SparkSession.builder.master("local[*]").appName("mapPartitionsTransformation").
    getOrCreate()
    val sc = spark.sparkContext
    val sqlContext = spark.sqlContext
    import spark.implicits._
    import spark.sql
    val x = sc.parallelize(Array(1,2,3),2)
    def f(i:Iterator[Int])={ (i.sum,42).productIterator}
    val y = x.mapPartitions(f)
    y.collect().foreach(println)
    spark.stop()
  }
}
```

12)mapPartitionsWithIndexTransformation

```
package SparkCore
```

```
/**
 * Created by Kalyan on 5/26/2017.
 */
import org.apache.spark.sql.SparkSession
```

```
object mapPartitionsWithIndexTransformation {
  def main(args: Array[String]) {
    //val spark = SparkSession.builder.master("local[*]").appName(
    "mapPartitionsWithIndexTransformation").config("spark.sql.warehouse.dir",
    "/home/hadoop/work/warehouse").enableHiveSupport().getOrCreate()
    val spark = SparkSession.builder.master("local[*]").appName(
    "mapPartitionsWithIndexTransformation").getOrCreate()
    val sc = spark.sparkContext
    val sqlContext = spark.sqlContext
    import spark.implicits._
    import spark.sql
    val x = sc.parallelize(Array(1,2,3),2)
    def f(partitionIndex:Int,i:Iterator[Int])={
      (partitionIndex,i.sum).productIterator
    }
    val y = x.mapPartitionsWithIndex(f)
    y.collect().foreach(println)
    spark.stop()
  }
}
```

13)mapTransformation

```
package SparkCore
```

```
/**
 * Created by Kalyan on 5/26/2017.
 */
import org.apache.spark.sql.SparkSession
```

```
object mapTransformation {
```

```
def main(args: Array[String]) {
  //val spark = SparkSession.builder.master("local[*]").appName("mapTransformation").config(
  "spark.sql.warehouse.dir", "/home/hadoop/work/warehouse").enableHiveSupport().getOrCreate()
  val spark = SparkSession.builder.master("local[*]").appName("mapTransformation").getOrCreate()
  val sc = spark.sparkContext
  val sqlContext = spark.sqlContext
  import spark.implicits._
  import spark.sql
  System.setProperty("hadoop.home.dir", "C:\\winutils");
  val dataRDD = sc.textFile("D:\\Sprak-Training\\spark_datasets\\animals.csv")
  val res = dataRDD.flatMap(x => x.split(","))
  res.collect().foreach(println)
  spark.stop()
}
}
```

14) reduceByKeyTransformation

package SparkCore

```
/**
 * Created by Kalyan on 5/26/2017.
 */
import org.apache.spark.sql.SparkSession

object reduceByKeyTransformation {
  def main(args: Array[String]) {
    //val spark = SparkSession.builder.master("local[*]").appName("reduceByKeyTransformation").
    config("spark.sql.warehouse.dir", "/home/hadoop/work/warehouse").enableHiveSupport().
    getOrCreate()
    val spark = SparkSession.builder.master("local[*]").appName("reduceByKeyTransformation").
    getOrCreate()
    val sc = spark.sparkContext
    val sqlContext = spark.sqlContext
    import spark.implicits._
    import spark.sql
    System.setProperty("hadoop.home.dir", "C:\\winutils")
    val data = sc.textFile("D:\\Sprak-Training\\spark_datasets\\sample_baby_names.csv")
    val header = data.first()
    val rows = data.filter(row => row != header).map(x=>x.split(","))
    val names = rows.map(x=>(x(1),x(4).toInt)).reduceByKey((v1,v2) => v1+v2)
    names.collect().foreach(println)
    spark.stop()
  }
}
```

15) repartitionAndSortWithinPartitionsTransformation

package SparkCore

```
/**
 * Created by Kalyan on 5/26/2017.
 */
import org.apache.spark.sql.SparkSession

object repartitionAndSortWithinPartitionsTransformation {
  def main(args: Array[String]) {
    //val spark = SparkSession.builder.master("local[*]").appName(
    "repartitionAndSortWithinPartitionsTransformation").config("spark.sql.warehouse.dir",
    "/home/hadoop/work/warehouse").enableHiveSupport().getOrCreate()
    val spark = SparkSession.builder.master("local[*]").appName(
    "repartitionAndSortWithinPartitionsTransformation").getOrCreate()
    val sc = spark.sparkContext
  }
}
```

```

    val sqlContext = spark.sqlContext
    import spark.implicitly._
    import spark.sql

    spark.stop()
  }
}

16) repartitionTransformation
-----
package SparkCore

/**
 * Created by Kalyan on 5/26/2017.
 */
import org.apache.spark.sql.SparkSession

object repartitionTransformation {
  def main(args: Array[String]) {
    //val spark = SparkSession.builder.master("local[*]").appName("repartitionTransformation").
    config("spark.sql.warehouse.dir", "/home/hadoop/work/warehouse").enableHiveSupport().
    getOrCreate()
    val spark = SparkSession.builder.master("local[*]").appName("repartitionTransformation").
    getOrCreate()
    val sc = spark.sparkContext
    val sqlContext = spark.sqlContext
    import spark.implicitly._
    import spark.sql
    val x = sc.parallelize(Array(1,2,3,4,5),3)
    val y = x.repartition(2)
    println(y.getNumPartitions)
    spark.stop()
  }
}

17) sampleTransformation
-----
package SparkCore

/**
 * Created by Kalyan on 5/26/2017.
 */
import org.apache.spark.sql.SparkSession

object sampleTransformation {
  def main(args: Array[String]) {
    //val spark = SparkSession.builder.master("local[*]").appName("sampleTransformation").config
    ("spark.sql.warehouse.dir", "/home/hadoop/work/warehouse").enableHiveSupport().getOrCreate()
    val spark = SparkSession.builder.master("local[*]").appName("sampleTransformation").
    getOrCreate()
    val sc = spark.sparkContext
    val sqlContext = spark.sqlContext
    import spark.implicitly._
    import spark.sql
    System.setProperty("hadoop.home.dir", "C:\\winutils");
    val dataset = sc.textFile("D:\\Sprak-Training\\spark_datasets\\sample_data.csv")
    val result = dataset.sample(true,0.5)//sample(withReplacement: Boolean, fraction: Double,
    seed: Long = Utils.random.nextLong)
    result.collect().foreach(println)
    spark.stop()
  }
}

18) sortByKeyTransformation
-----

```

```

package SparkCore

/**
 * Created by Kalyan on 5/26/2017.
 */
import org.apache.spark.sql.Session

object sortByKeyTransformation {
  def main(args: Array[String]) {
    //val spark = SparkSession.builder.master("local[*]").appName("sortByKeyTransformation").
    config("spark.sql.warehouse.dir", "/home/hadoop/work/warehouse").enableHiveSupport().
    getOrCreate()
    val spark = SparkSession.builder.master("local[*]").appName("sortByKeyTransformation").
    getOrCreate()
    val sc = spark.sparkContext
    val sqlContext = spark.sqlContext
    import spark.implicits._
    import spark.sql
    System.setProperty("hadoop.home.dir", "C:\\winutils")
    val data = sc.textFile("D:\\Sprak-Training\\spark_datasets\\sample_baby_names.csv")
    val header = data.first()
    val rows = data.filter(row => row != header).map(x => x.split(","))
    val names = rows.map(x => (x(1), x(4).toInt)).sortByKey(false) //false-descending order,true
    -Ascending order(default)
    names.collect().foreach(println)
    spark.stop()
  }
}

```

19)unionTransformation

```

-----
package SparkCore

/**
 * Created by Kalyan on 5/26/2017.
 */
import org.apache.spark.sql.Session

object unionTransformation {
  def main(args: Array[String]) {
    //val spark = SparkSession.builder.master("local[*]").appName("unionTransformation").config(
    "spark.sql.warehouse.dir", "/home/hadoop/work/warehouse").enableHiveSupport().getOrCreate()
    val spark = SparkSession.builder.master("local[*]").appName("unionTransformation").
    getOrCreate()
    val sc = spark.sparkContext
    val sqlContext = spark.sqlContext
    import spark.implicits._
    import spark.sql
    val data1 = sc.parallelize(1 to 10)
    val data2 = sc.parallelize(5 to 20)
    val result = data1.union(data2)
    result.collect().foreach(println)
    spark.stop()
  }
}

```

----- Spark Submit Commands

```

./bin/spark-submit --class org.apache.spark.examples.SparkPi \
  --master yarn \
  --deploy-mode cluster \
  --driver-memory 4g \
  --executor-memory 2g \
  --executor-cores 4 \

```

```
--queue default \
lib/spark-examples*.jar
```

The **key** configurations **to** run a Spark job **on** a YARN **cluster** are:

master - Determines how **to** run the job.

As we want **for** this blog review **to execute** Spark **in** YARN, the 'yarn' **value** has been selected **for** the example above. The other options available include 'mesos' **and** 'standalone'.

deploy-mode - We selected '**cluster**' **to** run the above SparkPi example within the **cluster**. **To** run the problem outside **of** the **cluster**, **then select** the 'client' **option**.

driver-memory - The amount memory available **for** the driver process. **In** a YARN **cluster** Spark configuration the Application Master runs the driver.

executor-memory - The amount **of** memory allocated **to** the executor process.

executor-cores - the total **number of** cores allocated **to** the executor process

queue - The YARN queue name **on** which this job will run. **If** you have **not** already defined queues **to** your **cluster**, it **is** best **to** utilize the '**default**' queue.

Broadcast **Variable** & Accumulators

1) Broadcast Variables

As part **of** this topic we will see details about Broadcast Variables

At times we need **to** pass (broadcast) **some** information **to all** the executors It can be done **by using** broadcast variables

A broadcast **variable** can be **of** preliminary **type or** it could be a hash **map**

Here **are** few examples

Single **value** - Common discount percent **for all** the products

Hash **map** - look up **or map** side **join**

When very **large data set** (fact) **is** tried **to join with** smaller **data set** (dimension), broadcasting dimension can have considerable performance improvement.

Broadcast variables have **to** be immutable.

Define problem **statement and** design the application

Get total revenue per department **for each day**

Again we will be **using** our retail_db database

Please refer **for data** model

Department name **is in** departments

To get department name we need **to join** these tables

order_items

products

categories

departments

We will **first join** products, categories **and** departments

get product_id **and** department_name

broadcast **data** which contain product_id **and** department_name

Then join orders **and** order_items

Perform simple **join between** orders **and** order_items

As part of join look up **into** hashmap **with** product_id **to get** department_name

Use reduceByKey **to** compute revenue **for each date and** department

Step-1) **//Read** products, categories **and** departments


```
val departments = sc.textFile(inputPath + "/departments")
val categories = sc.textFile(inputPath + "/categories")
val products = sc.textFile(inputPath + "/products")
```

```
Step-2) //Join products, categories and departments
val departmentsMap = departments.map(rec => (rec.split(",")(0).toInt, rec.split(",")(1)))
val categoriesMap = categories.map(rec => (rec.split(",")(0).toInt, rec.split(",")(1).toInt))
val productsMap = products.map(rec => (rec.split(",")(1).toInt, rec.split(",")(0).toInt))
val productCategories = productsMap.join(categoriesMap)
val productCategoriesMap = productCategories.map(rec => (rec._2._2, rec._2._1))
val productDepartments = productCategoriesMap.join(departmentsMap)

Step-3) //Build hash map and create broadcast variable
val productDepartmentsMap = productDepartments.map(rec => (rec._2._1, rec._2._2)).distinct
val bv = sc.broadcast(productDepartmentsMap.collectAsMap())

Step-4) //Read orders and order_items
val orders = sc.textFile(inputPath + "/orders")
val orderItems = sc.textFile(inputPath + "/order_items")

Step-5) //Filter for completed orders and extract required fields from orders
val ordersCompleted = orders.filter(rec => (rec.split(",")(3) == "COMPLETE")).map(rec => (
  rec.split(",")(0).toInt, rec.split(",")(1)))

Step-6) //Extract required fields from order_items while looking up into hash map to get
department name
val orderItemsMap = orderItems.map(rec =>
  (rec.split(",")(1).toInt, (bv.value.get(rec.split(",")(2).toInt).get, rec.split(",")(4).toFloat)))

Step-7) //Join orders and order_items and compute daily revenue for each product
val ordersJoin = ordersCompleted.join(orderItemsMap)
val revenuePerDayPerDepartment = ordersJoin.map(rec => ((rec._2._1, rec._2._2._1), rec._2._2._2)).reduceByKey((acc, value) => acc + value)

Step-8) //Save output to file system of your choice
revenuePerDayPerDepartment.sortByKey().saveAsTextFile(outputPath)
```

Accumulators

In this topic we will see details about accumulators

It is important to perform some counts as part of application for
unit testing
data quality

These counters cannot be global variables as part of the program
Instead we need to use accumulator which will be managed by spark
Accumulators will be passed to executors and scope is managed across
all the executors or executor tasks
We will see

how accumulators are implemented
issues with accumulators

Implementation of accumulators

Take our program AvgRevenueDaily
Add as many accumulators as you desire

Create variable like val ordersCompletedAccum = sc.accumulator(0, "ordersCompleted count")
Update the appropriate Spark API RDD function to increment accumulator (see the sample code below)
Compile the program

Build jar and ship it to the remote cluster

Run on remote cluster

Open spark history server and review the appropriate executor task in which accumulators are implemented

You will see counter as part of history server

```
Step-1) val ordersRDD = sc.textFile(inputPath + "/orders")
      val orderItemsRDD = sc.textFile(inputPath + "/order_items")
```

```
Step-2) val ordersCompletedAccum = sc.accumulator(0, "ordersCompleted count")
      val ordersFilterInvokedAccum = sc.accumulator(0, "orders filter invoked count")
```

```
Step-3) val ordersCompleted = ordersRDD.
      filter(rec => {
        ordersFilterInvokedAccum += 1
        if (rec.split(",")(3) == "COMPLETE") {
          ordersCompletedAccum += 1
        }
        rec.split(",")(3) == "COMPLETE"
      })

      val ordersAccum = sc.accumulator(0, "orders count")
      val orders = ordersCompleted.
      map(rec => {
        ordersAccum += 1
        (rec.split(",")(0).toInt, rec.split(",")(1))
      })

      val orderItemsMapAccum = sc.accumulator(0, "orderItemsMap count")
      val orderItemsMap = orderItemsRDD.
      map(rec => {
        orderItemsMapAccum += 1
        (rec.split(",")(1).toInt, rec.split(",")(4).toFloat)
      })

      val orderItemsValuesAccum = sc.accumulator(0, "reduceByKey values count")
      val orderItems = orderItemsMap.
      reduceByKey((acc, value) => {
        orderItemsValuesAccum += 1
        acc + value
      })
```

```
Step-4) val ordersJoin = orders.join(orderItems)
      val ordersJoinMap = ordersJoin.map(rec => (rec._2._1, rec._2._2))
```

```
Step-5) val revenuePerDay = ordersJoinMap.aggregateByKey((0.0, 0)) (
      (acc, value) => (acc._1 + value, acc._2 + 1),
      (total1, total2) => (total1._1 + total2._1, total1._2 + total2._2))

      val averageRevenuePerDay = revenuePerDay.
      map(rec => (rec._1, BigDecimal(rec._2._1 / rec._2._2).
        setScale(2, BigDecimal.RoundingMode.HALF_UP).toFloat))

      val averageRevenuePerDaySorted = averageRevenuePerDay.
      sortByKey()

      averageRevenuePerDaySorted.
      map(rec => rec._1 + "," + rec._2).
      saveAsTextFile(outputPath)
```

Cache()

```
ex: rdd.cache() //default StorageLevel is MEMORY_LEVEL
```

```
ex: rdd.uncache()
```

Persist()

ex : rdd.persist(StorageLevel.<type>)

1)MEMORY_ONLY

In this storage level, RDD is stored as deserialized Java object in the JVM. If the size of RDD is greater than memory, It will not cache some partition and recompute them next time whenever needed. In this level the space used for storage is very high, the CPU computation time is low, the data is stored in-memory. It does not make use of the disk.

2)MEMORY_AND_DISK

In this level, RDD is stored as deserialized Java object in the JVM. When the size of RDD is greater than the size of memory, it stores the excess partition on the disk, and retrieve from disk whenever required. In this level the space used for storage is high, the CPU computation time is medium, it makes use of both in-memory and on disk storage.

3)MEMORY_ONLY_SER

This level of Spark store the RDD as serialized Java object (one-byte array per partition). It is more space efficient as compared to deserialized objects, especially when it uses fast serializer. But it increases the overhead on CPU. In this level the storage space is low, the CPU computation time is high and the data is stored in-memory. It does not make use of the disk.

4)MEMORY_AND_DISK_SER

It is similar to MEMORY_ONLY_SER, but it drops the partition that does not fits into memory to disk, rather than recomputing each time it is needed. In this storage level, The space used for storage is low, the CPU computation time is high, it makes use of both in-memory and on disk storage.

5) DISK_ONLY

In this storage level, RDD is stored only on disk. The space used for storage is low, the CPU computation time is high and it makes use of on disk storage.

ex : rdd.unpersist()

---- Spark SQL

import org.apache.spark.sql.SparkSession

object csvExample_1 {

case class sfpd(auctionid: String, bid: Float, bidtime: Float, bidder: String, bidderrate: Integer, openbid: Float, price: Float, item: String, daystolive: Integer)

def main(args: Array[String]) {

```
//val spark = SparkSession.builder.master("local[*]").appName("csvExample_1").config(
  "spark.sql.warehouse.dir", "/home/hadoop/work/warehouse").enableHiveSupport().getOrCreate()
val spark = SparkSession.builder.master("local[*]").appName("csvExample_1").getOrCreate()
val sc = spark.sparkContext
val sqlContext = spark.sqlContext
import spark.implicits._
val df = sqlContext.read.format("com.databricks.spark.csv").option("header", "true").option(
  "inferSchema", "true").load(
  "C:\\Kalyan\\POC\\Spark\\spark_datasets\\Police_Department_Incidents.csv")
df.show()
df.select("Category").distinct.collect().foreach(println)
df.createOrReplaceTempView("sfpd")
sqlContext.sql("select Category from sfpd").collect().foreach(println)
```

// top 10 results

```

sqlContext.sql("SELECT Resolution , count(Resolution) as rescount FROM sfpd group by
Resolution order by rescount desc limit 10").collect().foreach(println)
val t = sqlContext.sql("select Category,count(Category) as catcount from sfpd group by
Category order by catcount desc limit 10")
t.show()
t.map(t=> "column 0: "+ t(0)).collect().foreach(println)
spark.stop()
}
}

=====
package SparkSQL.SparkSQL_CSV

/**
 * Created by Kalyan on 4/20/2017.
 */

import org.apache.spark.sql.SparkSession

object csvExample_2 {
  def main(args: Array[String]) {
    //val spark = SparkSession.builder.master("local[*]").appName("csvExample_2").config(
    "spark.sql.warehouse.dir", "/home/hadoop/work/warehouse").enableHiveSupport().getOrCreate()
    val spark = SparkSession.builder.master("local[*]").appName("csvExample_2").getOrCreate()
    val sc = spark.sparkContext
    val sqlContext = spark.sqlContext
    import spark.implicits._
    val df = sqlContext.read.format("com.databricks.spark.csv").option("header","false").option(
    "inferSchema","true")
      .load("C:\\Kalyan\\POC\\Spark\\spark_datasets\\ebay.csv")
    df.show()
    df.select($"_c0".alias("auctionid"), $"_c1".alias("bid"), $"_c2".alias("bidtime"), $"_c3".alias(
    "bidder"), $"_c4".alias("bidderrate")
    , $"_c5".alias("openbid"), $"_c6".alias("price"), $"_c7".alias("item"), $"_c8".alias(
    "daystolive")).show()
    // sqlContext.sql("select")
    spark.stop()
  }
}

=====
package SparkSQL.SparkSQL_CSV

/**
 * Created by Kalyan on 4/20/2017.
 */

import org.apache.spark.sql.SparkSession

object csvExample_3 {
  def main(args: Array[String]) {
    //val spark = SparkSession.builder.master("local[*]").appName("csvExample_3").config(
    "spark.sql.warehouse.dir", "/home/hadoop/work/warehouse").enableHiveSupport().getOrCreate()
    val spark = SparkSession.builder.master("local[*]").appName("csvExample_3").getOrCreate()
    val sc = spark.sparkContext
    val sqlContext = spark.sqlContext
    val df = sqlContext.read.format("com.databricks.spark.csv").option("header","true").option(
    "inferSchema","true")
      .load("C:\\Kalyan\\POC\\Spark\\spark_datasets\\cars.csv")
    df.show()
    spark.stop()
  }
}

```

```
package SparkSQL.SparkSQL_CSV
```

```
/**
```

```
 * Created by Kalyan on 4/20/2017.
```

```
 */
```

```
import org.apache.spark.sql.SparkSession
```

```
object csvExample_4 {
```

```
  def main(args: Array[String]) {
```

```
    //val spark = SparkSession.builder.master("local[*]").appName("csvExample_4").config(
    "spark.sql.warehouse.dir", "/home/hadoop/work/warehouse").enableHiveSupport().getOrCreate()
    val spark = SparkSession.builder.master("local[*]").appName("csvExample_4").getOrCreate()
```

```
    val sc = spark.sparkContext
```

```
    val sqlContext = spark.sqlContext
```

```
    val df = sqlContext.read.format("com.databricks.spark.csv").option("header", "true").option(
    "inferSchema", "true")
```

```
      .load(
      "C:\\Kalyan\\POC\\Spark\\spark_datasets\\csv_data\\spark-csv-master\\src\\test\\resources\\
      ages-alternative-malformed.csv")
```

```
    df.show()
```

```
    spark.stop()
```

```
  }
```

```
}
```

```
package SparkSQL.SparkSQL_CSV
```

```
/**
```

```
 * Created by Kalyan on 4/28/2017.
```

```
 */
```

```
import org.apache.spark.sql.SparkSession
```

```
object csvExample_5 {
```

```
  def main(args: Array[String]) {
```

```
    //val spark = SparkSession.builder.master("local[*]").appName("csvExample_5").config(
    "spark.sql.warehouse.dir", "/home/hadoop/work/warehouse").enableHiveSupport().getOrCreate()
    val spark = SparkSession.builder.master("local[*]").appName("csvExample_5").getOrCreate()
```

```
    val sc = spark.sparkContext
```

```
    val sqlContext = spark.sqlContext
```

```
    import spark.implicits._
```

```
    import spark.sql
```

```
    val df = sqlContext.read.format("com.databricks.spark.csv").option("header", "true").option(
    "inferSchema", "true").load("C:\\Users\\Kalyan\\Desktop\\sampledata.csv")
```

```
    df.show()
```

```
    spark.stop()
```

```
  }
```

```
}
```

```
package SparkSQL.SparkSQL_JSON
```

```
/**
```

```
 * Created by Kalyan on 4/20/2017.
```

```
 */
```

```
import org.apache.spark.sql.SparkSession
```

```
object jsonExample_1 {
```

```
  def main(args: Array[String]) {
```

```
    //val spark = SparkSession.builder.master("local[*]").appName("jsonExample_1").config(
    "spark.sql.warehouse.dir", "/home/hadoop/work/warehouse").enableHiveSupport().getOrCreate()
    val spark = SparkSession.builder.master("local[*]").appName("jsonExample_1").getOrCreate()
```

```
    val sc = spark.sparkContext
```

```
    val sqlContext = spark.sqlContext
```

```

import spark.implicitly._
import spark.sql
val df = sqlContext.read.json(
"C:\\Kalyan\\POC\\Spark\\spark_datasets\\json_data\\world_bank.json")
// df.printSchema()
df.createOrReplaceTempView("jsondata_one")
sqlContext.sql("select url,totalamt,abc.code,abc.name from jsondata_one " +
"lateral view explode(theme_namecode) as abc").show(5)
// result.write.format("com.databricks.spark.csv").option("header","true").save(
"C:\\Kalyan\\POC\\Spark\\spark_datasets\\json_data\\jsontocsv")
spark.stop()
}
}

```

```

=====
package SparkSQL.SparkSQL_XML

```

```

/**
 * Created by Kalyan on 4/20/2017.
 */

```

```

import org.apache.spark.sql.SparkSession

```

```

object xmlExample_1 {
  def main(args: Array[String]) {
    //val spark = SparkSession.builder.master("local[*]").appName("jsonExample_1").config(
    "spark.sql.warehouse.dir", "/home/hadoop/work/warehouse").enableHiveSupport().getOrCreate()
    val spark = SparkSession.builder.master("local[*]").appName("jsonExample_1").getOrCreate()
    val sc = spark.sparkContext
    val sqlContext = spark.sqlContext
    val df = sqlContext.read.format("com.databricks.spark.xml").option("rootTag","books").option(
    ("rowTag","book"))
    .load("C:\\Kalyan\\POC\\Spark\\spark_datasets\\books.xml")
    df.show()
    // df.printSchema()
    //df.createOrReplaceTempView("jsondata_one")
    //sqlContext.sql("select url,totalamt,abc.code,abc.name from jsondata_one " +
    //"lateral view explode(theme_namecode) as abc").show(5)
    // result.write.format("com.databricks.spark.csv").option("header","true").save(
    "C:\\Kalyan\\POC\\Spark\\spark_datasets\\json_data\\jsontocsv")
    spark.stop()
  }
}

```

```

=====
Spark SQL
=====

```

```

Example-1
=====

```

```

step 1)import org.apache.spark.sql.SQLContext
val sqlContext = new SQLContext(sc)
import sqlContext.implicitly._//to convert RDD's into DataFrames implicitly

step 2)case class Sample(a:Int,b:Int,c:Int)
val s1 = Sample(10,20,30)
val s2 = Sample(1,2,3)
val s3 = sample(100,200,300)
val s4 = Sample(1000,2000,3000)

val data = sc.parallelize(List(s1,s2,s3,s4))

```

```
data.collect
val x = data.map(v => v.a+v.b+v.c).collect
```

```
step 3) val df = data.toDF //if your rdd having schema its eligible for DataFrame.
df.show()
```

```
step 4) df.registerTempTable("sample")
```

```
step 5) val result = sqlContext.sql("select * from sample)
val result1 = sqlContext.sql("select a,b from sample")
result.show
result.printSchema()
result1.show()
val result2 = sqlContext.sql("select a,b,c,a+b+c as total from sample")
result2.show()
```

Example-2

=====

```
step 1) creating a file locally emp
```

```
101,aaa,10000,m,11
102,bbb,20000,f,12
103,ccc,30000,m,13
104,ddd,40000,f,11
105,eee,50000,m,12
106,fff,60000,f,13
```

```
step 2) creating a directory in hdfs
```

```
hdfs dfs -mkdir /user/cloudera/spark_sql/

hdfs dfs -put emp /user/cloudera/spark_sql/
```

```
step 3) loading the data into spark
```

```
val data = sc.textFile("/user/cloudera/spark_sql/emp")
data.count
data.take(1)
```

```
step 4) case class Emp(id:Int,name:String,sal:Int,gendar:String,dno:Int)
```

```
step 5) def toEmp(x:String)={
  val w = x.trim().split(",")
  val id = w(0).toInt
  val name = w(1)
  val sal = w(2).toInt
  val gendar = w(3)
  val dno = w(4).toInt
  val empdetails = Emp(id,name,sal,gendar,dno)
  empdetails
}
```

```
          //To check the function is working or not(toEmp)
val rec = "101,kalyan,70000,m,12"
val i = toEmp(rec)
i.name
i.sal
```

```
step 6) val einfo = data.map(x => toEmp(x)) //its RDD and having schema
einfo.foreach(println)
```

```
step 7) convert einfo to DF.
```

```
val edf = einfo.toDF
edf.show()
edf.printSchema()
```

```
step 8) create temp table
```

```
edf.registerTempTable("edf")
```

```
val result = sqlContext.sql("select dno,gendar,sum(sal)
                             from edf group by dno,gendar")
result.show()
```

```
step 9) //RDD style
einfo.filter(x => x.gendar == "m").collect
//sql style
val result = sqlContext.sql("select * from edf where gendar='m'")
result.show()

//select gendar,sum(sal) from emp
//RDD style
val pair = einfo.map(x=>(x.gendar,x.sal))
val result1 = pair.reduceByKey(_+_).collect
//sql style

val result2 = sqlContext.sql("select gendar,sum(sal) as total_sal from edf group by
gendar").show
```

Example 3)

=====

//RDD style of multiple aggregations

```
val pair = einfo.map(x=>(x.sex,x.sal))
val grp = pair.groupByKey()
val res = grp.map{x =>
  val gendar = x._1
  val cb = x._2
  val tot = cb.sum
  val cnt = cb.size
  val avg = tot/cnt
  val max = cb.max
  val min = cb.min
  (sex,tot,cnt,avg,max,min)
}
```

//sql style

```
val result = sqlContext.sql("select gendar,sum(sal) as tot,count(*) as cnt,avg(sal),max(sal) as
max,min(sal) as min from edf group by gendar").show()
//sql statement always returns DF.
```

```
val result1 = sqlContext.sql("select dno,gendar,sum(sal) as tot,count(*) as
cnt,avg(sal),max(sal) as max,min(sal) as min from edf group by gendar").show()
```

Example 4) //create a file emp2

```
201,kiran,14,m,90000
202,kalyani,12,f,100000
203,anushka,11,f,80000
204,kalyan,13,m,70000
```

step 1) hdfs dfs -put emp2 /user/cloudera/spark_sql/

step 2) val data2 = sc.textFile("/user/cloudera/spark_sql/emp2")

```
step 3) val emp2 = data2.map{x=>
  val w = x.split(",")
  Emp(w(0).toInt,w(1).toInt,w(2).toInt,w(3).toInt,w(4).toInt)
}
emp2.collect().foreach(println)
```

step 4) val edf2 = emp2.toDF


```
step 5) edf2.registerTempTable("edf2")
```

```
step 6) val result = sqlContext.sql("select * from edf union all select * from edf2").show()
      result.registerTempTable("emp_result")
      val result1 =
sqlContext.sql("select gendar,sum(sal) from emp_result group by gendar").show()
```

```
Example 5)//create dept file
11,marketing,hyd
12,hr,bang,
13,finance,hyd
14,sales,pune
```

```
step 1) hdfs dfs -put dept /user/cloudera/spark_sql/
```

```
step 2) val ddata = sc.textFile("/user/cloudera/spark_sql/dept")
```

```
step 3) case class Dept(dno:Int,dname:String,loc:String)
```

```
step 4) val dept = ddata.map{x=>
      val w = x.split(",")
      val dno = w(0).int
      val dname =w(1)
      val loc = w(2)
      Dept(dno,dname,loc)
    }
    dept.collect().foreach(println)
```

```
step 5) val ddf = dept.toDF
```

```
step 6) ddf.registerTempTable("dept")
```

```
step 7) val result = sqlContext.sql("select loc,sum(sal) as tot from edf e join dept d
on(e.dno= d.dno) group by loc").show()
```

```
=====using spark sql connect to
```

```
hive=====
```

```
Example 1)
```

```
step 1) import org.apache.spark.sql.hive.HiveContext
```

```
step 2) val hc = new HiveContext(sc)
```

```
step 3) val dbname = hc.sql("create database mysparkdb")
```

```
step 4) val usedbname = hc.sql("use mysparkdb")
```

```
step 5) val table = hc.sql("create table emp(eno int,name string,sal int,gendar string,dno int)
row format delimited fields terminated by ','")
```

```
step 6) val load= hc.sql("load data inpath "/user/cloudera/SparkSql/emp"
into table emp")
```

```
step 7) val result = hc.sql("select * from emp").show()
```

```
Spark SQL
```

```
=====
```

```
Example-1
```

```
=====
```

```
step 1) import org.apache.spark.sql.SQLContext
```

```
      val sqlContext = new SQLContext(sc)
```

```
      import sqlContext.implicits._//to convert RDD's into DataFrames implicitly
```

```
step 2) case class Sample(a:Int,b:Int,c:Int)
```

```
      val s1 = Sample(10,20,30)
```

```

val s2 = Sample(1,2,3)
val s3 = sample(100,200,300)
val s4 = Sample(1000,2000,3000)

val data = sc.parallelize(List(s1,s2,s3,s4))

data.collect
val x = data.map(v => v.a+v.b+v.c).collect

```

step 3) val df = data.toDF //if your rdd having schema its eligible for DataFrame.
df.show()

step 4) df.registerTempTable("sample")

step 5) val result = sqlContext.sql("select * from sample")
val result1 = sqlContext.sql("select a,b from sample")
result.show
result.printSchema()
result1.show()
val result2 = sqlContext.sql("select a,b,c,a+b+c as total from sample")
result2.show()

Example-2

=====

step 1) creating a file locally emp

```

101,aaa,10000,m,11
102,bbb,20000,f,12
103,ccc,30000,m,13
104,ddd,40000,f,11
105,eee,50000,m,12
106,fff,60000,f,13

```

step 2) creating a directory in hdfs

```
hdfs dfs -mkdir /user/cloudera/spark_sql/
```

```
hdfs dfs -put emp /user/cloudera/spark_sql/
```

step 3) loading the data into spark

```

val data = sc.textFile("/user/cloudera/spark_sql/emp")
data.count
data.take(1)

```

step 4) case class Emp(id:Int,name:String,sal:Int,gendar:String,dno:Int)

```

step 5) def toEmp(x:String)={
  val w = x.trim().split(",")
  val id = w(0).toInt
  val name = w(1)
  val sal = w(2).toInt
  val gendar = w(3)
  val dno = w(4).toInt
  val empdetails = Emp(id,name,sal,gendar,dno)
  empdetails
}

```

```

//To check the function is working or not(toEmp)
val rec = "101,kalyan,70000,m,12"
val i = toEmp(rec)
i.name
i.sal

```

step 6) val einfo = data.map(x => toEmp(x)) //its RDD and having schema
einfo.foreach(println)

step 7) convert einfo to DF.

```
val edf = einfo.toDF
edf.show()
edf.printSchema()
```

```
step 8) create temp table
edf.registerTempTable("edf")
```

```
val result = sqlContext.sql("select dno,gendar,sum(sal)
                             from edf group by dno,gendar")
result.show()
```

```
step 9) //RDD style
einfo.filter(x => x.gendar == "m").collect
//sql style
val result = sqlContext.sql("select * from edf where gendar='m'")
result.show()

//select gendar,sum(sal) from emp
//RDD style
val pair = einfo.map(x=>(x.gendar,x.sal))
val result1 = pair.reduceByKey(_+_).collect
//sql style

val result2 = sqlContext.sql("select gendar,sum(sal) as total_sal from edf group by
gendar").show
```

Example 3)

=====

```
//RDD style of multiple aggregations
val pair = einfo.map(x=>(x.sex,x.sal))
val grp = pair.groupByKey()
val res = grp.map{x =>
    val gendar = x._1
    val cb = x._2
    val tot = cb.sum
    val cnt = cb.size
    val avg = tot/cnt
    val max = cb.max
    val min = cb.min
    (sex,tot,cnt,avg,max,min)
}
```

```
//sql style
val result = sqlContext.sql("select gendar,sum(sal) as tot,count(*) as cnt,avg(sal),max(Sal) as
max,min(sal) as min from edf group by gendar").show()
//sql statement always returns DF.
```

```
val result1 = sqlContext.sql("select dno,gendar,sum(sal) as tot,count(*) as
cnt,avg(sal),max(Sal) as max,min(sal) as min from edf group by gendar").show()
```

Example 4) //create a file emp2

```
201,kiran,14,m,90000
202,kalyani,12,f,100000
203,anushka,11,f,80000
204,kalyan,13,m,70000
```

```
step 1) hdfs dfs -put emp2 /user/cloudera/spark_sql/
```

```
step 2) val data2 = sc.textFile("/user/cloudera/spark_sql/emp2")
```

```
step 3) val emp2 = data2.map{x=>
    val w = x.split(",")
```

```
Emp(w(0).toInt,w(1),w(4).toInt,w(3),w(2).toInt)
}
emp2.collect().foreach(println)
```

step 4) val edf2 = emp2.toDF

step 5) edf2.registerTempTable("edf2")

```
step 6) val result = sqlContext.sql("select * from edf union all select * from edf2").show()
result.registerTempTable("emp_result")
val result1 =
```

```
sqlContext.sql("select gendar,sum(sal) from emp_result group by gendar").show()
```

Example 5) //create dept file

11,marketing,hyd

12,hr,bang,

13,finance,hyd

14,sales,pune

step 1) hdfs dfs -put dept /user/cloudera/spark_sql/

step 2) val ddata = sc.textFile("/user/cloudera/spark_sql/dept")

step 3) case class Dept(dno:Int,dname:String,loc:String)

```
step 4) val dept = ddata.map{x=>
    val w = x.split(",")
    val dno = w(0).int
    val dname = w(1)
    val loc = w(2)
    Dept(dno,dname,loc)
}
dept.collect().foreach(println)
```

step 5) val ddf = dept.toDF

step 6) ddf.registerTempTable("dept")

```
step 7) val result = sqlContext.sql("select loc,sum(sal) as tot from edf e join dept d
on(e.dno= d.dno) group by loc").show()
```

=====using spark sql connect to

hive=====

Example 1)

step 1) import org.apache.spark.sql.hive.HiveContext

step 2) val hc = new HiveContext(sc)

step 3) val dbname = hc.sql("create database mysparkdb")

step 4) val usedbname = hc.sql("use mysparkdb")

```
step 5) val table = hc.sql("create table emp(eno int,name string,sal int,gendar string,dno int)
row format delimited fields terminated by ','")
```

```
step 6) val load= hc.sql("load data inpath "/user/cloudera/SparkSql/emp"
into table emp")
```

step 7) val result = hc.sql("select * from emp").show()

Example 2)

```
{"name":"kalyan","age":33}
```

```
{"name":"aparna","city":"hyd"}
```

```
{"name":"Akhil","age":5,"city":"nrt"}
```

```
{"name":"Akshaya","age":2}
```

```

step 3) create table jraw(line string);

step 4) load data local inpath 'sample_json' into table jraw;

step 5) select get_json_object(line,'$.name') from jraw;

        select get_json_object(line,'$.name'),
get_json_object(line,'$.age'),get_json_object(line,'$.city') from jraw;

step 6)select x.* from jraw
        lateral view json_tuple(line,'name','age','city') x as n,a,c;

```

```

//By using Spark SQL
step 1) hadoop fs -put sample_json /user/cloudera/spark_sql/

step 2) val df = sqlContext.read.json("/user/cloudera/Spark_Sql/sample_json")

step 3)df.show()

```

```

Example 3)sample_json2
{"name":"Kalyan","age":33,"wife":{"name":"aparna","age":28,"city":"hyd"},
"city":"hyd"}
{"name":"anil","age":32,"wife":{"name":"praveena","qual":"bsc","city":"cpt"},
"city":"nrt"}

```

```

//hive
create table jraw2(line string)
load data inpath '/user/cloudera/Spark_Sql/sample_json2' into table jraw2;

create table json2(name string,age int,wife string,city string);

insert into table json2
select x.* from jraw2
lateral view json_tuple(line,'name','age','wife','city') x as n,a,w,c;

select * from json2;

select name,get_json_object(wife,'$.name'),age,get_json_object(wife,'$.age'),
get_json_object(wife,'$.qual'),city,get_json_object(wife,'$.city') from json2;

```

```
//using Spark Sql
```

```
val data = sqlContext.read.json("/user/cloudera/Spark_Sql/sample_json2")
data.show()
data.collect.map(x=> x(3))
data.collect.map(x=> x(3)(1))
data.collect.map(x=> x(3)(2))
```

Example 4) Handling xml files using hive and spark sql

```
<rec><name>Kalyan</name><age>33</age></rec>
<rec><name>Aparna</name><gendar>F</gendar></rec>
<rec><name>Anil</name><age>33</age><gendar>M</gendar></rec>
```

```
hc.sql("use mysparkdb")
hc.sql("create table xraw(line string)")
hc.sql("load data inpath 'xml1' into table xraw")
hc.sql("create table xinfo(name string,age int,gendar string)
row format delimited fields terminated by ','")
hc.sql("insert into table xinfo
      select xpath_string(line,'rec/name'),xpath_int(line,'rec/age'),
      xpath_string(line,'rec/gendar') from xraw")
```

```
--- Spark To MYSQL
```

```
mysql to spark
```

```
-----
```

```
package SparkSQL.RDS
/**
 * Created by Kalyan on 5/31/2017.
 */
import org.apache.spark.sql.SparkSession
object mysqlConnect {
  def main(args: Array[String]) {
    //val spark =
    SparkSession.builder.master("local[*]").appName("coalesceTransformation").config("spark.sql.warehouse.dir", "/home/hadoop/work/warehouse").enableHiveSupport().getOrCreate()
    val spark =
    SparkSession.builder.master("local[*]").appName("coalesceTransformation").getOrCreate()
    val sc = spark.sparkContext
    val sqlContext = spark.sqlContext
    import java.util.Properties
    val mtable = "emp"
    val hostname="jdbc:mysql://mysql.c5ghyblpskst.ap-south-1.rds.amazonaws.com:3306/mysqlpdb"
    val mprp = new Properties()
    mprp.setProperty("driver","com.mysql.jdbc.Driver")
    mprp.setProperty("user","musername")
    mprp.setProperty("password","mpassword")
    val odf = spark.read.jdbc(hostname,mtable,mprp)
    odf.show()
    spark.stop()
  }
}
```

```

-----
oracle to spark
-----
package SparkSQL.RDS

/**
 * Created by Kalyan on 5/31/2017.
 */

import org.apache.spark.sql.SparkSession

object oracleConnect {
  def main(args: Array[String]) {
    //val spark =
    SparkSession.builder.master("local[*]").appName("coalesceTransformation").config("spark.sql.w
    arehouse.dir", "/home/hadoop/work/warehouse").enableHiveSupport().getOrCreate()
    val spark =
    SparkSession.builder.master("local[*]").appName("coalesceTransformation").getOrCreate()
    val sc = spark.sparkContext
    val sqlContext = spark.sqlContext
    import spark.implicits._
    import spark.sql
    import java.util.Properties
    val otable = "dept"
    val hostname="jdbc:oracle:
    thin://@oracle.c5ghyblpskst.ap-south-1.rds.amazonaws.com:1521/oracledb"
    val oprp = new Properties()
    oprp.setProperty("driver", "oracle.jdbc.OracleDriver")
    oprp.setProperty("user", "ouusername")
    oprp.setProperty("password", "opassword")

    val odf = spark.read.jdbc(hostname, otable, oprp)
    odf.show()
    //    val result = odf.where("DEPTNO>20")
    //    result.write.format("csv").option("header", "true").save("D:\\Sprak-Training\\oracleoutput")

    spark.stop()
  }
}

---- DATA SETS IN SPARK

ex 1)
val dataset = Seq(1, 2, 3).toDS()
dataset.show()

ex 2)
case class Person(name: String, age: Int)

val personDS = Seq(Person("Max", 33), Person("Adam", 32), Person("Muller", 62)).toDS()
personDS.show()

ex 3)
val rdd = sc.parallelize(Seq((1, "Spark"), (2, "Databricks")))
val integerDS = rdd.toDS()
integerDS.show()

ex 4)
case class Company(name: String, foundingYear: Int, numEmployees: Int)
val inputSeq = Seq(Company("ABC", 1998, 310), Company("XYZ", 1983, 904), Company("NOP", 2005,
83))
val df = sc.parallelize(inputSeq).toDF()

val companyDS = df.as[Company]
companyDS.show()

```

```

ex 5)
val rdd = sc.parallelize(Seq((1, "Spark"), (2, "Databricks"), (3, "Notebook")))
val df = rdd.toDF("Id", "Name")

val dataset = df.as[(Int, String)]
dataset.show()

ex 6)
val wordsDataset = sc.parallelize(Seq("Spark I am your father", "May the spark be with you",
"Spark I am your father")).toDS()
val groupedDataset = wordsDataset.flatMap(_.toLowerCase.split(" "))
                                .filter(_ != "")
                                .groupBy("value")

val countsDataset = groupedDataset.count()
countsDataset.show()

ex 7)
case class Employee(name: String, age: Int, departmentId: Int, salary: Double)
case class Department(id: Int, name: String)

case class Record(name: String, age: Int, salary: Double, departmentId: Int, departmentName:
String)
case class ResultSet(departmentId: Int, departmentName: String, avgSalary: Double)

val employeeDataSet1 = sc.parallelize(Seq(Employee("Max", 22, 1, 100000.0), Employee("Adam",
33, 2, 93000.0), Employee("Eve", 35, 2, 89999.0), Employee("Muller", 39, 3, 120000.0))).toDS()
val employeeDataSet2 = sc.parallelize(Seq(Employee("John", 26, 1, 990000.0), Employee("Joe",
38, 3, 115000.0))).toDS()
val departmentDataSet = sc.parallelize(Seq(Department(1, "Engineering"), Department(2,
"Marketing"), Department(3, "Sales"))).toDS()

val employeeDataset = employeeDataSet1.union(employeeDataSet2)

def averageSalary(key: (Int, String), iterator: Iterator[Record]): ResultSet = {
    val (total, count) = iterator.foldLeft(0.0, 0.0) {
        case ((total, count), x) => (total + x.salary, count + 1)
    }
    ResultSet(key._1, key._2, total/count)
}

val averageSalaryDataset = employeeDataset.joinWith(departmentDataSet, $"departmentId" ===
$"id", "inner")

                                .map(record => Record(record._1.name, record._1.age,
record._1.salary, record._1.departmentId,
record._2.name))
                                .filter(record => record.age > 25)
                                .groupBy($"departmentId", $"departmentName")
                                .avg()

ex 8)
import org.apache.spark.sql.functions._

val wordsDataset = sc.parallelize(Seq("Spark I am your father", "May the spark be with you",
"Spark I am your father")).toDS()
val result = wordsDataset
    .flatMap(_.split(" "))           // Split on whitespace
    .filter(_ != "")                 // Filter empty words
    .map(_.toLowerCase())
    .toDF()                          // Convert to DataFrame to perform
    aggregation / sorting
    .groupBy($"value")                // Count number of occurrences of each word
    .agg(count("*") as "numOccurrences")
    .orderBy($"numOccurrences" desc) // Show most common words first

result.show()

--- Spark Streaming

```



```

import java.util.Properties;
import java.sql.*;
import kafka.javaapi.producer.Producer;
import kafka.producer.KeyedMessage;
import kafka.producer.ProducerConfig;

public class JdbcProducer{
public static void main(String[] args) throws ClassNotFoundException,SQLException{
    Properties props = new Properties();
    props.put("zk.connect","localhost:2181");
    props.put("serializer.class","kafka.serializer.StringEncoder");
    props.put("metadata.broker.list","localhost:9092");
    ProducerConfig config = new ProducerConfig(props);
    Producer producer = new Producer(config);
    try{
        class.forName("com.mysql.jdbc.Driver")
        Connection con = DriverManager.getConnection(
            "jdbc:mysql://localhost:3306/test","root","root");
        Statement stmt = con.createStatement();
        ResultSet rs = stmt.executeQuery("select * from emp");
        while(rs.next())
            producer.send(new KeyedMessage("test",rs.getString(1)+"
            "+rs.getString(2)));
        con.close()
    }catch (Exception e){
        System.out.println(e)
    }
}
}

```

```
spark-shell --master local[2]
```

```

import org.apache.spark._
import org.apache.spark.streaming._
import org.apache.spark.streaming.StreamingContext._

```

```

val ssc = new StreamingContext(sc,Seconds(10))
val lines = ssc.socketTextStream("localhost",9999)
val words = lines.flatMap(x => x.split(" "))
val pairs = words.map(x => (x,1))
val res = pairs.reduceByKey(_+_ )
res.print()
ssc.start

```

```
nc -lk 9999
```

```
-----
kafka - Spark Streaming Example
```

- 1) please start zookeeper
./bin/zookeeper-server-start.sh config/zookeeper.properties &
- 2) please start kafka-Broker
./bin/kafka-server-start.sh config/server.properties &
- 3) spark-shell --master local[2]
- 4) create two topics mytopic,results
./bin/kafka-topics.sh --create --zookeeper localhost:2182 --partitions 1
--replication-factor 1 --topic mytopic

./bin/kafka-topics.sh --create --zookeeper localhost:2182 --partitions 1
--replication-factor 1 --topic results

./bin/kafka-topics.sh --list --zookeeper localhost:2182
- 5) please start console producer with mytopic
./bin/kafka-console-producer.sh --broker-list localhost:9092 --topic mytopic

```

6) create another to write the spark streaming output data(topic name result)

7) please start console consumer to read the data from result topic.
./bin/kafka-console-consumer.sh --bootstrap-server localhost:9092
--topic results --from-beginning

spark-shell --master local[2]
8) import org.apache.spark.streaming.StreamingContext
import org.apache.spark.streaming.Seconds
val ssc = new StreamingContext(sc, Seconds(5))
import org.apache.spark.streaming.kafka.KafkaUtils
val kafkaStream = KafkaUtils.createStream(ssc,"localhost:2182",
"spark-streaming-consumer-group",Map("mytopic" ->5))
val lines = kafkaStream.map(x => x._2.toUpperCase)
val words = lines.flatMap(x => x.split(" "))
val pairs = words.map(x => (x,1))
val res = pairs.reduceByKey(_+_ )
import org.apache.kafka.clients.producer.ProducerConfig
import java.util.HashMap
import org.apache.kafka.clients.producer.KafkaProducer
import org.apache.kafka.clients.producer.ProducerRecord
res.foreachRDD(rdd =>
  rdd.foreachPartition(partition =>
    partition.foreach{
      case(w:String,cnt:Int)=>{
        val x = w+"\t"+cnt
        val props = new HashMap[String,Object]()
        props.put(ProducerConfig.BOOTSTRAP_SERVERS_CONFIG,"localhost:9092")

        props.put(ProducerConfig.VALUE_SERIALIZER_CLASS_CONFIG,"org.apache.kafka.common.serialization.StringSerializer")

        props.put(ProducerConfig.KEY_SERIALIZER_CLASS_CONFIG,"org.apache.kafka.common.serialization.StringSerializer")

import org.apache.kafka.clients.producer.KafkaProducer
import org.apache.kafka.clients.producer.ProducerRecord
val producer = new KafkaProducer[String,String](props)
val message = new ProducerRecord[String,String]("results",null,x)
        producer.send(message)

      }
    }
  )

8)please start streaming context
ssc.start

```

Spark Streaming

```

Ex 1) spark-shell --master local[2]
import org.apache.spark._
import org.apache.spark.streaming._
import org.apache.spark.streaming.StreamingContext._
val ssc = new StreamingContext(sc,Seconds(10)) //10 seconds is microbatch period
val data = ssc.socketTextStream("localhost",9999)
val word = data.flatMap(x => x.split("\\W+"))
val pairs = word.map(x => (x,1))
val res = pairs.reduceByKey(_+_ )
res.print()
ssc.start //streaming started

```

```
Ex 2) Sliding and Windowing
pairs.persist()
pairs.reduceByKeyAndWindow(_+_ ,20,40)//20 seconds for sliding and 40 seconds for windowing
//sliding period is greater than equal to micro batch period.
```

```
--- SQOOP -----
```

```
sqoop.txt
Details
Activity
sqoop.txt
Sharing Info
k
```

```
J
R
```

```
+7
General Info
Type
Text
Size
10 KB (9,920 bytes)
Storage used
0 bytesOwned by someone else
Location
Sqoop
Owner
kumar K
Modified
18 Feb 2018 by kumar K
Opened
18:39 by me
Created
18 Feb 2018
Description
Add a description
Download permissions
Viewers can download
All selections cleared All selections cleared
```

```
Creating the tables in MySql DB
```

```
-----
CREATE TABLE cities (
id INTEGER UNSIGNED NOT NULL,
country VARCHAR(50),
city VARCHAR(150),
PRIMARY KEY (id)
);
```

```
INSERT INTO cities(id,country,city) VALUES (1, "USA", "Palo Alto");
INSERT INTO cities(id,country,city) VALUES (2, "Czech Republic", "Brno");
INSERT INTO cities(id,country,city) VALUES (3, "USA", "Sunnyvale");
```

```
CREATE TABLE countries(
country_id INTEGER UNSIGNED NOT NULL,
country VARCHAR(50),
PRIMARY KEY (country_id)
);
```

```
INSERT INTO countries(country_id,country) VALUES (1, "USA");
INSERT INTO countries(country_id,country) VALUES (2, "Czech Republic");
```

```
CREATE TABLE normcities (
id INTEGER UNSIGNED NOT NULL,
country_id INTEGER UNSIGNED NOT NULL,
city VARCHAR(150),
PRIMARY KEY (id),
FOREIGN KEY (country_id) REFERENCES countries(country_id)
);
```

```
INSERT INTO normcities(id,country_id,city) VALUES (1, 1, "Palo Alto");
INSERT INTO normcities(id,country_id,city) VALUES (2, 2, "Brno");
INSERT INTO normcities(id,country_id,city) VALUES (3, 1, "Sunnyvale");
```

```
CREATE TABLE visits (
id INTEGER UNSIGNED NOT NULL,
city VARCHAR(50),
last_update_date DATETIME NOT NULL,
PRIMARY KEY (id),
KEY (last_update_date)
);
```

```
INSERT INTO visits(id,city,last_update_date) VALUES(1, "Freemont", "1983-05-22 01:01:01");
INSERT INTO visits(id,city,last_update_date) VALUES(2, "Jicin", "1987-02-02 02:02:02");
```

---To find out user name MySql

```
1)select USER();
2)select CURRENT_USER();
```

Providing the permissions to HDFS home directory

```
-----
hdfs dfs -chmod -R 777 /user/hadoop;
```

Importing the data from MySql DB to HDFS,Hive and HBase

```
-----
```

1)Importing the data from mysql db to your HDFS if the table does not have a primary key and using warehouse-dir.

```
sqoop import -connect jdbc:mysql://localhost:3306/test -username root --table cities
--warehouse-dir /cities_output/ -m 1;
```

```
sqoop import -connect jdbc:mysql://localhost:3306/test -username root --table countries
--warehouse-dir /cities_output/ -m 1;
```

2)Importing the data from mysql db to your HDFS if the table does not have a primary key.

```
sqoop import -connect jdbc:mysql://localhost:3306/test -username root --table emp --target-dir
/sqoop -m 1;
```

3)Importing the data from mysql db to your HDFS if the table have primary key.

```
sqoop import -connect jdbc:mysql://localhost:3306/test -username root --table users
```

```
--target-dir /sqoop;
```

4) Importing the data from mysql db to your HDFS with compression.

```
sqoop import -connect jdbc:mysql://localhost:3306/test -username root --table cities
--target-dir /cities_compress --compress;
```

5) Importing the data from mysql db to your HDFS with where condition.

```
sqoop import -connect jdbc:mysql://localhost:3306/test -username root --table emp
--target-dir /some_emp --where "empname='XYZ'";
```

6) Importing the data from mysql db to your HDFS with direct condition.

```
sqoop import -connect jdbc:mysql://localhost:3306/test -username root --table users
--target-dir /fast_load --direct -m 1;
```

7) Importing the data from mysql db to your HDFS with overriding the type.

```
sqoop import -connect jdbc:mysql://localhost:3306/test -username root --table users
--target-dir /map_col --map-column-java user_id=Long;
```

8) Importing the data from mysql db to your HDFS applying parallelism

```
sqoop import -connect jdbc:mysql://localhost:3306/test -username root --table cities
--target-dir /sqoop_para --num-mappers 10;
```

9) Importing the data from mysql db to your HDFS --handling of null values

```
sqoop import -connect jdbc:mysql://localhost:3306/test -username root --table cities
--target-dir /handling_null --null-string '\\N' --null-non-string '\\N';
```

10) Importing the all tables from mysql db to your HDFS

```
sqoop import-all-tables -connect jdbc:mysql://localhost:3306/test -username root
--warehouse-dir /test_db;
```

11) Importing the some tables from mysql db to your HDFS

```
sqoop import-all-tables -connect jdbc:mysql://localhost:3306/test -username root
--warehouse-dir /test_dbase --exclude-tables emp,employee,users,empl,nyse;
```

12) Importing the data from 2 tables in mysql db to your HDFS.

```
sqoop import -connect jdbc:mysql://localhost:3306/test -username root --query 'SELECT normcities
.id,countries.country, normcities.city FROM normcities JOIN countries USING(country_id) WHERE $
CONDITIONS' --split-by id --target-dir /norm_cities;
```

```
--boundary-query "select min(id), max(id) from normcities"
```

13) Importing the data from 2 tables and applying boundary conditions in mysql db to your HDFS.

```
sqoop import -connect jdbc:mysql://localhost:3306/test -username root --query 'SELECT normcities
.id,countries.country, normcities.city FROM normcities JOIN countries USING(country_id) WHERE $
CONDITIONS' --split-by id --target-dir /norm_cities_boundary --boundary-query "select
min(id), max(id) from normcities";
```

14) Importing the data from mysql db to hive if the table have primary key.

```
sqoop import --connect jdbc:mysql://localhost:3306/test -username root --table users
--hive-table users --hive-import --target-dir /sqoop/hive/users;
```

15) Importing the data from mysql db to hive if the table not have primary key.

```
sqoop import --connect jdbc:mysql://localhost:3306/test -username root --table employee
```

```
--hive-table users --hive-import --target-dir /sqoop/hive/employee -m 1;
```

16) Importing the data from mysql db to hbase if the table not have primary key.

```
sqoop import --connect jdbc:mysql://localhost:3306/test --username root --table empl
--hbase-table employeeHBase --column-family info --hbase-row-key empid --hbase-create-table -m 1;
```

17) Importing the data from mysql db to hbase if the table have primary key.

```
sqoop import --connect jdbc:mysql://localhost:3306/test --username root --table empl
--hbase-table employeeHBase --column-family info --hbase-row-key empid --hbase-create-table;
```

18) If table have primary key and import only few columns of MySQL table into HBase table.

```
sqoop import --connect jdbc:mysql://localhost:3306/test --username root --table cities
--hbase-table empdet --columns id,country --column-family empf1 --hbase-create-table
```

Note : Column names specified in --columns attribute must contain the primary key column.

19) If table doesn't have **primary key** then choose one **column** as a **hbase-row-key**. Import **all** the **column** of MySQL table into HBase table.

```
sqoop import --connect jdbc:mysql://localhost:3306/test --username root --table emp
--hbase-table empwopk --column-family f1 --hbase-row-key empid --hbase-create-table -m 1;
```

20) If table doesn't have primary key then choose one column as a hbase-row-key. Import only few columns of MySQL table into HBase table.

```
sqoop import --connect jdbc:mysql://localhost:3306/test --username root --table emp
--hbase-table empwopkfc --columns empid,empname --column-family f1 --hbase-row-key empid
--hbase-create-table -m 1;
```

Note: Column name specified in hbase-row-key attribute must be in columns list. Otherwise command will execute successfully but no records are inserted into hbase.

21) Let us consider a MySQL table empdemo which have two columns name, address. The table empdemo doesn't have **primary key or unique key column**.

Records of empdemo:

```
create table empdemo(
-> name varchar(10),
-> address varchar(20)
-> );
```

```
insert into empdemo(name,address) values ('abc','123'),('sqw','345'),('abc','125'),('sdf',
'1234'),('aql','23dw');
```

name	address
abc	123
sqw	345
abc	125
sdf	1234
aql	23dw

```
sqoop import --connect jdbc:mysql://localhost:3306/test --username root --table empdemo
--hbase-table empdemo --column-family ed --hbase-row-key name --hbase-create-table -m 1;
```

Note:

Only 4 records are visible into HBase table instead of 5. In above example two rows have same value 'abc' of name column and value of this column is used as a HBase row key value. If record having value 'abc' of name column come then thoes record will inserted into HBase table. Next time, another record having the same value 'abc' of name column come then thoes column will

overwrite the **value** previous **column**.

Above problem also occurred **if table** have composite **primary key** because the one **column from** composite **key is** used **as** a HBase **row key**.

22) Saving the **output format in** a **sequence file**.

```
sqoop import --connect jdbc:mysql://localhost:3306/test -username root --table cities
--warehouse-dir /cities_output1/ --as-sequencefile -m 1;
```

23) Saving the **output format in** a **avrodata file**.

```
sqoop import --connect jdbc:mysql://localhost:3306/test -username root --table cities
--warehouse-dir /cities_output2/ --as-avrodatafile -m 1;
```

24) Protecting your password.

Sqoop execution that will **read** the password **from** standard **input**:

```
sqoop import --connect jdbc:mysql://localhost:3306/test -username root -P --table cities
--warehouse-dir /cities_output2/ --as-avrodatafile -m 1;
```

Reading the password **from** a **file**:

```
sqoop import --connect jdbc:mysql://localhost:3306/test -username root --password-file <path of
the file contains pwd> --table cities --warehouse-dir /cities_output2/ --as-avrodatafile -m 1;
```

Exporting the **data from** HDFS to MySQL DB

1) Exporting the **data from** HDFS db to your MySQL DB **if** the **table** does **not** have a **primary key**.

```
sqoop export --connect jdbc:mysql://localhost:3306/test --username root --table emp
--export-dir /sqoop/part-m-00000 --input-fields-terminated-by ',' -m 1;
```

2) Exporting the **data from** HDFS db to your MySQL DB **if** the **table** does **not** have a **primary key**.

```
sqoop export --connect jdbc:mysql://localhost:3306/test --username root --table emp
--export-dir /sqoop/part-m-00000 --input-fields-terminated-by ',';
```

SQOOP PRACTICE.

```
1) sqoop list-databases --connect "jdbc:mysql://quickstart.cloudera:3306" --username retail_dba
--password cloudera
```

```
2) sqoop list-databases --connect "jdbc:mysql://quickstart.cloudera:3306" --username retail_dba
-P
```

```
3) sqoop list-tables --connect "jdbc:mysql://quickstart.cloudera:3306/retail_db" --username
retail_dba --password cloudera
```

```
4) sqoop eval --connect "jdbc:mysql://quickstart.cloudera:3306/retail_db" --username retail_dba
--password cloudera --query "select count(1) from order_items"
```

```
5) sqoop import --all-tables -m 12 --connect "jdbc:mysql://quickstart.cloudera:3306/retail_db"
--username=retail_dba --password=cloudera --as-avrodatafile
--warehouse-dir=/user/hive/warehouse/retail_stage.db
```

```
6) sqoop import --connect "jdbc:mysql://quickstart.cloudera:3306/retail_db"
--username=retail_dba --password=cloudera --table departments --as-textfile
--target-dir=/user/cloudera/departments
```

```
7) sqoop import --connect "jdbc:mysql://quickstart.cloudera:3306/retail_db"
--username=retail_dba --password=cloudera --table departments --as-sequencefile
--target-dir=/user/cloudera/departments_seq
```

```

8) sqoop import --connect "jdbc:mysql://quickstart.cloudera:3306/retail_db"
--username=retail_dba --password=cloudera --table departments --as-avrodatafile
--target-dir=/user/cloudera/departments

9) sqoop import--all-tables --num-mappers 1 --connect
"jdbc:mysql://quickstart.cloudera:3306/retail_db" --username=retail_dba --password=cloudera
--hive-import --hive-overwrite --create-hive-table --compress --compression-codec
org.apache.hadoop.io.compress.SnappyCodec

10) sqoop import --connect "jdbc:mysql://quickstart.cloudera:3306/retail_db"
--username=retail_dba --password=cloudera --table departments --target-dir
/user/cloudera/departments -m 2 --boundary-query "select 2, 8 from departments"--columns
department_id,department_name

sqoop import --connect "jdbc:mysql://quickstart.cloudera:3306/retail_db"
--username=retail_dba --password=cloudera --table departments --target-dir
/user/cloudera/departments -m 2 --boundary-query "select
min(department_id),max(department_id) from departments"

sqoop import --connect "jdbc:mysql://quickstart.cloudera:3306/retail_db" --username=retail_dba
--password=cloudera --table departments --target-dir /user/cloudera/departments -m 2
--boundary-query "select 2,8 from departments where department_id <>1000 "

```

```

11) sqoop import --connect "jdbc:mysql://quickstart.cloudera:3306/retail_db"
--username=retail_dba --password=cloudera --query="select * from orders join order_items on
orders.order_id = order_items.order_item_order_id where \${CONDITIONS}" --target-dir
/user/cloudera/order_join --split-by order_id --num-mappers 4

```

```

12 )sqoop import --connect "jdbc:mysql://quickstart.cloudera:3306/retail_db"
--username=retail_dba --password=cloudera --table departments --target-dir
/user/hive/warehouse/retail_ods.db/departments --append --fields-terminated-by '|'
--lines-terminated-by '\n' --num-mappers 1

```

```

13) sqoop import --connect "jdbc:mysql://quickstart.cloudera:3306/retail_db"
--username=retail_dba --password=cloudera --table departments --target-dir
/user/hive/warehouse/retail_ods.db/departments --append --fields-terminated-by '|'
--lines-terminated-by '\n' --split-by department_id

```

Note :- Importing **table with out primary key using** multiple threads (split-by).When using split-by, using indexed **column is** highly desired.If the **column is not** indexed then performance will be bad because **of** full **table** scan **by each of** the thread.

```

14) sqoop import --connect "jdbc:mysql://quickstart.cloudera:3306/retail_db"
--username=retail_dba --password=cloudera --table departments --target-dir
/user/hive/warehouse/retail_ods.db/departments --append --fields-terminated-by '|'
--lines-terminated-by '\n' --split-by department_id --where "department_id > 7"

```

```

15) sqoop import --connect "jdbc:mysql://quickstart.cloudera:3306/retail_db"
--username=retail_dba --password=cloudera --table departments --target-dir
/user/hive/warehouse/retail_ods.db/departments --append --fields-terminated-by '|'
--lines-terminated-by '\n' --check-column "department_id" --incremental append --last-value 7

```

```

16) sqoop job --create sqoop_job -- import --connect "jdbc:
mysql://quickstart.cloudera:3306/retail_db" --username=retail_dba --password=cloudera --table
departments --target-dir /user/hive/warehouse/retail_ods.db/departments --append
--fields-terminated-by '|' --lines-terminated-by '\n' --check-column "department_id"
--incremental append --last-value 1000

```

```
sqoop job --list
```

```
sqoop job --exec sqoop_job
```

```

17) sqoop import --connect "jdbc:mysql://quickstart.cloudera:3306/retail_db"
--username=retail_dba --password=cloudera --table departments --fields-terminated-by '|'

```



```
--lines-terminated-by '\n' --hive-home /user/hive/warehouse/retail_ods.db --hive-import
--hive-overwrite --hive-table departments
```

```
18) sqoop import --connect "jdbc:mysql://quickstart.cloudera:3306/retail_db"
--username=retail_dba --password=cloudera --table departments --fields-terminated-by '|'
--lines-terminated-by '\n' --hive-home /user/hive/warehouse --hive-import --hive-table
mysparkdb.departments_test --create-hive-table
```

```
19) --Connect to mysql and create database for reporting database
```

```
--user:root, password:cloudera
```

```
mysql -u root -p
```

```
create database retail_rpt_db;
```

```
grant all on retail_rpt_db.* to retail_dba;
```

```
flush privileges;
```

```
use retail_rpt_db;
```

```
create table departments as select * from retail_db.departments where 1=2;
```

```
sqoop export --connect "jdbc:mysql://quickstart.cloudera:3306/retail_rpt_db" --username
retail_dba --password cloudera --table departments --export-dir
/user/hive/warehouse/retail_ods.db/departments --input-fields-terminated-by '|'
--input-lines-terminated-by '\n' --num-mappers 2 --batch
```

```
20) sqoop export --connect "jdbc:mysql://quickstart.cloudera:3306/retail_db" --username
retail_dba --password cloudera --table departments --export-dir
/user/hive/warehouse/retail_ods.db/departments --batch -m 1 --update-key department_id
--update-mode allowinsert
```

```
21) sqoop import --connect "jdbc:mysql://quickstart.cloudera:3306/retail_db"
--username=retail_dba --password=cloudera --table departments_test --target-dir
/user/hive/warehouse/retail_ods.db/departments_test --append --fields-terminated-by '|'
--lines-terminated-by '\n' --num-mappers 1 --null-string na --null-non-string -1
```

```
22) sqoop export --connect "jdbc:mysql://quickstart.cloudera:3306/retail_rpt_db" --username
root --password cloudera --table departments_test --export-dir
/user/hive/warehouse/retail_ods.db/departments_test --input-fields-terminated-by '|'
--input-lines-terminated-by '\n' --input-null-string na --input-null-non-string -1
```

```
23) sqoop import --connect "jdbc:mysql://quickstart.cloudera:3306/retail_db"
--username=retail_dba --password=cloudera --table departments --target-dir
/user/hive/warehouse/retail_ods.db/departments_new --append --fields-terminated-by '|'
--lines-terminated-by '\n' --num-mappers 1 --enclosed-by \"
```

```
24) sqoop import --connect "jdbc:mysql://quickstart.cloudera:3306/retail_db"
--username=retail_dba --password=cloudera --table departments_new --target-dir
/user/hive/warehouse/retail_ods.db/departments_new1 --append --fields-terminated-by '|'
--lines-terminated-by '\n' --num-mappers 1 --escaped-by \@
```

```
1) sqoop list-databases --connect jdbc:mysql://localhost:3306 --username root
```

```
2) sqoop list-tables --connect jdbc:mysql://localhost:3306/test --username root
```

```
3) sqoop eval --connect jdbc:mysql://localhost:3306/test --username root --query "select
count(*) from cities"
```

```
4) -- Reference:
```

```
http://www.cloudera.com/content/cloudera/en/developers/home/developer-admin-resources/get-started-with-hadoop-tutorial/exercise-1.html
```

```
sqoop import --all-tables -m 12 --connect jdbc:mysql://localhost:3306/test --username=root
--as-avrodatafile --warehouse-dir=/user/hive/warehouse/retail_stage.db
```

```
5) sqoop import --connect jdbc:mysql://localhost:3306/test --username root --table emp
--as-textfile --target-dir=/user/hadoop/emp_txt
```

```
6) sqoop import --connect jdbc:mysql://localhost:3306/test --username root --table emp
```

```
--as-sequencefile --target-dir=/user/hadoop/emp_seq
```

```
7) sqoop import --connect jdbc:mysql://localhost:3306/test --username root --table emp
--as-avrodatafile --target-dir=/user/hadoop/emp_avro
```

Note:- (This below comments are used for hive)

```
-- A file with extension avsc will be created under the directory from which sqoop import is
executed
-- Copy avsc file to HDFS location
-- Create hive table with LOCATION to /user/hadoop/emp_avro and TBLPROPERTIES pointing to avsc
file
hadoop fs -put sqoop_import_departments.avsc /user/cloudera
```

```
CREATE EXTERNAL TABLE departments
ROW FORMAT SERDE 'org.apache.hadoop.hive.serde2.avro.AvroSerDe'
STORED AS INPUTFORMAT 'org.apache.hadoop.hive.ql.io.avro.AvroContainerInputFormat'
OUTPUTFORMAT 'org.apache.hadoop.hive.ql.io.avro.AvroContainerOutputFormat'
LOCATION 'hdfs:///user/cloudera/departments'
TBLPROPERTIES ('avro.schema.url'='
'hdfs://quickstart.cloudera/user/cloudera/sqoop_import_departments.avsc');
```

```
-- It will create tables in default database in hive
-- Using snappy compression
-- As we have imported all tables before make sure you drop the directories
-- Launch hive drop all tables
drop table departments;
drop table categories;
drop table products;
drop table orders;
drop table order_items;
drop table customers;
```

```
-- Dropping directories, in case your hive database/tables in consistent state
hadoop fs -rm -R /user/hive/warehouse/departments
hadoop fs -rm -R /user/hive/warehouse/categories
hadoop fs -rm -R /user/hive/warehouse/products
hadoop fs -rm -R /user/hive/warehouse/orders
hadoop fs -rm -R /user/hive/warehouse/order_itmes
hadoop fs -rm -R /user/hive/warehouse/customers
```

```
8) sqoop import --all-tables --num-mappers 1 --connect jdbc:mysql://localhost:3306/test
--username root
--hive-import --hive-overwrite --create-hive-table --compress --compression-codec
org.apache.hadoop.io.compress.SnappyCodec
--outdir java_files
```

```
sudo -u hdfs hadoop fs -mkdir /user/cloudera/retail_stage
sudo -u hdfs hadoop fs -chmod +rw /user/cloudera/retail_stage
hadoop fs -copyFromLocal ~/.avsc /user/cloudera/retail_stage
```

```
-- Boundary Query and columns
9) sqoop import jdbc:mysql://localhost:3306/test --username root
--table departments --target-dir /user/hadoop/departments -m 2
--boundary-query "select 2, 8 from departments limit 1"
--columns department_id,department_name
```

```
-- query and split-by
sqoop import
--connect jdbc:mysql://localhost:3306/test --username root
--query="select * from orders join order_items on orders.order_id =
order_items.order_item_order_id where \${CONDITIONS}" \
--target-dir /user/hadoop/order_join --split-by order_id \
--num-mappers 4
```

```
-- Copying into existing table or directory (append)
-- Customizing number of threads (num-mappers)
-- Changing delimiter
sqoop import \
  --connect jdbc:mysql://localhost:3306/test --username root
  --table departments --target-dir /user/hive/warehouse/retail_ods.db/departments
  --append --fields-terminated-by '|' --lines-terminated-by '\n' --num-mappers 1
  --outdir java_files

-- Importing table with out primary key using multiple threads (split-by)
-- When using split-by, using indexed column is highly desired
-- If the column is not indexed then performance will be bad
-- because of full table scan by each of the thread
sqoop import \
  --connect jdbc:mysql://localhost:3306/test --username root
  --table departments --target-dir /user/hive/warehouse/retail_ods.db/departments \
  --append --fields-terminated-by '|' --lines-terminated-by '\n' --split-by department_id \
  --outdir java_files
```

Notes:-

--split-by : It is used to specify the column of the table used to generate splits for imports. This means that it specifies which column will be used to create the split while importing the data into your cluster. It can be used to enhance the import performance by achieving greater parallelism. Sqoop creates splits based on values in a particular column of the table which is specified by --split-by by the user through the import command. If it is not available, the primary key of the input table is used to create the splits.

Reason to use : Sometimes the primary key doesn't have an even distribution of values between the min and max values(which is used to create the splits if --split-by is not available). In such a situation you can specify some other column which has proper distribution of data to create splits for efficient imports.

--boundary-query : By default sqoop will use query select min(), max() from to find out boundaries for creating splits. In some cases this query is not the most optimal so you can specify any arbitrary query returning two numeric columns using --boundary-query argument.

Reason to use : If --split-by is not giving you the optimal performance you can use this to improve the performance further.

```
-- Getting delta (--where)
10) sqoop import --connect jdbc:mysql://localhost:3306/test --username root
  --table departments --target-dir /user/hive/warehouse/retail_ods.db/departments
  --append --fields-terminated-by '|' --lines-terminated-by '\n' --split-by department_id \
  --where "department_id > 7" --outdir java_files

-- Incremental load
11) sqoop import --connect jdbc:mysql://localhost:3306/test --username root
  --table departments --target-dir /user/hive/warehouse/retail_ods.db/departments
  --append --fields-terminated-by '|' --lines-terminated-by '\n' --check-column "department_id" \
  --incremental append --last-value 7 --outdir java_files

12) sqoop job --create sqoop_job
  --import --connect jdbc:mysql://localhost:3306/test --username root
  --table departments --target-dir /user/hive/warehouse/retail_ods.db/departments
  --append --fields-terminated-by '|' --lines-terminated-by '\n' --check-column "department_id" \
  --incremental append --last-value 7 --outdir java_files

sqoop job --list

sqoop job --show sqoop_job

sqoop job --exec sqoop_job

-- Hive related
-- Overwrite existing data associated with hive table (hive-overwrite)
13 ) sqoop import --connect jdbc:mysql://localhost:3306/test --username root
  --table departments --fields-terminated-by '|' --lines-terminated-by '\n'
  --hive-home /user/hive/warehouse/retail_ods.db --hive-import --hive-overwrite
  --hive-table departments --outdir java_files
```

```
--Create hive table example
14) sqoop import --connect dbc:mysql://localhost:3306/test --username root
   --table departments --fields-terminated-by '|' --lines-terminated-by '\n'
   --hive-home /user/hive/warehouse --hive-import --hive-table departments_test
   --create-hive-table \
   --outdir java_files

--Connect to mysql and create database for reporting database
--user:root, password:cloudera
mysql -u root -p
create database retail_rpt_db;
grant all on retail_rpt_db.* to retail_dba;
flush privileges;
use retail_rpt_db;
create table departments as select * from retail_db.departments where 1=2;
exit;

--For certification change database name retail_rpt_db to retail_db
15) sqoop export --connect dbc:mysql://localhost:3306/test --username root
   --table departments --export-dir /user/hive/warehouse/retail_ods.db/departments
   --input-fields-terminated-by '|' --input-lines-terminated-by '\n'
   --num-mappers 2 --batch --outdir java_files

16) sqoop export --connect jdbc:mysql://localhost:3306/test --username root
   --table departments --export-dir /user/hadoop/sqoop_import/departments_export
   --batch --outdir java_files -m 1 --update-key department_id --update-mode allowinsert

17) sqoop export --connect jdbc:mysql://localhost:3306/test --username root
   --table departments_test --export-dir /user/hive/warehouse/departments_test
   --input-fields-terminated-by '\001' --input-lines-terminated-by '\n'
   --num-mappers 2 --batch --outdir java_files --input-null-string nvl --input-null-non-string -1

--Merge process begins
hadoop fs -mkdir /user/cloudera/sqoop_merge

--Initial load
sqoop import \
  --connect "jdbc:mysql://quickstart.cloudera:3306/retail_db" \
  --username=retail_dba \
  --password=cloudera \
  --table departments \
  --as-textfile \
  --target-dir=/user/cloudera/sqoop_merge/departments

--Validate
sqoop eval --connect "jdbc:mysql://quickstart.cloudera:3306/retail_db" \
  --username retail_dba \
  --password cloudera \
  --query "select * from departments"

hadoop fs -cat /user/cloudera/sqoop_merge/departments/part*

--update
sqoop eval --connect "jdbc:mysql://quickstart.cloudera:3306/retail_db" \
  --username retail_dba \
  --password cloudera \
  --query "update departments set department_name='Testing Merge' where department_id = 9000"

--Insert
sqoop eval --connect "jdbc:mysql://quickstart.cloudera:3306/retail_db" \
  --username retail_dba \
  --password cloudera \
  --query "insert into departments values (10000, 'Inserting for merge')"

sqoop eval --connect "jdbc:mysql://quickstart.cloudera:3306/retail_db" \
```

```
--username retail_dba \
--password cloudera \
--query "select * from departments"

--New load
sqoop import \
  --connect "jdbc:mysql://quickstart.cloudera:3306/retail_db" \
  --username=retail_dba \
  --password=cloudera \
  --table departments \
  --as-textfile \
  --target-dir=/user/cloudera/sqoop_merge/departments_delta \
  --where "department_id >= 9000"

hadoop fs -cat /user/cloudera/sqoop_merge/departments_delta/part*

--Merge
sqoop merge --merge-key department_id \
  --new-data /user/cloudera/sqoop_merge/departments_delta \
  --onto /user/cloudera/sqoop_merge/departments \
  --target-dir /user/cloudera/sqoop_merge/departments_stage \
  --class-name departments \
  --jar-file <get_it_from_last_import>

hadoop fs -cat /user/cloudera/sqoop_merge/departments_stage/part*

--Delete old directory
hadoop fs -rm -R /user/cloudera/sqoop_merge/departments

--Move/rename stage directory to original directory
hadoop fs -mv /user/cloudera/sqoop_merge/departments_stage
/user/cloudera/sqoop_merge/departments

--Validate that original directory have merged data
hadoop fs -cat /user/cloudera/sqoop_merge/departments/part*

--Merge process ends
```

MONGO DB

1)Please download MongoDB by using below link

<http://www.mongodb.org/downloads>

2)place the MongoDB.taz into download directory

3)Extract the MongoDB.taz using below command

```
tar xvf mongodb-linux-i686-3.0.0.tgz
```

4)create a data directory under downloads

```
sudo mkdir -p /mongo/data
```

5)Then start the MongoDB server by using below command

```
sudo mongodb-linux-i686-3.0.0/bin/mongod --dbpath /mongo/data/
```

6)Then Start MongoDB client by using below command

```
mongodb-linux-i686-3.0.0/bin/mongo
```

=====MongoDB Commands=====

1)create a db in MongoDB

The below command creates a new db in MongoDB

```
>use sample
```

2)if you want check current db name

```
>db
```

3)list of dbs

```
>show dbs
```

4)Inserting a record into collection(Table)

syntax- >db.collection name.save({})

```
db.mycol.save({name:'gvipl',interest:'ericket'})
```

5)if you want looking for a record

```
>db.mycol.find()
```

6)if you want to looking for a one record

```
>db.mycol.findOne();
```

7)disply the list of tables in a db

```
>show collections;
```

8)Dropping a Database

```
db.dropDatabase();
```

9)Creating a collection

```
>db.createCollection("mycol", { capped : true, autoIndexID : true, size : 6142800, max : 10000 } );
```

Field	Type	Description
capped	Boolean	(Optional) If true, enables a capped collection. Capped collection is a collection fixed size colleccction that automatically overwrites its oldest entries when it reaches its maximum size. If you specify true, you need to specify size parameter also.
autoIndexID	Boolean	(Optional) If true, automatically create index on _id field.s Default value is false.
size	number	(Optional) Specifies a maximum size in bytes for a capped collection. If If capped is true, then you need to specify this field also.
max	number	(Optional) Specifies the maximum number of documents allowed in the capped collection.

10)Dropping a Collection

```
db.COLLECTION_NAME.drop()
```

```
>db.mycol.drop();
```

11)Inserting a Document

```
db.COLLECTION_NAME.insert(document)
```

```
db.post.insert([ {title:'MongoDB is no sql database', by:'Mongo DB',
url:'http://www.mongodb.org' ,tags:['mongodb','database','NoSQL'], likes:100 }]);
```

```
db.post.insert([ {title:'MongoDB is no sql database', by:'Mongo DB',
```

```

url:'http://www.mongodb.org' ,tags:['mongodb','database','NoSQL'], likes:20 }]);

db.post.insert([ {title:'MongoDB is no sql database', by:'Mongo DB',
url:'http://www.mongodb.org' ,tags:['mongodb','database','NoSQL'],
likes:50,comments:[{user:'user1',message: 'My first comment',dateCreated: new
Date(2013,11,10,2,35),like: 0}] }]);

12)db.post.find().pretty();

13)db.post.find({"by":"Mongo DB"}).pretty();

14)db.post.find({"likes":{$lt:50}}).pretty();

15)db.post.find({"likes":{$lte:50}}).pretty();

16)db.post.find({"likes":{$gt:50}}).pretty();

17)db.post.find({"likes":{$gte:50}}).pretty();

18)db.post.find({"likes":{$ne:50}}).pretty();

19)AND Condition
db.post.find({"by":"Mongo DB","title": "MongoDB Overview"}).pretty();

20)or condition
db.post.find({$or:[{"by":"Mongo DB"}, {"title": "MongoDB Overview"}]}).pretty();

21)And ,or both condition at a time
db.post.find("likes": {$gt:10}, $or: [{"by": "Mongo DB"}, {"title": "MongoDB Overview"}]
}).pretty()

22)Updating a collection

db.COLLECTION_NAME.update(SELECTIOIN_CRITERIA, UPDATED_DATA)

>db.post.update({'likes':100},{ $set: {'title': 'New MongoDB Tutorial'}})

23)Multiple documents update

>db.post.update({'title': 'MongoDB is no sql database'}, { $set: {'title': 'New MongoDB
Tutorial'} }, {multi:true})

24)save method

db.COLLECTION_NAME.save({_id:ObjectId(),NEW_DATA})

>db.post.save({"_id":ObjectId("550559e86a39df979dae4d80"),"title":"Mongo DB New Topic",
"by":"GVIPL Point"})

25)remove method

db.COLLECTION_NAME.remove(DELETION_CRITTERIA)

db.post.remove({'title': 'MongoDB Overview'})

db.COLLECTION_NAME.remove(DELETION_CRITERIA,1)

db.post.remove({'title': 'MongoDB Overview'},1);

if you want remove all documents

db.post.remove();

26)Projection of a fields

>db.post.find({}, {"title":1, "_id":0});

```

27) Limit ()

```
db.COLLECTION_NAME.find().limit(NUMBER)
```

```
>db.post.find({},{"title":1,"_id":0}).limit(2);
```

28) Skip ()

```
db.COLLECTION_NAME.find().limit(NUMBER).skip(NUMBER)
```

```
>db.post.find({},{"title":1,"_id":0}).limit(1).skip(1);
```

29) sort ()

```
db.COLLECTION_NAME.find().sort({KEY:1})
```

```
>db.post.find({},{"title":1,"_id":0}).sort({"title":-1});
```

DataTypes in Mongo DB

String : This is most commonly used datatype to store the data. String in mongodb must be UTF-8 valid.

Integer : This type is used to store a numerical value. Integer can be 32 bit or 64 bit depending upon your server.

Boolean : This type is used to store a boolean (true/ false) value.

Double : This type is used to store floating point values.

Min/ Max keys : This type is used to compare a value against the lowest and highest BSON elements.

Arrays : This type is used to store arrays or list or multiple values into one key.

Timestamp : ctimestamp. This can be handy for recording when a document has been modified or added.

Object : This datatype is used for embedded documents.

Null : This type is used to store a Null value.

Symbol : This datatype is used identically to a string however, it's generally reserved for languages that use a specific symbol type.

Date : This datatype is used to store the current date or time in UNIX time format. You can specify your own date time by creating object of Date and passing day, year into it.

Object ID : This datatype is used to store the document's ID.

Binary data : This datatype is used to store binay data.

Code : This datatype is used to store javascript code into document.

Regular expression : This datatype is used to store regular expression

1) Indexing in Mongo DB

```
ensureIndex()
```

```
db.COLLECTION_NAME.ensureIndex({KEY:1})
```

```
>db.post.ensureIndex({"title":1})
```



```
>db.mycol.ensureIndex({"title":1,"description":-1})

2) Aggregation

aggregate()

db.COLLECTION_NAME.aggregate(AGGREGATE_OPERATION)

group by($group)

>db.post.aggregate([{$group : {_id : "$by", num_tutorial : {$sum : 1}}}])

sum($sum)

>db.post.aggregate([{$group : {_id : "$by", num_tutorial : {$sum : "$likes"}}}])

avg($avg)

>db.post.aggregate([{$group : {_id : "$by", num_tutorial : {$avg : "$likes"}}}])

min($min)

>db.post.aggregate([{$group : {_id : "$by", num_tutorial : {$min : "$likes"}}}])

max($max)

>db.post.aggregate([{$group : {_id : "$by", num_tutorial : {$max : "$likes"}}}])

push($push)--Inserts the value to an array in the resulting document.

>db.post.aggregate([{$group : {_id : "$by", url : {$push: "$url"}}}])

addToSet($addToSet)-- Inserts the value to an array in the resulting document but does not
create duplicates.

>db.post.aggregate([{$group : {_id : "$by", url : {$addToSet: "$url"}}}])

first($first)

>db.post.aggregate([{$group : {_id : "$by", first_url : {$first: "$url"}}}])

last($last)

>db.post.aggregate([{$group : {_id : "$by", last_url : {$last: "$url"}}}])
```

Analyze the Query Performance

```
-----
>db.collection.explain();
>db.collection.explain().help()
```

Importing the data from a file

```
-----
>cd Downloads
>mongodb-linux-i686-3.0.0/bin/mongoimport --db sample --collection contacts --file test.json
>mongodb-linux-i686-3.0.0/bin/mongoimport --db sample --collection contacts --type csv
--headerline --contacts.csv
```

Aggregate Function Examples

```
-----
```

```
1) >db.zipcode.aggregate( [
    { $group: { _id: "$state", totalPop: { $sum: "$pop" } } },
    { $match: { totalPop: { $gte: 10*1000*1000 } } }
] )
```

```

2)>db.zipcode.aggregate( [
  { $group: { _id: { state: "$state", city: "$city" }, pop: { $sum: "$pop" } } },
  { $group: { _id: "$_id.state", avgCityPop: { $avg: "$pop" } } }
] )

3)>db.zipcode.aggregate( [
  { $group:
    {
      _id: { state: "$state", city: "$city" },
      pop: { $sum: "$pop" }
    }
  },
  { $sort: { pop: 1 } },
  { $group:
    {
      _id : "$_id.state",
      biggestCity: { $last: "$_id.city" },
      biggestPop:   { $last: "$pop" },
      smallestCity: { $first: "$_id.city" },
      smallestPop:  { $first: "$pop" }
    }
  },
  { $project:
    { _id: 0,
      state: "$_id",
      biggestCity: { name: "$biggestCity", pop: "$biggestPop" },
      smallestCity: { name: "$smallestCity", pop: "$smallestPop" }
    }
  }
] )

```

Mongo DB Relationships

- 1) Embaded Relationship
- 2) Referenced Relationship

```

db.user.find().pretty();
{
  "_id": ObjectId("52ffc33cd85242f436000001"),
  "name": "gvipl",
  "contact": "987654321",
  "dob": "01-01-1991"
}

```

```

db.address.find().pretty()
{
  "_id": ObjectId("52ffc4a5d85242602e000000"),
  "building": "GV IPL",
  "pincode": 123456,
  "city": "Hyderabad",
  "state": "AP"
}

```

Embaded Relationship

```

([
{
  "_id": ObjectId("52ffc33cd85242f436000001"),
  "contact": "987654321",
  "dob": "01-01-1991",
  "name": "gvipl",
  "address": [
    {
      "building": "GV IPL",

```

```

        "pincode": 123456,
        "city": "Hyderabad",
        "state": "AP"
    },
    {
        "building": "GVIPL2",
        "pincode": 456789,
        "city": "Hyderabad",
        "state": "TS"
    }
]
)

>db.users.findOne({"name":"gvipl"}, {"address":1})

```

Referenced Relationships

```

-----
{
  "_id":ObjectId("52ffc33cd85242f436000001"),
  "contact": "987654321",
  "dob": "01-01-1991",
  "name": "gvipl",
  "address_ids": [
    ObjectId("52ffc4a5d85242602e000000"),
    ObjectId("52ffc4a5d85242602e000001")
  ]
}

>var result = db.users.findOne({"name":"gvipl"}, {"address_ids":1})
>var addresses = db.address.find({"_id":{"$in":result["address_ids"]}})

```

1)Backup and Restore

```

-----
Backup
-----
>mongodump
>mongodump --host localhost --port 27017
>mongodump --out /data/backup/
>mongodump --collection myCollection --db test
>mongodump --host mongodbl.example.net --port 3017 --username user --password pass --out
/opt/backup/mongodump-2013-10-24

```

Restore

```

-----
syntax:-mongorestore --port <port number> <path to the backup>

>mongorestore dump-2013-10-25/
>mongorestore --host mongodbl.example.net --port 3017 --username user --password pass
/opt/backup/mongodump-2013-10-24

```

```
=====
```

```
=====
                                HBASE
=====
```

```

[hadoop@localhost ~]$ sudo /sbin/service zookeeper-server start;
[sudo] password for hadoop
[hadoop@localhost ~]$ sudo /sbin/service hbase-master start;
[hadoop@localhost ~]$ sudo /sbin/service hbase-master status;

[hadoop@localhost ~]$ sudo service hbase-regionserver start;

[hadoop@localhost ~]$ hbase shell

```

```

-----
<property>
  <name>hbase.rootdir</name>
  <value>hdfs://localhost:8020/hbase</value>
</property>
<property>
<name>hbase.zookeeper.quorum</name>
<value>localhost:2181</value>
</property>
<property>
<name>hbase.cluster.distributed</name>
<value>true</value>
</property>
</configuration>

```

// Below command gives the all list of all hbase shell commands

1.hbase> help

//General hbase shell commands

1.status:-how cluster status. Can be 'summary', 'simple', or 'detailed'. The default is 'summary'.

```

hbase> status
hbase> status 'simple'
hbase> status 'summary'
hbase> status 'detailed'

```

2.version:-Output this HBase versionUsage

```

hbase>version

```

3.whoami:shows the current hbase user.

```

hbase>whoami

```

//Tables Management commands (DDL)

1.create:-Create table; pass table name, a dictionary of specifications per column family, and optionally a dictionary of table configuration.

```

hbase> create 't1', {NAME => 'f1', VERSIONS => 5}
hbase> create 't1', {NAME => 'f1'}, {NAME => 'f2'}, {NAME => 'f3'}
hbase> # The above in shorthand would be the following:
hbase> create 't1', 'f1', 'f2', 'f3'
hbase> create 't1', {NAME => 'f1', VERSIONS => 1, TTL => 2592000, BLOCKCACHE => true}
hbase> create 't1', {NAME => 'f1', CONFIGURATION => {'hbase.hstore.blockingStoreFiles' =>
1101}}

```

2.describe:-this command describe the hbase table

```

hbase>describe '<table name>'

```

3.disable:-disable the specified table

```

hbase>disable '<table name>'

```

if you set a disable condition on a table that table cannot scanned.

```

hbase>scan '<disabled table>'

```

4.disable_all:-Disable all of tables matching the given regex

```

hbase>disable_all 't.*'

```

5.is_disabled:-verifies Is named table disabled and it returns true/false

```
hbase>is_disabled '<table name>'
```

6.drop:-Drop the named table. Table must first be disabled

```
hbase>drop 't1'
```

7.drop_all:-Drop all of the tables matching the given regex

```
hbase>drop_all 't.*'
```

8.is_enabled:-verifies Is named table enabled and returns true/false

```
hbase>is_enabled '<table name>'
```

9.exists:-Does the named table exist and returns true/false

```
hbase>exists '<table name>'
```

10.list:-List all tables in hbase. Optional regular expression parameter could be used to filter the output

```
hbase>list
```

```
hbase>list 'abc.*'
```

11.show_filters:-Show all the filters in hbase

```
hbase>show_filters
```

12.alter:-alter column family schema; pass table name and a dictionary specifying new column family schema.

Dictionaries are described on the main help command output.

Dictionary must include name of column family to alter.

For example, to change or add the 'f1' column family in table 't1' from current value to keep a maximum of 5 cell VERSIONS, do:

First we need to disable the table

```
hbase>alter 't1',NAME=>'f1',VERSIONS=>5
```

13.alter_status:-get the status of the alter command. Indicates the number of regions of the table that have received the updated schema Pass table name

```
hbase>alter_status<table_name>
```

//Data Manipulation commands (DML)

1.put:-Put a cell value at specified table/row/column and optionally timestamp coordinates.

To put a cell value into table 't1' at

row 'r1' under column 'c1' marked with the time 'ts1', do

```
hbase>create 'emp',{NAME=>'address'},{NAME=>'expinfo'}
```

```
put <Table Name>,<row_id>,<family:col name>,<value>
```

```
hbase>put 'emp','1','address:city','hyd'
```

```
hbase>put 'emp','1','address:state','ap'
```

```
hbase>put 'emp','1','address:country','india'
```

```
hbase>put 'emp','1','expinfo:doj','02022012'
```

```
hbase>put 'emp','1','expinfo:dol',''
```

2.scan:-Scan a table; pass table name and optionally a dictionary of scanner specifications. Scanner specifications may include one or more of:

TIMERANGE, FILTER, LIMIT, STARTROW, STOPROW, TIMESTAMP, MAXLENGTH,

or COLUMNS, CACHEIf no columns are specified, all columns will be scanned.

To scan all members of a column family, leave the qualifier empty as in 'col_family:'.The filter can be specified in two ways:

1. Using a filterString ñ more information on this is available in the Filter Language document attached to the HBASE-4176 JIRA
2. Using the entire package name of the filter. Some examples: hbase> scan ñ.META.1

```
hbase>scan 'emp'
```

- 3.count:-count the no.of rows in a table.

```
hbase>count 'em'
```

- 4.delete:-Put a delete cell value at specified table/row/column and optionally timestamp coordinates. Deletes must match the deleted cellís coordinates exactly. When scanning, a delete cell suppresses older versions. To delete a cell from ñt1' at row ñr1' under column ñc1' marked with the time ñts1', do:

```
hbase>delete '<Table Name>','<row_id>','<family:colname>'
hbase>delete 'emp','1','expinfo:dol'
```

- 5.deleteall:-Delete all cells in a given row; pass a table name, row, and optionally a column and timestamp.

```
hbase> deleteall ñ<Table Name>','<row_id>'
```

- 6.get:-Get row or cell contents; pass table name, row, and optionally a dictionary of column(s), timestamp, timerange and versions.

```
hbase>get '<Table Name>','<row_id>'
hbase>get 'emp','1'
```

- 7.incr:-Increments a cell ñvalueí at specified table/row/column coordinates. To increment a cell value in table ñt1' at row ñr1' under column ñc1' by 1 (can be omitted) or 10 do

```
hbase>incr '<Table_Name>','<row_id>','<family:colname>','<incrvalue>'
hbase>incr 'emp','1','expinfo:doj','1'
```

- 8.truncate:-Disables, drops and recreates the specified table

```
hbase>truncate <Table_Name>
hbase>truncate 'Employee'
```

-----SPARK INTELLIJ-----

```
*****WordCount
```

```
package SparkCore
```

```
import org.apache.spark.sql.SparkSession
```

```
object WordCount {
  def main(args: Array[String]): Unit = {
    val spark = SparkSession.builder.master("local[*]").appName("word count").getOrCreate()
    val sc = spark.sparkContext
    val sqlContext = spark.sqlContext
    import spark.implicits._
    import spark.sql
    println("Success")
  }
}
```

```
val rawData =
sc.textFile("C:\\Users\\NaraVish\\IdeaProjects\\SparkPractice\\Data\\wordcount.txt")

println("Data SUCCESSFULL")
//convert the lines into words using flatMap operation
```

```

val words = rawData.flatMap(line => line.split(" "))

//count the individual words using map and reduceByKey operation
val wordCount = words.map(word => (word, 1)).reduceByKey(_ + _)

println(wordCount.count())

wordCount.foreach(println)

// wordCount.map(x=> println(x))

wordCount.saveAsTextFile("C:\\Users\\NaraVish\\IdeaProjects\\SparkPractice\\Data\\output")

spark.stop()

}
}

-----

***MapvsMapPartitions

package SparkCore

/*
Map Transformation iterates each and every element and creates an RDD at each element level
In Case of Map Partition RDD is created per Partition

Map Partition contain a iterator object , it iterates each and every record from a given data
set.
And Finally creates an RDD per partitions

Map Partition with Index is similar to map partition and it takes one extra parameter called it
as index number.
If we are using map partition with index we are loosing spark functionality of spliiting the
data.

*/

import org.apache.spark.sql.SparkSession

object MapvsMapPartitions {

  def main(args: Array[String]): Unit = {

    val spark = SparkSession.builder().master("local[*]").getOrCreate()
    val sc = spark.sparkContext

    // Map
    val m = sc.parallelize(Array(1,2,3))

    println(m.getNumPartitions)

    // Map Partitions
    val x = sc.parallelize(Array(1,2,3),1)

    println(x.getNumPartitions)

```

```
// Mapr Partitions with Indix
val y = sc.parallelize(Array(1,2,3),2)

def f(partitionIndex:Int,i:Iterator[Int])={
  (partitionIndex,i.sum,42).productIterator
}

val z = y.mapPartitionsWithIndex(f)
z.collect().foreach(println)

}
}
```

****JOIN

package SparkCore

import org.apache.spark.sql.SparkSession

object Join {

def main(args: Array[String]): Unit = {

val spark= SparkSession.builder().master("local[*]").appName("Join").getOrCreate()
val sc=spark.sparkContext

val edata = sc.textFile("C:\\Users\\NaraVish\\IdeaProjects\\SparkPractice\\Data\\emp.txt")
val ddata = sc.textFile("C:\\Users\\NaraVish\\IdeaProjects\\SparkPractice\\Data\\dept.txt")

val edata_pair = edata.map{ x =>
 val w = x.split(",")
 val eno = w(0).toInt
 val ename = w(1)
 val sal = w(2).toInt
 val gendar = w(3)
 val dno = w(4).toInt
 (dno,(eno,ename,sal,gendar))
}

val ddata_pair = ddata.map { x=>
 val w = x.split(",")
 val dno = w(0).toInt
 val dname = w(1)
 val dloc = w(2)
 (dno,(dname,dloc))
}

val edata_pair_join_ddata_pair = edata_pair.leftOuterJoin(ddata_pair)

edata_pair_join_ddata_pair.foreach(println)

// edata_pair_join_ddata_pair.saveAsTextFile(
"C:\\Users\\NaraVish\\IdeaProjects\\SparkPractice\\Data\\output")
edata_pair_join_ddata_pair.toDebugString
spark.stop()

}

}

```
---RepartitionVsCoalesce
```

```
package SparkCore
```

```
import org.apache.spark.sql.SparkSession
```

```
object RepartitionVsCoalesce {
```

```
  /*
```

1. Both coalesce and repartition enables the re assinging of the partitions at run time.
2. coalesce by defaultly shuffling is false
3. Re-partition by defaultly shuffling is true

```
we can swith off shuffling in repartition that will behave as coalesce
```

```
we can not switch on shuffling in coalesce
```

```
repartition is not avaialble in apache storm
```

```
coalesce is re-commendable
```

```
  */
```

```
  def main(args: Array[String]): Unit = {
```

```
    val spark= SparkSession.builder().master("local[*]").getOrCreate()
```

```
    val sc= spark.sparkContext
```

```
  //    Example coalesce :
```

```
    val x = sc.parallelize(Array(1,2,3,4,5),3)
```

```
    val y = x.coalesce(2,false)
```

```
    println(y.getNumPartitions)
```

```
  //Example repartition :
```

```
    val repartList = sc.parallelize(Array(1,2,3,4,5),3)
```

```
    val repartListOutput = repartList.repartition(2)
```

```
    println(repartListOutput.getNumPartitions)
```

```
  }
```

```
}
```

```
-----
```

```
*****DifferentWaysToCreateRDD
```

```
package SparkCore
```

```
import org.apache.spark.sql.SparkSession
```

```
object DifferentWaysToCreateRDD {
```

```
  def main(args: Array[String]): Unit = {
```

```
    val spark = SparkSession.builder.master("local[*]").appName("word count").getOrCreate()
```

```
    val sc = spark.sparkContext
```

```

    /// Through Serialized Method
    val l =List(10,20,30)

    val rdd = sc.parallelize(l)

    rdd.foreach(println)

    /// From Existing RDD

    val newRdd = rdd.filter(x=> x>10)

    newRdd.foreach(println)

    // External Sources sc.TextFile

    val emp = sc.textFile("C:\\Users\\NaraVish\\IdeaProjects\\SparkPractice\\Data\\emp.txt")

    sc.stop()

}

}

-----

*****CSVProcessing

package SparkPackage

import org.apache.spark.sql.SparkSession

object CSVProcessing {

    def main(args: Array[String]): Unit = {

        val spark = SparkSession.builder.master("local[*]").appName("SparkCsv").getOrCreate()
        //local represents local machine & star represents no of resources i.e utilising all resources
        //getOrCreate => creating a application or using already existed application
        //appName => creating an application with name "csvExample_1"

        //creating 2 contex sparkContext and sqlContext
        val sc = spark.sparkContext
        val sqlContext = spark.sqlContext

        import spark.implicits._
        //Creating DataFrame using SQLContext Reading data)
        val df = sqlContext.read.format("com.databricks.spark.csv").option("header","true").option(
            "inferSchema","true").load(
            "C:\\Users\\NaraVish\\IdeaProjects\\SparkPractice\\Data\\Police_Department_Incidents.csv")
        //option("header","true") => Use first line of all files as Header
        //option("inferSchema","true") => Automatically infer data types

        df.show()
        //df.select("Category").distinct.collect().foreach(println)
        df.createOrReplaceTempView("sfpd") //Creating Temp table
        // sqlContext.sql("select Category from sfpd").collect().foreach(println)
        //top 10 results
        //sqlContext.sql("SELECT Resolution , count(Resolution) as rescount FROM sfpd group by
        Resolution order by rescount desc limit 10").collect().foreach(println)
        val t = sqlContext.sql("select Category,count(Category) as catcount from sfpd group by
        Category order by catcount desc limit 10")
        t.show()
        t.map(t=> "column 0: "+ t(0)).collect().foreach(println)
        spark.stop()

    }
}

```

}

*****ExcelExample

package SparkPackage

import org.apache.spark.sql.Session

object ExcelExample {

def main(args: Array[String]): Unit = {

val spark= SparkSession.builder().master("local[*]").appName("ExcelExample").getOrCreate()

val sc=spark.sparkContext

val sqlContext=spark.sqlContext

val df = sqlContext.read

.format("com.crealytics.spark.excel")

.option("useHeader", "true")

.option("treatEmptyValuesAsNulls", "false")

.option("inferSchema", "false")

.option("startColumn", 0) // Optional, default: 0

.option("endColumn", 99) // Optional, default: Int.MaxValue

.option("timestampFormat", "dd-MON-YY HH:mm:ss") // Optional, default: yyyy-mm-dd hh:mm:ss
[.ffffffffff]

.option("location", "C:\\Users\\NaraVish\\IdeaProjects\\SparkPractice\\Data\\EMP.xlsx")

.option("addColorColumns", "False")

.option("sheetName", "EMP")

.load("C:\\Users\\NaraVish\\IdeaProjects\\SparkPractice\\Data\\EMP.xlsx")

//

/*

df.write

.format("com.crealytics.spark.excel")

.option("sheetName", "Daily")

.option("useHeader", "true")

.option("dateFormat", "yy-mm-dd") // Optional, default: yy-m-d h:mm

.option("timestampFormat", "mm-dd-yyyy hh:mm:ss") // Optional, default: yyyy-mm-dd hh:mm:ss.000

.mode("overwrite")

.save("Worktime2.xlsx")

* */

df.show(10)

}

}

*****JsonProcessing

package SparkPackage

import org.apache.spark.sql.Session

object JsonProcessing {

def main(args: Array[String]): Unit = {

val spark = SparkSession.builder.master("local[*]").appName("JsonExample").getOrCreate()

```

val sc = spark.sparkContext
val sqlContext = spark.sqlContext

//Converting RDD to Data Frame
import spark.implicits._
import spark.sql
val df = sqlContext.read.json(
"C:\\Users\\NaraVish\\IdeaProjects\\SparkPractice\\Data\\world_bank.json")
df.printSchema() //printing schema
df.createOrReplaceTempView("jsontdata_one")
df.show()

val result = sqlContext.sql("select url,totalamt,abc.* from jsontdata_one " + "lateral
view explode(theme_namecode) as abc")
val result2 = sqlContext.sql("select _id from jsontdata_one")
result.show(10)
result2.show(10)
// println(result)

//result.write.format("com.databricks.spark.csv").option("header","true").save(
"C:\\Users\\sonirai\\Desktop\\Hadoop GV\\Spark\\SparkSQL\\datasets\\jsontocsv")
spark.stop()
}

```

*****RDS_FromOracle

package SparkPackage

import org.apache.spark.sql.SparkSession

object RDS_FromOracle {

def main(args: **Array**[String]): Unit = {

val spark = SparkSession.builder().master("local[*]").appName("Oracle Example").getOrCreate()

val sc = spark.sparkContext

val sqlContext = spark.sqlContext

val empDF = spark.read.format("jdbc").option("url",
"jdbc:oracle:thin:scoot/tiger@//HYRDSLVM0028.es.ad.adp.com:1521/cri02hyd").option("dbtable",
"EMP")
.option("user", "scott").option("password", "tiger").option("driver",
"oracle.jdbc.driver.OracleDriver").load()

empDF.show()

empDF.registerTempTable("emp")

val empDF2 = sqlContext.sql("select job,count(0) as jobcount from emp group by job order by
job")

empDF2.show(10)

}

}

*****XMLProcessing

package SparkPackage

import org.apache.spark.sql.SparkSession

object XMLProcessing {

def main(args: Array[String]): Unit = {

```
val spark = SparkSession.builder.master("local[*]").appName("xmlExample_1").getOrCreate()
val sc = spark.sparkContext
val sqlContext = spark.sqlContext
val df = sqlContext.read.format("com.databricks.spark.xml").option("rootTag", "books").option(
  "rowTag", "book").load("C:\\Users\\NaraVish\\IdeaProjects\\SparkPractice\\Data\\books.xml")
df.show()
// df.printSchema()
//df.createOrReplaceTempView("jsontdata_one")
//sqlContext.sql("select url, totalamt, abc.code, abc.name from jsontdata_one " +
//"lateral view explode(theme_namecode) as abc").show(5)
// result.write.format("com.databricks.spark.csv").option("header", "true").save(
  "C:\\Kalyan\\POC\\Spark\\spark_datasets\\json_data\\jsontocsv")
spark.stop()
```

}

}

SparkSQL

package SparkSQL

case class Sample(a: Int, b: Int, c: Int)

import org.apache.spark.sql.SparkSession

object SparkSQLExample1 {

def main(args: Array[String]): Unit = {

```
val spark = SparkSession.builder().master("local[*]").appName("SparkSQLExample").getOrCreate()
val sc = spark.sparkContext
val sqlContext = spark.sqlContext
```

```
import sqlContext.implicits._
```

```
val s1 = Sample(10, 20, 30)
val s2 = Sample(1, 2, 3)
val s3 = Sample(100, 200, 300)
val s4 = Sample(1000, 2000, 3000)
```

```
val data = sc.parallelize(List(s1, s2, s3, s4))
```

```
data.collect
```

```
val x = data.map(v => v.a + v.b + v.c).collect
x.foreach(println)
```

```

val df = data.toDF //if your rdd having schema its eligible for DataFrame.
df.show()

df.registerTempTable("sample")

val result = sqlContext.sql("select * from sample")
val result1 = sqlContext.sql("select a,b from sample")
result.show()
result.printSchema()
result1.show()
val result2 = sqlContext.sql("select a,b,c,a+b+c as total from sample")
result2.show()

val result4 = sqlContext.sql("select a,b,c,a+b+c as total from sample")
result4.show(10)

val groupbySql = sqlContext.sql("select a,sum(b) from sample group by a")
groupbySql.show(10)

}

}

```

FirstDataFrame

package SparkSQL

import org.apache.spark.sql.SparkSession

case class Employee(eno: Int, ename: String, sal: Int, gendar: String, dno: Int)

case class Department(dno: Int, dname: String, dloc: String)

object FirstDataFrame {

def main(args: Array[String]): Unit = {

val spark = SparkSession.builder().master("local[*]").appName("FirstDataFrame").getOrCreate()

val sc = spark.sparkContext

val sqlContext=spark.sqlContext

```

/*
select dno,loc,avg(sal),max(sal),min(sal) from emp e join dept d
where e.dno=d.dno
group by dno,dloc
*/

```

val emp = sc.textFile("C:\\Users\\NaraVish\\IdeaProjects\\SparkPractice\\Data\\emp.txt")

val dept = sc.textFile("C:\\Users\\NaraVish\\IdeaProjects\\SparkPractice\\Data\\dept.txt")

val edata = emp.map { x =>

val w = x.split(",")

val eno = w(0).toInt

val ename = w(1)

val sal = w(2).toInt

val gendar = w(3)

val dno = w(4).toInt

```

    Employee(eno, ename, sal, gendar, dno)
  }

val ddata = dept.map { x =>
  val w = x.split(",")
  val dno = w(0).toInt
  val dname = w(1)
  val dloc = w(2)
  Department(dno, dname, dloc)
}

import spark.implicits._
import spark.sqlContext

val empDataFrame = edata.toDF()
val deptDataFrame = ddata.toDF()

empDataFrame.createGlobalTempView("empview")
deptDataFrame.createGlobalTempView("deptview")

val empdf= sqlContext.sql("select * from empview").toDF()
val deptdf= sqlContext.sql("select * from deptview").toDF()
val avg_sal_max_sal= sqlContext.sql("select d.dno,d.dloc,avg(sal) as AVG_SAL ,max(sal)
MAX_SAL ,min(sal) MIN_SAL,count(*) COUNT_SAL from empview e join deptview d on e.dno=d.dno
group by d.dno,d.dloc").toDF()

val join= sqlContext.sql("select d.dno,d.dloc from empview e join deptview d on e.dno=d.dno"
).toDF()

empdf.show()
deptdf.show()
join.show()
avg_sal_max_sal.show()

spark.stop()

}

}

```

FirstDataSet

package SparkSQL

import org.apache.spark.sql.SparkSession

//case class Employee(eno: **Int**, ename: String, sal: **Int**, gendar: String, dno: **Int**)
//case class Department(dno: **Int**, dname: String, dloc: String)

object FirstDataSet {

def main(args: **Array**[String]): Unit = {

val spark = SparkSession.builder().master("local[*]").appName("FirstDataFrame").getOrCreate()
val sc = spark.sparkContext
val sqlContext=spark.sqlContext

/*
select dno,loc,avg(sal),max(sal),min(sal) from emp e join dept d

```

where e.dno=d.dno
group by dno,dloc

```

```

*/

```

```

val emp = sc.textFile("C:\\Users\\NaraVish\\IdeaProjects\\SparkPractice\\Data\\emp.txt")
val dept = sc.textFile("C:\\Users\\NaraVish\\IdeaProjects\\SparkPractice\\Data\\dept.txt")

```

```

val edata = emp.map { x =>
  val w = x.split(",")
  val eno = w(0).toInt
  val ename = w(1)
  val sal = w(2).toInt
  val gendar = w(3)
  val dno = w(4).toInt
  Employee(eno, ename, sal, gendar, dno)
}

```

```

val ddata = dept.map { x =>
  val w = x.split(",")
  val dno = w(0).toInt
  val dname = w(1)
  val dloc = w(2)
  Department(dno, dname, dloc)
}

```

```

import spark.implicit._
import spark.sqlContext

```

```

val empDataSet = edata.toDS
val deptDataSet = ddata.toDS

```

```

empDataSet.show(10)
deptDataSet.show(10)

```

```

empDataSet.select($"eno",$"ename",$"sal",$"gendar",$"dno").show()

```

```

empDataSet.filter("eno=101").show(10)

```

```

empDataSet.select($"(sal)").show()

```

```

// Pending hwo to get Max
//empDataSet.select($max"(sal)").show()

```

```

spark.stop()

```

```

}

```

```

}

```

```

package SparkSQL

```

```

case class Person(name: String, age: Int)
case class Company(name: String, foundingYear: Int, numEmployees: Int)

```



```

case class Employee(name: String, age: Int, departmentId: Int, salary: Double)
case class Department(id: Int, name: String)

case class Record(name: String, age: Int, salary: Double, departmentId: Int, departmentName: String)
case class ResultSet(departmentId: Int, departmentName: String, avgSalary: Double)

import org.apache.spark.sql.Session
import org.apache.spark.sql.functions._

object AnotherDataSets {

  def main(args: Array[String]): Unit = {

    val spark = SparkSession.builder().master("local[*]").appName("AnotherDataSets").getOrCreate()

    val sc = spark.sparkContext
    val sqlContext = spark.sqlContext

    import spark.implicits._

    val dataset = sc.parallelize(Seq(1, 2, 3))
    dataset.toDS().show()
    /* -- Not working in 2.0
    val dataset = sc.parallelize(Seq(1, 2, 3))

    val personDS = Seq(Person("Max", 33), Person("Adam", 32), Person("Muller", 62)).toDS()
    personDS.show()

    val rdd = sc.parallelize(Seq((1, "Spark"), (2, "Databricks")))
    val integerDS = rdd.toDS()
    integerDS.show()

    val inputSeq = Seq(Company("ABC", 1998, 310), Company("XYZ", 1983, 904), Company("NOP", 2005, 83))
    val df = sc.parallelize(inputSeq).toDF()

    val companyDS = df.as[Company]
    companyDS.show()

    val rdd = sc.parallelize(Seq((1, "Spark"), (2, "Databricks"), (3, "Notebook")))
    val df = rdd.toDF("Id", "Name")

    val dataset = df.as[(Int, String)]
    dataset.show()

    val wordsDataset = sc.parallelize(Seq("Spark I am your father", "May the spark be with you", "Spark I am your father")).toDS()
    val groupedDataset = wordsDataset.flatMap(_.toLowerCase.split(" "))
      .filter(_ != "")
      .groupBy("value")
    val countsDataset = groupedDataset.count()
    countsDataset.show()
  }
}

```

```

val employeeDataSet1 = sc.parallelize(Seq(Employee("Max", 22, 1, 100000.0),
Employee("Adam", 33, 2, 93000.0), Employee("Eve", 35, 2, 89999.0), Employee("Muller", 39,
3, 120000.0))).toDS()
val employeeDataSet2 = sc.parallelize(Seq(Employee("John", 26, 1, 990000.0),
Employee("Joe", 38, 3, 115000.0))).toDS()
val departmentDataSet = sc.parallelize(Seq(Department(1, "Engineering"), Department(2,
"Marketing"), Department(3, "Sales"))).toDS()

val employeeDataset = employeeDataSet1.union(employeeDataSet2)

def averageSalary(key: (Int, String), iterator: Iterator[Record]): ResultSet = {
  val (total, count) = iterator.foldLeft(0.0, 0.0) {
    case ((total, count), x) => (total + x.salary, count + 1)
  }
  ResultSet(key._1, key._2, total / count)
}

val averageSalaryDataset = employeeDataset.joinWith(departmentDataSet, $"departmentId" ===
$"id", "inner")
  .map(record => Record(record._1.name, record._1.age, record._1.salary,
record._1.departmentId, record._2.name))
  .filter(record => record.age > 25)
  .groupBy($"departmentId", $"departmentName")
  .avg()

val wordsDataset = sc.parallelize(Seq("Spark I am your father", "May the spark be with
you", "Spark I am your father")).toDS()
val result = wordsDataset
  .flatMap(_.split(" ")) // Split on whitespace
  .filter(_ != "") // Filter empty words
  .map(_.toLowerCase())
  .toDF() // Convert to DataFrame to perform aggregation / sorting
  .groupBy($"value") // Count number of occurrences of each word
  .agg(count("*") as "numOccurrences")
  .orderBy($"numOccurrences" desc) // Show most common words first
result.show()
*/
}
}

```
