

## 'General Questions'

## 1. Tell About Yourself

I would like to thanks for giving me an opportunity.

My Name is Vishal Nara. I am from mumbai. I m currently working in ADP Pvt limited from last 17 Months as Hadoop/Spark developers.

I have completed in Bachelor in Engineering from Mumbai University.

My strengths are a positive attitude, quick learning.

Thats all about me.

## 2. Rate Yourself in

Spark 7/10

Scala 7/10

Hive 8/10

Sql 8/10

Unix 6/10

Java 5/10

## 3. Tell About your current Project

I m currently working in Data Cloud Insights projects.

We pull the Client Data from Oracle Datawarehouse using Spark Jobs. Generate Cube by comparing left and right value (For Example Turnover Rate in 2014 vs Turnover Rate in 2015) and generate Difference.

We can do this for 3 to 4 dimentions like time, job , location.

From these geneated values we filter most usefull differences and exported to Oracle which is used to send mobile notification to Managers.

Ex : Turnover Rate of your company is 4.0% More than previous year  
Turnover Rate of your company is 2% more in New york location

Z-Score ?

## 4. How many years of experience in Spark and Big data Ecosystem

Initially I started working in Oracle and PL/SQL and In Techmahindra I got opportunity to work in Hadoop and Spark.

And After moving to ADP being in Data Cloud Project I got more opportunity to explore and worked in Hive,Sqoop,Spark,Kafka

## 5. What are roles and Responsibility of you in your team

I work as Data Engineer in my project. My Roles and Responsibility are

1. Writing Optimized Oracle Query to get required details from data Warehouse
2. Data Extraction
3. Data Cleansing to remove anamolies
4. Creating Hive Table
5. Data Processing in Spark using Scala
6. Data Export to Oracle

## 6. Explain your Dev and Production cluster

Dev cluseter :

1. 15 Node Cluster
2. 50 TB Hard Disk
3. 2 TB of RAM

Prod **Cluster** :

1. 30 Node **Cluster**
2. 100 TB Hard Disk
3. 4 TB **of** RAM

7. What version that your **using for**

Spark  
Hive  
Scala  
Hadoop

'Spark '

1. What **is** RDDs **and** why they **are** immutable

RDD (Resilient Distributed Dataset) **is** the fundamental **data structure of** Apache Spark which **are** an immutable collection **of** objects which computes **on** the different node **of** the **cluster**. **Each and every** dataset **in** Spark RDD **is** logically partitioned across many servers so that they can be computed **on** different nodes **of** the **cluster**.

RDD stands **for** "Resilient Distributed Dataset". It **is** the fundamental **data structure of** Apache Spark. RDD **in** Apache Spark **is** an immutable collection **of** objects which computes **on** the different node **of** the **cluster**.

Decomposing the name RDD:

Resilient, i.e. fault-tolerant **with** the help **of** RDD lineage graph(DAG) **and** so able **to** recompute missing **or** damaged partitions due **to** node failures.  
Distributed, since **Data** resides **on** multiple nodes.  
Dataset represents records **of** the **data** you **work with**. The **user** can load the **data set** externally which can be either JSON **file**, CSV **file**, text **file or** database via JDBC **with no specific data structure**.

There **are** three ways **to create** RDDs

1. Sc.parallelize()
2. **From** Other RDD
3. **From Data Sets like** csv,json,xmlls

2. What **is** Data Frame

DataFrame appeared **in** Spark **Release 1.3.0**. We can term DataFrame **as** Dataset organized **into** named columns. DataFrames **are** similar **to** the **table in** a relational database **or data frame in** R /Python. It can be said **as** a relational **table with** good optimization technique.

The idea behind DataFrame **is** it allows processing **of** a **large** amount **of** structured **data**. DataFrame **contains rows with Schema**. The **schema is** the illustration **of** the **structure of data**.

DataFrame **in** Apache Spark prevails over RDD but **contains** the features **of** RDD **as** well. The features common **to** RDD **and** DataFrame **are** immutability, **in-memory**, resilient, distributed computing capability. It allows the **user to** impose the **structure** onto a distributed collection **of data**. Thus provides higher **level** abstraction.

We can build DataFrame **from** different **data** sources. **For** Example structured **data file**, tables **in** Hive, **external** databases **or** existing RDDs. The Application Programming **Interface (APIs)** **of** DataFrame **is** available **in** various languages. Examples include Scala, **Java**, Python, **and** R.

It makes **large data set** processing even easier. **Data** Frame also allows developers **to** impose a **structure** onto a distributed collection **of data**. **As a result**, it allows higher-level abstraction.

**Data** frame **is both space and** performance efficient.

It can deal **with both** structured **and** unstructured **data** formats, **for** example, Avro, CSV etc .

And also storage systems like HDFS, HIVE tables, MySQL, etc.  
 The DataFrame API's are available in various programming languages. For example Java, Scala, Python, and R.  
 It provides Hive compatibility. As a result, we can run unmodified Hive queries on existing Hive warehouse.  
 Catalyst tree transformation uses DataFrame in four phases: a) Analyze logical plan to solve references. b) Logical plan optimization c) Physical planning d) Code generation to compile part of the query to Java bytecode.  
 It can scale from kilobytes of data on the single laptop to petabytes of data on the large cluster.

### 3. What is Data Set

Dataset is a data structure in SparkSQL which is strongly typed and is a map to a relational schema. It represents structured queries with encoders. It is an extension to dataframe API. Spark Dataset provides both type safety and object-oriented programming interface. We encounter the release of the dataset in Spark 1.6.

The encoder is primary concept in serialization and deserialization (SerDe) framework in Spark SQL. Encoders translate between JVM objects and Spark's internal binary format. Spark has built-in encoders which are very advanced. They generate bytecode to interact with off-heap data.

An encoder provides on-demand access to individual attributes without having to de-serialize an entire object. To make input output time and space efficient, Spark SQL uses SerDe framework. Since encoder knows the schema of record, it can achieve serialization and deserialization.

Spark Dataset is structured and lazy query expression that triggers on the action. Internally dataset represents logical plan. The logical plan tells the computational query that we need to produce the data. the logical plan is a base catalyst query plan for the logical operator to form a logical query plan. When we analyze this and resolve we can form a physical query plan.

Dataset clubs the features of RDD and DataFrame. It provides:

The convenience of RDD.

Performance optimization of DataFrame.

Static type-safety of Scala.

Thus, Datasets provides a more functional programming interface to work with structured data.

### 4. Difference between RDDs and Data Frame and Data sets

(<https://data-flair.training/blogs/apache-spark-rdd-vs-dataframe-vs-dataset/>)

Spark RDD APIs - An RDD stands for Resilient Distributed Datasets. It is Read-only partition collection of records. RDD is the fundamental data structure of Spark. It allows a programmer to perform in-memory computations on large clusters in a fault-tolerant manner. Thus, speed up the task. Follow this link to learn Spark RDD in great detail.

Spark Dataframe APIs - Unlike an RDD, data organized into named columns. For example a table in a relational database. It is an immutable distributed collection of data. DataFrame in Spark allows developers to impose a structure onto a distributed collection of data, allowing higher-level abstraction. Follow this link to learn Spark DataFrame in detail.

Spark Dataset APIs - Datasets in Apache Spark are an extension of DataFrame API which provides type-safe, object-oriented programming interface. Dataset takes advantage of Spark's Catalyst optimizer by exposing expressions and data fields to a query planner.

'Spark Release'

RDD - The RDD APIs have been on Spark since the 1.0 release.

DataFrames - Spark introduced DataFrames in Spark 1.3 release.

DataSet - Spark introduced Dataset in Spark 1.6 release.

'Data Representation'

RDD - RDD is a distributed collection of data elements spread across many machines in

the **cluster**. RDDs are a **set of Java or Scala** objects representing **data**.  
 DataFrame - A DataFrame **is** a distributed collection **of data** organized **into** named columns. It **is** conceptually equal **to** a **table in** a relational database.  
 DataSet - It **is** an extension **of** DataFrame API that provides the functionality **of - type-safe, object-oriented programming interface of** the RDD API **and** performance benefits **of** the Catalyst query optimizer **and off heap** storage mechanism **of** a DataFrame API.

#### 'Data Formats'

RDD - It can easily **and** efficiently process **data** which **is** structured **as well as** unstructured. But **like** Dataframe **and** DataSets, RDD does **not** infer the **schema of** the ingested **data and** requires the **user to** specify it.  
 DataFrame - It can process structured **and** unstructured **data** efficiently. It organizes the **data in** the named **column**. DataFrames allow the Spark **to** manage **schema**.  
 DataSet - It also efficiently processes structured **and** unstructured **data**. It represents **data in** the form **of** JVM objects **of row or** a collection **of row object**. Which **is** represented **in** tabular forms through encoders.

#### 'Data Sources API'

RDD - **Data** source API allows that an RDD could come **from any data** source e.g. text **file,** a database via JDBC etc. **and** easily handle **data with no** predefined **structure**.  
 DataFrame - **Data** source API allows **Data** processing **in** different formats (AVRO, CSV, JSON, **and** storage system HDFS, HIVE tables, MySQL). It can **read and write from** various **data** sources that **are** mentioned above.  
 DataSet - Dataset API **of** spark also support **data from** different sources.

#### 'Immutability and Interoperability'

RDD - RDDs **contains** the collection **of** records which **are** partitioned. The basic unit **of** parallelism **in** an RDD **is** called **partition**. **Each partition is** one logical division **of data** which **is** immutable **and** created through **some** transformation **on** existing partitions. Immutability helps **to** achieve consistency **in** computations. We can move **from** RDD **to** DataFrame (**If** RDD **is in** tabular **format**) **by** toDF() method **or** we can **do** the **reverse by** the .rdd method. Learn various RDD Transformations **and** Actions APIs **with** examples.  
 DataFrame - **After** transforming **into** DataFrame one cannot regenerate a **domain object**. **For** example, **if** you generate testDF **from** testRDD, **then** you won't be able **to** recover the original RDD **of** the test **class**.  
 DataSet - It overcomes the limitation **of** DataFrame **to** regenerate the RDD **from** Dataframe. Datasets allow you **to convert** your existing RDD **and** DataFrames **into** Datasets.

#### 'Compile-time type safety'

RDD - RDD provides a familiar **object-oriented** programming style **with** compile-time **type** safety.  
 DataFrame - **If** you **are** trying **to access** the **column** which does **not** exist **in** the **table in** such **case** Dataframe APIs does **not** support compile-time error. It detects attribute error **only at** runtime.  
 DataSet - It provides compile-time **type** safety.  
 Learn: Apache Spark vs. Hadoop MapReduce

#### 'Optimization'

RDD - **No** inbuilt optimization engine **is** available **in** RDD. **When** working **with** structured **data**, RDDs cannot take advantages **of** sparks advance optimizers. **For** example, catalyst optimizer **and** Tungsten execution engine. Developers optimise **each** RDD **on** the basis **of** its attributes.  
 DataFrame - Optimization takes place **using** catalyst optimizer. Dataframes **use** catalyst tree transformation framework **in** four phases: a) Analyzing a logical plan **to** resolve **references**. b) Logical plan optimization. c) Physical planning. d) Code generation **to** compile parts **of** the query **to Java** bytecode. The brief overview **of** optimization phase **is** also given **in** the below figure:  
 Spark-SQL-Optimization  
 Dataset - It includes the concept **of** Dataframe Catalyst optimizer **for** optimizing query plan.

#### 'Serialization'

RDD - **Whenever** Spark needs **to** distribute the **data** within the **cluster or write** the **data** **to** disk, it does so **use Java** serialization. The overhead **of** serializing individual **Java**

and Scala objects **is** expensive **and** requires sending **both data and structure between** nodes.

DataFrame - Spark DataFrame Can serialize the **data into off-heap** storage (**in memory**) **in binary format and then** perform many transformations directly **on** this **off heap** memory because spark understands the **schema**. There **is no** need **to use java** serialization **to encode** the **data**. It provides a Tungsten physical execution backend which explicitly manages memory **and** dynamically generates bytecode **for** expression evaluation.

DataSet - **When** it comes **to** serializing **data**, the Dataset API **in** Spark has the concept **of** an encoder which handles conversion **between** JVM objects **to** tabular representation. It stores tabular representation **using** spark internal Tungsten **binary format**. Dataset allows performing the **operation on** serialized **data and** improving memory **use**. It allows **on-demand access to** individual attribute **without** deserializing the entire **object**.

#### 'Garbage Collection'

RDD - There **is** overhead **for** garbage collection that results **from** creating **and** destroying individual objects.

DataFrame - Avoids the garbage collection costs **in** constructing individual objects **for each row in** the dataset.

DataSet - There **is** also **no** need **for** the garbage collector **to destroy object** because serialization takes place through Tungsten. That uses **off heap data** serialization.

#### 'Efficiency/Memory use'

RDD - Efficiency **is** decreased **when** serialization **is** performed individually **on** a **java and scala object** which takes lots **of time**.

DataFrame - **Use of off heap** memory **for** serialization reduces the overhead. It generates byte code dynamically so that many operations can be performed **on** that serialized **data**. **No** need **for** deserialization **for** small operations.

DataSet - It allows performing an **operation on** serialized **data and** improving memory **use**. Thus it allows **on-demand access to** individual attribute **without** deserializing the entire **object**.

#### 'Lazy Evolution'

RDD - Spark evaluates RDDs lazily. They **do not** compute their **result right** away. Instead, they just remember the transformation applied **to some** base **data set**. Spark compute Transformations **only when** an action needs a **result to** sent **to** the driver program. Refer this guide **if** you **are new to** the Lazy Evaluation feature **of** Spark.  
Apache Spark Lazy Evaluation Feature.

DataFrame - Spark evaluates DataFrame lazily, that means computation happens **only when** action appears (**like** display **result**, save **output**).

DataSet - It also evaluates lazily **as** RDD **and** Dataset.

### 5. Difference between Spark 1.0 and Spark 2.0

### 6. Difference Between Repartitions and coalesce

1. **Both** coalesce **and** repartition enables the re assigning **of** the partitions **at run time**.
2. coalesce **by** defaultly shuffling **is false**
3. Re-partition **by** defaultly shuffling **is true**

we can switch **off** shuffling **in** repartition that will behave **as** coalesce  
we can **not** switch **on** shuffling **in** coalesce

repartition **is not** available **in** apache storm

coalesce **is** re-recommendable

Example coalesce :

```
val x = sc.parallelize(Array(1,2,3,4,5),3)
val y = x.coalesce(2,false)
```

```
println(y.getNumPartitions)
```

Example repartition :

```
val x = sc.parallelize(Array(1,2,3,4,5),3)
val y = x.repartition(2)
println(y.getNumPartitions)
```

## 7. Different kinds of Transformation and Different types of Transformation

'Types Transformation in Spark'

- ```
-----
```
1. Map --> Realtime
  2. FlatMap --> Realtime
  3. Filter --> Realtime
  4. join --> Realtime
  5. groupByKey --> Realtime
  6. reduceByKey --> Realtime
  7. aggregateByKey
  8. mapPartition
  9. mapPartitionWithIndex
  10. coalesce --> Realtime
  11. repartition --> Realtime
  12. cogroup
  13. union
  14. union all
  15. distinct
  16. sortBy
  17. intersect
  18. cartesian

### Key-value

1. aggregateByKey
2. reduceByKey
3. groupByKey
4. sortByKey
5. join
6. cogroup

## 8. Different Actions

```
count()
collect()
take(n)
top()
countByValue()
reduce()
fold()
aggregate()
foreach()
```

## 9. Features of RDD

In-memory computation  
 Lazy Evaluation  
 Fault Tolerance  
 Immutability  
 Persistence  
 Partitioning  
 Parallel  
 Location-Stickiness

Coarse-grained **Operation**  
 Typed  
**No** limitation

## 10. Performance Tuning in Spark

### 11. Difference between Persist vs cache

Spark RDD persistence **is** an optimization technique **in** which saves the **result of** RDD evaluation. **Using** this we save the intermediate **result** so that we can **use** it further **if** required. It reduces the computation overhead.

We can make persisted RDD through `cache()` and `persist()` methods. **When** we **use** the `cache()` method we can store **all** the RDD **in-memory**. We can persist the RDD **in** memory **and use** it efficiently across parallel operations.

The **difference between** `cache()` and `persist()` **is** that **using** `cache()` the **default** storage **level** **is** `MEMORY_ONLY` **while using** `persist()` we can **use** various storage levels (described below). It **is** a **key** tool **for** an interactive algorithm. Because, **when** we persist RDD **each** node stores **any partition of** it that it computes **in** memory **and** makes it reusable **for** future **use**. This process speeds up the further computation ten times.

There **are some** advantages **of** RDD caching **and** persistence mechanism **in** spark. It makes the whole system

**Time** efficient  
 Cost efficient  
 Lessen the execution **time**.

**Using** `persist()` we can **use** various storage levels **to** Store Persisted RDDs **in** Apache Spark.

1. `MEMORY_ONLY`
2. `MEMORY_AND_DISK`
3. `MEMORY_ONLY_SER`
4. `MEMORY_AND_DISK_SER`
5. `DISK_ONLY`
6. `MEMORY_ONLY_2`
7. `MEMORY_AND_DISK_2`

`RDD.unpersist()` --> To Unpersisit

## 12. What is Spark SQL

Apache Spark **SQL** **is** a **module for** structured **data** processing **in** Spark. **Using** the **interface** provided **by** Spark **SQL** we **get** more information about the **structure of** the **data and** the computation performed. **With** this extra information, one can achieve extra optimization **in** Apache Spark. We can interact **with** Spark **SQL** **in** various ways **like** `DataFrame` **and** the `Dataset` API. The Same execution engine **is** used **while** computing a **result**, irrespective **of** which API/**language** we **use to** express the computation. Thus, the **user** can easily switch back **and** forth **between** different APIs, it provides the most **natural** way **to** express a given transformation.

**In** Apache Spark **SQL** we can **use** structured **and** semi-structured **data in** three ways:

**To** simplify working **with** structured **data** it provides `DataFrame` abstraction **in** Python, **Java**, **and** Scala. `DataFrame` **is** a distributed collection **of** **data** organized **into** named columns. It provides a good optimization technique.

The **data** can be **read and** written **in** a variety **of** structured formats. **For** example, JSON, Hive Tables, **and** Parquet.

**Using** **SQL** we can query **data**, **both from** inside a Spark program **and from external** tools. The **external** tool connects through standard database connectors (`JDBC/ODBC`) **to** Spark **SQL**.

The best way **to use** Spark **SQL** **is** inside a Spark application. This empowers us **to** load **data and** query it **with** **SQL**. **At** the same **time**, we can also combine it **with** "regular" program code **in** Python, **Java** or Scala.

There were **some** limitations **with** RDDs. **When** working **with** structured **data**, there was **no**

inbuilt optimization engine. On the basis of attributes, the developer optimized each RDD. Also, there was no provision to handle structured data. The DataFrame in Spark SQL overcomes these limitations of RDD. Spark DataFrame is Spark 1.3 release. It is a distributed collection of data ordered into named columns. Concept wise it is equal to the table in a relational database or a data frame in R/Python. We can create DataFrame using:

Structured data files  
Tables in Hive  
External databases  
Using existing RDD

Spark SQL Datasets  
Spark Dataset is an interface added in version Spark 1.6. it is a distributed collection of data. Dataset provides the benefits of RDDs along with the benefits of Apache Spark SQL's optimized execution engine. Here an encoder is a concept that does conversion between JVM objects and tabular representation.

A Dataset can be made using JVM objects and after that, it can be manipulated using functional transformations (map, filter etc.). The Dataset API is accessible in Scala and Java. Dataset API is not supported by Python. But because of the dynamic nature of Python, many benefits of Dataset API are available. The same is the case with R. Using a Dataset of rows we represent DataFrame in Scala and Java. Follow this comparison guide to learn the comparison between Java vs Scala.

Spark Catalyst Optimizer  
The optimizer used by Spark SQL is Catalyst optimizer. It optimizes all the queries written in Spark SQL and DataFrame DSL. The optimizer helps us to run queries much faster than their counter RDD part. This increases the performance of the system.

Spark Catalyst is a library built as a rule-based system. And each rule focusses on the specific optimization. For example, ConstantFolding focus on eliminating constant expression from the query.

Uses of Apache Spark SQL  
It executes SQL queries.  
We can read data from existing Hive installation using SparkSQL.  
When we run SQL within another programming language we will get the result as Dataset/DataFrame

Advantages of Spark SQL

1. Integrated
2. Unified Data Access
3. High compatibility
4. Standard Connectivity
5. Performance Optimization
6. For batch processing of Hive tables

Disadvantages :

- a. Unsupportive Union type
- b. No error for oversize of varchar type
- c. No support for transactional table
- d. Unsupportive Char type
- e. No support for time-stamp in Avro table.

13. How Fault tolerant achieved in Spark

The basic fault-tolerant semantic of Spark are:  
Since all RDD is an immutable data set. Each RDD keeps track of the lineage of the deterministic operation that employee on fault-tolerant input dataset to create it.

If any partition of an RDD is lost due to a worker node failure, then that partition can be re-computed from the original fault-tolerant dataset using the lineage of operations.

Assuming that all of the RDD transformations are deterministic, the data in the final transformed RDD will always be the same irrespective of failures in the Spark cluster.



To achieve fault tolerance for all the generated RDDs, the achieved data replicates among multiple Spark executors in worker node in the cluster. This result in two types of data that should recover in the event of failure:

Data received and replicated - In this, the data replicates on one of the other nodes. Thus we can retrieve data when a failure occurs.

Data received but buffered for replication - the data does not replicate. Thus the only way to recover fault is by retrieving it again from the source.

Failure can also occur in worker and driver nodes.

Failure of worker node - The node which runs the application code on the cluster is worker node. These are the slave nodes. Any of the worker nodes running executor can fail, thus resulting in loss of in-memory data. If any receivers were running on failed nodes, then their buffer data will vanish.

Failure of driver node - If the driver node running the Spark Streaming application fails, then there is the loss of SparkContent. All executors along with their in-memory data vanishes.

14. What version you are using in Spark

2.1

15. Code Sample 1.x and 2.x

1.x

```
import org.apache.spark.SparkContext
import org.apache.spark.SparkConf

object Wordcount {
  def main(args: Array[String]) {

    val conf = new SparkConf().setAppName("WordCount")
    val sc = new SparkContext(conf)
    if (args.length < 2) {
      println("Usage: ScalaWordCount <input> <output>")
      System.exit(1)
    }
    val rawData = sc.textFile(args(0))
    val words = rawData.flatMap(line => line.split(" "))
    val wordCount = words.map(word => (word, 1)).reduceByKey(_ + _)
    wordCount.saveAsTextFile(args(1))
    sc.stop
  }
}
```

2.x

```
import org.apache.spark.sql.SparkSession

object WordCount {
  def main(args: Array[String]): Unit = {
    val spark = SparkSession.builder.master("local[*]").appName("word count").getOrCreate()
    val sc = spark.sparkContext
    val sqlContext = spark.sqlContext
    import spark.implicits._
    import spark.sql
    println("Success")
    val rawData = sc.textFile(
      "C:\\Users\\NaraVish\\IdeaProjects\\SparkPractice\\Data\\wordcount.txt")
    println("Data SUCCESSFULL")
    val words = rawData.flatMap(line => line.split(" "))
    val wordCount = words.map(word => (word, 1)).reduceByKey(_ + _)
    println(wordCount.count())
  }
}
```

```

wordCount.foreach(println)
wordCount.saveAsTextFile(
"C:\\Users\\NaraVish\\IdeaProjects\\SparkPractice\\Data\\output")
spark.stop()

}
}

```

## 16. What is Lineage Graph in Spark and how does it help in fault tolerant

The basic fault-tolerant semantic of Spark are:

Since all RDD is an immutable data set. Each RDD keeps track of the lineage of the deterministic operation that employee on fault-tolerant input dataset to create it.

If any partition of an RDD is lost due to a worker node failure, then that partition can be re-computed from the original fault-tolerant dataset using the lineage of operations.

Assuming that all of the RDD transformations are deterministic, the data in the final transformed RDD will always be the same irrespective of failures in the Spark cluster.

To achieve fault tolerance for all the generated RDDs, the achieved data replicates among multiple Spark executors in worker node in the cluster. This result in two types of data that should recover in the event of failure:

**Data received and replicated** - In this, the data replicates on one of the other nodes. Thus we can retrieve data when a failure occurs.

**Data received but buffered for replication** - the data does not replicate. Thus the only way to recover fault is by retrieving it again from the source.

Failure can also occur in worker and driver nodes.

**Failure of worker node** - The node which runs the application code on the cluster is worker node. These are the slave nodes. Any of the worker nodes running executor can fail, thus resulting in loss of in-memory data. If any receivers were running on failed nodes, then their buffer data will vanish.

**Failure of driver node** - If the driver node running the Spark Streaming application fails, then there is the loss of SparkContent. All executors along with their in-memory data vanishes.

## 17. Why Data Set are faster than Data Frame

Along with all the above benefits, you cannot overlook the space efficiency and performance gains in using DataFrames and Dataset APIs for two reasons.

**First**, because DataFrame and Dataset APIs are built on top of the Spark SQL engine, it uses Catalyst to generate an optimized logical and physical query plan. Across R, Java, Scala, or Python DataFrame/Dataset APIs, all relation type queries undergo the same code optimizer, providing the space and speed efficiency. Whereas the Dataset[T] typed API is optimized for data engineering tasks, the untyped Dataset[Row] (an alias of DataFrame) is even faster and suitable for interactive analysis.

**Second**, since Spark as a compiler understands your Dataset type JVM object, it maps your type-specific JVM object to Tungsten's internal memory representation using Encoders. As a result, Tungsten Encoders can efficiently serialize/deserialize JVM objects as well as generate compact bytecode that can execute at superior speeds.

**When should I use DataFrames or Datasets?**

If you want rich semantics, high-level abstractions, and domain specific APIs, use DataFrame or Dataset.

If your processing demands high-level expressions, filters, maps, aggregation, averages, sum, SQL queries, columnar access and use of lambda functions on semi-structured data, use DataFrame or Dataset.

If you want higher degree of type-safety at compile time, want typed JVM objects, take advantage of Catalyst optimization, and benefit from Tungsten's efficient code generation, use Dataset.

If you want unification and simplification of APIs across Spark Libraries, use DataFrame or Dataset.  
 If you are a R user, use DataFrames.  
 If you are a Python user, use DataFrames and resort back to RDDs if you need more control.  
 Note that you can always seamlessly interoperate or convert from DataFrame and/or Dataset to an RDD, by simple method call `.rdd`. For instance,

```
// select specific fields from the Dataset, apply a predicate
// using the where() method, convert to an RDD, and show first 10
// RDD rows
val deviceEventsDS = ds.select($"device_name", $"cca3", $"c02_level").where($"c02_level" > 1300)
// convert to RDDs and take the first 10 rows
val eventsRDD = deviceEventsDS.rdd.take(10)
```

Bringing It All Together

In summation, the choice of when to use RDD or DataFrame and/or Dataset seems obvious. While the former offers you low-level functionality and control, the latter allows custom view and structure, offers high-level and domain specific operations, saves space, and executes at superior speeds.

As we examined the lessons we learned from early releases of Spark—how to simplify Spark for developers, how to optimize and make it performant—we decided to elevate the low-level RDD APIs to a high-level abstraction as DataFrame and Dataset and to build this unified data abstraction across libraries atop Catalyst optimizer and Tungsten.

Pick one—DataFrames and/or Dataset or RDDs APIs—that meets your needs and use-case, but I would not be surprised if you fall into the camp of most developers who work with structure and semi-structured data.

## 18. Role of Encoder and working of Encoder

Project Tungsten will be the largest change to Spark's execution engine since the project's inception. It focuses on substantially improving the efficiency of memory and CPU for Spark applications, to push performance closer to the limits of modern hardware. This effort includes three initiatives:

Memory Management and Binary Processing: leveraging application semantics to manage memory explicitly and eliminate the overhead of JVM object model and garbage collection  
 Cache-aware computation: algorithms and data structures to exploit memory hierarchy  
 Code generation: using code generation to exploit modern compilers and CPUs  
 The focus on CPU efficiency is motivated by the fact that Spark workloads are increasingly bottlenecked by CPU and memory use rather than IO and network communication. This trend is shown by recent research on the performance of big data workloads (Ousterhout et al) and we've arrived at similar findings as part of our ongoing tuning and optimization efforts for Databricks Cloud customers.

Why is CPU the new bottleneck? There are many reasons for this. One is that hardware configurations offer increasingly large aggregate IO bandwidth, such as 10Gbps links in networks and high bandwidth SSD's or striped HDD arrays for storage. From a software perspective, Spark's optimizer now allows many workloads to avoid significant disk IO by pruning input data that is not needed in a given job. In Spark's shuffle subsystem, serialization and hashing (which are CPU bound) have been shown to be key bottlenecks, rather than raw network throughput of underlying hardware. All these trends mean that Spark today is often constrained by CPU efficiency and memory pressure rather than IO.

## 19. How Spark is Better than Hadoop

Apache Spark is lightening fast cluster computing tool. It is up to 100 times faster than Hadoop MapReduce due to its very fast in-memory data analytics processing power.

Apache Spark is a Big Data Framework. Apache Spark is a general purpose data processing engine and is generally used on top of HDFS. Apache Spark is suitable for the variety of data processing requirements ranging from Batch Processing to Data Streaming.

Hadoop is an open source framework which processes data stored in HDFS. Hadoop can process structured, unstructured or semi-structured data. Hadoop MapReduce can process the data only in

Batch mode.

Apache Spark surpasses Hadoop in many cases such as

1. Processing the data in memory which is not possible in Hadoop
2. Processing the data that is in batch, iterative, interactive & streaming i.e. Real Time mode. Whereas Hadoop processes only in batch mode.
3. Spark is faster because it reduces the number of disk read-write operations due to its virtue of storing intermediate data in memory. Whereas in Hadoop MapReduce intermediate output which is output of Map() is always written on local hard disk
4. Apache Spark is easy to program as it has hundreds of high-level operators with RDD (Resilient Distributed Dataset)
5. Apache Spark code is compact due compared to Hadoop MapReduce. Use of Scala makes it very short, reduces programming efforts. Also, Spark provides rich APIs in various languages such as Java, Scala, Python, and R.
6. Spark & Hadoop are both highly fault-tolerant.
7. Spark application running in Hadoop clusters is up to 10 times faster on disk than Hadoop MapReduce.

## 20. Explain Spark Architecture and Spark Ecosystem

Spark Core - Spark Core is the foundation of the whole project. All the functionality that is in Spark, is present on the top of Spark Core.

Spark Streaming - It allows fault-tolerant streaming of live data streams. It is an add-on to core Spark API. Here it makes use of micro-batching for real-time streaming. Thus it packages live data into small batches and delivers to the batch system for processing.

Spark SQL - Spark SQL component is distributed framework for structured data processing. Using Spark SQL Spark gets more information about the structure of data and the computation being performed. As a result, by using this information Spark can perform extra optimization.

Spark MLlib - MLlib is a scalable learning library that discusses both: High-quality algorithm, High speed. The motive behind MLlib creation is to make machine learning scalable and easy. Thus it contains machine learning libraries that have an implementation of various machine learning algorithms.

Spark GraphX - GraphX is API for graphs and graph parallel execution. In order to support graph computation, graphX contains set of fundamental operators like sub graph, joinvertices and an optimized variant of Pregel API. Also, clustering, classification, traversal, searching, and pathfinding is possible in graphX.

SparkR - SparkR is Apache Spark 1.4 release. The key component of SparkR is SparkR DataFrame. Data frames are a fundamental data structure for data processing in R and the concept of data frames extends to other languages with libraries like Pandas etc.

## 21. What is Main Abstraction of Spark

whenever the term basic abstraction in Apache Spark arises, the only name strikes in mind is .. RDD.., RDD stands for "Resilient Distributed Dataset". It is the fundamental abstraction in Apache Spark. It is the basic data structure. RDD in Apache Spark is an immutable collection of objects which computes on the different node of the cluster.

RDD stands for "Resilient Distributed Dataset". It is the fundamental data structure of Apache Spark. RDD in Apache Spark is an immutable collection of objects which computes on the different node of the cluster.

Resilient, i.e. fault-tolerant with the help of RDD lineage graph(DAG) and so able to recompute missing or damaged partitions due to node failures.  
Distributed, since Data resides on multiple nodes.  
Dataset represents records of the data you work with. The user can load the data set externally which can be either JSON file, CSV file, text file or database via JDBC with no specific data structure

## 22. How to Integrate Hive and Spark ? And What are its advantages

Spark **SQL** supports Apache Hive **using** HiveContext. It uses the Spark **SQL** execution engine **to work with data** stored **in** Hive.

HiveContext **is** a specialized SQLContext **to work with** Hive.

Import org.apache.spark.sql.hive **package to use** HiveContext

log4j.logger.org.apache.spark.sql.hive.HiveContext=DEBUG

SQLContext.**sql** (or simply **sql**) allows you **to** interact **with** Hive.+

You can **use** show functions **to** learn about the Hive functions supported through the Hive integration.

'How to enable Hive context in Spark 2.x' --\*\*\*\*\* V Imp

Method :

enablehivesupport

Warehouse Directory :

cp /usr/lib/hive/conf/hive-site.xml /usr/lib/spark/

```
val spark= SparkSession
    .builder()
    .appName("Spark Hive Example")
    .config("spark.sql.warehouse.dir","warehouseLocation")
    .enablehivesupport
    .getOrCreate()
```

'Why DataFrames are very Powerful'

DataFrame = RDD + Catalyst Optimizer + DAG + **in** Memory

DataSet = RDDs + Catalyst Optimizer + CPU Caches

CBO = Cost Based Optimizer

Catalyst Optimizer internally uses CBO

### 23. Pair RDD **and** Different Transformation

Paired RDDs **are** the RDD-containing **key-value** pair. A **key-value** pair (KVP) **contains** two linked **data** item. Here **Key is** the identifier **and Value are** the **data corresponding to the key value**.

Transformations :

#### **Key-value**

1. aggregateByKey
2. reduceByKey
3. groupByKey
4. sortByKey
5. **join**
6. cogroup

### 24. lazy Evaluation **in** Spark **and** its benefits

The lazy evaluation known **as call-by-need is** a strategy that delays the execution until one requires a **value**. The transformation **in** Spark **is** lazy **in** nature. Spark evaluate them lazily.

**When** we **call some operation in** RDD it does **not execute** immediately; Spark maintains the

graph **of** which **operation** it demands. We can **execute** the **operation at any** instance **by** calling the action **on** the **data**. The **data** does **not** loads until it **is** necessary.

**Read** about Spark Lazy Evaluation **in** detail.

**Q.21)** What **are** the benefits **of** lazy evaluation?

**Using** lazy evaluation we can:

Increase the manageability **of** the program.

Saves computation overhead **and** increases the speed **of** the system.

Reduces the **time and space** complexity.

provides the optimization **by** reducing the **number of** queries.

## 25. Json **in** Hive **and** Spark

JSON **In** Hive

```
create table json_table(str String);
load data local inpath '' into table json_table;
select get_json_object(str,'$.ecode') as ecode, get_json_object(str,'$.ename') as ename ,
get_json_object(str,'$.sal') as salary from json_guru;
```

```
import org.apache.spark.sql.SparkSession
```

**object** JsonExample {

```
def main(args: Array[String]): Unit = {
```

```
val spark = SparkSession.builder.master("local[*]").appName("JsonExample").getOrCreate()
val sc = spark.sparkContext
val sqlContext = spark.sqlContext
```

```
//Converting RDD to Data Frame
```

```
import spark.implicits._
```

```
import spark.sql
```

```
val df = sqlContext.read.json("C:\\Users\\NaraVish\\Desktop\\#Personal\\#Interview Documents\\filformatsinspark\\world_bank.json")
```

```
df.printSchema() //printing schema
```

```
df.createOrReplaceTempView("jsondata_one")
```

```
df.show()
```

```
//val result = sqlContext.sql("select url,totalamt,abc.* from jsondata_one " + "lateral view explode(theme_namecode) as abc")
```

```
val result = sqlContext.sql("select _id from jsondata_one")
```

```
result.show(10)
```

```
// println(result)
```

```
//result.write.format("com.databricks.spark.csv").option("header","true").save(
"C:\\Users\\sonirai\\Desktop\\Hadoop GV\\Spark\\SparkSQL\\datasets\\jsontocsv")
spark.stop()
```

```
}
```

```
}
```

## 26. Join Example **using** Spark Core **and** Spark SQL

```
val edata = sc.textFile("file:///home/cloudera/emp.txt")
```

```
val ddata = sc.textFile("file:///home/cloudera/dept.txt")
```

```
val edata_pair = edata.map{ x =>
```

```
val w = x.split(",")
```

```

val eno = w(0).toInt
val ename = w(1)
val sal = w(2).toInt
val gendar = w(3)
val dno = w(4).toInt
(dno, (eno, ename, sal, gendar))
}

val ddata_pair = ddata.map { x =>
val w = x.split(",")
val dno = w(0).toInt
val dname = w(1)
val dloc = w(2)
(dno, (dname, dloc))
}

val edata_pair_join_ddata_pair = edata_pair.join(ddata_pair)

select dno, loc, avg(sal), max(sal), min(sal) from emp e join dept d
where e.dno=d.dno
group by dno, dloc

import org.apache.spark.sql.SQLContext

val sqlContext = new SQLContext(sc)

val emp = sc.textFile("file:///home/cloudera/emp.txt")
val dept = sc.textFile("file:///home/cloudera/dept.txt")

case class Employee (eno:Int, ename:String, sal:Int, gendar:String, dno:Int)
case class Department (dno:Int, dname:String, dloc:String)

val edata = emp.map{ x =>
val w = x.split(",")
val eno = w(0).toInt
val ename=w(1)
val sal = w(2).toInt
val gendar = w(3)
val dno =w(4).toInt
Employee (eno, ename, sal, gendar, dno)
}

val ddata = dept.map{ x =>
val w = x.split(",")
val dno = w(0).toInt
val dname=w(1)
val dloc = w(2)
Department (dno, dname, dloc)
}

--//converting RDD to DataFrame

import sqlContext.implicits._

val edf = edata.toDF
val ddf = ddata.toDF

edf.show()
ddf.show()

edf.registerTempTable("empview")
ddf.registerTempTable("deptview")

```

```
val eresult= sqlContext.sql("select d.dno,d.dloc,avg(sal) as AVG_SAL ,max(sal) MAX_SAL
,min(sal) MIN_SAL,count(*) COUNT_SAL from empview e join deptview d on e.dno=d.dno group by
d.dno,d.dloc")
```

```
val eresult= sqlContext.sql("select d.dno,d.dloc from empview e join deptview d on e.dno=d.dno")
```

## 27. What is Project Tungsten in Spark

Project Tungsten will be the largest change to Spark's execution engine since the project's inception. It focuses on substantially improving the efficiency of memory and CPU for Spark applications, to push performance closer to the limits of modern hardware. This effort includes three initiatives:

Memory Management and Binary Processing: leveraging application semantics to manage memory explicitly and eliminate the overhead of JVM object model and garbage collection  
 Cache-aware computation: algorithms and data structures to exploit memory hierarchy  
 Code generation: using code generation to exploit modern compilers and CPUs  
 The focus on CPU efficiency is motivated by the fact that Spark workloads are increasingly bottlenecked by CPU and memory use rather than IO and network communication. This trend is shown by recent research on the performance of big data workloads (Ousterhout et al) and we've arrived at similar findings as part of our ongoing tuning and optimization efforts for Databricks Cloud customers.

Why is CPU the new bottleneck? There are many reasons for this. One is that hardware configurations offer increasingly large aggregate IO bandwidth, such as 10Gbps links in networks and high bandwidth SSD's or striped HDD arrays for storage. From a software perspective, Spark's optimizer now allows many workloads to avoid significant disk IO by pruning input data that is not needed in a given job. In Spark's shuffle subsystem, serialization and hashing (which are CPU bound) have been shown to be key bottlenecks, rather than raw network throughput of underlying hardware. All these trends mean that Spark today is often constrained by CPU efficiency and memory pressure rather than IO.

## 28. Why we won't use collect() in production code

When a collect operation is issued on a RDD, the dataset is copied to the driver, i.e. the master node. A memory exception will be thrown if the dataset is too large to fit in memory; takeOrTakeSample can be used to retrieve only a capped number of elements instead.

## 29. Does Spark Requires Hadoop or not ? Explain

Spark is an in-memory distributed computing engine.  
 Hadoop is a framework for distributed storage (HDFS) and distributed processing (YARN).  
 Spark can run with or without Hadoop components (HDFS/YARN)

Distributed Storage:

Since Spark does not have its own distributed storage system, it has to depend on one of these storage systems for distributed computing.

S3 - Non-urgent batch jobs. S3 fits very specific use cases when data locality isn't critical.

Cassandra - Perfect for streaming data analysis and an overkill for batch jobs.

HDFS - Great fit for batch jobs without compromising on data locality.

Distributed processing:

You can run Spark in three different modes: Standalone, YARN and Mesos

## 30. What is Broadcast Variable and Accumulators and What are its usage

Broadcast variables allow the programmer to keep a read-only variable cached on each machine rather than shipping a copy of it with tasks. They can be used, for example, to give every



node a copy of a large input dataset in an efficient manner. Spark also attempts to distribute broadcast variables using efficient broadcast algorithms to reduce communication cost.

Spark actions are executed through a set of stages, separated by distributed "shuffle" operations. Spark automatically broadcasts the common data needed by tasks within each stage. The data broadcasted this way is cached in serialized form and deserialized before running each task. This means that explicitly creating broadcast variables is only useful when tasks across multiple stages need the same data or when caching the data in deserialized form is important.

Broadcast variables are created from a variable v by calling SparkContext.broadcast(v). The broadcast variable is a wrapper around v, and its value can be accessed by calling the value method. The code below shows this:

Accumulators are variables that are only "added" to through an associative and commutative operation and can therefore be efficiently supported in parallel. They can be used to implement counters (as in MapReduce) or sums. Spark natively supports accumulators of numeric types, and programmers can add support for new types.

As a user, you can create named or unnamed accumulators. As seen in the image below, a named accumulator (in this instance counter) will display in the web UI for the stage that modifies that accumulator. Spark displays the value for each accumulator modified by a task in the "Tasks" table.

### 31. Where you used Apache Spark in your Project

We used Spark to generate Insight Cube to generate all values for all dimensions. And also NRT Streaming to read kafka topics

### 32. Explain Catalyst Framework

Spark SQL is one of the newest and most technically involved components of Spark. It powers both SQL queries and the new DataFrame API. At the core of Spark SQL is the Catalyst optimizer, which leverages advanced programming language features (e.g. Scala's pattern matching and quasiquotes) in a novel way to build an extensible query optimizer.

We recently published a paper on Spark SQL that will appear in SIGMOD 2015 (co-authored with Davies Liu, Joseph K. Bradley, Xiangrui Meng, Tomer Kaftan, Michael J. Franklin, and Ali Ghodsi). In this blog post we are republishing a section in the paper that explains the internals of the Catalyst optimizer for broader consumption.

To implement Spark SQL, we designed a new extensible optimizer, Catalyst, based on functional programming constructs in Scala. Catalyst's extensible design had two purposes. First, we wanted to make it easy to add new optimization techniques and features to Spark SQL, especially for the purpose of tackling various problems we were seeing with big data (e.g., semistructured data and advanced analytics). Second, we wanted to enable external developers to extend the optimizer – for example, by adding data source specific rules that can push filtering or aggregation into external storage systems, or support for new data types. Catalyst supports both rule-based and cost-based optimization.

At its core, Catalyst contains a general library for representing trees and applying rules to manipulate them. On top of this framework, we have built libraries specific to relational query processing (e.g., expressions, logical query plans), and several sets of rules that handle different phases of query execution: analysis, logical optimization, physical planning, and code generation to compile parts of queries to Java bytecode. For the latter, we use another Scala feature, quasiquotes, that makes it easy to generate code at runtime from composable expressions. Finally, Catalyst offers several public extension points, including external data sources and user-defined types.

We use Catalyst's general tree transformation framework in four phases, as shown below: (1) analyzing a logical plan to resolve references, (2) logical plan optimization, (3) physical planning, and (4) code generation to compile parts of the query to Java bytecode. In the physical planning phase, Catalyst may generate multiple plans and compare them based on cost

. All other phases are purely rule-based. Each phase uses different types of tree nodes; Catalyst includes libraries of nodes for expressions, data types, and logical and physical operators. We now describe each of these phases.

### 33. What are advantages of Parquet File format ? Difference between Avro and Parquet file

Parquet is an open source file format for Hadoop. Parquet stores nested data structures in a flat columnar format compared to a traditional approach where data is stored in row-oriented approach, parquet is more efficient in terms of storage and performance.

There are several advantages to columnar formats:

- 1) Organizing by column allows for better compression, as data is more homogeneous. The space savings are very noticeable at the scale of a Hadoop cluster.
- 2) I/O will be reduced as we can efficiently scan only a subset of the columns while reading the data. Better compression also reduces the bandwidth required to read the input.
- 3) As we store data of the same type in each column, we can use encoding better suited to the modern processors' pipeline by making instruction branching more predictable.

Parquet vs Avro

Avro is a row-based storage format for Hadoop.

Parquet is a column-based storage format for Hadoop.

If your use case typically scans or retrieves all of the fields in a row in each query, Avro is usually the best choice.

If your dataset has many columns, and your use case typically involves working with a subset of those columns rather than entire records, Parquet is optimized for that kind of work.

### 34. Why kairo Serialization is better the Default Java Serialization

Data Serialization

Serialization plays an important role in the performance of any distributed application. Formats that are slow to serialize objects into, or consume a large number of bytes, will greatly slow down the computation. Often, this will be the first thing you should tune to optimize a Spark application. Spark aims to strike a balance between convenience (allowing you to work with any Java type in your operations) and performance. It provides two serialization libraries:

**Java serialization:** By default, Spark serializes objects using Java's ObjectOutputStream framework, and can work with any class you create that implements java.io.Serializable. You can also control the performance of your serialization more closely by extending java.io.Externalizable. Java serialization is flexible but often quite slow, and leads to large serialized formats for many classes.

**Kryo serialization:** Spark can also use the Kryo library (version 2) to serialize objects more quickly. Kryo is significantly faster and more compact than Java serialization (often as much as 10x), but does not support all Serializable types and requires you to register the classes you'll use in the program in advance for best performance.

### 35. Checkpointing in Spark

As an Apache Spark application developer, memory management is one of the most essential tasks, but the difference between caching and checkpointing can cause confusion. Both operations are essential in preventing Spark from having to lazily recompute a resilient distributed dataset (RDD) every time it is referenced, but there are also key differences between the two.

Caching computes and materializes an RDD in memory while keeping track of its lineage (dependencies). There are many levels of persistence supported that allow you to make space and compute cost tradeoffs, and specify the behavior of the RDD when it runs out of memory. Since caching remembers an RDD's lineage, Spark can recompute lost partitions in the event

of node failures. Lastly, an RDD that is cached lives within the context of the running application, and once the application terminates, cached RDDs are deleted as well.

Checkpointing saves an RDD to a reliable storage system (e.g. HDFS, S3) while forgetting the RDD's lineage completely. Truncating dependencies becomes relevant especially when the RDD's lineage starts getting long. Checkpointing an RDD is similar to how Hadoop stores intermediate computation values to disk, trading off execution latency with ease of recovering from failures. Since an RDD is checkpointed in an external storage system, it can be reused by other applications.

Now the bigger question is how caching and checkpointing interplay. Let's trace through the compute path of an RDD to find out more.

At the core of Spark's engine is the DAGScheduler that breaks down a job (generated by a Spark action) into a DAG of stages. Each of these shuffle or result stages is further broken down into individual tasks that run on a partition of an RDD. An RDD's iterator method is the entry point for a task to access the underlying data partition. We can see from this method that if the storage level is set, indicating that the RDD may be cached, it first attempts to getOrCompute the partition from the block manager. If the block manager does not have the RDD's partition, it falls back to computeOrReadCheckpoint. As you can guess, computeOrReadCheckpoint retrieves checkpointed values if it exists, and if not, only then is the data partition computed.

All that being said, it is up to you to decide which of the two match your use case at different points in your job. It takes longer to read and write a checkpointed RDD simply because it has to be persisted to an external storage system, but Spark worker failures need not result in a recomputation (assuming the data is intact in the external storage system). On the other hand, cached RDD's will not permanently take up storage space, but recomputation is necessary on worker failure. In general, the length of time it takes to do a computation is a good indicator to use one or the other.

### 36. MLib in your Project ?

Paycode classification

### 37. Fold Operation in Spark

Fold is a very powerful operation in spark which allows you to calculate many important values in  $O(n)$  time. If you are familiar with Scala collection it will be like using fold operation on collection. Even if you not used fold in Scala, this post will make you comfortable in using fold.

Syntax

```
def fold[T](acc:T)((acc,value) => acc)
```

The above is kind of high level view of fold api. It has following three things

T is the data type of RDD

acc is accumulator of type T which will be return value of the fold operation

A function, which will be called for each element in rdd with previous accumulator.

Let's see some examples of fold

Finding max in a given RDD

Let's first build a RDD

```
val sparkContext = new SparkContext("local", "functional")
val employeeData = List(("Jack",1000.0),("Bob",2000.0),("Carl",7000.0))
val employeeRDD = sparkContext.makeRDD(employeeData)
```

Now we want to find an employee, with maximum salary. We can do that using fold.

To use fold we need a start value. The following code defines a dummy employee as starting accumulator.

```
val dummyEmployee = ("dummy",0.0);
```

Now using fold, we can find the employee with maximum salary.

```
val maxSalaryEmployee = employeeRDD.fold(dummyEmployee)((acc,employee) => {
  if(acc._2 < employee._2) employee else acc})
println("employee with maximum salary is"+maxSalaryEmployee)
Fold by key
```

In Map/Reduce **key** plays a role of grouping values. We can use **foldByKey** operation to aggregate values based on keys.

In this example, employees are grouped by department name. If you want to find the maximum salaries in a given department we can use following code.

```
val deptEmployees = List(
  ("cs", ("jack", 1000.0)),
  ("cs", ("bron", 1200.0)),
  ("phy", ("sam", 2200.0)),
  ("phy", ("ronaldo", 500.0))
)
val employeeRDD = sparkContext.makeRDD(deptEmployees)

val maxByDept = employeeRDD.foldByKey(("dummy", 0.0))
((acc, element) => if(acc._2 > element._2) acc else element)

println("maximum salaries in each dept" + maxByDept.collect().toList)
```

38. How Spark Can you be used for Data Extraction from RDBMS,  
How it is better than Sqoop

```
/spark-2.1.0-bin-hadoop2.7/bin/pyspark
--jars "/home/jars/ojdbc6.jar"
--master yarn-client
--num-executors 10
--driver-memory 16g
--executor-memory 8g
```

```
empDF = spark.read \
  .format("jdbc") \
  .option("url", "jdbc:oracle:thin:username/password@//hostname:portnumber/SID") \
  .option("dbtable", "hr.emp") \
  .option("user", "db_user_name") \
  .option("password", "password") \
  .option("driver", "oracle.jdbc.driver.OracleDriver") \
  .load()
```

```
empDF.printSchema()
```

```
empDF.show()
```

The reason Spark is Faster than Sqoop is Spark works with In-Memory.  
Sqoop rights the data to Disk that increases I/O Operation.

39. Roles and Responsibility of

1. Driver
2. Executor
3. Worker Node

Spark Driver - Master Node of a Spark Application

It is the central point and the entry point of the Spark Shell (Scala, Python, and R). The driver program runs the main () function of the application and is the place where the Spark Context is created. Spark Driver contains various components - DAGScheduler, TaskScheduler, BackendScheduler and BlockManager responsible for the translation of spark user code into

actual spark jobs executed **on** the **cluster**.

The driver program that runs **on** the master node **of** the spark **cluster** schedules the job execution **and** negotiates **with** the **cluster** manager.  
 It translates the RDD's **into** the execution graph **and** splits the graph **into** multiple stages.  
 Driver stores the metadata about **all** the Resilient Distributed Databases **and** their partitions.  
 Cockpits **of** Jobs **and** Tasks Execution -Driver program converts a **user** application **into** smaller execution units known **as** tasks. Tasks **are then** executed **by** the executors i.e. the worker processes which run individual tasks.  
 Driver exposes the information about the running spark application through a Web UI **at** port **4040**.  
 Role **of** Executor **in** Spark Architecture

Executor **is** a distributed agent responsible **for** the execution **of** tasks. **Every** spark applications has its own executor process. Executors usually run **for** the entire lifetime **of** a Spark application **and** this phenomenon **is** known **as** "**Static Allocation of Executors**". However, users can also opt **for dynamic** allocations **of** executors wherein they can **add or** remove spark executors dynamically **to match with** the overall workload.

Executor performs **all** the **data** processing.  
**Reads from and Writes data to external** sources.  
 Executor stores the computation results **data in-memory**, cache **or on** hard disk drives.  
 Interacts **with** the storage systems.  
 Role **of Cluster Manager in** Spark Architecture

An **external** service responsible **for** acquiring resources **on** the spark **cluster and** allocating them **to** a spark job. There **are 3** different types **of cluster** managers a Spark application can leverage **for** the allocation **and** deallocation **of** various physical resources such **as** memory **for** client spark jobs, CPU memory, etc. Hadoop YARN, Apache Mesos **or** the simple standalone spark **cluster** manager either **of** them can be launched **on-premise or in** the cloud **for** a spark application **to** run.

Choosing a **cluster** manager **for any** spark application depends **on** the goals **of** the application because **all cluster** managers provide different **set of** scheduling capabilities. **To get** started **with** apache spark, the standalone **cluster** manager **is** the easiest one **to use when** developing a **new** spark application.

#### 40. Spark Submit Job Command

```
execute spark-submit --executor-cores 8 --num-executors 16 \
  --executor-memory 35g --master yarn \
  --driver-memory 20g \
  --deploy-mode cluster \
  --name {env}-annual_benchmarks \
  --files /app/dsenv-{env}/dsmain-benchmarks/cook/builder/annual_comp_benchmarks.xml \
  --conf spark.yarn.executor.memoryOverhead=12000 \
  --conf spark.core.connection.ack.wait.timeout=200s \
  --conf spark.yarn.driver.memoryOverhead=12000 \
  /app/dsenv-{env}/dscommon-benchmarkstudio/benchmark_builder/benchmark_builder.py -f
  annual_comp_benchmarks.xml -p {201706}
```

#### 41. Explain Apache Streaming and How it is Achieved

Spark Streaming **is** an extension **of** the core Spark API that allows enables scalable, high-throughput, fault-tolerant stream processing **of** live **data** streams. **Data** can be ingested **from** many sources **like** Kafka, Flume, Twitter, ZeroMQ, Kinesis **or** plain old TCP sockets **and** be processed **using** complex algorithms expressed **with** high-level functions **like** map, reduce, join and window. Finally, processed **data** can be pushed **out to** filesystems, databases, and live dashboards. **In** fact, you can apply Spark's machine learning algorithms, **and** graph processing algorithms **on data** streams.

Internally, it works **as** follows. Spark Streaming receives live **input data** streams **and** divides the **data into** batches, which **are then** processed **by** the Spark engine **to** generate the final stream **of** results **in** batches.

Spark Streaming provides a high-level abstraction called discretized stream **or** DStream, which represents a continuous stream **of data**. DStreams can be created either **from input data** stream **from** sources such **as** Kafka, Flume, **and** Kinesis, **or by** applying high-level operations **on** other DStreams. Internally, a DStream **is** represented **as** a **sequence of** RDDs.

This guide shows you how **to start** writing Spark Streaming programs **with** DStreams. You can **write** Spark Streaming programs **in** Scala **or** Java, **both of which are** presented **in** this guide. You will find tabs throughout this guide that let you choose **between** Scala **and** Java code snippets.

```
import org.apache.spark._
import org.apache.spark.streaming._
import org.apache.spark.streaming.StreamingContext._

// Create a local StreamingContext with two working thread and batch interval of 1 second
val conf = new SparkConf().setMaster("local[2]").setAppName("NetworkWordCount")
val ssc = new StreamingContext(conf, Seconds(1))

val lines = ssc.socketTextStream("localhost", 9999)

val words = lines.flatMap(_.split(" "))

import org.apache.spark.streaming.StreamingContext._
// Count each word in each batch
val pairs = words.map(word => (word, 1))
val wordCounts = pairs.reduceByKey(_ + _)

// Print the first ten elements of each RDD generated in this DStream to the console
wordCounts.print()
```

## 42. Explain D-Stream

A Discretized Stream (DStream), it's the fundamental abstraction in Spark Streaming, is a continuous sequence of RDDs of constant kind representing a steady/nonstop stream of information. DStreams may be created from live data like information/data from TCP sockets, Kafka, Flume, etc employing a StreamingContext or it may be generated by working on existing DStreams exploitation functions like map, window, and reduceByKeyAndWindow. Periodically DStream create an RDD which is generated by a parent DStream.

This category contains the fundamental operations offered on all DStreams, like map, filter, and window. additionally, PairDStreamFunctions contains operations offered solely on DStreams of key-value pairs, like groupByKeyAndWindow and be a part of. Through implicit conversions, these operations are offered on any DStream of pairs (e.g., DStream[(Int, Int)]). DStreams internally is characterized by basic properties: - a listing of alternative DStreams depends on - An amount at that the DStream generates an RDD - operate that's want **to** generate an RDD once **on every** occasion **interval**

Discretized Stream may be a **sequence of** Resilient Distributed Databases that represent a stream **of** information. DStreams may be created **from** varied sources **like** Apache Kafka, HDFS, **and** Apache Flume

## 43. What **is** Speculative Execution **in** Spark

Speculative Execution **of** Tasks

Speculative tasks (also speculatable tasks **or** task strugglers) **are** tasks that run slower **than** most (FIXME the setting) **of** the **all** tasks **in** a job. Speculative execution **of** tasks **is** a health-check procedure that checks **for** tasks **to** be speculated, i.e. running slower **in** a stage **than** the median **of all** successfully completed tasks **in** a taskset (FIXME the setting). Such slow tasks will be re-submitted **to** another worker. It

will **not** stop the slow tasks, but run a **new** copy **in** parallel.

The thread starts **as** TaskSchedulerImpl starts **in** clustered deployment modes **with** spark.speculation enabled. It executes periodically **every** spark.speculation.interval **after** the **initial** spark.speculation.interval passes.

**When** enabled, you should see the following INFO message **in** the logs:

INFO TaskSchedulerImpl: Starting speculative execution thread

It works **as** task-scheduler-speculation daemon thread pool **using** j.u.c.

ScheduledThreadPoolExecutor **with** core pool **size** 1.

The job **with** speculatable tasks should finish **while** speculative tasks **are** running, **and** it will leave these tasks running - **no** KILL command yet.

It uses checkSpeculatableTasks method that asks rootPool **to check for** speculatable tasks. **If** there **are any**, SchedulerBackend **is** called **for** reviveOffers.

#### 44. What **are** the Machine Learning algorithm **is** possible **in** Spark

Moreover, it provides following ML Algorithms:

- Basic statistics
- Classification **and** Regression
- Clustering
- Collaborative filtering

#### 45. Difference between Spark Session and Spark Context

Spark Context:

**Prior to** Spark 2.0.0 sparkContext was used **as** a channel **to access all** spark functionality. The spark driver program uses spark context **to connect to** the **cluster** through a **resource** manager (YARN or Mesos..).

sparkConf **is** required **to create** the spark context **object**, which stores configuration **parameter like** appName (**to identify your spark driver**), application, **number of core and memory size of** executor running **on** worker node.

**In order to use** APIs of SQL, HIVE, and Streaming, **separate** contexts need **to** be created.

Example:

creating sparkConf :

```
val conf = new SparkConf().setAppName("RetailDataAnalysis").setMaster("spark://master:7077").set("spark.executor.memory", "2g")
```

creation of sparkContext:

```
val sc = new SparkContext(conf)
```

Spark **Session**:

SPARK 2.0.0 onwards, SparkSession provides a single point **of** entry **to** interact **with** underlying Spark functionality **and** allows programming Spark **with** DataFrame **and** Dataset APIs. **All** the functionality available **with** sparkContext **are** also available **in** sparkSession.

**In order to use** APIs of SQL, HIVE, and Streaming, **no** need **to create separate** contexts **as** sparkSession includes **all** the APIs.

Once the SparkSession **is** instantiated, we can configure Spark's run-time config properties.

Example:

Creating Spark **session**:

```
val spark = SparkSession
  .builder
  .appName("WorldBankIndex")
  .getOrCreate()
```

Configuring properties:

```
spark.conf.set("spark.sql.shuffle.partitions", 6)
```



```
spark.conf.set("spark.executor.memory", "2g")
```

Spark 2.0.0 onwards, it is better to use sparkSession as it provides access to all the spark functionalities that sparkContext does. Also, it provides APIs to work on DataFrames and Datasets.

46. How do you do logging in Spark Job and how to retrieve

Log4j in Apache Spark

Spark uses log4j as the standard library for its own logging. Everything that happens inside Spark gets logged to the shell console and to the configured underlying storage. Spark also provides a template for app writers so we could use the same log4j libraries to add whatever messages we want to the existing and in place implementation of logging in Spark.

Example : ?

47. Difference Between

- SoryByKey vs distributeByKey
- Map vs Map Partition
- Map Partition vs Map Partition with Index

**map(func)** What does it do? Pass each element of the RDD through the supplied function; i.e. func

**flatMap(func)** "Similar to map, but each input item can be mapped to 0 or more output items (so func should return a Seq rather than a single item)."

**mapPartitions(func)** Consider mapPartitions a tool for performance optimization. It won't do much for you when running examples on your local machine compared to running across a cluster. It's the same as map, but works with Spark RDD partitions. Remember the first D in RDD is "Distributed" - Resilient Distributed Datasets. Or, put another way, you could say it is distributed over partitions. enter image description here

**mapPartitionsWithIndex(func)** Similar to mapPartitions, but also provides a function with an Int value to indicate the index position of the partition. enter image description here

d. Repartitions vs coalesce

The repartition algorithm does a full shuffle and creates new partitions with data that is distributed evenly. Let's create a DataFrame with the numbers from 1 to 12.

```
val x = (1 to 12).toList
val numbersDf = x.toDF("number")
numbersDf contains 4 partitions on my machine.
```

```
numbersDf.rdd.partitions.size // => 4
Here is how the data is divided on the partitions:
```

```
Partition 00000: 1, 2, 3
Partition 00001: 4, 5, 6
Partition 00002: 7, 8, 9
Partition 00003: 10, 11, 12
```

Let's do a full-shuffle with the repartition method and get this data on two nodes.

```
val numbersDfR = numbersDf.repartition(2)
Here is how the numbersDfR data is partitioned on my machine:
```

```
Partition A: 1, 3, 4, 6, 7, 9, 10, 12
Partition B: 2, 5, 8, 11
```

The repartition method makes new partitions and evenly distributes the data in the new partitions (the data distribution is more even for larger data sets).



## Difference between coalesce and repartition

coalesce uses existing partitions to minimize the amount of data that's shuffled. repartition creates **new** partitions **and** does a full shuffle. **coalesce** results **in** partitions **with** different amounts **of data** (sometimes partitions that have much different sizes) **and** repartition results **in** roughly equal sized partitions.

## Is coalesce or repartition faster?

**coalesce** may run faster **than** repartition, but unequal sized partitions **are** generally slower **to work with than** equal sized partitions. You'll usually need to repartition datasets after filtering a large data set. I've **found** repartition **to** be faster overall because Spark **is** built **to work with** equal sized partitions.

1. **Both** coalesce **and** repartition enables the re assinging **of** the partitions **at** run **time**.
2. coalesce **by** defaultly shuffling **is false**
3. Re-partition **by** defaultly shuffling **is true**

we can swith **off** shuffling **in** repartition that will behave **as** coalesce  
we can **not** switch **on** shuffling **in** coalesce

repartition **is not** avaiable **in** apache storm

coalesce **is** re-commendable

48. How **to** Identify shuffling **in** spark

**By** Looking **at** DAG graph , **if** it shows extra stage that means shuffling happened

49. Common Mistake developers make **when** it comparately

People often **do** mistakes **in** DAG controlling. So **in order to** avoid such mistakes. We should **do** the following:

Always try **to use** reducebykey instead **of** groupbykey : The ReduceByKey **and** GroupByKey can perform almost similar functions, but GroupByKey **contains large data**. Hence, try **to use** ReduceByKey **to** the most. Make sure you stay away **from** shuffles **as much as** possible:  
Always try **to lower** the side **of** maps **as much as** possibleTry **not to** waste more **time in** PartitioningTry **not to** shuffle moreTry **to** keep away **from** Skews **as well as** partitions  
tooReduce should be lesser **than** TreeReduce: Always **use** TreeReduce instead **of** Reduce, Because TreeReduce does much more **work in** comparison **to** the Reduce **on** the executors.

50. **Difference between** Spark **SQL and** Hive

Apache Hive:

-----  
Primarily, its database model **is** Relational DBMS.

It supports an additional database model, i.e. **Key-value** store

Basically, it supports **all** Operating Systems **with a Java** VM.

It has predefined **data** types. **For** example, **float or date**.

It possesses **SQL-like** DML **and** DDL statements.

Apache Hive supports JDBC, ODBC, **and** Thrift.

We can **use** several programming languages **in** Hive. **For** example C++, **Java**, **PHP**, **and** **Python**.

It uses **data** sharding method **for** storing **data on** different nodes.

There **is** a selectable replication factor **for** redundantly storing **data on** multiple nodes.

Basically, hive supports concurrent manipulation **of data**.

Basically, it supports **for** making **data** persistent.

There **are** access rights **for** users, groups **as well as** roles.

**Schema** flexibility **and** evolution.

Also, can portion **and** bucket, tables **in** Apache Hive.

**As** JDBC/ODBC drivers **are** available **in** Hive, we can **use** it.

It does **not** offer **real-time** queries **and** **row level** updates.

Also provides acceptable latency **for** interactive **data** browsing.

Hive does **not** support **online transaction** processing.

In Apache Hive, latency **for** queries **is** generally very high.

Spark **SQL**:

Primarily, its database model **is** also Relational DBMS

**As** similar **as** Hive, it also supports **Key-value** store **as** additional database model.

It supports several operating systems. **For** example Linux OS, X, **and** Windows.

**As** similar **to** Spark **SQL**, it also has predefined **data** types. **For** Example, **float or date**.

**Like** Apache Hive, it also possesses **SQL-like** DML **and** DDL statements.

Spark **SQL** supports **only** JDBC **and** ODBC.

We can **use** several programming languages **in** Spark **SQL**. **For** example **Java**, Python, R, **and** Scala.

This creates **difference between** SparkSQL **and** Hive.

It uses spark core **for** storing **data on** different nodes.

Basically, **for** redundantly storing **data on** multiple nodes, there **is** a **no** replication factor **in** Spark **SQL**.

Whereas, spark **SQL** also supports concurrent manipulation **of data**.

**As** same **as** Hive, Spark **SQL** also support **for** making **data** persistent.

There **are no** access rights **for** users.

Basically, it performs **SQL** queries.

Through Spark **SQL**, it **is** possible **to read data from** existing Hive installation.

We **get** the **result as** Dataset/DataFrame **if** we run Spark **SQL** **with** another programming **language**.

It does **not** support **union type**

Although, **no** provision **of** error **for** oversize **of** varchar **type**

It does **not** support transactional **table**

However, **no** support **for** Char **type**

It does **not** support **time-stamp in** Avro **table**.

## 51. Explain sliding window operations

Spark Streaming also provides windowed computations, which allow you **to** apply transformations over a sliding window **of data**. The following figure illustrates this sliding window.

Spark Streaming

**As** shown **in** the figure, **every time** the window slides over a source DStream, the source RDDs that fall within the window **are** combined **and** operated upon **to** produce the RDDs **of** the windowed DStream. **In** this **specific case**, the **operation is** applied over the **last 3 time** units **of data**, **and** slides **by 2 time** units. This shows that **any window operation** needs **to** specify two **parameters**.

window **length** - The duration **of** the window (**3 in** the figure).

sliding **interval** - The **interval at** which the window **operation is** performed (**2 in** the figure).

These two **parameters** must be multiples **of** the batch **interval of** the source DStream (**1 in** the figure).

Let's illustrate the window operations **with** an example. Say, you want **to** extend the earlier example **by** generating word counts over the **last 30 seconds of data**, **every 10 seconds**. **To do** this, we have **to** apply the reduceByKey **operation on** the pairs DStream **of** (word, 1) pairs over the **last 30 seconds of data**. This **is** done **using** the **operation** reduceByKeyAndWindow.

Sliding Window controls transmission **of data** packets **between** various computer networks. Spark Streaming library provides windowed computations **where** the transformations **on** RDDs **are** applied over a sliding window **of data**. **Whenever** the window slides, the RDDs that fall within the particular window **are** combined **and** operated upon **to** produce **new** RDDs **of** the windowed DStream.

## 52. Why there **are no** indexes **in** spark **Sql**

Indexing **input data**

The fundamental reason why indexing over **external data** sources **is not in** the Spark **scope is** that Spark **is not** a **data** management system but a batch **data** processing engine. Since it doesn't own

the data it is using it cannot reliably monitor changes and as a consequence cannot maintain indices.

If data source supports indexing it can be indirectly utilized by Spark through mechanisms like predicate pushdown.

Indexing Distributed Data Structures:

standard indexing techniques require persistent and well defined data distribution but data in Spark is typically ephemeral and its exact distribution is nondeterministic.

high level data layout achieved by proper partitioning combined with columnar storage and compression can provide very efficient distributed access without an overhead of creating, storing and maintaining indices. This is a common pattern used by different in-memory columnar systems.

That being said some forms of indexed structures do exist in Spark ecosystem. Most notably Databricks provides Data Skipping Index on its platform.

Other projects, like Succinct (mostly inactive today) take different approach and use advanced compression techniques with with random ac'

### 53. How Memory Handled in Data Sets

### 54. What is Data Piping

A data pipeline is a software that consolidates data from multiple sources and makes it available to be used strategically.

The data pipeline architecture consists of several layers:-

- 1) Data Ingestion
- 2) Data Collector
- 3) Data Processing
- 4) Data Storage
- 5) Data Query
- 6) Data Visualization

Let's get into details of each layer & understand how we can build a real-time data pipeline.

### 55. How Data Security Achieved in Spark

Spark currently supports authentication via a shared secret. Authentication can be configured to be on via the spark.authenticate configuration parameter. This parameter controls whether the Spark communication protocols do authentication using the shared secret. This authentication is a basic handshake to make sure both sides have the same shared secret and are allowed to communicate. If the shared secret is not identical they will not be allowed to communicate. The shared secret is created as follows:

For Spark on YARN deployments, configuring spark.authenticate to true will automatically handle generating and distributing the shared secret. Each application will use a unique shared secret. For other types of Spark deployments, the Spark parameter spark.authenticate.secret should be configured on each of the nodes. This secret will be used by all the Master/Workers and applications.

### 56. Explain Kerberos Security

One of the more confusing topics in Hadoop is how authorization and authentication work in the system. The first and most important thing to recognize is the subtle, yet extremely important, differentiation between authorization and authentication, so let's define these terms first:

Authentication **is** the process **of** determining whether someone **is** who they claim **to** be.

Authorization **is** the **function of** specifying **access** rights **to** resources.

In simpler terms, authentication **is** a way **of** proving who I am, **and authorization is** a way **of** determining what I can **do**.

#### Authentication

If Hadoop **is** configured **with all of** its defaults, Hadoop doesn't **do any** authentication **of** users. This **is** an important realization **to** make, because it can have serious implications **in** a corporate **data** center. Let's look **at** an example **of** this.

Let's say Joe **User** has **access to** a Hadoop **cluster**. The **cluster** does **not** have **any** Hadoop security features enabled, which means that there **are no** attempts made **to** verify the identities **of** users who interact **with** the **cluster**. The **cluster's** superuser **is** hdfs, **and** Joe doesn't have the password **for** the hdfs **user on any of** the **cluster** servers. However, Joe happens **to** have a client machine which has a **set of** configurations that will allow Joe **to access** the Hadoop **cluster**, **and** Joe **is** very disgruntled. He runs these commands:

```
sudo useradd hdfs
sudo -u hdfs hadoop fs -rmr /
```

1  
2

```
sudo useradd hdfs
sudo -u hdfs hadoop fs -rmr /
```

The **cluster** goes **off and** does **some work**, **and** comes back **and** says "Ok, hdfs, I deleted everything!".

So what happened here? Well, **in** an insecure **cluster**, the NameNode **and** the JobTracker don't require **any** authentication. **If** you make a request, **and** say you're hdfs **or** mapred, the NN/JT will **both** say "ok, I believe that," **and** allow you **to do** whatever the hdfs **or** mapred users have the ability **to do**.

Hadoop has the ability **to** require authentication, **in** the form **of** Kerberos principals. Kerberos **is** an authentication protocol which uses "tickets" **to** allow nodes **to** identify themselves. **If** you need a more **in depth** introduction **to** Kerberos, I strongly recommend checking **out** the Wikipedia page.

Hadoop can **use** the Kerberos protocol **to** ensure that **when** someone makes a request, they really **are** who they say they **are**. This mechanism **is** used throughout the **cluster**. **In** a secure Hadoop configuration, **all of** the Hadoop daemons **use** Kerberos **to** perform mutual authentication, which means that **when** two daemons talk **to each** other, they **each** make sure that the other daemon **is** who it says it **is**. Additionally, this allows the NameNode **and** JobTracker **to** ensure that **any** HDFS **or** MR requests **are** being executed **with** the appropriate **authorization level**.

#### Authorization

Authorization **is** a much different beast **than** authentication. **Authorization** tells us what **any** given **user** can **or** cannot **do** within a Hadoop **cluster**, **after** the **user** has been successfully authenticated. **In** HDFS this **is** primarily governed **by file** permissions.

HDFS **file** permissions **are** very similar **to** BSD **file** permissions. **If** you've ever run `ls -l` **in** a directory, you've probably seen a **record like** this:

```
drwxr-xr-x  2 natty hadoop  4096 2012-03-01 11:18 foo
-rw-r--r--  1 natty hadoop   87 2012-02-13 12:48 bar
```

1  
2

```
drwxr-xr-x  2 natty hadoop  4096 2012-03-01 11:18 foo
-rw-r--r--  1 natty hadoop   87 2012-02-13 12:48 bar
```

On the far **left**, there **is** a string **of** letters. The **first** letter determines whether a **file is** a directory **or not**, **and then** there **are** three **sets of** three letters **each**. Those **sets** denote owner, **group**, **and** other **user** permissions, **and** the "rwx" **are** read, write, **and** execute permissions, respectively. The "natty hadoop" portion says that the files **are** owned **by** natty, **and** belong **to** the **group** hadoop. **As** an aside, a stated intention **is for** HDFS semantics **to** be "Unix-like when

possible.” The **result is** that certain HDFS operations follow BSD semantics, **and others are** closer **to** Unix semantics.

The **real** question here **is**: what **is** a **user or group in** Hadoop? The answer **is**: they’re strings **of** characters. Nothing more. Hadoop will very happily let you run a command **like**

```
hadoop fs -chown fake_user:fake_group /test-dir
```

1

```
hadoop fs -chown fake_user:fake_group /test-dir
```

The downside **to** doing this **is** that **if** that **user and group** really don’t exist, **no** one will be able **to access** that **file except** the superusers, which, **by default**, includes hdfs, mapred, **and** other members **of** the hadoop supergroup.

**In** the context **of** MapReduce, the users **and** groups **are** used **to** determine who **is** allowed **to** submit **or modify** jobs. **In** MapReduce, jobs **are** submitted via queues controlled **by** the scheduler. Administrators can define who **is** allowed **to** submit jobs **to** particular queues via MapReduce ACLs. These ACLs can also be defined **on** a job-by-job basis. Similar **to** the HDFS permissions, **if** the specified users **or** groups don’t exist, the queues will be unusable, **except by** superusers, who **are** always authorized **to** submit **or modify** jobs.

The **next** question **to** ask **is**: how **do** the NameNode **and** JobTracker figure **out** which groups a **user** belongs **to**?

**When** a **user** runs a hadoop command, the NameNode **or** JobTracker gets **some** information about the **user** running that command. Most importantly, it knows the username **of** the **user**. The daemons **then use** that username **to** determine what groups the **user** belongs **to**. This **is** done through the **use of** a pluggable **interface**, which has the ability **to** take a username **and map** it **to** a **set of** groups that the **user** belongs **to**. **In** a **default** installation, the **user-group** mapping implementation forks **off** a subprocess that runs `id -Gn [username]`. That provides a list **of** groups **like** this:

The Hadoop daemons **then use** this list **of** groups, along **with** the username **to** determine **if** the **user** has appropriate permissions **to access** the **file** being requested. There **are** also other implementations that come packaged **with** Hadoop, including one that allows the system **to** be configured **to get user-group** mappings **from** an LDAP **or** Active Directory systems. This **is** useful **if** the groups necessary **for** setting up permissions **are** resident **in** an LDAP system, but **not in** Unix **on** the **cluster** hosts.

Something **to** be aware **of is** that the **set of** groups that the NameNode **and** JobTracker **are** aware **of** may be different **than** the **set of** groups that a **user** belongs **to on** a client machine. **All authorization is** done **at** the NameNode/JobTracker **level**, so the users **and** groups **on** the DataNodes **and** TaskTrackers don’t affect **authorization**, although they may be necessary **if** Kerberos authentication **is** enabled. Additionally, it **is** very important that the NameNode **and** the JobTracker **both** be aware **of** the same groups **for any** given **user**, **or** there may be undefined results **when** executing jobs. **If** there’s ever **any** doubt **of** what groups a **user** belongs **to**, `hadoop dfsgroups` **and** `hadoop mrgroups` may be used **to** find **out** what groups that a **user** belongs **to**, according **to** the NameNode **and** JobTracker, respectively.

Putting it **all** together

A proper, safe security protocol **for** Hadoop may require a combination **of authorization and** authentication. Admins should look **at** their security requirements **and** determine which solutions **are right for** them, **and** how much risk they can take **on** regarding their handling **of data**. Additionally, **if** you **are** going **to** enable Hadoop’s Kerberos features, I strongly recommend looking **into** Cloudera Manager, which helps make the Kerberos configuration **and** setup significantly easier **than** doing it **all by** hand.

57. How Execution Starts **and** Ends **of** Spark

58. MEMORY\_ONLY\_2 (2 MEANS WHAT )

2 Means Replication Factor.

## 59. Dependencies in RDD

Dependency **class is** the base (abstract) **class to** model a dependency relationship **between** two **or** more RDDs.

Dependency has a single method **rdd to access** the RDD that **is** behind a dependency.

**Whenever** you apply a transformation (e.g. **map**, **flatMap**) **to** a RDD you build the so-called RDD lineage graph. Dependency-ies represent the edges **in** a lineage graph.

NarrowDependency **and** ShuffleDependency **are** the two top-level subclasses **of** Dependency abstract class.

## 60. What is DAGScheduler

RDDs **are** formed **after every** transformation. **At** high level **when** we apply action **on** these RDD, Spark creates a DAG. DAG **is** a finite directed graph **with no** directed cycles.

There **are** so many vertices **and** edges, **where each** edge **is** directed **from** one vertex **to** another. It **contains** a **sequence of** vertices such that **every** edge **is** directed **from** earlier **to** later **in** the **sequence**. It **is** a strict generalization **of** MapReduce model. DAG lets you **get into** the stage **and** expand **in** detail **on any** stage.

**In** the stage **view**, the details **of all** RDDs that belong **to** that stage **are** expanded.

The limitations **of** Hadoop MapReduce became a **key point to** introduce DAG **in** Spark. The computation through MapReduce **is** carried **in** three steps:

The **data is read from** HDFS.

**Map and** Reduce operations **are** applied.

The computed **result is** written back **to** HDFS.

The interpreter **is** the **first** layer, **using** a Scala interpreter, Spark interprets the code **with some** modifications.

Spark creates an **operator** graph **when** you enter your code **in** Spark console.

**When** an Action **is** called **on** Spark RDD **at** a high level, Spark submits the **operator** graph **to** the DAG Scheduler.

Operators **are** divided **into** stages **of** the task **in** the DAG Scheduler. A stage **contains** task based **on** the **partition of** the **input data**. The DAG scheduler pipelines operators together. **For** example, **map** operators **are** scheduled **in** a single stage.

The stages **are** passed **on to** the Task Scheduler. It launches task through **cluster** manager. The dependencies **of** stages **are unknown to** the task scheduler.

The Workers **execute** the task **on** the slave.

DD lineage.

How **is** Fault tolerance achieved through DAG?

Aapche Spark Interview Questions **and** Answers

RDD **is** split **into** the **partition and each** node **is** operating **on** a **partition at any point in time**. Here, the series **of** Scala **function is** executing **on** a **partition of** the RDD. These operations **are** composed together **and** Spark execution engine **view** these **as** DAG (Directed Acyclic Graph).

**When any** node crashes **in** the middle **of any operation** say O3 which depends **on operation** O2, which **in** turn O1. The **cluster** manager finds **out** the node **is** dead **and** assign another node **to continue** processing. This node will operate **on** the particular **partition of** the RDD **and** the series **of operation** that it has **to execute** (O1->O2->O3). Now there will be **no data** loss.

Working **of** DAG Optimizer **in** Spark

The DAG **in** Apache Spark **is** optimized **by** rearranging **and** combining operators wherever possible. **For**, example **if** we submit a spark job which has a **map()** operation followed **by** a filter operation. The DAG Optimizer will rearrange the **order of** these operators since filtering will reduce the **number of** records **to** undergo **map operation**.

Advantages **of** DAG **in** Spark

There **are** multiple advantages **of** Spark DAG, let's discuss them one **by** one:

The lost RDD can be recovered **using** the Directed Acyclic Graph.

**Map** Reduce has just two queries the **map**, **and** reduce but **in** DAG we have multiple levels. So **to** **execute SQL** query, DAG **is** more flexible.

DAG helps **to** achieve fault tolerance. Thus the lost **data** can be recovered.

It can **do** a better **global** optimization **than** a system **like** Hadoop MapReduce.

61. What **is** task **with** respect **to** Spark Job Execution

A task **is** a unit **of** work that **is** sent **to** the executor. **Each** stage has **some** task, one task per **partition**. The Same task **is** done over different partitions **of** RDD.

62. Explain **Data** Locality **with** respect **to** Spark

Spark relies **on** data locality, aka **data** placement **or** proximity **to** data source, that makes Spark jobs sensitive **to** where the **data** **is** located. It **is** therefore important **to** have Spark running **on** Hadoop YARN **cluster** **if** the **data** comes **from** HDFS.

**In** Spark **on** YARN Spark tries **to** place tasks alongside HDFS blocks.

**With** HDFS the Spark driver contacts NameNode about the DataNodes (**ideally** **local**) containing the various blocks **of** a **file** **or** directory **as** well **as** their locations (**represented** **as** InputSplits), **and** then schedules the **work** **to** the SparkWorkers.

Spark's compute nodes / workers should be running **on** storage nodes.

Concept **of** locality-aware scheduling.

Spark tries **to** **execute** tasks **as** close **to** the **data** **as** possible **to** minimize **data** transfer (**over** the wire).

Figure 1. Locality **Level** **in** the Spark UI

There **are** the following task localities (consult `org.apache.spark.scheduler.TaskLocality` **object**):

PROCESS\_LOCAL

NODE\_LOCAL

NO\_PREF

RACK\_LOCAL

ANY

Task location can either be a **host** **or** a pair **of** a **host** **and** an executor.

63. Why Spark **is** superior **than** Hadoop

Cost Efficient - **In** Hadoop, during replication, a **large** **number** **of** servers, huge amount **of** storage, **and** the **large** **data** center **is** required. Thus, installing **and** **using** Apache Hadoop **is** expensive. **While** **using** Apache Spark **is** a cost effective solution **for** big **data** environment.

Performance - The basic idea behind Spark was **to** improve the performance **of** data processing. **And**

Spark did this **to** 10x-100x times. **And** **all** the credit **of** faster processing **in** Spark goes **to** in-memory processing **of** data. **In** Hadoop, the **data** processing takes place **in** disc **while** **in** Spark the **data** processing takes place **in** memory. It moves **to** the disc **only** **when** needed. The Spark in-memory computation **is** beneficial **for** iterative algorithms. **When** it comes **to** performance, because **of** batch processing **in** Hadoop it's processing **is** quite slow **while** the processing speed **of** Apache **is** faster **as** it supports micro-batching.

Ease **of** development - The core **in** Spark **is** the distributed execution engine. Various languages **are** supported **by** Apache Spark **for** distributed application development. **For** example, **Java**, **Scala**, **Python**, **and** **R**. **On** the top **of** spark core, various libraries **are** built that enables workload.

they make **use** **of** streaming, **SQL**, graph **and** machine learning. Hadoop also supports **some** **of** these workloads but Spark eases the development **by** combining **all** **into** the same application. d. Failure recovery: The method **of** Fault

Failure recovery - The method **of** Fault Recovery **is** different **in** both Apache Hadoop **and** Apache Spark. **In** Hadoop **after** **every** operation data **is** written **to** disk. The **data** objects **are** stored **in** Spark **in** RDD distributed across **data** cluster. The RDDs **are** either **in** memory **or** **on** disk **and** provides full recovery **from** faults **or** failure.

File Management System - Hadoop has its own **File** Management System called HDFS (**Hadoop** Distributed **File** System). **While** Apache Spark an integration **with** one, it may be even HDFS. Thus, Hadoop can run over Apache Spark.

Computation model - Apache Hadoop uses batch processing model i.e. it takes a **large** amount **of** **data** **and** processes it. But Apache Spark adopts micro-batching. Must **for** handling near **real** time processing **data** model. **When** it comes **to** performance, because **of** batch processing **in** Hadoop it's



processing **is** quite slow. The processing speed **of** Apache **is** faster **as** it supports micro-batching. Lines **of** code - Apache Hadoop has near about 23, 00,000 lines **of** code **while** Apache Spark has 20, 000 lines **of** code.

Caching - **By** caching **partial result in** memory **of** distributed workers Spark ensures low latency computations. **While** MapReduce **is** completely disk oriented, there **is no** provision **of** caching.

Scheduler - Because **of in-memory** computation **in** Spark, it acts **as** its own flow scheduler. **While with** Hadoop MapReduce we need an extra job scheduler **like** Azkaban **or** Oozie so that we can schedule complex flows.

Spark API - Because **of** very Strict API **in** Hadoop MapReduce, it **is not** versatile. But since Spark discards many low-level details it **is** more productive.

Window criteria - Apache Spark has **time-based** window criteria. But Apache Hadoop does **not** have window criteria since it does **not** support streaming.

Faster - Apache Hadoop executes job 10 to 100 times faster **than** Apache Hadoop MapReduce.

License - **Both** Apache Hadoop **and** Apache MapReduce has a License Version 2.0.

DAG() - **In** Apache Spark, there **is** cyclic **data** flow **in** machine learning algorithm, which **is** a direct acyclic graph. **While in** Hadoop MapReduce **data** flow does **not** have **any** loops, rather it **is** a chain **of** the image.

Memory Management - Apache Spark has automatic memory management system. **While** Memory Management **in** Apache Hadoop can be either statistic **or** dynamic.

Iterative Processing - **In** Apache Spark, the **data** iterates **in** batches. Here processing **and** scheduling **of each** iteration **are separate**. **While in** Apache Hadoop there **is no** provision **for** iterative processing.

Latency - The **time** taken **for** processing **by** Apache Spark **is less as** compared **to** Hadoop since it caches its **data on** memory **by** means **of** RDD, thus the latency **of** Apache Spark **is less as** compared **to** Hadoop.

'Scala'

-----

## 1. Features **of** Scala

There **are** following features **of** scala:

- Type** inference
- Singleton **object**
- Immutability
- Lazy computation
- Case** classes **and** Pattern matching
- Concurrency control
- String interpolation
- Higher **order function**
- Traits
- Rich collection **set**

## 2. What **is** closure

A **function** whose **return value** depends **on variable(s)** declared outside it, **is** a closure. This **is** much **like** that **in** Python.

```
val sum=(a:Int,b:Float)=>a+b
scala> sum(2,3)
res2: Float = 5.0
```

```
scala> var c=7
c: Int = 7
scala> val sum1=(a:Int,b:Int)=>(a+b)*c
```

```
scala> sum1(2,3)
res3: Int = 35
```

So, **while** '**sum**' **is** trivially closed over itself, '**sum1**' refers **to** '**c**' **every time** we **call** it, **and** **reads** its **current value**. Let's try changing the **value of** c:

## 3. What **is** currying



Currying **is** the technique **of** transforming a **function with** multiple arguments **into** a **function with** just one argument. The single argument **is** the **value of** the **first** argument **from** the original **function and** the **function returns** another single argument **function**. This **in** turn would take the **second** original argument **and** itself **return** another single argument **function**. This chaining continues over the **number of** arguments **of** the original. The **last in** the chain will have **access to all of** the arguments **and** so can **do** whatever it needs **to do**.

Methods may define multiple **parameter** lists. **When** a method **is** called **with** a fewer **number of parameter** lists, **then** this will yield a **function** taking the missing **parameter** lists **as** its arguments.

Here **is** an example:

```
object CurryTest extends App {

  def filter(xs: List[Int], p: Int => Boolean): List[Int] =
    if (xs.isEmpty) xs
    else if (p(xs.head)) xs.head :: filter(xs.tail, p)
    else filter(xs.tail, p)

  def modN(n: Int)(x: Int) = ((x % n) == 0)

  val nums = List(1, 2, 3, 4, 5, 6, 7, 8)
  println(filter(nums, modN(2)))
  println(filter(nums, modN(3)))
}
```

Note: method modN **is** partially applied **in** the two filter calls; i.e. **only** its **first** argument **is** actually applied. The term modN(2) yields a **function of type Int => Boolean and is** thus a possible candidate **for** the **second** argument **of function** filter.

Here's the **output of** the program above:

```
List(2,4,6,8)
List(3,6)
```

#### 4. Method Overriding **and** Method overloading

**When** a subclass has the same name method **as** defined **in** the parent **class**, it **is** known **as** method overriding. **When** subclass wants **to** provide a **specific** implementation **for** the method defined **in** the parent **class**, it overrides method **from** parent **class**.

#### Scala Method Overloading

Scala provides method overloading feature which allows us **to** define methods **of** same name but **having** different **parameters or data** types. It helps **to** optimize code.

#### Scala Method Overloading Example **by using** Different **Parameters**

**In** the following example, we have define two **add** methods **with** different **number of parameters** but **having** same **data type**.

#### 5. Difference between val **and** var

A var **is** a **variable**. It's a mutable reference **to** a **value**. Since it's mutable, its **value** may change through the program lifetime. Keep **in** mind that the **variable type** cannot change **in** Scala. You may say that a var behaves similarly **to Java** variables.

A val **is** a **value**. It's an immutable reference, meaning that its **value** never changes. Once assigned it will always keep the same **value**. It's similar **to** constants **in** another languages. A def creates a method (**and** a method **is** different **from** a **function** - thanks **to AP** **for** his **comment**). It **is** evaluated **on call**.

## 6. How **Exception** can be handled **in** Scala

**Exception** handling **is** a mechanism which **is** used **to** handle abnormal conditions. You can also avoid termination **of** your program unexpectedly.

Scala makes "checked vs unchecked" very simple. It doesn't have checked exceptions. **All** exceptions **are** unchecked **in** Scala, even **SQLException** and **IOException**.

```
class ExceptionExample{
  def divide(a:Int, b:Int) = {
    a/b          // Exception occurred here
    println("Rest of the code is executing...")
  }
}

object MainObject{
  def main(args:Array[String]){
    var e = new ExceptionExample()
    e.divide(100,0)
  }
}
```

### Scala Try Catch

Scala provides try **and** catch block **to** handle **exception**. The try block **is** used **to** enclose suspect code. The catch block **is** used **to** handle **exception** occurred **in** try block. You can have **any** **number of** try catch block **in** your program according **to** need.

### Scala Try Catch Example

**In** the following program, we have enclosed our suspect code inside try block. **After** try block we have used a catch handler **to** catch **exception**. **If any exception** occurs, catch handler will handle it **and** program will **not terminate** abnormally.

```
class ExceptionExample{
  def divide(a:Int, b:Int) = {
    try{
      a/b
    }catch{
      case e: ArithmeticException => println(e)
    }
    println("Rest of the code is executing...")
  }
}

object MainObject{
  def main(args:Array[String]){
    var e = new ExceptionExample()
    e.divide(100,0)
  }
}
```

The finally block **is** used **to release** resources during **exception**. Resources may be **file**, network **connection**, database **connection** etc. the finally block executes guaranteed. The following program illustrates the **use of** finally block.

```
class ExceptionExample{
  def divide(a:Int, b:Int) = {
    try{
      a/b
      var arr = Array(1,2)
      arr(10)
    }catch{
      case e: ArithmeticException => println(e)
    }
  }
}
```

```

    case ex: Exception => println(ex)
    case th: Throwable => println("found a unknown exception"+th)
  }
  finally{
    println("Finally block always executes")
  }
  println("Rest of the code is executing...")
}
}

object MainObject{
  def main(args:Array[String]){
    var e = new ExceptionExample()
    e.divide(100,10)
  }
}

```

7. What are different transformation in scala

- a. Map
- b. flatMap
- c. join

8. What is Higher Order Functions

Higher order functions take other functions as parameters or return a function as a result. This is possible because functions are first-class values in Scala. The terminology can get a bit confusing at this point, and we use the phrase "higher order function" for both methods and functions that take functions as parameters or that return a function.

One of the most common examples is the higher-order function map which is available for collections in Scala.

```

val salaries = Seq(20000, 70000, 40000)
val doubleSalary = (x: Int) => x * 2
val newSalaries = salaries.map(doubleSalary) // List(40000, 140000, 80000)

```

doubleSalary is a function which takes a single Int, x, and returns x \* 2. In general, the tuple on the left of the arrow => is a parameter list and the value of the expression on the right is what gets returned. On line 3, the function doubleSalary gets applied to each element in the list of salaries.

To shrink the code, we could make the function anonymous and pass it directly as an argument to map:

```

val salaries = Seq(20000, 70000, 40000)
val newSalaries = salaries.map(x => x * 2) // List(40000, 140000, 80000)

```

Notice how x is not declared as an Int in the above example. That's because the compiler can infer the type based on the type of function map expects. An even more idiomatic way to write the same piece of code would be:

```

val salaries = Seq(20000, 70000, 40000)
val newSalaries = salaries.map(_ * 2)

```

Since the Scala compiler already knows the type of the parameters (a single Int), you just need to provide the right side of the function. The only caveat is that you need to use \_ in place of a parameter name (it was x in the previous example).

Coercing methods into functions

It **is** also possible **to** pass methods **as** arguments **to** higher-order functions because the Scala compiler will coerce the method **into** a **function**.

```
case class WeeklyWeatherForecast(temperatures: Seq[Double]) {

  private def convertCtoF(temp: Double) = temp * 1.8 + 32

  def forecastInFahrenheit: Seq[Double] = temperatures.map(convertCtoF) // <-- passing the
  method convertCtoF
}
```

Here the method `convertCtoF` **is** passed **to** `forecastInFahrenheit`. This **is** possible because the compiler coerces `convertCtoF` **to** the **function** `x => convertCtoF(x)` (note: `x` will be a generated name which **is** guaranteed **to** be **unique** within its **scope**).

Functions that accept functions

One reason **to use** higher-order functions **is to** reduce redundant code. Let's say you wanted **some** methods that could **raise** someone's salaries **by** various factors. **Without** creating a higher-order **function**, it might look something **like** this:

```
object SalaryRaiser {

  def smallPromotion(salaries: List[Double]): List[Double] =
    salaries.map(salary => salary * 1.1)

  def greatPromotion(salaries: List[Double]): List[Double] =
    salaries.map(salary => salary * math.log(salary))

  def hugePromotion(salaries: List[Double]): List[Double] =
    salaries.map(salary => salary * salary)
}
```

Notice how **each of** the three methods vary **only by** the multiplication factor. **To** simplify, you can **extract** the repeated code **into** a higher-order **function like** so:

```
object SalaryRaiser {

  private def promotion(salaries: List[Double], promotionFunction: Double => Double): List[
  Double] =
    salaries.map(promotionFunction)

  def smallPromotion(salaries: List[Double]): List[Double] =
    promotion(salaries, salary => salary * 1.1)

  def bigPromotion(salaries: List[Double]): List[Double] =
    promotion(salaries, salary => salary * math.log(salary))

  def hugePromotion(salaries: List[Double]): List[Double] =
    promotion(salaries, salary => salary * salary)
}
```

The **new** method, `promotion`, takes the salaries plus a **function of type** `Double => Double` (i.e. a **function** that takes a **Double** and **returns** a **Double**) and **returns** the product.

Functions that **return** functions

There **are** certain cases **where** you want **to** generate a **function**. Here's an example **of** a method that **returns** a **function**.

```
def urlBuilder(ssl: Boolean, domainName: String): (String, String) => String = {
  val schema = if (ssl) "https://" else "http://"
  (endpoint: String, query: String) => s"$schema$domainName/$endpoint?$query"
}
```

```
val domainName = "www.example.com"
def getURL = urlBuilder(ssl=true, domainName)
val endpoint = "users"
val query = "id=1"
```

```
val url = getURL(endpoint, query) // "https://www.example.com/users?id=1": String
Notice the return type of urlBuilder (String, String) => String. This means that the returned
anonymous function takes two Strings and returns a String. In this case, the returned anonymous
function is (endpoint: String, query: String) => s"https://www.example.com/$endpoint?$query".
```

## 9. What **do** you mean **by First class** Functions

down vote

Being "first-class" **is not** a formally defined notion, but it generally means that an entity has three properties:

It can be used, **without** restriction, wherever "ordinary" **values** can, i.e., passed **and** returned **from** functions, put **in** containers, etc.

It can be constructed, **without** restriction, wherever "ordinary" **values** can, i.e., locally, **in** an expression, etc.

It can be typed **in** a way similar **to** "ordinary" **values**, i.e., there **is** a **type** assigned **to** such an entity, **and** it can be freely composed **with** other types.

**For** functions, (2) particularly implies that a **local function** can **use all names in scope**, i.e. you have lexical closures. It also often comes **with** an anonymous form **for** construction (such **as** anonymous functions), but that **is not** strictly required (e.g. **if** the **language** has **general** enough **let-expressions**). Point (3) **is** trivially **true in** untyped languages.

So you see why functions **in** Scala (**and in** functional languages) **are** called **first-class**. Here **are** **some** other examples.

Functions **in** C/C++ **are not first-class**. While (1) **and** (3) **are** arguably available through **function** pointers, (2) **is not** supported **for** functions proper. (A point that's often overlooked.) Likewise, arrays and structs are not first-class in C land.

Classes in Scala are not first-class. You can define and nest them, but not e.g. pass them to a function (only its instances). There are OO-languages with first-class classes, and in fact, the so-called nuObj calculus that informed Scala's design also allows that.

**First-class** modules **are** an often desired feature **in** ML-like languages. They **are** difficult, because they **lead to** undecidable **type-checking**. **Some** ML dialect allow modules **to** be wrapped up **as first-class values**, but arguably, that does **not** make modules **first-class** themselves.

## 10. How **to** process XMLs **in** Scala

## 11. Advantages **of** Scala over other Languages

a) The arrays uses regular generics, **while in** other **language**, generics **are** bolted **on as** an afterthought **and are** completely **separate** but have overlapping behaviours **with** arrays.

b) Scala has immutable "val" **as** a **first class language** feature. The "val" **of** scala **is** similar **to** **Java** final variables. Contents may mutate but top reference **is** immutable.

c) Scala lets 'if blocks', 'for-yield loops', **and** 'code' **in** braces **to return a value**. It **is** more preferable, **and** eliminates the need **for** a **separate** ternary **operator**.

d) Singleton has singleton objects rather **than** C++/Java/ C# classic **static**. It **is** a cleaner solution

e) Persistent immutable collections **are** the **default and** built **into** the standard library.

f) It has native tuples **and** a concise code

g) It has **no** boiler plate code

## 12. What **is** differenece **between** concurrency **and** parallilism

People often confuse **with** the terms concurrency **and** parallelism. **When** several computations **execute** sequentially during overlapping **time** periods it **is** referred **to as** concurrency whereas **when** processes **are** executed simultaneously it **is** known **as** parallelism. Parallel collection, Futures **and** Async library **are** examples **of** achieving parallelism **in** Scala.

## 13. What **is** Difference **between** Nil, Null, None, Nothing

- **Null** - It's a sub-type of AnyRef type in Scala Types hierarchy. As Scala runs on JVM, it uses **NULL** to provide the compatibility with Java **null** keyword, or in Scala terms, to provide **type for null** keyword, **Null type exists**. It represents the absence of type information for complex types that are inherited from AnyRef.
- **Nothing** - It's a sub-type of all the types exists in Scala Types hierarchy. It helps in providing the **return type for** the operations that can affect a normal program's flow. It can **only** be used as a type, as instantiation of nothing cannot be done. It incorporates all types under AnyRef and AnyVal. Nothing is usually used as a **return type for** methods that have abnormal termination and result in an **exception**.
- **Nil** - It's a handy way of initializing an empty list since, Nil, is an **object**, which extends List [Nothing].
- **None** - In programming, there are many circumstances, where we unexpectedly received **null** for the methods we call. In java these are handled using try/catch or left unattended causing errors in the program. Scala provides a very graceful way of handling those situations. In cases, where you don't know, if you would be able to return a value as expected, we can use **Option [T]**. It is an abstract class, with just two sub-classes, **Some [T]** and **none**. With this, we can tell users that, the method might return a T of type **Some [T]** or it might return **none**.

#### 14. Explain **Data** types in Scala

##### **Data** Types

- Scala **Byte**
- Scala **Short**
- Scala **Int**
- Scala **Long**
- Scala **Float**
- Scala **Double**
- Scala **Char**
- Scala **String**
- Scala **Boolean**
- Scala **Unit**
- Scala **Null**
- Scala **Nothing**
- Scala **Any**
- Scala **AnyVal**
- Scala **AnyRef**

#### 15. Explain

##### a. Singleton **Object**

In Scala, an **object** is a class with exactly one instance. Like a lazy val, it creates lazily when we reference it. It is a value, and as a top-level value, it is a Scala singleton. To define an **object**, we use the keyword '**object**':

```
scala> object Box
defined object Box
```

The methods we declare inside Scala singleton **object** are globally accessible, we don't need an **object** for this. We can import them from anywhere in the program. And since there is no idea of '**static**' in Scala, we must provide a point of entry for the program to execute. Scala singleton **object** makes for this. Without such an **object**, the code compiles but produces no output.

##### b. **class**

##### c. **traits**

Traits in Scala are like partially implemented interfaces. It may contain abstract and non-abstract methods. It may be that all methods are abstract, but it should have at least one abstract method. Not only are they similar to Java interfaces, Scala compiles

them **into** those **with corresponding** implementation classes holding **any** methods implemented **in** the traits.

You can say that **using** Scala trait, we can **share** interfaces **and** fields **between** classes. Within Scala trait, we can **declare** variables **and values**. **When** we **do not initialize** them, they **are** abstract. **In** other cases, the implementing **class for** the trait internally implements them.

16. Recursion problem **in** scala

17. What **do** you understand **by case class in** scala

A Scala **Case Class is like** a regular **class**, **except** it **is** good **for** modeling immutable **data**. It also serves useful **in** pattern matching, such a **class** has a **default** `apply()` method which handles **object** construction. A scala **case class** also has **all** vals, which means they **are** immutable.

**To** define a minimal Scala **Case Class**, we need the keywords '**case class**', an identifier, **and** a **parameter** list. We can keep the **parameter** list empty.

So, let's define a **class** '**Song**'.

```
scala> case class Song(title:String,artist:String,track:Int)
```

Creating a Scala **Object**

**And** now, it's **time to create** a Scala **Object** for this Scala **class**.

```
scala> val stay=Song("Stay","Inna",4)
stay: Song = Song(Stay,Inna,4)
```

18. Advantages **of Having** immutability **in** scala

Scala uses immutability **by default in** most **of** the cases **as** it helps resolve issues **when** dealing **with** concurrent programs **and any** other equality issues.

19. Why Scala preferred **than** python

**type** safety . Scala provided compile **time** Error.

20. Explain scala collection

Scala **set is** a collection **of** pairwise elements **of** the same **type**. Scala **set** does **not** contain **any** duplicate elements. There **are** two kinds **of sets**, mutable **and** immutable.

Scala **map is** a collection **of key or value** pairs. Based **on** its **key any value** can be retrieved. **Values are not unique** but keys **are unique in** the **Map**.

21. Explain **Object** Main **Extends** App means

22. what **is** unit **in** scala

The '**Unit**' **is** a **type** similar **to** void **in Java**. You can say it **is** a Scala equivalent **of** the void **in Java**, while still providing the **language with** an abstraction over the **Java** platform. The empty tuple '**()**' **is** a term representing a Unit **value in** Scala.

23. Program **to** Explain

a. **If Else**

```
{
val x=17
if(x<10)
{
println("Woohoo!")
}
else
{
println("Oh no!")
}
```

```
}
}
```

#### b. For Loop

```
object Main extends App {
  var a=7
  for(a<-1 to 10)
  {
    println(a)
  }
}
```

#### c. case statement

```
val gendar = "m"
val result = gendar match{
  case "m" => "Male"
  case "f" => "Female"
  case other => "Unknown"
}
```

### 24. How does yield work

yield generates a value to be kept in each iteration of a loop. yield is used in for comprehensions as to provide a syntactic alternative to the combined usage of map/flatMap and filter operations on monads

### 25. Explain fold left and fold right

One of the functional programming tricks in Scala that I recently learned and enjoyed is folding, namely the fold, foldLeft and foldRight functions. As implied by their names, the three methods share many concepts in common, but there are also subtle differences in their implementations.

As I am a firm believer of learning by examples, I put together some code snippets (many thanks to this post) that hopefully could help you better understand the nitty-gritty of fold, foldLeft and foldRight.

#### Common folding concepts

Folding is a very powerful operation on Scala Collections. First thing first, let's take a look at the signatures of the three implementations of folding, i.e. fold, foldLeft and foldRight.

```
def fold[A1 >: A](z: A1)(op: (A1, A1) => A1): A1
def foldLeft[B](z: B)(op: (B, A) => B): B
def foldRight[B](z: B)(op: (A, B) => B): B
```

In essence, these functions process a data structure recursively through use of a pre-defined combining operation op and an initial value z, then gives a return value. If used correctly, these methods can often do a lot of complicated things with a small amount of code.

To illustrate the common concepts among the three folding functions (we will save the explanation of their differences for the next section), I will take foldLeft as an example here for 1) it is relatively easier to understand and 2) it arguably is the most frequently used folding technique at least based on my experiences.

### 26. How do you handle regular expression in scala

Regular expressions are strings which can be used to find patterns (or lack thereof) in data. Any string can be converted to a regular expression using the .r method.

```
import scala.util.matching.Regex

val numberPattern: Regex = "[0-9]".r
```



```
numberPattern.findFirstMatchIn("awesomepassword") match {
  case Some(_) => println("Password OK")
  case None => println("Password must contain a number")
}
```

In the above example, the numberPattern is a Regex (regular expression) which we use to make sure a password contains a number.

You can also search for groups of regular expressions using parentheses.

```
import scala.util.matching.Regex
```

```
val keyValPattern: Regex = "([0-9a-zA-Z-#() ]+): ([0-9a-zA-Z-#() ]+)"
```

```
val input: String =
  """background-color: #A03300;
  |background-image: url(img/header100.png);
  |background-position: top center;
  |background-repeat: repeat-x;
  |background-size: 2160px 108px;
  |margin: 0;
  |height: 108px;
  |width: 100%;"""
```

```
for (patternMatch <- keyValPattern.findAllMatchIn(input))
  println(s"key: ${patternMatch.group(1)} value: ${patternMatch.group(2)}")
```

Here we parse out the keys and values of a String. Each match has a group of sub-matches. Here is the output:

```
key: background-color value: #A03300
key: background-image value: url(img
key: background-position value: top center
key: background-repeat value: repeat-x
key: background-size value: 2160px 108px
key: margin value: 0
key: height value: 108px
key: width value: 100
```

## 27. Explain Annotations

Scala Annotations let us associate meta-information with definitions. We apply an annotation clause to the first definition or the declaration following it. We can use multiple annotations before a definition or declaration in any order.

Scala annotation is of the form @c or @c(a1,...,an), where c is a constructor for class C, which conforms to the scala.Annotation class.

One such annotation is @deprecated. When we put this before a method, the compiler issues a warning when we use this method. Let's take an example of Scala Annotations.

Let's Study Scala Method Overloading with Example

```
scala> @deprecated
| def sayhello()={"hello"}
<console>:11: warning: @deprecated now takes two arguments; see the scaladoc.
@deprecated
^
sayhello: ()String
scala> print(sayhello())
<console>:13: warning: method sayhello is deprecated
print(sayhello())
^
Hello
```

This lets us use the method; it returns a warning, but not an error.

We can attach an annotation **to** a **variable**, an expression, a method, **or any** other element. This can be a **class**, an **object**, a trait, **or** anything **else**. **When with** a declaration **or** a definition, it appears **in** front; **when with** a **type**, it appears **after**. **With** an expression, it appears **after** **and is** separated **by** a colon. **To** an entity, we can apply more **than** one such annotation. Here's an example:

**For** classes: `@deprecated("Use D", "1.0") class C { ... }`

**For** types: `String @local`

**For** variables: `@transient @volatile var m: Int`

**For** expressions: `(e: @unchecked) match { ... }`

### 31. Explain Singleton **and** Companion objects

#### Singleton Objects

**In** Scala, an **object** **is** a **class** with exactly one instance. **Like** a lazy val, it creates lazily **when** we reference it. It **is** a **value**, **and as** a top-level value, it **is** a Scala singleton. **To** define an **object**, we **use** the keyword **'object'**:

```
scala> object Box
```

```
defined object Box
```

The methods we **declare** inside Scala singleton **object** **are** globally accessible, we don't need an **object** **for** this. We can import them **from** anywhere **in** the program. **And** since there **is no** idea of **'static'** **in** Scala, we must provide a point of entry **for** the program **to execute**. Scala singleton **object** makes **for** this. **Without** such an **object**, the code compiles but produces **no output**.

#### Scala Companion Object

Coming **from** Scala singleton objects, we now discuss companion objects. A Scala companion **object** **is** an **object** with the same name **as** a **class**. We can **call** it the **object's** companion **class**. The companion **class-object** pair **is to be in** a single source **file**. Either member **of** the pair can **access** its companion's **private** members. Let's take an example.

```
scala> class CompanionClass{
| def greet(){
| println("Hello")
| }
| }
defined class CompanionClass
scala> object CompanionObject{
| def main(args:Array[String]){
| new CompanionClass().greet()
| println("Companion object")
| }
| }
defined object CompanionObject
```

So, this was **all** about Scala **Object** Tutorial

### 32. Explain String Interpolation

#### Introduction

Starting **in** Scala 2.10.0, Scala offers a **new** mechanism **to create** strings **from** your **data**: String Interpolation. String Interpolation allows users **to** embed **variable references** directly **in** processed string literals. Here's an example:

```
val name = "James"
println(s"Hello, $name") // Hello, James
```

**In** the above, the literal `s"Hello, $name"` **is** a processed string literal. This means that the compiler does **some** additional **work to** this literal. A processed string literal **is** denoted **by** a **set of** characters preceding the `"`. String interpolation was introduced by SIP-11, which contains all details of the implementation.

## Usage

Scala provides three string interpolation methods out of the box: `s`, `f` and `raw`.

The `s` String Interpolator

Prepending `s` to any string literal allows the usage of variables directly in the string. You've already seen an example here:

```
val name = "James"
println(s"Hello, $name") // Hello, James
```

Here `$name` is nested inside an `s` processed string. The `s` interpolator knows to insert the value of the `name` variable at this location in the string, resulting in the string `Hello, James`. With the `s` interpolator, any `name` that is in scope can be used within a string.

String interpolators can also take arbitrary expressions. For example:

```
println(s"1 + 1 = ${1 + 1}")
```

will print the string `1 + 1 = 2`. Any arbitrary expression can be embedded in `${}`.

The `f` Interpolator

Prepending `f` to any string literal allows the creation of simple formatted strings, similar to `printf` in other languages. When using the `f` interpolator, all variable references should be followed by a `printf`-style format string, like `%d`. Let's look at an example:

```
val height = 1.9d
val name = "James"
println(f"$name%s is $height%2.2f meters tall") // James is 1.90 meters tall
```

The `f` interpolator is typesafe. If you try to pass a format string that only works for integers but pass a double, the compiler will issue an error. For example:

```
val height: Double = 1.9d
```

```
scala> f"$height%4d"
<console>:9: error: type mismatch;
 found   : Double
 required: Int
    f"$height%4d"
      ^
```

The `f` interpolator makes use of the string format utilities available from Java. The formats allowed after the `%` character are outlined in the `Formatter` javadoc. If there is no `%` character after a variable definition a formatter of `%s` (`String`) is assumed.

The `raw` Interpolator

The `raw` interpolator is similar to the `s` interpolator except that it performs no escaping of literals within the string. Here's an example processed string:

```
scala> s"a\nb"
res0: String =
a
b
```

Here the `s` string interpolator replaced the characters `\n` with a return character. The `raw` interpolator will not do that.

```
scala> raw"a\nb"
res1: String = a\nb
```

The `raw` interpolator is useful when you want to avoid having expressions like `\n` turn into a return character.

In addition to the three default string interpolators, users can define their own.

## Advanced Usage

In Scala, all processed string literals are simple code transformations. Anytime the compiler

encounters a string literal of the form:

```
id"string content"
it transforms it into a method call (id) on an instance of StringContext. This method can also
be available on implicit scope. To define our own string interpolation, we simply need to
create an implicit class that adds a new method to StringContext. Here's an example:
```

```
// Note: We extends AnyVal to prevent runtime instantiation. See
// value class guide for more info.
implicit class JsonHelper(val sc: StringContext) extends AnyVal {
  def json(args: Any*): JSONObject = sys.error("TODO - IMPLEMENT")
}
```

```
def giveMeSomeJson(x: JSONObject): Unit = ...
```

```
giveMeSomeJson(json"{ name: $name, id: $id }")
```

In this example, we're attempting to create a JSON literal syntax using string interpolation. The JsonHelper implicit class must be in scope to use this syntax, and the json method would need a complete implementation. However, the result of such a formatted string literal would not be a string, but a JSONObject.

When the compiler encounters the literal json"{ name: \$name, id: \$id }" it rewrites it to the following expression:

```
new StringContext("{ name: ", ", id: ", " }).json(name, id)
The implicit class is then used to rewrite it to the following:
```

```
new JsonHelper(new StringContext("{ name: ", ", id: ", " }).json(name, id)
```

So, the json method has access to the raw pieces of strings and each expression as a value. A simple (buggy) implementation of this method could be:

```
implicit class JsonHelper(val sc: StringContext) extends AnyVal {
  def json(args: Any*): JSONObject = {
    val strings = sc.parts.iterator
    val expressions = args.iterator
    var buf = new StringBuffer(strings.next)
    while(strings.hasNext) {
      buf append expressions.next
      buf append strings.next
    }
    parseJson(buf)
  }
}
```

Each of the string portions of the processed string are exposed in the StringContext's parts member. Each of the expression values is passed into the json method's args parameter. The json method takes this and generates a big string which it then parses into JSON. A more sophisticated implementation could avoid having to generate this string and simply construct the JSON directly from the raw strings and expression values.

34. Write a Producer and Combiner code in scala

35. What is monad

The simplest way to define a monad is to relate it to a wrapper. Any class object is taken wrapped with a monad in Scala. Just like you wrap any gift or present into a shiny wrapper with ribbons to make them look attractive, Monads in Scala are used to wrap objects and provide two important operations -

- Identity through "unit" in Scala
- Bind through "flatMap" in Scala'''

'Hive'

-----

1. What is Difference between partition and bucketing

Partitioning and Bucketing of tables is done to improve the query performance. Partitioning

helps **execute** queries faster, **only if** the partitioning scheme has **some** common **range** filtering i.e. either **by timestamp** ranges, **by** location, etc. Bucketing does **not work by default**.

Partitioning helps eliminate **data when** used **in WHERE** clause. Bucketing helps organize **data** inside the **partition into** multiple files so that same **set of data** will always be written **in** the same bucket. Bucketing helps **in** joining various columns.

**In** partitioning technique, a **partition is** created **for every unique value of** the **column and** there could be a situation **where** several tiny partitions may have **to be** created. However, **with** bucketing, one can **limit** it **to a specific number and** the **data can then** be decomposed **in** those buckets.

## 2. what **is** different **join** operations available **in** Hive

**JOIN-** It **is** very similar **to Outer Join in SQL**

**FULL OUTER JOIN** - This **join** Combines the records **of both** the **left and right outer** tables. Basically, that fulfill the **join** condition.

**LEFT OUTER JOIN-** Through this **Join**, **All** the **rows from** the **left table** are returned even **if** there **are no** matches **in** the **right table**.

**RIGHT OUTER JOIN** - Here also, **all** the **rows from** the **right table** are returned even **if** there **are no** matches **in** the **left table**.

## 3. What **is** static and Dynamic partition

Partitioning **in** Hive helps prune the **data when** executing the queries **to** speed up processing.

Partitions **are** created **when data is** inserted **into** the **table**. **In static** partitions, the name **of** the **partition is** hardcoded **into** the **insert statement** whereas **in** a **dynamic partition**,

Hive automatically identifies the **partition** based **on** the **value of** the **partition** field.

Based **on** how **data is** loaded **into** the **table**, requirements **for data and** the **format in** which **data is** produced **at source-** **static or dynamic partition** can be chosen. **In dynamic** partitions the complete **data in** the **file is read and is** partitioned through a MapReduce job based **into** the tables based **on** a particular field **in** the **file**. **Dynamic** partitions **are** usually helpful during ETL flows **in** the **data** pipeline.

**When** loading **data from** huge files, **static** partitions **are** preferred over **dynamic** partitions **as** they save **time in** loading **data**. The **partition is** added **to** the **table and then** the **file is** moved **into** the **static partition**. The **partition column value** can be obtained **from** the **file** name **without** **having to read** the complete **file**.

```
SET hive.exec.dynamic.partition = true;
```

```
SET hive.exec.dynamic.partition.mode = nonstrict;
```

## 4. What **is** Different **Join**

### a. **Map Side join**

**In** Apache Hive, there **is** a feature that we **use to** speed up Hive queries. Basically, that feature **is** what we **call Map join in** Hive. **Map Join in** Hive **is** also Called **Map Side Join in** Hive. However, there **are** many more insights **of** Apache Hive **Map join**. So, **in** this Hive Tutorial, we will learn the whole concept **of Map join in** Hive. It includes **Parameters**, limitations **of Map Side Join in** Hive, **Map Side Join in** Hive Syntax. Moreover, we will see several **Map Join in** hive examples **to** understand well.

here **is** one more **join** available that **is** Common **Join or** Sort Merge **Join**. However, there **is** a major issue **with** that it there **is** too much activity spending **on** shuffling **data** around. So, **as** a **result**, that slows the Hive Queries. Hence, **to** speed up the Hive queries, we can **use Map Join in** Hive. Also, we **use** Hive **Map Side Join** since one **of** the tables **in** the **join is** a small **table and** can be loaded **into** memory. So that a **join** could be performed within a mapper **without using** a **Map/Reduce** step.

**Parameters of** Hive **Map Side Join**

a. hive.auto.convert.join

b. Hive.auto.convert.join.noconditionaltask

Limitations **of Map Join in** Hive

Below are some limitations of Map Side join in Hive:

At First, the major restriction is, we can never convert Full outer joins to map-side joins. However, it is possible to convert a left-outer join to a map side join in hive. However, only possible since the right table that is to the right side of the join conditions, is lesser than 25 MB in size.

Also, we can convert a right-outer join to a map side join in hive. Similarly, only possible if the left table size is lesser than 25 MB.

#### 5. How to Identify Hive Map Join

Basically, we will see Hive Map Side Join Operator just below Map Operator Tree while using EXPLAIN command.

Other

Although, we can use the hint to specify the query using Map Join in Hive. Hence, below an example shows that smaller table is the one put in the hint, and force to cache table B manually.

```
Select /*+ MAPJOIN(b) */ a.key, a.value from a join b on a.key = b.key
```

For Example,

```
hive> set hive.auto.convert.join=true;
hive> set hive.auto.convert.join.noconditionaltask=true;
hive> set hive.auto.convert.join.noconditionaltask.size=20971520
hive> set hive.auto.convert.join.use.nonstaged=true;
hive> set hive.mapjoin.smalltable.filesize = 30000000;
```

#### b. Bucket Map Join

Basically, while the tables are large and all the tables used in the join are bucketed on the join columns we use a Bucket Map Join in Hive. In this article, we will cover the whole concept of Apache Hive Bucket Map Join. It also includes use cases, disadvantages, and Bucket Map Join example which will enhance our knowledge.

In Apache Hive, while the tables are large and all the tables used in the join are bucketed on the join columns we use Hive Bucket Map Join feature. Moreover, one table should have buckets in multiples of the number of buckets in another table in this type of join.

Basically, Join is done in Mapper only. However, let's understand it in this way, the mapper processing bucket 1 for table A will only fetch bucket 1 of table B.

#### Use Case

To be more specific we use this feature with several scenarios. Like:

- i. While all tables are large.
- ii. Also, while all tables are bucketed using the join columns.
- iii. Moreover, while The number of buckets in one table is a multiple of the number of buckets in the other table.
- iii. Also, when all tables are not sorted.

#### Disadvantages of Bucket Map Join in Hive

The major disadvantage of using Bucket Map Join is, here tables need to be bucketed in the same way how the SQL joins. That implies we can not use it for other types of SQLs.

#### c. Skew Join

Basically, when there is a table with skew data in the joining column, we use skew join feature. On defining what is skewed table, it is a table that is having values that are present in large numbers in the table compared to other data. However, while the rest of the data is stored in a separate file Skew data is stored in a separate file.

#### Parameter

However, to be set for a Hive skew join we need the following parameter:

```
set hive.optimize.skewjoin=true;
set hive.skewjoin.key=100000;
```

### How Hive Skew Join Works

However, let's assume if table A join B, and A has skew data "1" in joining column. At First store, the rows with key 1 in an in-memory hash table and read B. Further to read A run a set of mappers. Afterward, do the following:

Make sure use the hashed version of B to compute the result since it has key 1. Then, send all other keys to a reducer which does the join. Basically, from a mapper, this reducer will get rows of B also.

Learn Hadoop from Industry Experts

Hence, as a result, we end up reading only B twice. Basically, that implies that the skewed keys in A are only read and processed by the Mapper. Also, they are not sent to the reducer. Moreover, remaining keys in A go through only a single Map/Reduce. However, the assumption is that B has few rows with keys which are skewed in A. Hence, in this way these rows can be loaded into the memory.

### Skew Join - Use Case

Basically, on the joining column, one table has huge skew values.

Let Explore Joins in Hive with Examples

### Disadvantages of Skew Join in Hive

Here, are some Limitations of Hive Skew Join are discussed:

So, the major disadvantage of it is One table is read twice here.

Moreover, it is necessary that users should be aware of the skew key.

### d. Sort Merge Bucket Join

In Hive, while each mapper reads a bucket from the first table and the corresponding bucket from the second table, in SMB join. Basically, then we perform a merge sort join feature. Moreover, we mainly use it when there is no limit on file or partition or table join. Also, when the tables are large we can use Hive Sort Merge Bucket join. However, using the join columns, all join the columns are bucketed and sorted in SMB. Although, make sure in SMB join all tables should have the same number of buckets.

### How Hive SMB Works

Basically, in Mapper, only Join is done. Moreover, all the buckets are joined with each other at the mapper which are corresponding.

### Use Case of Sort Merge Bucket Join in Hive

There are several scenarios when we can use Hive Sort Merge Bucket Join:

While all tables are Large.

Also, while all tables are bucketed using the join columns.

While by using the join columns, Sorted.

Also, when the number of buckets is same as the number of all tables.

Read about Map Join in Hive | Map Side Join

### Disadvantages of Sort Merge Bucket Join in Hive

Following are the limitations of Hive Sort Merge Bucket Join:

However, in the same way how the SQL joins Tables need to be bucketed. Hence, for other types of SQL, it cannot be used.

Also, it is possible that Partition tables might slow down here.

## 5. Difference between order by , sort by , distribute by , cluster by

**SORT BY** - Data is ordered at each of 'N' reducers where the reducers can have overlapping range of data.

**ORDER BY**- This is similar to the ORDER BY in SQL where total ordering of data takes place by passing it to a single reducer.

**DISTRIBUTE BY** - It is used to distribute the rows among the reducers. Rows that have the same distribute by columns will go to the same reducer.

**CLUSTER BY**- It is a combination of DISTRIBUTE BY and SORT BY where each of the N reducers gets

non overlapping **range of data** which **is then** sorted **by** those ranges **at** the respective reducers.

6. How **do** we intergrate Hive **with** Spark

7. **Difference between** Managed Tables **and** External Tables

Ans. Managed **table**,

The metadata information along **with** the **table data is** deleted **from** the Hive warehouse directory **if** one drops a managed **table**.

**External table**,

Hive just deletes the metadata information regarding the **table**. Further, it leaves the **table data** present **in** HDFS untouched.

8. Different indexes **in** Hive

An **Index** acts **as** a reference **to** the records. Instead **of** searching **all** the records, we can refer **to** the **index to search for** a particular **record**. Indexes maintain the reference **of** the records. So that it **is** easy **to search for** a **record with** minimum overhead. Indexes also speed up the searching **of data**.

Types **of** Indexes **in** Hive

Compact Indexing

Bitmap Indexing

**Bit map** indexing was introduced **in** Hive **0.8** and **is** commonly used **for** columns **with distinct values**.

Differences **between** Compact **and** Bitmap Indexing

The main **difference is** the storing **of** the mapped **values of** the **rows in** the different blocks.

**When** the **data** inside a Hive **table is** stored **by default in** the HDFS, they **are** distributed across the nodes **in** a **cluster**. There needs **to** be a proper identification **of** the **data, like** the **data in** block indexing. This **data** will be able **to identity** which **row is** present **in** which block, so that **when** a query **is** triggered it can **go** directly **into** that block. So, **while** performing a query, it will **first check** the **index and then go** directly **into** that block.

Compact indexing stores the pair **of** indexed **column's value and** its blockid.

Bitmap indexing stores the combination **of** indexed **column value and** list **of rows as** a bitmap.

9. How **to create** a **Schema for** the **Data in** Hive

```
create table hive_table
(
)
fields terminated by ','
lines terminated by '\n'
;
```

10. What **are** different **Data** types **in** Hive

1. **Numeric** Types

**TINYINT** (1-byte signed **integer, from -128 to 127**)

**SMALLINT** (2-byte signed **integer, from -32,768 to 32,767**)

**INT** (4-byte signed **integer, from -2,147,483,648 to 2,147,483,647**)

**BIGINT** (8-byte signed **integer, from -9,223,372,036,854,775,808 to 9,223,372,036,854,775,807**)

**FLOAT** (4-byte single **precision** floating point **number**)

**DOUBLE** (8-byte **double precision** floating point **number**)

**DECIMAL** (Hive 0.13.0 introduced **user** definable **precision and** scale)

2. **Date/Time** Types

**TIMESTAMP**



**DATE**

## 3. String Types

STRING

**VARCHAR****CHAR**

## 4. Misc Types

**BOOLEAN****BINARY**

## 5. Complex Types

arrays: **ARRAY**<data\_type>maps: **MAP**<primitive\_type, data\_type>structs: **STRUCT**<col\_name : data\_type [**COMMENT** col\_comment], ...>**union**: **UNIONTYPE**<data\_type, data\_type, ...>

## 11. How to Select Complex Data Types in Hive

**--arrays****select** ename,subordinates[0] **from** employees;**--maps****select** ename,deductions["Federal Taxes"] **from** employees;**--structs****select** ename,address.state,address.city,address.street,address.zip **from** employees;

## 12. How to create Partition Table for Date column

## 13. Why Hive is not suitable for OLTP Applications

Hive is not suitable for OLTP systems because it does not provide insert and update function at the row level.

## 14. What is Metastore in Hive &amp; What is the Metastore in that you used. And How do you configure

Basically, to store the metadata information in Hive we use Metastore. Though, it is possible by using RDBMS and an open source ORM (Object Relational Model) layer called Data Nucleus. That converts the object representation into the relational schema and vice versa.

**Local Metastore:**

It is the metastore service runs in the same JVM in which the Hive service is running and connects to a database running in a separate JVM. Either on the same machine or on a remote machine.

**Remote Metastore:**

In this configuration, the metastore service runs on its own separate JVM and not in the Hive service JVM.

Basically, hive-site.xml file has to be configured with the below property, to configure metastore in Hive -

hive.metastore.uris

thrift: //node1 (or IP Address):9083

IP address and port of the metastore host

## 15. When you should use Sort by instead of Order by

Despite ORDER BY we should use SORT BY. Especially while we have to sort huge datasets. The reason is SORT BY clause sorts the data using multiple reducers. However, ORDER BY sorts all

of the data together using a single reducer. Hence, using ORDER BY will take a lot of time to execute a large number of inputs.

#### 16. What is Partitioning and when do you perform Partitioning

Basically, for the purpose of grouping similar type of data together on the basis of column or partition key, Hive organizes tables into partitions. Moreover, to identify a particular partition Each Table can have one or more partition keys. On defining Partition, in other words, it is a sub-directory in the table directory.

However, in a Hive table, Partitioning provides granularity. Hence, by scanning only relevant partitioned data instead of the whole dataset it reduces the query latency.

Dynamic partitioning values for partition columns are known in the runtime. In other words, it is known during loading of the data into a Hive table.

Usage:

While we Load data from an existing non-partitioned table, in order to improve the sampling. Thus it decreases the query latency.

Also, while we do not know all the values of the partitions beforehand. Thus, finding these partition values manually from a huge dataset is a tedious task.

#### 17. What is bucketing and when do you use bucketing

Basically, for performing bucketing to a partition there are two main reasons:

A map side join requires the data belonging to a unique join key to be present in the same partition.

It allows us to decrease the query time. Also, makes the sampling process more efficient.

However, by using the formula:  $\text{hash\_function}(\text{bucketing\_column}) \bmod \text{num\_of\_buckets}$  Hive determines the bucket number for a row. Basically, hash\_function depends on the column data type. Although, hash\_function for integer data type will be:

$\text{hash\_function}(\text{int\_type\_column}) = \text{value of int\_type\_column}$

#### 18. Explain Hive Indexing

An Index acts as a reference to the records. Instead of searching all the records, we can refer to the index to search for a particular record. Indexes maintain the reference of the records. So that it is easy to search for a record with minimum overhead. Indexes also speed up the searching of data.

Hive is a data warehousing tool present on the top of Hadoop, which provides the SQL kind of interface to perform queries on large data sets. Since Hive deals with Big Data, the size of files is naturally large and can span up to Terabytes and Petabytes. Now if we want to perform any operation or a query on this huge amount of data it will take large amount of time.

In a Hive table, there are many numbers of rows and columns. If we want to perform queries only on some columns without indexing, it will take large amount of time because queries will be executed on all the columns present in the table.

The major advantage of using indexing is; whenever we perform a query on a table that has an index, there is no need for the query to scan all the rows in the table. Further, it checks the index first and then goes to the particular column and performs the operation.

So if we maintain indexes, it will be easier for Hive query to look into the indexes first and then perform the needed operations within less amount of time.

When to use Indexing?

Indexing can be use under the following circumstances:

If the dataset is very large.

If the query execution is more amount of time than you expected.

If a speedy query execution is required.

When building a data model.

Indexes are maintained in a separate table in Hive so that it won't affect the data inside the table, which contains the data. Another major advantage for indexing in Hive is that indexes can also be partitioned depending on the size of the data we have.

Types of Indexes in Hive

Compact Indexing

Bitmap Indexing

Bit map indexing was introduced in Hive 0.8 and is commonly used for columns with distinct values.

Differences between Compact and Bitmap Indexing

The main difference is the storing of the mapped values of the rows in the different blocks. When the data inside a Hive table is stored by default in the HDFS, they are distributed across the nodes in a cluster. There needs to be a proper identification of the data, like the data in block indexing. This data will be able to identify which row is present in which block, so that when a query is triggered it can go directly into that block. So, while performing a query, it will first check the index and then go directly into that block.

Compact indexing stores the pair of indexed column's value and its blockid.

Bitmap indexing stores the combination of indexed column value and list of rows as a bitmap.

Let's now understand what is bitmap?

A bitmap is a type of memory organization or image file format used to store digital images so with this meaning of bitmap, we can redefine bitmap indexing as given below.

"Bitmap index stores the combination of value and list of rows as a digital image."

The following are the different operations that can be performed on Hive indexes:

Creating index

Showing index

Alter index

Dropping index

Creating Index in Hive

Syntax for creating a compact index in Hive is as follows:

```
CREATE INDEX index_name
ON TABLE table_name (columns,...)
AS 'org.apache.hadoop.hive.ql.index.compact.CompactIndexHandler'
WITH DEFERRED REBUILD;
```

Here, in the place of index\_name we can give any name of our choice, which will be the table's INDEX NAME.

In the ON TABLE line, we can give the table\_name for which we are creating the index and the names of the columns in brackets for which the indexes are to be created. We should specify the columns which are available only in the table.

The org.apache.hadoop.hive.ql.index.compact.CompactIndexHandler' line specifies that a built in CompactIndexHandler will act on the created index, which means we are creating a compact index for the table.

The WITH DEFERRED REBUILD statement should be present in the created index because we need to alter the index in later stages using this statement.

This syntax will create an index for our table, but to complete the creation, we need to complete the REBUILD statement. For this to happen, we need to add one more alter statement. A MapReduce job will be launched and the index creation is now completed.

Hadoop

```
ALTER INDEX index_name ON table_name REBUILD;
```

This ALTER statement will complete our REBUILT index creation for the table.

Examples - Creating Index

In this section we will first execute the hive query on non-indexed table and will note down the time taken by query to fetch the result.

In the second part, we will be performing the same query on indexed table and then will compare the time taken by query to fetch the result with the earlier case.

We will be demonstrating this difference of time with practical examples.

In first scenario we are performing operations on non-indexed table.

Let's create a normal managed table to contain the olympic dataset first.

#### Table Creation

```
create table olympic(athlete STRING,age INT,country STRING,year STRING,closing STRING,
sport STRING,gold INT,silver INT,bronze INT,total INT) row format delimited fields
terminated by '\t' stored as textfile;
```

Here we are creating a table with name 'olympic'. The schema of the table is as specified and the data inside the input file is delimited by tab space.

At the end of the line, we have specified 'stored as textfile', which means we are using a TEXTFILE format.

You can check the schema of your created table using the command 'describe olympic;'

We can load data into the created table as follows:

```
load data local inpath 'path of your file' into table olympic;
```

We have successfully loaded the input file data into the table which is in the TEXTFILE format.

Let's perform an Average operation on this 'olympic' data. Let's calculate the average age of the athletes using the following command:

```
SELECT AVG(age) from olympic;
```

Here we can see the average age of the athletes to be 26.405433646812956 and the time for performing this operation is 21.08 seconds.

Now, let's create the index for this table:

```
CREATE INDEX olympic_index
ON TABLE olympic (age)
AS 'org.apache.hadoop.hive.ql.index.compact.CompactIndexHandler'
WITH DEFERRED REBUILD;
```

```
ALTER INDEX olympic_index on olympic REBUILD;
```

Here we have created an index for the 'olympic' table on the age column. We can view the indexes created for the table by using the below command:

```
show formatted index on olympic;
```

We can see the indexes available for the 'olympic' table in the above image.

Now, let's perform the same Average operation on the same table.

We have now got the average age as 26.405433646812956, which is same as the above, but now the time taken for performing this operation is 17.26 seconds, which is less than the above case.

Now we know that by using indexes we can reduce the time of performing the queries.

Can we have different indexes for the same table?

Yes! We can have any number of indexes for a particular table and any type of indexes as well.

Let's create a Bitmap index for the same table:

```
CREATE INDEX olympic_index_bitmap
ON TABLE olympic (age)
AS 'BITMAP'
WITH DEFERRED REBUILD;

ALTER INDEX olympic_index_bitmap on olympic REBUILD;
```

19. Explain Different types of Joins in Hive -- Duplicate
20. Explain --Duplicate
  - a. Bucket Map Join
  - b. Skew Join
  - c. Sort Merge Bucket Join
21. Explain SORT BY, ORDER BY, DISTRIBUTE BY and CLUSTER BY with Example
22. How do process query for
  - a. XML

```
1) create table xmlsample_guru(str string);
2) load data local inpath '/home/hduser/test.xml' overwrite into table xmlsample_guru;
3) select xpath(str,'emp/ename/text()'), xpath(str,'emp/esal/text()') from
xmlsample_guru;
```

b. Json

```
1) create table json_guru(str string);
2) load data inpath 'home/hduser/test.json' into table json_guru;
3) select * from jsonl;
4) select get_json_object(str,'$.ecode') as ecode, get_json_object(str,'$.ename') as
ename ,get_json_object(str,'$.sal') as salary from json_guru;
```

c. CSV

```
CREATE TABLE AllstarFull (playerID string,yearID string,gameNum string,gameID string,
teamID string,lgID string,GP string,startingPos string) row format delimited fields
terminated by ',' stored as textfile;
```

```
LOAD DATA INPATH '/user/bigdataproject/AllstarFull.csv' OVERWRITE INTO TABLE AllstarFull;

SELECT * FROM AllstarFull;
```

23. What are complex data types and how do you query Hive Collections

```
hive> create table employees ( ename string,salary float,subordinates array<string>,
deductions map<string,float>,address struct<street:string,city:string,state:string,zip:int>)
> row format delimited
> fields terminated by '\001'
> collection items terminated by '\002'
> map keys terminated by '\003';
```

```
load data local inpath '/home/cloudera/Downloads/employees.txt'
OVERWRITE into table employees;
```

```
--arrays
select ename,subordinates[0] from employees;

--maps
select ename,deductions["Federal Taxes"] from employees;

--structs
select ename,address.state,address.city,address.street,address.zip from employees;
```

24. Explain What **are** the Optimization Technique Availalble **in** Hive

Types **of** Query Optimization Techniques **in** Hive

- a. Tez-Execution Engine **in** Hive
- b. **Usage of** Suitable **File Format in** Hive
- c. Hive Partitioning
- d. Bucketing **in** Hive
- e. Vectorization **In** Hive
- f. Cost-Based Optimization **in** Hive (CBO)
- g. Hive Indexing

25. Explain Views **in** Hive

A **view** allows a query **to** be saved **and** treated **like** a **table**. It **is** a logical construct, **as** it does **not** store **data like** a **table**. **In** other words, materialized views **are not** currently supported **by** Hive.

**When** a query **references** a **view**, the information **in** its definition **is** combined **with** the rest **of** the query **by** Hive's query planner. Logically, you can imagine that Hive executes the **view** **and then** uses the results **in** the rest **of** the query.

Views **to** Reduce Query Complexity

**When** a query becomes **long or** complicated, a **view** may be used **to** hide the complexity **by** dividing the query **into** smaller, more manageable pieces; similar **to** writing a **function in** a programming **language or** the concept **of** layered design **in** software. Encapsulating the complexity makes it easier **for end** users **to** construct complex queries **from** reusable parts. **For** example, consider the following query **with** a nested subquery:

It **is** common **for** Hive queries **to** have many levels **of** nesting. **In** the following example, the nested portion **of** the query **is** turned **into** a **view**:

```
CREATE VIEW shorter_join AS
SELECT * FROM people JOIN cart
ON (cart.people_id=people.id) WHERE firstname='john';
Now the view is used like any other table. In this query we added a WHERE clause to the SELECT statement. This exactly emulates the original query:
```

```
SELECT lastname FROM shorter_join WHERE id=3;
```

26. Did you used UDFs **in** Hive

-- Yes (But its not written by Me) we used it for Percentile calculations

27. What **is** Beeline

HiveServer2 supports a command shell Beeline that works **with** HiveServer2. It's a JDBC client that is based on the SQLLine CLI (<http://sqlline.sourceforge.net/>). There's detailed documentation of SQLLine which is applicable to Beeline as well.

Replacing the Implementation of Hive CLI Using Beeline

The Beeline shell works in both embedded mode as well as remote mode. In the embedded mode, it runs an embedded Hive (similar to Hive CLI) whereas remote mode is for connecting to a separate HiveServer2 process over Thrift. Starting in Hive 0.14, when Beeline is used with HiveServer2, it also prints the log messages from HiveServer2 for queries it executes to STDERR. Remote HiveServer2 mode is recommended for production use, as it is more secure and doesn't require direct HDFS/metastore **access to** be granted **for** users.

28. What version of Hive you used in your organization

29. What is Impala --Not used

30. Explain Different SET Operations in Hive

```
set hive.cli.current.print.current.db=true
set hive.auto.convert.join=true
set hive.exec.dynamic.partition=true
set hive.exec.dynamic.partition.mode=nonstrict;
set mapred.reduce.tasks=50
set hive.exec.reducers.max=50
```

31. Why do you drop a External Table

-- Needs answer

32. Explain Serde in Hive

33. What are File Formats supported by Hive -- Check 44

34. Explain variables in Hive

```
hive> set CURRENT_DATE='2012-09-16';
hive> select * from foo where day >= '${hiveconf:CURRENT_DATE}'
% hive -hiveconf CURRENT_DATE='2012-09-16' -f test.hql
```

35. Explain How do you insert Date in Hive Table

```
insert into table_name values ();
```

```
insert into new_tables
select * from table_name;
```

36. Explain Analytical functions in Hive

1. count
2. sum
3. lead
4. lag
5. FIRST\_VALUE
6. ROW\_NUMBER
7. Rank
8. Dense Rank

37. How do you delete Duplicates in Hive

```
insert overwrite table dynpart select distinct * from dynpart;
```

- 1) Create a new table from old table (with same structure).
- 2) Copy distinct rows in new table from existing table.  

```
select col1,col2,col3,col4,max(<duplicate column>) as <name of duplicate column> from <table name> group by col1,col2,col3,col4;
```
- 3) Delete old table.
- 4) Rename new table to old one.

38. Explain Architecture of Hive

There are several components of Hive Architecture. Such as -

**User Interface** - Basically, it calls the **execute interface** to the driver. Further, driver creates a **session** handle to the query. Then sends the query to the compiler to generate an execution plan for it.

**Metastore** - It is used to Send the metadata to the compiler. Basically, for the execution of the query on receiving the send MetaData request.

**Compiler**- However, it generates the execution plan. Especially, that is a DAG of stages where

each stage is either a metadata operation, a map or reduce job or an operation on HDFS.

39. What is Apache HCatalog

Hcatalog can be used to share data structures with external systems. Hcatalog provides access to hive metastore to users of other tools on Hadoop so that they can read and write data to hive's data warehouse.

40. What is Hive Current Version and What is Hive stable Version

41. Difference between SQL and HQL

SQL : It supports DML

HQL : It doesn't support DML

42. How do you pull the Oracle data into Hive

```
scoop import \
--connect jdbc:mysql://localhost/dualcore \
--username training \
--password training \
--m 1 \
--target-dir /queryresult \
--table employees \
--hive-import
```

43. How to integrate Hive with Spark

```
val sparkSession = SparkSession.builder
.master("local")
.appName("demo")
.enableHiveSupport()
.getOrCreate()
```

```
sparkSession.sqlContext.sql("INSERT INTO TABLE students VALUES ('Rahul','Kumar'), ('abc','xyz')")
```

44. Hive File types

File formats in Hive

- a. Row Based File format
- b. column based file format

- a. Row based file format
  1. Text file Format
  2. Sequence File Format
  3. Avro File Format
- b. Column Based File format
  1. RC File
  2. ORC File
  3. Parquet

a.1. Text file Format :

It is the default format of Hive, It is a human readable file format.

Text file format doesnot allow compression technique. It has interoperable to HDFS and non HDFS. It takes huge space.

Example : tab separated file, comma separated , space separated

How to create text file format :



```
create table emp (
    empno int,ename string,salary float)
stored as TextFile;
```

```
describe formatted table_name : org.apache.hadoop.hive.ql.io.
HiveIgnoreKeyTextOutputFormat.
```

#### a.2 Sequence File Format :

It is the binary file format and it is a row based file format. It is not a human readable format.

Sequence file format binary or images. Sequence file format supports compression . performance wise very good.

Drawback : Poor interoperable. It supports only HDFs .

```
create table emp_sequence
(eno int,
ename string,
salary int,
gendar string,
dno int
) stored as SEQUENCEFILE;
```

you cannot extract the schema

#### a.3. AVRO

AVRO file format is sequence file format and on top of the data machine creates a schema.

AVRO tool creates a schema on top of the data.

#### Step 1:

```
sqoop import-all-tables \
--connect "jdbc:mysql://quickstart.cloudera:3306/retail_db" \
--username retail_dba \
--password cloudera \
--warehouse-dir /user/hive/warehouse/retail_stage.db \
--compress \
--compression-codec snappy \
--as-avrodatafile
-m 1;
```

#### Step 2:

```
hadoop fs -get /user/hive/warehouse/retail_stage.db/orders/part-m-00000.avro
avro-tools getschema part-m-00000.avro > orders.avsc
hadoop fs -mkdir /user/hive/schemas
hadoop fs -ls /user/hive/schemas/order
hadoop fs -copyFromLocal orders.avsc /user/hive/schemas/order
```

Launch HIVE using 'hive' command in a separate terminal

Below HIVE command will create a table pointing to the avro data file for orders data

```
create external table orders_sqoop
STORED AS AVRO
LOCATION '/user/hive/warehouse/retail_stage.db/orders'
```

Column Based format .

### 1. RC File format (Row column) --

The Default column based file format is RC File format .

If we are using RC file format. All the column formats by defaultly string. It occupies more space.

RC File format supports compression

Drawback is : Poor performance and less interoperability

Example :

```
create table emp_rc (
no int,
ename string,
salary int,
gender string,
dno int
) stored as RCFILE;
```

### 2. ORC File Format :

Horton works introduced improved version of RC File.

Based on Input data is stored as that data type instead of all string.

Drawback is it only support Hadoop.

### 3. Parquet

Check more about parquet

1. If use case is more about reading the data -- AVRO

2. If use case is more about writing the data of single column -- PARQUET

3. If use case is more about writing data to many columns -- AVRO

Mostly we use AVRO in realtime

Spark default file format is parquet.

'Sqoop'

### 1. How to Import Query data into HDFS

```
sqoop import \
--connect jdbc:oracle:thin:@//localhost:1521/xe \
--username scott \
--password tiger \
--table EMP \
--warehouse-dir /user/cloudera/sqoop_dir or --target-dir /etl/input/cities
```

### 2. How to Import Data from Oracle to Hive Table or Hive Partitions

```
sqoop import \
--connect jdbc:oracle:thin:@//HYRDSLVM0028.es.ad.adp.com:1521/cri02hyd \
--username scott \
--password tiger \
--table EMP \
--hive-import \
--create-hive-table \
--hive-table naravishdb.EMP \
--null-string '\\N' \
--null-non-string '\\N' \
```

### 3. How to do incremental import using sqoop

```

sqoop import \
--connect jdbc:oracle:thin:@//localhost:1521/xe \
--username scott \
--password tiger \
--table SALGRADE \
--incremental append \
--check-column GRADE \
--last-value 5 \
--warehouse-dir /user/cloudera/sqoop_dir/

```

4. How to create job or store the last value and retrieve in sqoop

```

sqoop job \
--create salgrade \
-- \
import \
--connect jdbc:oracle:thin:@//localhost:1521/xe \
--username scott \
--password tiger \
--table SALGRADE \
--incremental append \
--check-column GRADE \
--last-value 6

```

5. How to set the boundary in sqoop

```
--boundary clause
```

6. How to import data into HBase

7. Boundary Query

```

-- Boundary Condition
sqoop import \
--connect jdbc:oracle:thin:@//localhost:1521/xe \
--username scott \
--password tiger \
--query 'SELECT normcities.id, \
countries.country, \
normcities.city \
FROM normcities \
JOIN countries USING(country_id) \
WHERE $CONDITIONS' \
--split-by id \
--target-dir cities \
--boundary-query "select min(id), max(id) from normcities"

```

8. \$CONDITIONS

Sqoop performs highly efficient data transfers by inheriting Hadoop's parallelism.

To help Sqoop split your query into multiple chunks that can be transferred in parallel, you need to include the \$CONDITIONS placeholder in the where clause of your query.

To help Sqoop split your query into multiple chunks that can be transferred in parallel, you need to include the \$CONDITIONS placeholder in the where clause of your query.

Sqoop will automatically substitute this placeholder with the generated conditions specifying which slice of data should be transferred by each individual task.

Sqoop will automatically substitute this placeholder with the generated conditions specifying which slice of data should be transferred by each individual task.

While you could skip \$CONDITIONS by forcing Sqoop to run only one job using the --num-mappers 1 parameter, such a limitation would have a severe performance impact.

While you could skip \$CONDITIONS by forcing Sqoop to run only one job using the

```
--num-mappers 1 parameter, such a limitation would have a severe performance impact.
```

#### 9. --where

To specify conditions **while** import **or** export

10. Append **and** overwrite Directo  
ry (overwrite doesnot exist, we need **to** handle separatlely **in** shell)

#### 11. How **to do** Incremental load **or** delta load

```
--Using Last value
```

#### 12. **Insert/update in** Sqoop Incremental Why **update not work in** sqoop

**to check**

#### 13. Integration **of** Hive **with** Sqoop

#### 14. How you query **using** sqoop

```
sqoop eval --connect --query "select count(0) from emp"
```

#### 15. How **to** pull **all** the tables **using** sqoop

```
--import-all-tables
```

#### 16. What **are file** formats supported **by** sqoop

```
sqoop su
```

Newer Version **of** Sqoop support **file** formats **like**

1. **sequence file format**
2. Avro **File format**
3. Parquet **file format**

#### 17. Does Sqoop supports **CLOB** Columns

```
sqoop import \  
-Dmapred.job.queue.name=default \  
-connect jdbc:oracle:thin:@hostname:port/port \  
-username XXXXXX \  
-password XXXXXX \  
-query "SELECT * FROM tablename WHERE \${CONDITIONS}" \  
-hive-drop-import-delims \  
-map-column-java column1=String,column2=String \  
-m 8 \  
-hive-import \  
-hive-table tablename \  
-target-dir /user/hdfs/ \  
-fields-terminated-by '01' \  
-split-by id;
```

#### 18. Different Options avaiable **in** sqoop --> Same as 38

#### 19. What **is** better sqoop **or** Spark pull

Spark **is** better **than** sqoop **for Data** Extraction **as** spark works **on In-memory**.  
Sqoop works **with** I/o

#### 20. How you **do** incremental pull **using** sqoop job

#### 21. How **to** Handle **Null in** sqoop import

```
--null-string          --> Null String  
--null-non-string      --> Null for non strings
```

22. Explain `--append` option in sqoop

23. Explain free form query `in` sqoop

Use `--query`

24. Difference between `--target-dir` `--warehouse-dir`

`--target-dir` Mainly used Importing a Single Table into HDFS  
`--always-target-dir` looking for a new directory in HDFS

`--warehouse-dir`

If you want to import all the tables of schema we use

25. How to store and use last value in sqoop job

.If an incremental import is run from the command line, the value which should be specified as `--last-value` in a subsequent incremental import will be printed to the screen for your reference.

If an incremental import is run from a saved job, this value will be retained in the saved job. Subsequent runs of sqoop job will continue to import only newer rows than those previously imported.

26. How to use password file

```
sqoop import --connect jdbc:mysql://localhost:3306/db --username bhavesh --password-file /pwd --table t1 --target-dir '/erp/test'
```

27. where you should copy the jars

```
cp /usr/lib/hive/lib/mysql*.jar /usr/lib/hadoop/lib
```

28. How to exclude table in import all

```
--exclude-table table_list
```

29. How to increase number of mappers

```
--m 10
--num-of_mapper
```

30. how to do compression

```
--compress \
```

31. Is it possible to update record using sqoop

With insert mode, records exported by Sqoop are appended to the end of the target table. Sqoop also provides an update mode that you can use by providing the `-update-key <column(s)>` command line argument. This action causes Sqoop to generate a SQL UPDATE statement to run on the RDBMS or data warehouse.

Assume that you want to update a three-column table with data stored in the HDFS file `/user/my-hdfs-file`. The file contains this data:

100, 1000, 2000

The following abbreviated Sqoop export command generates the corresponding SQL UPDATE statement on your database system:

```
$ sqoop export (Generic Arguments)
--table target-relational-table
```

```
--update-key column1
--export-dir /user/my-hdfs-file
...
```

Generates => **UPDATE** target-relational-table **SET**  
                     column2=1000,column3=2000  
           **WHERE** column1=100;

### 32. Export and Import Data from and to Oracle

```
sqoop import \
--connect jdbc:oracle:thin:@//HYRDSLVM0028.es.ad.adp.com:1521/cri02hyd \
--username hr \
--password hr \
--table JOBS

sqoop export \
--connect jdbc:oracle:thin:@//localhost:1521/xe \
--username scott \
--password tiger \
--table EMP_DEPT \
--export-dir '/user/cloudera/emp_dept.txt'
```

### 33. Export and Import Data from and to Hive

```
sqoop import \
--connect jdbc:mysql://mysql.example.com/sqoop \
--username sqoop \
--password sqoop \
--table cities \
--hive-import \
--hive-partition-key day \
--hive-partition-value "2013-05-22"

sqoop export \
--connect jdbc:oracle:thin:@//HYRDSLVM0028.es.ad.adp.com:1521/cri02hyd \
--username scott \
--password tiger \
--table EMP \
--null-string '\\N' --input-null-string '\\N' \
--export-dir 'naravish.db/emp/part*' \
--fields-terminated-by '\\001'
```

### 34. Export and Import Data from and to Hbase

```
sqoop import \
--connect jdbc:mysql://mysql.example.com/sqoop \
--username sqoop \
--password sqoop \
--table cities \
--hbase-table cities \
--column-family world

sqoop export \
--connect jdbc:oracle:thin:@//HYRDSLVM0028.es.ad.adp.com:1521/cri02hyd \
--username scott \
--password tiger \
--table EMP2 \
--input-null-string '\\N' \
--input-null-non-string '\\N' \
--export-dir 'naravish.db/emp_permanent_hbase/*' \
--fields-terminated-by ','
```

### 36. The nine functions of Sqoop?

A. Full Load

- B. Incremental Load
- C. Parallel import/export
- D. Import results of SQL query
- E. Compression
- F. Connectors for all major RDBMS Databases
- G. Kerberos Security Integration
- H. Load data directly into Hive/Hbase
- I. Support for Accumulo

### 37. Default number of parallel jobs

Defaultly sqoop does 4 parallel jobs

### 38. Explain

```
--append                --> creates a new part file
--as-avrodatafile        --> Saves data as avro Data file
--as-sequencefile        --> Saves data as s
--as-textfile            --> Saves as Textfile
--boundary-query         --> Query to specify min and max value
--columns                --> To pull only certain columns
--direct                 --> --direct - Use direct import fast path
--direct-split-size      -->
--inline-lob-limit       -->
---m                     --> Num of Mappers
--e,--query              --> Customer Query
--split-by               --> column value that should be used to min and max value incase of
no PK
--table                  --> Table Name
--target-dir             --> Target Directory
--warehouse-dir          --> Incase of importing import-all tables
--where                  --> Where clause to import data
--compress               --> Compression
--compression-codec      --> Dont know
--null-string            --> Null String
--null-non-string        --> Null for non strings
```

'HDFS'

### 1. What is Data Locality

Data locality refers to the ability to move the computation close to where the actual data resides on the node, instead of moving large data to computation. This minimizes network congestion and increases the overall throughput of the system.

### 2. Difference between 1.0 vs 2.0

Hadoop V.1.x Components

Apache Hadoop V.1.x has the following two major Components

HDFS (HDFS V1)

MapReduce (MR V1)

In Hadoop V.1.x, these two are also know as Two Pillars of Hadoop.

hadoop1.x-components

Hadoop V.2.x Components

Apache Hadoop V.2.x has the following three major Components

HDFS V.2

YARN (MR V2)

MapReduce (MR V1)

In Hadoop V.2.x, these two are also know as Three Pillars of Hadoop.

hadoop2.x-components

## Hadoop 1.x Limitations

Hadoop 1.x has many limitations or drawbacks. Main drawback of Hadoop 1.x is that MapReduce Component in it's Architecture. That means it supports only MapReduce-based Batch/Data Processing Applications.

Hadoop 1.x has the following Limitations/Drawbacks:

It is only suitable for Batch Processing of Huge amount of Data, which is already in Hadoop System.

It is not suitable for Real-time Data Processing.

It is not suitable for Data Streaming.

It supports upto 4000 Nodes per Cluster.

It has a single component : JobTracker to perform many activities like Resource Management, Job Scheduling, Job Monitoring, Re-scheduling Jobs etc.

JobTracker is the single point of failure.

It does not support Multi-tenancy Support.

It supports only one Name Node and One Namespace per Cluster.

It does not support Horizontal Scalability.

It runs only Map/Reduce jobs.

It follows Slots concept in HDFS to allocate Resources (Memory, RAM, CPU). It has static Map and Reduce Slots. That means once it assigns resources to Map/Reduce jobs, it cannot re-use them even though some slots are idle.

For Example:- Suppose, 10 Map and 10 Reduce Jobs are running with 10 + 10 Slots to perform a computation. All Map Jobs are doing their tasks but all Reduce jobs are idle. We cannot use these Idle jobs for other purpose.

NOTE:- In Summary, Hadoop 1.x System is a Single Purpose System. We can use it only for MapReduce Based Applications.

## Differences between Hadoop 1.x and Hadoop 2.x

If we observe the components of Hadoop 1.x and 2.x, Hadoop 2.x Architecture has one extra and new component that is : YARN (Yet Another Resource Negotiator).

It is the game changing component for BigData Hadoop System.

### New Components and API

As shown in the below diagram, Hadoop 1.x is re-architected and introduced new component to solve Hadoop 1.x Limitations.

### hadoop1\_vs\_hadoop2

#### Hadoop 1.x Job Tracker

As shown in the below diagram, Hadoop 1.x Job Tracker component is divided into two components:

#### Resource Manager:-

To manage resources in cluster

#### Application Master:-

To manage applications like MapReduce, Spark etc.

### hadoop1\_jobtracker\_hadoop2

Hadoop 1.x supports only one namespace for managing HDFS filesystem whereas Hadoop 2.x supports multiple namespaces.

Hadoop 1.x supports one and only one programming model: MapReduce. Hadoop 2.x supports multiple programming models with YARN Component like MapReduce, Interactive, Streaming, Graph, Spark, Storm etc.

Hadoop 1.x has lot of limitations in Scalability. Hadoop 2.x has overcome that limitation with new architecture.

Hadoop 2.x has Multi-tenancy Support, but Hadoop 1.x doesn't.

Hadoop 1.x HDFS uses fixed-size Slots mechanism for storage purpose whereas Hadoop 2.x uses variable-sized Containers.



Hadoop 1.x supports maximum 4,000 nodes per cluster where Hadoop 2.x supports more than 10,000 nodes per cluster.

How Hadoop 2.x solves Hadoop 1.x Limitations

Hadoop 2.x has resolved most of the Hadoop 1.x limitations by using new architecture.

By decoupling MapReduce component responsibilities into different components.

By Introducing new YARN component for Resource management.

By decoupling component's responsibilities, it supports multiple namespace, Multi-tenancy, Higher Availability and Higher Scalability.

Hadoop 2.x YARN Benefits

Hadoop 2.x YARN has the following benefits.

Highly Scalability

Highly Availability

Supports Multiple Programming Models

Supports Multi-Tenancy

Supports Multiple Namespaces

Improved Cluster Utilization

Supports Horizontal Scalability

### 3. Explain the Architecture of 2.0

Hadoop 2.x Architecture

Apache Hadoop 2.x or later versions are using the following Hadoop Architecture. It is a Hadoop 2.x High-level Architecture. We will discuss in-detailed Low-level Architecture in coming sections.

hadoop architecture

Hadoop Common Module is a Hadoop Base API (A Jar file) for all Hadoop Components. All other components works on top of this module.

HDFS stands for Hadoop Distributed File System. It is also know as HDFS V2 as it is part of Hadoop 2.x with some enhanced features. It is used as a Distributed Storage System in Hadoop Architecture.

YARN stands for Yet Another Resource Negotiator. It is new Component in Hadoop 2.x Architecture. It is also know as "MR V2".

MapReduce is a Batch Processing or Distributed Data Processing Module. It is also know as "MR V1" as it is part of Hadoop 1.x with some updated features.

Remaining all Hadoop Ecosystem components work on top of these three major components: HDFS, YARN and MapReduce. We will discuss all Hadoop Ecosystem components in-detail in my coming posts.

When compared to Hadoop 1.x, Hadoop 2.x Architecture is designed completely different. It has added one new component : YARN and also updated HDFS and MapReduce component's Responsibilities.

Hadoop 2.x Major Components

Hadoop 2.x has the following three Major Components:

HDFS

YARN

MapReduce

These three are also known as Three Pillars of Hadoop 2. Here major key component change is YARN. It is really game changing component in BigData Hadoop System.

How Hadoop 2.x Major Components Works

Hadoop 2.x components follow this architecture to interact each other and to work parallel in a reliable, highly available and fault-tolerant manner.

Hadoop 2.x Components High-Level Architecture

hadoop 2 architecture diagram

All Master Nodes and Slave Nodes contains both MapReduce and HDFS Components.

One Master Node has two components:

Resource Manager (YARN or MapReduce v2)

HDFS

It's HDFS component is also known as NameNode. It's NameNode is used to store Meta Data.

In Hadoop 2.x, some more Nodes act as Master Nodes as shown in the above diagram. Each this 2nd level Master Node has 3 components:

Node Manager

Application Master

Data Node

Each this 2nd level Master Node again contains one or more Slave Nodes as shown in the above diagram.

These Slave Nodes have two components:

Node Manager

HDFS

It's HDFS component is also known as Data Node. It's Data Node component is used to store actual our application Big Data. These nodes do not contain Application Master component.

Hadoop 2.x Components In-detail Architecture

hadoop components and architecture

Hadoop 2.x Architecture Description

Resource Manager:

Resource Manager is a Per-Cluster Level Component.

Resource Manager is again divided into two components:

Scheduler

Application Manager

Resource Manager's Scheduler is :

Responsible to schedule required resources to Applications (that is Per-Application Master).

It does only scheduling.

It does care about monitoring or tracking of those Applications.

Application Master:

Application Master is a per-application level component. It is responsible for:

Managing assigned Application Life cycle.

It interacts with both Resource Manager's Scheduler and Node Manager

It interacts with Scheduler to acquire required resources.

It interacts with Node Manager to execute assigned tasks and monitor those task's status.

Node Manager:

Node Manager is a Per-Node Level component.

It is responsible for:

Managing the life-cycle of the Container.

Monitoring each Container's Resources utilization.

Container:

Each Master Node or Slave Node contains set of Containers. In this diagram, Main Node's Name Node is not showing the Containers. However, it also contains a set of Containers. Container is a portion of Memory in HDFS (Either Name Node or Data Node).

In Hadoop 2.x, Container is similar to Data Slots in Hadoop 1.x. We will see the major differences between these two Components: Slots Vs Containers in my coming posts.

NOTE:-

Resource Manager is Per-Cluster component where as Application Master is per-application component.

Both Hadoop 1.x and Hadoop 2.x Architectures follow Master-Slave Architecture Model.

4. Explain the role of YARN

Apache Yarn - "Yet Another Resource Negotiator" is the resource management layer of Hadoop. The Yarn was introduced in Hadoop 2.x. Yarn allows different data processing engines like graph processing, interactive processing, stream processing as well as batch processing to run and process data stored in HDFS (Hadoop Distributed File System). Apart from resource management, Yarn also does job Scheduling. Yarn extends the power of Hadoop to other evolving technologies, so they can take the advantages of HDFS (most reliable and popular storage system on the planet) and economic cluster. To learn installation of Apache Hadoop 2 with Yarn follows this quick installation guide.

Apache yarn is also a data operating system for Hadoop 2.x. This architecture of Hadoop 2.x provides a general purpose data processing platform which is not just limited to the MapReduce. It enables Hadoop to process other purpose-built data processing system other than MapReduce. It allows running several different frameworks on the same hardware where Hadoop is deployed.

In this section of Hadoop Yarn tutorial, we will discuss the complete architecture of Yarn. Apache Yarn Framework consists of a master daemon known as "Resource Manager", slave daemon called node manager (one per slave node) and Application Master (one per application).

Resource Manager has two Main components

Scheduler  
Application manager

Node Manager (NM)  
Application Master (AM)  
Resource Manager Restart  
Yarn Resource Manager High availability

## 5. What is the Issue with Hadoop 1.0.

The NameNode is the single point of failure in Hadoop 1.0.

Each cluster has a single NameNode and if that machine is not available, the whole cluster will be not available.

This impacts the total availability of HDFS in two ways:

For any unplanned event such as machine crashes, the whole cluster is not available until the Name node is brought up manually.

For planned maintenance such as Hardware or Software upgrades on NameNode would result in cluster unavailability.

In Hadoop 2.0, HDFS High Availability feature addresses the above problem, by providing an option to run two NameNodes in the same cluster in an Active/Passive configuration with a hot standby.

This allows fast Failover to a new NameNode for any machine crashes or administrator initiated fail-over for any planned maintenance activities.

## 6. How Name node single point of failure is rectified in Hadoop 2.0

HDFS High Availability of Namenode is introduced with Hadoop 2. In this two separate machines are getting configured as NameNodes, where one NameNode always in working state and another is in standby. Working Name node handling all clients request in the cluster where standby is behaving as the slave and maintaining enough state to provide a fast failover on Working Name node.

## 7. Why block size is 128 KB in Hadoop

HDFS blocks are large compared to disk blocks, and the reason is to minimize the cost of seeks. By making a block large enough, the time to transfer the data from the disk can be significantly longer than the time to seek to the start of the block. Thus the time to transfer a large file made of multiple blocks operates at the disk transfer rate.

A quick calculation shows that if the seek time is around 10 ms and the transfer rate is 100 MB/s, to make the seek time 1% of the transfer time, we need to make the block size around 100 MB. The default is actually 64 MB, although many HDFS installations use 128 MB blocks. This figure will continue to be revised upward as transfer speeds grow with new generations of disk drives.

This argument shouldn't be taken too far, however. Map tasks in MapReduce normally operate on one block at a time, so if you have too few tasks (fewer than nodes in the cluster), your jobs will run slower than they could otherwise.

#### 8. Explain

- a. Edit logs
- b. FSImage

#### 9. Explain how fault tolerant is achieved in Hadoop

Using Replication factor of 3

#### 10. Why Hadoop

Hadoop Provides 2 Main Important things

1. HDFS For Distributed storage
2. Map Reduce for Distributed Processing

Hadoop Make it possible

- a. Expense is low (Commodity Hardware)
- b. Fast Processing because of Data locality
- c. Fault tolerant (Replication Factor)
- d. Possible Horizontal Scalability

#### 11. Explain Heartbeat in Hadoop

Heartbeat is the signal that is sent by the datanode to the namenode after the regular interval to time to indicate its presence, i.e. it is alive. NameNode and DataNode do communicate using Heartbeat. If after the certain time of heartbeat, NameNode do not receive any response from DataNode, then that Node is dead.

#### 12. Explain the replication factor in Hadoop

For example, if the replication factor was set to 3 (default value in HDFS) there would be one original block and two replicas. hdfs-site.xml is used to configure HDFS. ... The block size setting is used by HDFS to divide files into blocks and then distribute those blocks across the cluster. Feb 12, 2018

If we write a data to block Hadoop takes responsibility to copy to other 2 data nodes. And it is rack aware.

#### 13. Explain Safe mode in Hadoop

Safemode in Apache Hadoop is a maintenance state of NameNode. During which NameNode doesn't allow any modifications to the file system. During Safemode, HDFS cluster is in read only and doesn't replicate or delete block.

In SafeMode NameNode performs collection of block reports from datanodes. NameNode enters safemode automatically during its start up. NameNode leaves Safemode after the DataNodes have reported that most Blocks are available.

To know the status of Safemode, use command: `hadoop dfsadmin -safemode get`

To enter Safemode, use command: `bin/hadoop dfsadmin -safemode enter`  
To come out of Safemode, use command: `hadoop dfsadmin -safemode`

#### 14. Explain Small file problem in Hadoop

Problems with small files and HDFS

A small file is one which is significantly smaller than the HDFS block size (default 64MB). If you're storing small files, then you probably have lots of them (otherwise you wouldn't turn to Hadoop), and the problem is that HDFS can't handle lots of files.

Every file, directory and block in HDFS is represented as an object in the namenode's memory, each of which occupies 150 bytes, as a rule of thumb. So 10 million files, each using a block, would use about 3 gigabytes of memory. Scaling up much beyond this level is a problem with current hardware. Certainly a billion files is not feasible.

Furthermore, HDFS is not geared up to efficiently accessing small files: it is primarily designed for streaming access of large files. Reading through small files normally causes lots of seeks and lots of hopping from datanode to datanode to retrieve each small file, all of which is an inefficient data access pattern.

Problems with small files and MapReduce

Map tasks usually process a block of input at a time (using the default FileInputFormat). If the file is very small and there are a lot of them, then each map task processes very little input, and there are a lot more map tasks, each of which imposes extra bookkeeping overhead. Compare a 1GB file broken into 16 64MB blocks, and 10,000 or so 100KB files. The 10,000 files use one map each, and the job time can be tens or hundreds of times slower than the equivalent one with a single input file.

There are a couple of features to help alleviate the bookkeeping overhead: task JVM reuse for running multiple map tasks in one JVM, thereby avoiding some JVM startup overhead (see the `mapred.job.reuse.jvm.num.tasks` property), and `MultiFileInputSplit` which can run more than one split per map.

#### 15. Why Hadoop is less costly

Because hadoop does not need high end computing servers, it relies on large number of commodity hardware.

#### 16. Explain Rack Awareness in Hadoop

Data Replications happen in different racks. For example if there are 2 racks then Hadoop Framework makes sure one replication is at least available in each rack.

#### 17. Explain the Daemons of Hadoop

Various daemons of Hadoop are:

**NameNode**- It is also known as Master in Hadoop cluster. It stores meta-data i.e. number of blocks, their location, replicas and other details. NameNode maintains and manages the slave nodes, and assigns tasks to them. It should be deployed on reliable hardware as it is the centerpiece of HDFS.

**Secondary NameNode**- It downloads the `FsImage` and `EditLogs` from the NameNode. Then it merges `EditLogs` with the `FsImage` periodically. It keeps edits log size within a limit. Then it stores the modified `FsImage` into persistent storage. So we can use `FsImage` in case of NameNode failure.

**DataNode**- It is also known as Slave node. It is responsible for storing actual data in HDFS.

DataNode performs read and write operations as per request for the clients.

**Node Manager**- It is the per-machine/per-node framework agent. It is responsible for containers, monitoring their resource usage and reporting the same to the Resource manager.

**ResourceManager**- The YARN Resource Manager Service (RM) is the central controlling authority for resource management and makes allocation decisions. ResourceManager has two main components: Scheduler and ApplicationsManager.

#### 18. What are 4 configuration files in Hadoop

<https://www.edureka.co/blog/hadoop-cluster-configuration-files/>

- `hadoop-env.sh`
- `core-site.xml`
- `hdfs-site.xml`
- `mapred-site.xml`
- `masters`
- `slaves`

`HADOOP_HOME` directory (the extracted directory etc) is called as `HADOOP_HOME`. e.g. `hadoop-2.6.0-cdh5.5.1` contain all the libraries, scripts, configuration files, etc.

hadoop-env.sh

1. This **file** specifies environment variables that affect the JDK used **by** Hadoop Daemon (**bin/hadoop**).

**As** Hadoop framework **is** written **in Java** and uses **Java** Runtime environment, one **of** the important environment variables **for** Hadoop daemon **is** \$JAVA\_HOME **in** hadoop-env.sh.

2. This **variable** directs Hadoop daemon **to** the **Java path in** the system

Actual: export JAVA\_HOME=<path-to-the-root-of-your-Java-installation>

Change: export JAVA\_HOME=</usr/lib/jvm/java-8-oracle/>

core-site.sh

3. This **file** informs Hadoop daemon **where** NameNode runs **in** the **cluster**. It **contains** the configuration settings **for** Hadoop Core such **as** I/O settings that **are** common **to** HDFS **and** MapReduce.

```
<property>
<name>fs.defaultFS</name>
<value>hdfs://localhost:9000</value>
</property>
<property>
<name>hadoop.tmp.dir</name>
<value>/home/dataflair/Hadmin</value>
</property>
```

☐ Location **of** namenode **is** specified **by** fs.defaultFS property

☐ namenode running **at** 9000 port **on** localhost.

☐ hadoop.tmp.dir property **to** specify the location **where** temporary **as** well **as** permanent **data of** Hadoop will be stored.

☐ "/home/dataflair/hadmin" **is** my location; here you need **to** specify a location **where** you have **Read Write privileges**.

hdfs-site.sh

☐ we need **to** make changes **in** Hadoop configuration **file** hdfs-site.xml (which **is** located **in** HADOOP\_HOME/etc/hadoop) **by** executing the below command:

Hdata@ubuntu:~/hadoop-2.6.0-cdh5.5.1/etc/hadoop\$ nano hdfs-site.xml

Replication factor

```
<property>
<name>dfs.replication</name>
<value>1</value>
</property>
```

☐ Replication factor **is** specified **by** dfs.replication property;

☐ **as** it **is** a single node **cluster** hence we will **set** replication **to** 1.

mapred-site.xml

☐ we need **to** make changes **in** Hadoop configuration **file** mapred-site.xml (which **is** located **in** HADOOP\_HOME/etc/hadoop)

☐ Note: **In order to** edit mapred-site.xml **file** we need **to first create** a copy **of file**

mapred-site.xml.template. A copy **of** this **file** can be created **using** the following command:

Hdata@ubuntu:~/hadoop-2.6.0-cdh5.5.1/etc/hadoop\$ cp mapred-site.xml.template mapred-site.xml

☐ We will now edit the mapred-site.xml **file by using** the following command:

Hdata@ubuntu:~/hadoop-2.6.0-cdh5.5.1/etc/hadoop\$ nano mapred-site.xml

Changes

```
<property>
<name>mapreduce.framework.name</name>
<value>yarn</value>
</property>
```

**In order to** specify which framework should be used **for** MapReduce, we **use** mapreduce. framework.name property, yarn **is** used here.

yarn-site.xml

Changes

```

<property>
<name>yarn.nodemanager.aux-services</name>
<value>mapreduce_shuffle</value>
</property>
<property>
<name>yarn.nodemanager.aux-services.mapreduce.
shuffle.class</name>
<value>org.apache.hadoop.mapred.ShuffleHandler</value>
</property>

```

□ In order to specify auxiliary service need to run with nodemanager "yarn.nodemanager.aux-services" property is used.

□ Here Shuffling is used as auxiliary service. And in order to know the class that should be used for shuffling we use "yarn.nodemanager.aux-services.mapreduce.shuffle.class"

#### 19. Commands

- a. copyFromLocal
- b. moveFromLocal
- c. put
- d. **get**
- e. copyToLocal
- f. moveToLocal
- g. **get**
- h. put
- i. mkdir
- j. ls
- h. append
- i. setrep
- j. mv
- k. put
- l. rm
- m. fsck

#### 20. What do you know about Speculative Execution

In MapReduce, jobs are broken into tasks and the tasks are run in parallel to make the overall job execution time smaller than it would otherwise be if the tasks ran sequentially. Now among the divided tasks, if one of the tasks take more time than desired, then the overall execution time of job increases.

Tasks may be slow for various reasons:

Including hardware degradation or software misconfiguration, but the causes may be hard to detect since the tasks may be completed successfully, could be after a longer time than expected. Apache Hadoop does not fix or diagnose slow-running tasks. Instead, it tries to detect when a task is running slower than expected and launches another, equivalent task as a backup (the backup task is called as speculative task). This process is called Speculative execution in MapReduce.

Speculative execution in Hadoop does not imply that launching duplicate tasks at the same time so they can race. As this will result in wastage of resources in the cluster. Rather, a speculative task is launched only after a task runs for the significant amount of time and framework detects it running slow as compared to other tasks, running for the same job.

When a task successfully completes, then duplicate tasks that are running are killed since they are no longer needed.

If the speculative task after the original task, then kill the speculative task. on the other hand, if the speculative task finishes first, then the original one is killed. Speculative execution in Hadoop is just an optimization, it is not a feature to make jobs run more reliably.

So if, I summarize:

The speed of MapReduce job is dominated by the slowest task. MapReduce first detects slow tasks. Then, run redundant (speculative) tasks. This will optimistically commit before the corresponding stragglers. This process is known as speculative execution. Only one copy of a straggler is allowed to be speculated. Whichever copy (among the two copies) of a task commits first, it becomes the definitive copy, and the other copy is killed by the framework.



- What **is default size of** block?
- **64** MB
- what **is default** replication factor?
- **3**
- Shell we Increase the hdfs block **size**?
- Yes, Multiples **of 64** MB
- Where** will store namenode meta **data**?
- Master(editlogs,fsimage)
- What **are** the daemon processes will be there **in** hdfs?
- NameNode, DataNode **and** SecondaryNameNode
- Different modes **of** Hadoop ?
- **Local**, pseudo distribution **and** distribution
- Hdfs command **to format** namenode?
- Hadoop namenode **-format**
- What **is** namenode **and** jobtracker UI port?
- Namenode UI - **50070** , Jobtracker - **50030**
  
- Where the main hdfs configuration files?
- core-site.xml, hdfs-site.xml, mapred-site.xml
- Where** we will configure blocksize, replication?
- hdfs-site.xml
- If** we **get any** errors **while** loading hdfs **data where** we need **to check**?
- **In** logfiles
- What **are** Hadoop Basic components?
- HDFS **and** MapReduce
- Main Features **of** Hadoop?
- Horizontal Scalability, Distributed Storage, Fault Tolerance
- How much storage needed **for 1TB data with** replication **3** ?
- **3TB**
- How many blocks will **create for 10 GB file with** block **size 64** MB?
- **160**
- what **is** the command **for file** checking
- Hadoop fsck

'MR'

1. **In Map** Reduce ideally how many mappers should be configured **on** a slave

You cannot change the **no of** Mappers via **new Java** API, because we **are using** Job **class in** MapReduce configuration core. **In old** API(deprecated), we can **set no of** mappers **using** setNumMapTasks(**int** n) methods via the JobConf **object**. Ideally, this **is not** the best way **to** **set/change** the **no of** mappers.

**By default**, **no of** mappers **are 2** **on each** slave-node. We can **set/change** this **value using** mapreduce.tasktracker.map.tasks.maximum **parameter**. You need **to set** this **parameter in** mapred-site.xml **file**. We should **not** directly **select** random **value to set** the **no of** mappers.

Ideally **for each** logical InputSplit, a independent mapper **or map dynamic** container will **get** invoked. **If** we **go with default case**, **on each** particular slave node, Node-manager can run **only** two mappers **or map dynamic** containers parallelly irrespective **of** logical **input** splits. **Initially** two **input** split **are** assigned **to two map dynamic** containers **on** slave-node1. **then** the remaining **input** split might be **in** a queue. **In some** cases, these **input** splits might got traveled **to some** other

slave-node(Let's say SN2) which is having map dynamic container sitting idle. This mapper can process the traveled input-split on this slave-node (SN2).

Even though if you specified 2 value(No of mappers) in configuration file. Node-manager doesn't invoke **all** mappers parallelly. This decision **is** taken care **by Resource** Manager based **on** the **input** split(s) available **on** a particular slave-node. But that slave-node can run maximum **2 map dynamic** container parallelly.

Please **go** through below one, so that you can come **to** know how many **no of** maximum mapper we need **to set in order to** get optimize solution **on** a particular slave node.

**When** you **are** setting up the **cluster**, **at that time** you should decide how many maximum **no of** mappers that should be configured/run parallelly **on all** slaves-nodes. Basically, **no of** mappers **are** decided based **on** the below two factors, that **is**,

- 1) **No of** cores
- 2) Ram memory



Lets say we have 10 cores on your system. we can have 10 mappers (One mapper = one core) if go with one core for each mapper. Each mapper/map dynamic container can run on one core. This case might not be true in all cases.

Let's say you have 10 cores on your slave-node, and ram memory is 25GB. Your job need 5GB of memory, so every map tasks requires 5GB of ram. You will have 5 cores on each slave-node. So that we can run maximum 5 mappers parallely. On slave-node, it doesn't have enough memory to run more than 5 mappers parallely even though we have more no of cores available on slave-node. In this case, maximum no of mappers are limited by amount of ram available in your systems. It is not limited by cores available in your system.

If your job required, every map tasks to be loaded with 5GB of memory, then you are wasting cores if you are having 10 cores on each slave-nodes. Here we are using only 5 cores on each slave-node, remaining 5 cores are not utilized. Either go with "10 cores with 50GB memory" or "5 cores with 25GB of ram memory". This will gives the optimal usage of resources.

In general, for each mapper, we will go with 1 to 1.5 core processor. If the usage/processing is very small/light, then go with 1 core processor for each map dynamic container. If the usage/processing is very heavy, then go with 1.5 core processor for each map dynamic container. And also you should the keep above two factors in mind to serve the optimized solution.

## 2. How to set no of Mappers in Map Reduce

In a Single word, no we cannot change the number of Mappers in MapReduce job but we can configure Reducers as per our requirement.

And the number of Mappers depends upon the number of InputSplit and this InputSplit depends on Size of your files and Block size. For example- If we have the block size of 128MB, then the number of mappers will be approximately 4.

## 3. Where is output of Mappers Stored

disk

The output of the maps jobs is stored in the local disk of the mappers. Once the map job finishes these local outputs are then transferred to reducers. You can check your \$ HADOOP\_HOME/conf/mapred-site.xml to check where mapper outputs are stored.

## 4. What is Partitioner and Combiner

A partitioner divides the intermediate key-value pairs produced by map tasks into partition. The total number of partition is equal to the number of reducers where each partition is processed by the corresponding reducer. The partitioning is done using the hash function based on a single key or group of keys. The default partitioner available in Hadoop is HashPartitioner.

In Hadoop MapReduce concept, we have a class in between Mapper and Reducer, called Combiner. When a MapReduce (MR) job is run on a large dataset, Map task generates huge chunks of intermediate data, which is passed on to Reduce task. During this phase, the output from Mapper has to travel over the network to the node where Reducer is running. This data movement may cause network congestion if the data is huge. To reduce this network congestion, MR framework provides a function called 'Combiner', which is also called as 'Mini-Reducer'. The role of Combiner is to take the output of Mapper as it's input, process it and sends its output to the reducer. Combiner reads each key-value pair, combines all the values for the same key, and sends this as input to the reducer, which reduces the data movement in the network. Combiner works along with each Mapper. Combiner uses same class as Reducer.

## 5. Explain shuffling and sorting

Now, the output is Shuffled to the reduce node (which is a normal slave node but reduce phase will run here hence called as reducer node). The shuffling is the physical movement of the data which is done over the network. Once all the mappers are finished and their output

is shuffled on the reducer nodes, then this intermediate output is merged and sorted, which is then provided as input to reduce phase. Follow this comprehensive guide to read more about Shuffling and Sorting in Hadoop MapReduce.

## 6. Explain input split

It is created by InputFormat, logically represent the data which will be processed by an individual Mapper (We will understand mapper below). One map task is created for each split; thus the number of map tasks will be equal to the number of InputSplits. The split is divided into records and each record will be processed by the mapper. Learn MapReduce InputSplit in detail.

## 7. Explain Record Reader

It communicates with the InputSplit in Hadoop MapReduce and converts the data into key-value pairs suitable for reading by the mapper. By default, it uses TextInputFormat for converting data into a key-value pair. RecordReader communicates with the InputSplit until the file reading is not completed. It assigns byte offset (unique number) to each line present in the file. Further, these key-value pairs are sent to the mapper for further processing.

## 8. Explain Reducer

It takes the set of intermediate key-value pairs produced by the mappers as the input and then runs a reducer function on each of them to generate the output. The output of the reducer is the final output, which is stored in HDFS. Follow this link to learn about Reducer in detail.

## 9. Is map only job possible

For Example Sqoop Job

In Hadoop, Map-Only job is the process in which mapper does all task, no task is done by the reducer and mapper's output is the final output. In this tutorial on Map only job in Hadoop MapReduce, we will learn about MapReduce process, the need of map only job in Hadoop, how to set a number of reducers to 0 for Hadoop map only job. We will also learn what are the advantages of Map Only job in Hadoop MapReduce, processing in Hadoop without reducer along with MapReduce example with no reducer. Learn how to install Hadoop 2 with Yarn on pseudo distributed mode

MapReduce is a software framework for easily writing applications that process the vast amount of structured and unstructured data stored in the Hadoop Distributed Filesystem (HDFS). Two important tasks done by MapReduce algorithm are: Map task and Reduce task. Hadoop Map phase takes a set of data and converts it into another set of data, where individual element are broken down into tuples (key/value pairs). Hadoop Reduce phase takes the output from the map as input and combines those data tuples based on the key and accordingly modifies the value of the key.

From the above word-count example, we can say that there are two sets of parallel process, map and reduce; in map process, the first input is split to distribute the work among all the map nodes as shown in a figure, and then each word is identified and mapped to the number 1. Thus the pairs called tuples (key-value) pairs. In the first mapper node three words lion, tiger, and river are passed. Thus the output of the node will be three key-value pairs with three different keys and value set to 1 and the same process repeated for all nodes. These tuples are then passed to the reducer nodes and partitioner comes into action. It carries out shuffling so that all tuples with the same key are sent to the same node. Thus, in reduce process basically what happens is an aggregation of values or rather an operation on values that share the same key. Now, let us consider a scenario where we just need to perform the operation and no aggregation required, in such case, we will prefer 'Map-Only' job in Hadoop. In Hadoop Map-Only job, the map does all task with its InputSplit and no job is done by the reducer. Here map output is the final output. Refer this guide to learn Hadoop features and design principles.

Advantages of Map only job in Hadoop In between map and reduces phases there is key, sort and shuffle phase. Sort and shuffle are responsible for sorting the keys in ascending order

and then grouping values based on same keys. This phase is very expensive and if reduce phase is not required we should avoid it, as avoiding reduce phase would eliminate sort and shuffle phase as well. This also saves network congestion as in shuffling, an output of mapper travels to reducer and when data size is huge, large data needs to travel to the reducer. Learn more about shuffling and sorting in Hadoop MapReduce. The output of mapper is written to local disk before sending to reducer but in map only job, this output is directly written to HDFS. This further saves time and reduces cost as well. Also, there is no need of partitioner and combiner in Hadoop Map Only job that makes the process fast.

#### 10. Explain Distributed Cache

In Hadoop, data chunks process independently in parallel among DataNodes, using a program written by the user. If we want to access some files from all the DataNodes, then we will put that file to Distributed Cache.

MapReduce framework provides a service called Distributed Cache to caches files needed by the applications. It can cache read-only text files, archives, jar files etc.

First of all, an application which need to use distributed cache to distribute a file:

Should make sure that the file is available.

And also make sure that file can accessed via urls. Urls can be either hdfs: // or http:// or it can be file://

Now, if the file is present on the mentioned urls. The user mentions it to be a cache file to the distributed cache. MapReduce job will copy the cache file on all the nodes before starting of tasks on those nodes.

Process as Follows:

Copy the requisite file to the HDFS: \$ hdfs dfs put /user/dataflair/lib/jar\_file.jar

Setup the application's JobConf: DistributedCache.addFileToClasspath(new Path ("/user/dataflair/lib/jar-file.jar"), conf).

Add it in Driver class.

#### 11. Write a word count problem in Map reduce

```
package co.edureka.mapreduce;
import java.io.IOException;
import java.util.StringTokenizer;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Mapper;
import org.apache.hadoop.mapreduce.Reducer;
import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.lib.input.TextInputFormat;
import org.apache.hadoop.mapreduce.lib.output.TextOutputFormat;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
import org.apache.hadoop.fs.Path;

public class WordCount
{
    public static class Map extends Mapper<LongWritable,Text,Text,IntWritable> {
        public void map(LongWritable key, Text value,Context context) throws IOException,
        InterruptedException{
            String line = value.toString();
            StringTokenizer tokenizer = new StringTokenizer(line);
            while (tokenizer.hasMoreTokens()) {
                value.set(tokenizer.nextToken());
                context.write(value, new IntWritable(1));
            }
        }
    }
}
```

```

}
}

```

```

public static class Reduce extends Reducer<Text,IntWritable,Text,IntWritable> {
public void reduce(Text key, Iterable<IntWritable> values,Context context) throws IOException,
InterruptedException {
int sum=0;
for(IntWritable x: values)
{
sum+=x.get();
}
context.write(key, new IntWritable(sum));
}
}

public static void main(String[] args) throws Exception {

Configuration conf= new Configuration();
Job job = new Job(conf,"My Word Count Program");
job.setJarByClass(WordCount.class);
job.setMapperClass(Map.class);
job.setReducerClass(Reduce.class);
job.setOutputKeyClass(Text.class);
job.setOutputValueClass(IntWritable.class);
job.setInputFormatClass(TextInputFormat.class);
job.setOutputFormatClass(TextOutputFormat.class);
Path outputPath = new Path(args[1]);
//Configuring the input/output path from the filesystem into the job
FileInputFormat.addInputPath(job, new Path(args[0]));
FileOutputFormat.setOutputPath(job, new Path(args[1]));
//deleting the output path automatically from hdfs so that we don't have to delete it explicitly
outputPath.getFileSystem(conf).delete(outputPath);
//exiting the job only if the flag value becomes false
System.exit(job.waitForCompletion(true) ? 0 : 1);
}
}
'

```

'KAFKA' <https://mindmajix.com/apache-kafka-interview-questions>

<https://data-flair.training/blogs/kafka-interview-questions/>

## 1. Explain Different components of KAFKA

Kafka Interview Questions- Components of Kafka

Topic -

Kafka Topic is the bunch or a collection of messages.

Producer -

In Kafka, Producers issue communications as well as publishes messages to a Kafka topic.

Consumer -

Kafka Consumers subscribes to a topic(s) and also reads and processes messages from the topic(s).

Brokers -

While it comes to manage storage of messages in the topic(s) we use Kafka Brokers.

## 2. Explain role of offset in Kafka

There is a sequential ID number given to the messages in the partitions what we call, an offset. So, to identify each message in the partition uniquely, we use these offsets.

### 3. Explain consumer group

The concept of Consumer Groups is exclusive to Apache Kafka. Basically, every Kafka consumer group consists of one or more consumers that jointly consume a set of subscribed topics.

### 4. Explain role of zookeeper

Apache Kafka is a distributed system is built to use Zookeeper. Although, Zookeeper's main role here is to build coordination between different nodes in a cluster. However, we also use Zookeeper to recover from previously committed offset if any node fails because it works as periodically commit offset

### 5. Explain the term of leader and follower in Kafka Environment

In every partition of Kafka, there is one server which acts as the Leader, and none or more servers plays the role as a Followers.

### 6. Why Replications are important in Kafka

Because of Replication, we can be sure that published messages are not lost and can be consumed in the event of any machine error, program error or frequent software upgrades.

### 7. Explain Kafka Architecture

Apache Kafka APIs

#### a. Producer API

In order to publish a stream of records to one or more Kafka topics, the Producer API allows an application.

#### b. Consumer API

This API permits an application to subscribe to one or more topics and also to process the stream of records produced to them.

#### c. Streams API

Moreover, to act as a stream processor, consuming an input stream from one or more topics and producing an output stream to one or more output topics, effectively transforming the input streams to output streams, the streams API permits an application.

#### d. Connector API

While it comes to building and running reusable producers or consumers that connect Kafka topics to existing applications or data systems, we use the Connector API. For example, a connector to a relational database might capture every change to a table.

#### a. Kafka Broker

Basically, to maintain load balance Kafka cluster typically consists of multiple brokers. However, these are stateless, hence for maintaining the cluster state they use ZooKeeper. Although, one Kafka Broker instance can handle hundreds of thousands of reads and writes per second. Whereas, without performance impact, each broker can handle TB of messages. In addition, make sure ZooKeeper performs Kafka broker leader election.

#### b. Kafka - ZooKeeper

For the purpose of managing and coordinating, Kafka broker uses ZooKeeper. Also, uses it to notify producer and consumer about the presence of any new broker in the Kafka system or failure of the broker in the Kafka system. As soon as Zookeeper send the notification regarding presence or failure of the broker then producer and consumer, take the decision and starts coordinating their task with some other broker.

#### c. Kafka Producers

Further, Producers in Kafka push data to brokers. Also, all the producers search it and automatically sends a message to that new broker, exactly when the new broker starts. However, keep in mind that producer sends messages as fast as the broker can handle, it doesn't wait for acknowledgments from the broker.

#### d. Kafka Consumers

Basically, **by using partition** offset the Kafka Consumer maintains that how many messages have been consumed because Kafka brokers **are** stateless. Moreover, you can assure that the consumer has consumed **all prior** messages once the consumer acknowledges a particular message offset. Also, **in order to** have a buffer **of** bytes ready **to** consume, the consumer issues an asynchronous pull request **to** the broker. **Then** simply **by** supplying an offset **value**, consumers can rewind **or** skip **to any point in** a **partition**. **In** addition, ZooKeeper notifies Consumer offset **value**.

#### 4. Kafka Fundamental Concepts

Here, we **are** listing **some of** the fundamental concepts **of** Kafka Architecture that you must know:

##### a. Kafka Topics

The topic **is** a logical channel **to** which producers publish message **and from** which the consumers receive messages.

A topic defines the stream **of** a particular **type/classification of data**, **in** Kafka. Moreover, here messages **are** structured **or** organized. A particular **type of** messages **is** published **on** a particular topic.

Basically, **at first**, a producer writes its messages **to** the topics. **Then** consumers **read** those messages **from** topics.

**In** a Kafka **cluster**, a topic **is identified by** its name **and** must be **unique**.

There can be **any number of** topics, there **is no** limitation.

We can **not** change **or** update data, **as soon as** it gets published.

##### b. Partitions **in** Kafka

**In** a Kafka **cluster**, Topics **are** split **into** Partitions **and** also replicated across brokers.

However, **to** which **partition** a published message will be written, there **is no** guarantee about that.

Also, we can **add** a **key** **to** a message. Basically, we will **get** ensured that **all** these messages (**with** the same **key**) will **end up in** the same **partition if** a producer publishes a message **with** a **key**. Due **to** this feature, Kafka offers message sequencing guarantee.

Though, unless a **key is** added **to** it, **data is** written **to** partitions randomly.

Moreover, **in** one **partition**, messages **are** stored **in** the sequenced fashion.

**In** a **partition**, **each** message **is** assigned an incremental id, also called offset.

However, **only** within the **partition**, these offsets **are** meaningful. Moreover, **in** a topic, it does **not** have **any value** across partitions.

There can be **any number of** Partitions, there **is no** limitation.

##### c. Topic Replication Factor **in** Kafka

**While** designing a Kafka system, it's always a wise decision **to** factor **in** topic replication. **As** a **result**, its topics' replicas **from** another broker can solve the crisis, **if** a broker goes down. **For** example, we have **3** brokers **and** **3** topics. Broker1 has Topic **1** **and** Partition **0**, its replica **is in** Broker2, so **on and** so forth. It has got a replication factor **of** **2**; it means it will have one additional copy other **than** the **primary** one. Below **is** the image **of** Topic Replication Factor:

**Some key points -**

Replication takes place **in** the **partition level only**.

**For** a given **partition**, **only** one broker can be a leader, **at** a **time**. Meanwhile, other brokers will have **in-sync** replica; what we **call** ISR.

It **is not** possible **to** have the **number of** replication factor more **than** the **number of** available brokers.

##### d. Consumer **Group**

It can have multiple consumer process/instance running.

Basically, one consumer **group** will have one **unique group-id**.

Moreover, exactly one consumer instance **reads** the **data from** one **partition in** one consumer **group**, **at** the **time of** reading.

Since, there **is** more **than** one consumer **group**, **in** that **case**, one instance **from each of**

these groups can **read from** one single **partition**.

However, there will be **some** inactive consumers, **if** the **number of** consumers exceeds the **number of** partitions. Let's understand it **with** an example **if** there **are** **8** consumers **and** **6** partitions **in** a single consumer **group**, that means there will be **2** inactive consumers.  
**Read Apache Kafka + Spark Streaming Integration**

So, this was **all** about Apache Kafka Architecture. Hope you **like** our explanation.

## 8. Explain Partitioning **Key**

Kafka topics **are** divided **into** a **number of** partitions, which **contains** messages **in** an unchangeable **sequence**. **Each** message **in** a **partition is** assigned **and identified by** its **unique** offset. A topic can also have multiple **partition** logs **like** the click-topic has **in** the image **to** the **right**. This allows **for** multiple consumers **to read from** a topic **in** parallel.

**In** Kafka, replication **is** implemented **at** the **partition level**. The redundant unit **of** a topic **partition is** called a replica. **Each partition** usually has one **or** more replicas meaning that partitions contain messages that **are** replicated over a few Kafka brokers **in** the **cluster**. **As** we can see **in** the pictures - the click-topic **is** replicated **to** Kafka node **2 and** Kafka node **3**.

### Apache Kafka **Partition**

It's possible for the producer to attach a key to the messages and tell which partition the message should go to. All messages with the same key will arrive at the same partition.'

## 9. Advantages **of** Kafka

### High-throughput

We **do not** need **any large** hardware **in** Kafka, because it **is** capable **of** handling high-velocity **and** high-volume **data**. Moreover, it can also support message throughput **of** thousands **of** messages per **second**.

### Low Latency

Kafka can easily handle these messages **with** the very low latency **of** the **range of** milliseconds, demanded **by** most **of** the **new use** cases.

### Fault-Tolerant

Kafka **is** resistant **to** node/machine failure within a **cluster**.

### Durability

**As** Kafka supports messages replication, so, messages **are** never lost. It **is** one **of** the reasons behind durability.

### Scalability

Kafka can be scaled-out, **without** incurring **any** downtime **on** the fly **by** adding additional nodes.

## 10. Explain (Same as 1)

- Producer
- Consumer
- Broker
- topic
- partition**

## 11. Main components **where** the **data is** processed seamlessly **in** kakka

## 12. **Difference between** Kafka **and** flume **and** Why Kafka **is** better **than** flume

Flume **and** Kakfa **both** can act **as** the event backbone **for real-time** event processing. **Some** features **are** overlapping **between** the two **and** there **are some** confusions about what should be used **in** what **use** cases. This post tries **to** elaborate **on** the pros **and** cons **of both** products **and** the **use** cases that they fit the best.

Flume **and** Kafka **are** actually two quite different products. Kafka **is** a **general** purpose publish-subscribe model messaging system, which offers strong durability, scalability **and** fault-tolerance support. It **is not** specifically designed **for** Hadoop. Hadoop ecosystem **is**



just be one of its possible consumers.

Flume is a distributed, reliable, and available system for efficiently collecting, aggregating, and moving large amounts of data from many different sources to a centralized data store, such as HDFS or HBase. It is more tightly integrated with Hadoop ecosystem. For example, the flume HDFS sink integrates with the HDFS security very well. So its common use case is to act as a data pipeline to ingest data into Hadoop.

Kafka is very scalable. One of the key benefits of Kafka is that it is very easy to add large number of consumers without affecting performance and without down time. That's because Kafka does not track which messages in the topic have been consumed by consumers. It simply keeps all messages in the topic within a configurable period. It is the consumers' responsibility to do the tracking through offset. In contrast, adding more consumers to Flume means changing the topology of Flume pipeline design, replicating the channel to deliver the messages to a new sink. It is not really a scalable solution when you have huge number of consumers. Also since the flume topology needs to be changed, it requires some down time.

Kafka's scalability is also demonstrated by its ability to handle spike of the events. This is where Kafka truly shines because it acts as a "shock absorber" between the producers and consumers. Kafka can handle events at 100k+ per second rate coming from producers. Because Kafka consumers pull data from the topic, different consumers can consume the messages at different pace. Kafka also supports different consumption model. You can have one consumer processing the messages at real-time and another consumer processing the messages in batch mode. On the contrary, Flume sink supports push model. When event producers suddenly generate a flood of messages, even though flume channel somewhat acts as a buffer between source and sink, the sink endpoints might still be overwhelmed by the write operations.

Message durability is also an important consideration. Flume supports both ephemeral memory-based channel and durable file-based channel. Even when you use a durable file-based channel, any event stored in a channel not yet written to a sink will be unavailable until the agent is recovered. Moreover, the file-based channel does not replicate event data to a different node. It totally depends on the durability of the storage it writes upon. If message durability is crucial, it is recommended to use SAN or RAID. Kafka supports both synchronous and asynchronous replication based on your durability requirement and it uses commodity hard drive.

Flume does have some features that makes it attractive to be a data ingestion and simple event processing framework. The key benefit of Flume is that it supports many built-in sources and sinks, which you can use out of box. If you use Kafka, most likely you have to write your own producer and consumer. Of course, as Kafka becomes more and more popular, other frameworks are constantly adding integration support for Kafka. For example, Apache Storm added Kafka Spout in release 0.9.2, allowing Storm topology to consume data from Kafka 0.8.x directly.

Kafka does not provide native support for message processing. So mostly likely it needs to integrate with other event processing frameworks such as Apache Storm to complete the job. In contrast, Flume supports different data flow models and interceptors chaining, which makes event filtering and transforming very easy. For example, you can filter out messages that you are not interested in the pipeline first before sending it through the network for obvious performance reason. However, It is not suitable for complex event processing, which I will address in a future post.

The good news is that the latest trend is to use both together to get the best of both worlds. For example, Flume in CDH 5.2 starts to accept data from Kafka via the KafkaSource and push to Kafka using the KafkaSink. Also CDH 5.3 (the latest release) adds Kafka Channel support, which addresses the event durability issue mentioned above.

,

#### 14. ISR in Kafka

Basically, a list of nodes that replicate the log is Replicas. Especially, for a particular partition. However, they are irrespective of whether they play the role of the Leader.

In addition, ISR refers to In-Sync Replicas. On defining ISR, it is a set of message replicas that are synced to the leaders.



## 15. Key advantages of Kafka

Advantages of Kafka

Apache Kafka is selected for its strengths in the space of messaging. The following are some of the advantages which Kafka possess, making it ideal for our Data Lake implementation:

High-throughput: Kafka is capable of handling high-velocity and high-volume data using not so large hardware. It is capable of supporting message throughput of thousands of messages per second.

Low latency: Kafka is able to handle these messages with very low latency of the range of milliseconds, demanded by most of new use cases.

Fault tolerant: The inherent capability of Kafka to be resistant to node/machine failure within a cluster.

Durability: The data/messages are persistent on disk, making it durable and messages are also replicated ...

'

## 16. How to create a topic in kafka

```
./kafka-topics.sh --create --zookeeper localhost:2182 --partitions 2 --replication-factor 1
--topic test_20180613
```

## 17. how to start zookeeper

```
bin/zookeeper-server-start.sh config/zookeeper.properties
```

Next, to start the Kafka server: > bin/kafka-server-start.sh config/server.properties

## 18. What is default retention period of Kafka Broker

160 Hours

## 19. How do intergrate Spark Streaming with Kafka

## 20. How to make RDBMS or Producer and RDBMS as consumer

'PIG'

-----

## 1. Difference between PIG and Hive

Language	Pig Latin	SQL-like
Application	Programming purposes	Report creation
Operation	Client Side	Server side
Data support	Semi-structured	Structured
Connectivity	Can be called by other applications	JDBC & BI tool integration

## 2. Explain ( ILLUSTRATE, DESCRIBE, EXPLAIN, Define)

**DUMP** : It helps to display the results on screen.

**DESCRIBE** : It helps to display the schema of a particular relation.

**ILLUSTRATE** : It helps to display step by step execution of a sequence of pig statements

**EXPLAIN** : It helps to display the execution plan for Pig Latin statements.

## 3. What are the Data types available in PIG

Int  
Long  
Float  
Double  
Char array  
Byte array

Complex :  
Bag  
Map  
Tuple

## 4. Explain What are the transformation available in PIG

- a. Distinct
- b. filter
- c. for each
- d. order by
- e. group
- f. cogroup
- g. Join
  - join
  - left outer Join
  - Right outer Join
  - Full outer join
  - cross
- h. limit
- i. Union
- j. split

## 5. Explain Data types available in PIG --Same as 3

## 6. Explain Flatten in PIG

Sometimes there is data in a tuple or a bag and if we want to remove the level of nesting from that data, then Flatten modifier in Pig can be used. Flatten un-nests bags and tuples. For tuples, the Flatten operator will substitute the fields of a tuple in place of a tuple, whereas un-nesting bags is a little complex because it requires creating new tuples.

## 7. How do you process below formats using PIG

- a. JSON

```
ins_json = LOAD 'PIG_SCRIPTS/ins_json' USING JsonLoader
(
  'this:float,
  that:float,
  insight: (
    diff : float,
    percentage_diff : float,
    normalised_diff : float,
    normalised_percentage_diff : float,
    zscore_diff : float,
    zscore_percentage_diff : float,
    normalised_zscore_diff : float,
    normalised_zscore_percentage_diff : float,
```

```
percentile_rank : float)'
);
```

```
json_insign = foreach ins_json generate this,that,insight.diff,insight.percentage_diff,insight.
normalised_diff,insight.normalised_percentage_diff,insight.zscore_diff,insight.
zscore_percentage_diff,insight.normalised_zscore_diff,insight.normalised_zscore_percentage_diff,
insight.percentile_rank;
```

```
dump json_insign;
```

```
store json_insign;
```

b. CSV

```
A = LOAD '/tmp/test.csv' USING PigStorage(',') AS (a:chararray, b:chararray, c:chararray, d:
chararray, e:chararray);
```

```
DUMP A;
```

c. XML

```
hdfs dfs -copyFromLocal customers_data.xml PIG_SCRIPTS/customers_data.xml
```

```
CUSTOMERS_DATA = load 'PIG_SCRIPTS/customers_data.xml' using org.apache.pig.piggybank.storage.
XMLLoader('customer') as (customer:chararray);
```

```
grunt> CUSTOMERS_DATA = load 'PIG_SCRIPTS/customers_data.xml' using org.apache.pig.piggybank.
storage.XMLLoader('customer') as (customer:chararray);
grunt> dump CUSTOMERS_DATA;
```

## 8. Scenerios that we can you PIG

MapReduce **is** a powerful programming model based **on** the principle parallel processing **or** computation **of data**. Hadoop MapReduce gives the programmers the ability **to filter and aggregate data from** HDFS **to** gain business insights **from** big **data**. MapReduce programming can be implemented **using** many conventional programming languages **like Java, Python, C** etc.

**On** the other hand, Apache Pig **is** a platform **for** analyzing **large data sets** containing high-level **language for** expressing **data** analysis programs, coupled **with** infrastructure **for** evaluating these programs. It gives ease **of** programming **to** the developers **by** enabling complex programmatical challenges **to** be written **in** simple **data** flow **sequence and less** complex textual **language**.

Most **of** the jobs can be run **using** Pig **and** Hive but **to** make **use of** the advanced application programming interfaces, developers may look up **to** MapReduce alternatives. **In** certain situations we need MapReduce alternative over Pig **like** below:

- 1) **When** Hadoop developers need definite driver program control **then** they should make **use of** Hadoop MapReduce instead **of** Pig **and** Hive.
- 2) **When** Hadoop developer needs implementing a custom partitioner they choose MapReduce over Pig **and** Hive.
- 3) **If** there already **exists** pre-defined library **of Java** Mappers **or for** a job **then** it **is** a wise decision **to use** Hadoop MapReduce instead **of** Pig **and** Hive.
- 4) Hadoop MapReduce can prove **to** be a better coding approach over Pig **and** Hive **if** the job requires optimization **at** a particular stage **of** processing **by** making the best **use of** tricks **like in-mapper** combining.
- 5) **If** the job has **some** tricky **usage of** Distributed cache (replicated **join**), **cross** products, groupings **or** joins **then** Hadoop MapReduce **is** a better programming approach over Pig

## 9. Explain Tuple ,Bag and Map

Tuple

An ordered **set of** fields **is** what we **call** a tuple.

**For** Example: (Ankit, 32)

Bag  
A collection of tuples is what we call a bag.  
For Example: {(Ankit,32),(Neha,30)}

Map  
A set of key-value pairs is what we call a Map.  
For Example: [ 'name' #'Ankit', 'age' #32]

#### 10. Is PIG case sensitive

PIG key-words are case insensitive but all other elements are case sensitive.

#### 11. Explain Architecture of PIG

##### Pig Latin Scripts

Initially as illustrated in the above image, we submit Pig scripts to the Apache Pig execution environment which can be written in Pig Latin using built-in operators.

There are three ways to execute the Pig script:

Grunt Shell: This is Pig's interactive shell provided to execute all Pig Scripts.

Script File: Write all the Pig commands in a script file and execute the Pig script file. This is executed by the Pig Server.

Embedded Script: If some functions are unavailable in built-in operators, we can programmatically create User Defined Functions to bring that functionalities using other languages like Java, Python, Ruby, etc. and embed it in Pig Latin Script file. Then, execute that script file.

##### Parser

From the above image you can see, after passing through Grunt or Pig Server, Pig Scripts are passed to the Parser. The Parser does type checking and checks the syntax of the script. The parser outputs a DAG (directed acyclic graph). DAG represents the Pig Latin statements and logical operators. The logical operators are represented as the nodes and the data flows are represented as edges.

##### Optimizer

Then the DAG is submitted to the optimizer. The Optimizer performs the optimization activities like split, merge, transform, and reorder operators etc. This optimizer provides the automatic optimization feature to Apache Pig. The optimizer basically aims to reduce the amount of data in the pipeline at any instance of time while processing the extracted data, and for that it performs functions like:

PushUpFilter: If there are multiple conditions in the filter and the filter can be split, Pig splits the conditions and pushes up each condition separately. Selecting these conditions earlier, helps in reducing the number of records remaining in the pipeline.

PushDownForEachFlatten: Applying flatten, which produces a cross product between a complex type such as a tuple or a bag and the other fields in the record, as late as possible in the plan. This keeps the number of records low in the pipeline.

ColumnPruner: Omitting columns that are never used or no longer needed, reducing the size of the record. This can be applied after each operator, so that fields can be pruned as aggressively as possible.

MapKeyPruner: Omitting map keys that are never used, reducing the size of the record.

LimitOptimizer: If the limit operator is immediately applied after a load or sort operator, Pig converts the load or sort operator into a limit-sensitive implementation, which does not require processing the whole data set. Applying the limit earlier, reduces the number of records. This is just a flavor of the optimization process. Over that it also performs Join, Order By and Group By functions.

To shutdown, automatic optimization, you can execute this command:

```
pig -optimizer_off [opt_rule | all ]
Compiler
```

After the optimization process, the compiler compiles the optimized code into a series of MapReduce jobs. The compiler is the one who is responsible for converting Pig jobs automatically into MapReduce jobs.

Execution engine

Finally, as shown in the figure, these MapReduce jobs are submitted for execution to the execution engine. Then the MapReduce jobs are executed and gives the required result. The result can be displayed on the screen using "DUMP" statement and can be stored in the HDFS using "STORE" statement.

After understanding the Architecture, now in this Apache Pig tutorial, I will explain you the Pig Latins's Data Model.

## 12. Use Cases of PIG

MapReduce is a powerful programming model based on the principle parallel processing or computation of data. Hadoop MapReduce gives the programmers the ability to filter and aggregate data from HDFS to gain business insights from big data. MapReduce programming can be implemented using many conventional programming languages like Java, Python, C etc.

On the other hand, Apache Pig is a platform for analyzing large data sets containing high-level language for expressing data analysis programs, coupled with infrastructure for evaluating these programs. It gives ease of programming to the developers by enabling complex programmatical challenges to be written in simple data flow sequence and less complex textual language.

Most of the jobs can be run using Pig and Hive but to make use of the advanced application programming interfaces, developers may look up to MapReduce alternatives. In certain situations we need MapReduce alternative over Pig like below:

- 1) When Hadoop developers need definite driver program control then they should make use of Hadoop MapReduce instead of Pig and Hive.
- 2) When Hadoop developer needs implementing a custom partitioner they choose MapReduce over Pig and Hive.
- 3) If there already exists pre-defined library of Java Mappers or for a job then it is a wise decision to use Hadoop MapReduce instead of Pig and Hive.
- 4) Hadoop MapReduce can prove to be a better coding approach over Pig and Hive if the job requires optimization at a particular stage of processing by making the best use of tricks like in-mapper combining.
- 5) If the job has some tricky usage of Distributed cache (replicated join), cross products, groupings or joins then Hadoop MapReduce is a better programming approach over Pig

## 13. How files are referenced in PIG when schema is not available

```
salgrade = load '/user/cloudera/pig/hr/HR/salgrade' using PigStorage(',');
```

## 14. What are Different in-built functions available in PIG

AVG  
 CONCAT  
 COUNT  
 COUNT\_STAR  
 DIFF  
 IsEmpty  
 MAX  
 MIN  
 SIZE  
 SUM  
 TOKENIZE  
 Load/Store Functions  
 Handling Compression  
 BinStorage  
 PigDump  
 PigStorage  
 TextLoader

## Math Functions

ABS  
 ACOS  
 ASIN  
 ATAN  
 CBRT  
 CEIL  
 COS  
 COSH  
 EXP  
 FLOOR  
 LOG  
 LOG10  
 RANDOM  
 ROUND  
 SIN  
 SINH  
 SQRT  
 TAN  
 TANH

## String Functions

INDEXOF  
 LAST\_INDEX\_OF  
 LCFIRST  
 LOWER  
 REGEX\_EXTRACT  
 REGEX\_EXTRACT\_ALL  
 REPLACE  
 STRSPLIT  
 SUBSTRING  
 TRIM  
 UCFIRST  
 UPPER

## Bag and Tuple Functions

TOBAG  
 TOP  
 TOTUPLE

## 15. Difference between group and cogroup

Group and Cogroup operators are identical. For readability, GROUP is used in statements involving one relation and COGROUP is used in statements involving two or more relations. Group operator collects all records with the same key. Cogroup is a combination of group and join, it is a generalization of a group instead of collecting records of one input depends on a key, it collects records of n inputs based on a key. At a time, we can Cogroup up to 127 relations.

```
cogroup_data = COGROUP emp by DEPTNO, dept by DEPTNO;
```

```
pos = foreach baseball_limit generate name, flatten(position) as position;
bypos = group pos by position;
```

## 16. How to get the metadata

```
describe
```

## 17. UDFx in Pig

## 18. How do you create pig script and run

```
pig -x mapreduce hdfs://localhost:9000/pig_data/Sample_script.pig
```

## 19. How to read and store the data

```
emp = load 'emp.txt' using PigStorage(',');
store emp into 'emp';
```

20. How **do** you store processed **data in** Hive

```
dw_data_set = LOAD 'default.customers' USING org.apache.hive.hcatalog.pig.HCatLoader();
final_data_set= foreach dw_data_set generate customer_id,customer_fname;
create table customers2 (customer_id int ,customer_fname string);
STORE final_data_set INTO 'default.customers2' USING org.apache.hive.hcatalog.pig.HCatStorer();
```

'SQL Questions '

1. What **is** Different types **of SQL Statement**
2. What **are** the different Database objects you know
3. What **is View ? Types ? and** how it **is** different **from Table**
4. What **is** Materialized **view and** What **are** the types **of** refreshed method
5. **Difference between view and MV**
6. What **is Partition** and what **are** different types **of** partion can be added **to table**
7. Explain advantage **of Using** Partitioning **in Oracle**
8. Exaplain **use of** Indexes **and** Different types **of** Indexes
9. **Difference between B-tree and Bitmap Index**
10. What **do** you mean **by local and global index**
11. What **is** **Synonym** and what **are** the types **of** synonoyms
12. What you mean **by** DB-link
13. What **are** the **Data Dictionary** tables avaiialble **in Oracle**
14. What **are** the Different **constraints** available **in Oracle**
15. What **is** different **between Table level and column level constraint**
16. **Use of** Sequences
17. What **is** the Oracle version that you **are** currently **using**
18. Explain
  - a. DDL
  - b. DML
  - c. DRL
  - d. DCL
19. What **are** the pre-defined **data** types avaiialble **in oracle**
  - a. **Character**
  - b. Numeric
  - c. **Date**
  - d. What **are** aggregate function
20. Explain working **of**
  - a. Co-related sub queries
  - b. **group by** query
21. Explain Different types **of** Joins available **in Oracle**
22. How **do** you **delete** duplicates **from** the **table**
23. Explain Locking mechanism **in oracle**
24. Explain **Use of Global Temporary table (GTT)**
25. **Difference between Rank() and Dense\_Rank()**
26. Explain **Use of** RowNumber() **and Rowid**
27. Practice Hierarchiel queries
28. **Use of LISTAGG() Queries -- Practice 3 Queries**
29. **Difference between RowNumber() and rownum**
30. Explain the working **for** B-tree
31. **Difference between Delete, Truncate and Drop**
32. Explain ACID properties
33. Explain **use of Decode() and case**
34. **Difference between SGA and PGA**
35. Explain Complete flow **of**

```
select * from emp ;
```
36. Explain complete working **of**

```
update emp set ename='VISHAL' where empno=7900;
```
37. Explain Merge **Operation in Oracle.**
38. Explain **Current of Operation in Oracle**
39. Explain types **of** Sub-Query **in Oracle**
40. Explain **On Delete null and On delete cascade.**
41. **Difference between varchar vs varchar2 vs Nvarchar2**
42. Explain Pseudo Columns **in Oracle**
43. Explain Sub-partitioning **in Oracle.**

44. Explain
  - a. Hard Parse
  - b. soft parse
45. Explain **with** respect **to** oracle Architecture
  - a. Blocks
  - b. segments
  - c. Extents
  - d. **Data** Files
  - e. Tablespace
46. Various Hints **in** Oracle
47. Page **no 148 to 185**
48. How **do** you **create table** faster **in** Oracle
49. Basic checks you **do to** improve performance **of** query
50. Normalization **and** its Types.
51. Nth Highest Paid Employee
52. Employees **with** Maximum salary **in Each** Department
53. Explain
  - a. **Union**
  - b. **Union all**
  - c. Intersection
  - d. **Minus**
54. **Difference** between user\_\*, all\_\* **and** dba\_\* **data Dictionary** objects
55. Explain **Difference** Keys **in** Oracle

## PL/SQL

-----

1. What **is** the **Use of** PL/SQL ? What **are** the Advantages
2. **Write** an annonyms blocks **to update** an Employee
3. What **are**
  - a. **Procedure**
  - b. Functions
  - c. Packages**and** what **are** scenarios that above **are** used
4. **Difference between** Functions **and** Procedure
5. What **is** context switching
6. What **is Bulk collect and Bulk Exception**  
**And when it is** used **and** what **is** its significance
7. What **is Trigger** and what **are** the different types **of** triggers
8. What **is** mutating **table** error
9. Can we **use commit in trigger** ? Justify the Answer
10. What **is Cursor** and its types
11. Explain Parameterized **cursor**
12. What **is Ref-Cursor**
13. What **are** Excpetion ? List pre-defined **Exception**
14. Explain **Raise** vs **Raise** Application Error
15. **Use of SQLCODE , SQLERRM**
16. How **do** you find the line **no** Error **in** PL/SQL -->DBMS\_SQLBACKTRACE
17. Collections **in** PL/SQL
18. Explain **Pragma** Autonomous **Transaction**
19. **Use of Pragma** Exception\_INT
20. Modes **of** Paramter
  - a. **In**
  - b. **In-out**
  - c. **out**
21. Types **of** Notations
22. Explain Overloaing Procedurs
23. Explain Dynmaic **SQL in** PL/SQL
24. How **do** you perform DDL **in** PL/SQL
25. **Check SQL%ROW\_COUNT Usage in** PL/SQL
26. What **are** PL/SQL Datatypes
27. **Difference between %ROWTYPE AND %TYPE**  
**AND** Explain **both**
28. Practice Example
  - a. **Function**



- b. Procedure
  - c. Package
  - d. Bulk Collect
  - e. Bulk collect with Exception
  - f. Collectiosn
  - g. Cursor
  - h. Excpetion
  - g. Autonomous Transaction
  - h. Dynamic SQL
  - i. IF , IF-ELSE
  - J. for loop
29. Check Error logging mechanism in Exception from Steven Feuerstein.
30. DBMS Scheduler Jobs in Oracle
31. Doing Activities Fast , Read more on it
- a. Create table with parallel 32 and nologging
  - b. Insert /\*+ Append\*/
  - c. create index with parallel 32 and nologging
  - d. Disable any triggers while loading any data into table
  - e. Parallel session using Shell script and primary key columns

'Data Warehouse'

-----

1. What is Surrogate Key
2. What is Normalization and its types
3. What is SCD ? Type 1 and Type 2 Dimention
4. Explain Star Schema
5. Explain Snowflake Schema
6. Explain
  - a. Junk Dimentions
  - b. Confimed Dimensions
  - c. Denerated Dimensions
7. What is ETL